

"Iqra' bismi rabbika allathee khalaq, khalaqal insana min 'alaq."

Simple Guide to PostgreSQL Installation, Replication and Backup

Assumptions:

- PostgreSQL Version: 16 or newer
- Primary Server: 192.168.20.21 (hostname: pg-primary)
- Replica Server: 192.168.20.22 (hostname: pg-replica)
- OS: Ubuntu 24.04LTS (adjustable for RHEL/CentOS/Almalinux)
- Replication User: replicator
- Data Directory: /custom_pg_data
- Network: Both servers can communicate with port 5432
- Disk Setup: External disks /dev/sdb and /dev/sdc for LVM

Step 0: Prepare the Environment (Both Servers)

1. Set Hostnames

On Primary:

```
hostnamectl set-hostname pg-primary; bash
```

On Replica:

```
hostnamectl set-hostname pg-replica; bash
```

2. Update /etc/hosts

On both servers:

```
cat >> /etc/hosts <<EOF
```

```
192.168.20.21      primary.yourdomain.com      pg-primary
```

```
192.168.20.22      replica.yourdomain.com      pg-replica
```

```
EOF
```

3. Set Up LVM for Custom Data Directory

a. Partition External Disks:

```
fdisk /dev/sdb
```

- Press 'n' for new partition, 'p' for primary, accept defaults, then 'w' to write.

```
fdisk /dev/sdc (repeat steps)
```

b. Verify Partitions:

```
root@pg-primary:~# lsblk
```

Expected output:

NAME	MAJ:MIN	RM	SIZE	RO	TYPE	MOUNTPOINTS
sda	8:0	0	50G	0	disk	
└─sda1	8:1	0	1M	0	part	
└─sda2	8:2	0	2G	0	part	/boot
└─sda3	8:3	0	48G	0	part	
└─ubuntu--vg-ubuntu--lv 252:0		0	48G	0	lvm	/
sdb	8:16	0	50G	0	disk	
└─sdb1	8:17	0	50G	0	part	
sdsc	8:32	0	50G	0	disk	
└─sdsc1	8:33	0	50G	0	part	
sr0	11:0	1	2.6G	0	rom	

c. Create LVM:

```
sudo pvcreate /dev/sdb1 /dev/sdsc1
```

```
sudo vgcreate pg_vg /dev/sdb1 /dev/sdsc1
```

```
sudo lvcreate -n pg_lv -L 50G pg_vg
```

d. Format and Mount Logical Volume:

```
sudo mkfs.ext4 /dev/pg_vg/pg_lv
```

```
sudo mkdir -p /custom_pg_data
```

```
sudo mount /dev/pg_vg/pg_lv /custom_pg_data/
```

```
echo '/dev/pg_vg/pg_lv /custom_pg_data ext4 defaults 0 0' | sudo tee -a /etc/fstab
```

e. Verify LVM:

```
sudo lvscan
```

```
sudo blkid /dev/pg_vg/pg_lv
```

f. Activate Logical Volume (if inactive):

```
sudo lvchange -ay /dev/pg_vg/pg_lv
```

```
root@pg-primary:~# lsblk
```

NAME	MAJ:MIN	RM	SIZE	RO	TYPE	MOUNTPOINTS
sda	8:0	0	50G	0	disk	
└─sda1	8:1	0	1M	0	part	
└─sda2	8:2	0	2G	0	part	/boot
└─sda3	8:3	0	48G	0	part	
└─ubuntu--vg-ubuntu--lv 252:0		0	48G	0	lvm	/
sdb	8:16	0	50G	0	disk	
└─sdb1	8:17	0	50G	0	part	
└─pg_vg-pg_lv	252:1	0	50G	0	lvm	/custom_pg_data
sdsc	8:32	0	50G	0	disk	
└─sdsc1	8:33	0	50G	0	part	
└─pg_vg-pg_lv	252:1	0	50G	0	lvm	/custom_pg_data
sr0	11:0	1	2.6G	0	rom	

Step 1: Install PostgreSQL (Both Servers)

1. Update and Install PostgreSQL:

```
sudo apt update -y  
sudo apt upgrade -y  
sudo apt install postgresql postgresql-contrib postgresql-client -y
```

2. Verify Installation:

```
pg_config --version  
  
Expected output: PostgreSQL 16.x
```

3. Stop PostgreSQL:

```
sudo systemctl stop postgresql  
sudo systemctl disable postgresql.service
```

4. Set Up Custom Data Directory:

```
sudo chown postgres:postgres /custom_pg_data  
sudo chmod -R 700 /custom_pg_data  
sudo rm -rf /custom_pg_data/*  
sudo -u postgres /usr/lib/postgresql/16/bin/initdb -D /custom_pg_data
```

5. Update PostgreSQL Configuration:

```
sudo sed -i 's|^#data_directory = .*|data_directory = '/custom_pg_data'|"  
/etc/postgresql/16/main/postgresql.conf
```

6. Modify Authentication:

Edit /custom_pg_data/pg_hba.conf:

Change:

local	all	all	trust
-------	-----	-----	-------

To:

local	all	all	md5
-------	-----	-----	-----

7. Start PostgreSQL and Verify:

```
sudo systemctl start postgresql  
sudo -u postgres psql -c "SHOW data_directory;"
```

Expected output:

```
data_directory  
-----  
/custom_pg_data
```

Step 2: Configure the Primary Server

1. Edit postgresql.conf:

```
sudo vi /etc/postgresql/16/main/postgresql.conf
```

Ensure:

```
listen_addresses = '*'
```

```
wal_level = replica
```

```
max_wal_senders = 10
```

```
wal_keep_size = 512MB
```

```
max_replication_slots = 5
```

```
hot_standby = on
```

```
archive_mode = on
```

```
archive_command = 'cp %p /var/lib/postgresql/wal_archive/%f'
```

Or use sed:

```
sudo sed -i 's/^#listen_addresses = */listen_addresses = "*"/' /etc/postgresql/16/main/postgresql.conf
```

```
sudo sed -i 's/^#wal_level = */wal_level = replica/' /etc/postgresql/16/main/postgresql.conf
```

```
sudo sed -i 's/^#max_wal_senders = */max_wal_senders = 10/' /etc/postgresql/16/main/postgresql.conf
```

```
sudo sed -i 's/^#wal_keep_size = */wal_keep_size = 512MB/' /etc/postgresql/16/main/postgresql.conf
```

```
sudo sed -i 's/^#max_replication_slots = */max_replication_slots = 5/'
```

```
etc/postgresql/16/main/postgresql.conf
```

```
sudo sed -i 's/^#hot_standby = */hot_standby = on/' /etc/postgresql/16/main/postgresql.conf
```

```
sudo sed -i 's/^#archive_mode = */archive_mode = on/' /etc/postgresql/16/main/postgresql.conf
```

```
sudo sed -i 's|^#archive_command = .*|archive_command = 'cp %p
```

```
/var/lib/postgresql/wal_archive/%f'|" /etc/postgresql/16/main/postgresql.conf
```

2. Create WAL Archive Directory:

```
sudo mkdir -p /var/lib/postgresql/wal_archive
```

```
sudo chown postgres:postgres /var/lib/postgresql/wal_archive
```

```
sudo chmod 700 /var/lib/postgresql/wal_archive
```

3. Edit pg_hba.conf:

```
sudo vi /etc/postgresql/16/main/pg_hba.conf
```

Add:

```
host          replication    replicator    192.168.20.22/32    md5
```

Or:

```
echo "host    replication    replicator    192.168.20.22/32    md5" | sudo tee -a  
/etc/postgresql/16/main/pg_hba.conf
```

4. Reload Configuration:

```
sudo systemctl reload postgresql
```

5. Create Replication User:

```
sudo -u postgres psql -c "CREATE USER replicator WITH REPLICATION ENCRYPTED PASSWORD 'replica_pass';"
```

6. Restart PostgreSQL:

```
sudo systemctl restart postgresql
```

7. Verify Replication Setup:

```
sudo -u postgres psql -c "SELECT * FROM pg_stat_replication;"
```

(No output expected yet)

Step 3: Set Up the Replica Server

1. Stop PostgreSQL:

```
sudo systemctl stop postgresql
```

2. Clear Data Directory:

```
sudo rm -rf /custom_pg_data/*
```

3. Take Base Backup from Primary:

```
PGPASSWORD='replica_pass' pg_basebackup -h 192.168.20.21 -D /custom_pg_data -U replicator -Fp -Xs -P -R
```

-R: Creates standby.signal and sets up primary_conninfo

-Xs: Includes WAL files

-Fp: Plain format

4. Set Permissions:

```
sudo chown -R postgres:postgres /custom_pg_data
sudo chown -R postgres:postgres /custom_pg_data/*
sudo chmod -R 700 /custom_pg_data
sudo chmod 600 /custom_pg_data/postgresql.auto.conf
```

5. Verify postgresql.auto.conf:

```
cat /custom_pg_data/postgresql.auto.conf
```

Ensure:

```
primary_conninfo = 'host=192.168.20.21 port=5432 user=replicator password=replica_pass'
```

6. Create WAL Archive Directory:

```
sudo mkdir -p /custom_pg_data/archive
```

```
sudo chown postgres:postgres /custom_pg_data/archive
sudo chmod 700 /custom_pg_data/archive
```

7. Update postgresql.conf:

```
sudo vi /custom_pg_data/postgresql.conf
```

Add/verify:

```
primary_conninfo = 'host=192.168.20.21 port=5432 user=replicator password=replica_pass'
restore_command = 'cp /custom_pg_data/archive/%f %p'
hot_standby = on
```

8. Start PostgreSQL:

```
sudo systemctl start postgresql
sudo systemctl enable postgresql
```

9. Check Logs:

```
sudo journalctl -u postgresql -f
```

Step 4: Monitoring and Verification

On Primary:

```
sudo -u postgres psql -c "SELECT pid, username, client_addr, state, sync_state FROM
pg_stat_replication;"
```

(Expected: Replica (192.168.20.22) with state=streaming)

pid	username	client_addr	state	sync_state
40179	replicator	192.168.20.41	streaming	async

(1 row)

On Replica:

```
sudo -u postgres psql -c "SELECT status, receive_start_lsn, replay_lsn FROM pg_stat_wal_receiver;"
```

(Expected: status=streaming)

status	receive_start_lsn
streaming	0/5000000

(1 row)

Step 5: Test Replication

1. On Primary:

```
sudo -u postgres psql -c "CREATE DATABASE testdb;"
sudo -u postgres psql -d testdb -c "CREATE TABLE test_replication (id SERIAL, name TEXT);"
sudo -u postgres psql -d testdb -c "INSERT INTO test_replication (name) VALUES ('replicated\!');"
```

2. On Replica:

```
sudo -u postgres psql -d testdb -c "SELECT * FROM test_replication;"
```

Expected:

```
id | name
----+-----
  1 | replicated\!
(1 row)
```

Step 6: Configure Replication Slot (Prevents WAL loss)

1. On Primary:

```
sudo -u postgres psql -c "SELECT * FROM pg_create_physical_replication_slot('replica_slot');"
```

2. On Replica:

Edit /custom_pg_data/postgresql.conf:

```
primary_slot_name = 'replica_slot'
```

3. Restart PostgreSQL on Replica:

```
sudo systemctl restart postgresql
```

4. Verify Slot Usage on Primary:

```
sudo -u postgres psql -c "SELECT slot_name, active FROM pg_replication_slots;"
```

Expected: replica_slot with active=true

Final Step: Database Backup using bash Script

```
#!/bin/bash

# === CONFIGURATION ===
DB_NAME="yourdatabase"
BACKUP_DIR="/yourdirectory"
DATE=$(date +"%Y-%m-%d")
BACKUP_FILE="${BACKUP_DIR}/${DB_NAME}_backup_${DATE}.sql.gz" # Adding compression
RETENTION_DAYS=15

# === CHECK IF BACKUP DIRECTORY EXISTS ===
if [ ! -d "$BACKUP_DIR" ]; then
    echo "[ERROR] Backup directory $BACKUP_DIR does not exist: $BACKUP_DIR" >&2
    exit 1
fi

# === CREATE BACKUP ===
echo "[INFO] $(date) - Starting backup for database: $DB_NAME"
su - postgres -c "pg_dump -d $DB_NAME -Fc | gzip > $BACKUP_FILE" # Compress the backup with gzip

if [ $? -eq 0 ]; then
    echo "[INFO] $(date) - Backup successful: $BACKUP_FILE"
else
    echo "[ERROR] $(date) - Backup failed" >&2
    exit 1
fi

# === DELETE OLD BACKUPS (older than $RETENTION_DAYS days) ===
echo "[INFO] $(date) - Deleting backups older than $RETENTION_DAYS days from $BACKUP_DIR"
find "$BACKUP_DIR" -name "${DB_NAME}_backup_*.sql.gz" -type f -mtime +$RETENTION_DAYS -exec
rm -f {} \;

if [ $? -eq 0 ]; then
    echo "[INFO] $(date) - Cleanup complete."
else
    echo "[ERROR] $(date) - Failed to clean up old backups" >&2
    exit 1
fi

#Crontab Entry:
sudo nano /etc/cron.d/db_backup
0 2 * * * postgres /bin/bash /path/to/your/backup_script.sh >> /path/to/backup/logs/backup_log.log
2>&1
```