

# PostgreSQL parallel backups with pg\_dump

Backing up your PostgreSQL database efficiently is very important, especially for large databases. One way to speed up backups is by using...

Backing up your PostgreSQL database efficiently is very important, especially for large databases. One way to speed up backups is by using `pg_dump` in directory format with parallel jobs. This method stores each table in a separate file and allows multiple tables to be dumped simultaneously.

## 1. Checking your data directory

Before starting, let's check where our PostgreSQL data is stored. Typically, the data directory contains multiple subdirectories and configuration files like:

```
[postgres@node1 pgdata]$ ls -rlt /pgdata/

total 72

drwx----- 8 postgres postgres 76 Sep 19 11:28 base
drwx----- 2 postgres postgres 4096 Sep 19 11:30 global
drwx----- 3 postgres postgres 4096 Sep 19 11:31 pg_wal-rw----- 1

postgres postgres 29671 Jul 15 11:01 postgresql.conf-rw----- 1

postgres postgres 5779 Jul 15 10:34 pg_hba.conf

[postgres@node1 pgdata]$
[postgres@node1 pgdata]$ ls
base                pg_commit_ts       pg_logical          pg_serial          pg_subtrans
current_logfiles    pg_dynshmem         pg_multixact        pg_snapshots       pg_tblspc
global              pg_hba.conf         pg_notify           pg_stat            pg_twophase
log                 pg_ident.conf       pg_replslot         pg_stat_tmp        PG_VERSION
```

- 'base/' — contains the actual database files
- 'pg\_wal/' — transaction logs (Write-Ahead Logs)
- 'global/' — cluster-wide metadata

- Configuration files like 'postgresql.conf', 'pg\_hba.conf', and 'PG\_VERSION'

## 2. Creating a parallel directory ackup

To create a backup in directory format with parallel jobs, use the -Fd option with the -j option.

```
Copy[postgres@node1 pgdata]$  
[postgres@node1 pgdata]$ ls -rlt /pgbackup/  
total 0  
[postgres@node1 pgdata]$  
[postgres@node1 pgdata]$  
[postgres@node1 pgdata]$  
[postgres@node1 pgdata]$ pg_dump -U postgres -Fd -j 4 -f /pgbackup/  
postgres -v  
  
pg_dump: last built-in OID is 16383  
pg_dump: reading extensions  
pg_dump: identifying extension members  
pg_dump: reading schemas  
pg_dump: reading user-defined tables  
pg_dump: reading user-defined functions  
pg_dump: reading user-defined types  
pg_dump: reading procedural languages  
pg_dump: reading user-defined aggregate functions  
pg_dump: reading user-defined operators  
pg_dump: reading user-defined access methods  
pg_dump: reading user-defined operator classes  
pg_dump: reading user-defined operator families  
pg_dump: reading user-defined text search parsers  
pg_dump: reading user-defined text search templates  
pg_dump: reading user-defined text search dictionaries  
pg_dump: reading user-defined text search configurations  
pg_dump: reading user-defined foreign-data wrappers  
pg_dump: reading user-defined foreign servers  
pg_dump: reading default privileges  
pg_dump: reading user-defined collations  
pg_dump: reading user-defined conversions  
pg_dump: reading type casts  
pg_dump: reading transforms  
pg_dump: reading table inheritance information  
pg_dump: reading event triggers  
pg_dump: finding extension tables  
pg_dump: finding inheritance relationships  
pg_dump: reading column info for interesting tables  
pg_dump: finding table default expressions  
pg_dump: flagging inherited columns in subtables  
pg_dump: reading partitioning data  
pg_dump: reading indexes  
pg_dump: flagging indexes in partitioned tables  
pg_dump: reading extended statistics  
pg_dump: reading constraints
```

```
pg_dump: reading triggers
pg_dump: reading rewrite rules
pg_dump: reading policies
pg_dump: reading row-level security policies
pg_dump: reading publications
pg_dump: reading publication membership of tables
pg_dump: reading publication membership of schemas
pg_dump: reading subscriptions
pg_dump: reading large objects
pg_dump: reading dependency data
pg_dump: saving encoding = UTF8
pg_dump: saving standard_conforming_strings = on
pg_dump: saving search_path =
pg_dump: saving database definition
pg_dump: dumping contents of table "public.all_university"
pg_dump: dumping contents of table "public.olympic_history"
pg_dump: finished item 4345 TABLE DATA all_university
pg_dump: dumping contents of table "public.google_employees"
pg_dump: finished item 4347 TABLE DATA google_employees
pg_dump: dumping contents of table "public.test_unique"
pg_dump: finished item 4350 TABLE DATA test_unique
pg_dump: dumping contents of table "public.products"
pg_dump: finished item 4355 TABLE DATA products
pg_dump: dumping contents of table "public.groups"
pg_dump: finished item 4357 TABLE DATA groups
pg_dump: dumping contents of table "public.students"
pg_dump: finished item 4359 TABLE DATA students
pg_dump: dumping contents of table "public.netflix_shows"
pg_dump: dumping contents of table "public.fifa_players"
pg_dump: finished item 4348 TABLE DATA fifa_players
pg_dump: finished item 4351 TABLE DATA netflix_shows
pg_dump: finished item 4352 TABLE DATA olympic_history
[postgres@node1 pgdata]$
```

### **pg\_dump options:**

-U postgres → the PostgreSQL user -Fd → directory format (stores each table as a separate file) -j 4 → use 4 parallel jobs to speed up the backup -f /pgbackup/ → backup directory - postgres → the database name

This will create /pgbackup/ containing files for all tables and other database objects. Parallel jobs make the backup faster for large databases.

```
[postgres@node1 pgdata]$
[postgres@node1 pgdata]$ ls -rlt /pgbackup/
total 7216
-rw-r--r-- 1 postgres postgres 16265 Sep 26 14:16 toc.dat
-rw-r--r-- 1 postgres postgres 165844 Sep 26 14:16 4345.dat.gz
-rw-r--r-- 1 postgres postgres 29388 Sep 26 14:16 4347.dat.gz
-rw-r--r-- 1 postgres postgres 75 Sep 26 14:16 4359.dat.gz
-rw-r--r-- 1 postgres postgres 63 Sep 26 14:16 4357.dat.gz
-rw-r--r-- 1 postgres postgres 76 Sep 26 14:16 4355.dat.gz
-rw-r--r-- 1 postgres postgres 43 Sep 26 14:16 4350.dat.gz
-rw-r--r-- 1 postgres postgres 375325 Sep 26 14:16 4348.dat.gz
-rw-r--r-- 1 postgres postgres 1369421 Sep 26 14:16 4351.dat.gz
-rw-r--r-- 1 postgres postgres 5403069 Sep 26 14:16 4352.dat.gz
[postgres@node1 pgdata]$
[postgres@node1 pgdata]$
```

## 2. Restoring from a Parallel Directory Backup

Count postgres DB object list:

```
SELECT
    CASE UPPER(relkind)
        WHEN 'R' THEN 'TABLE'
        WHEN 'P' THEN 'PARTITION TABLE'
        WHEN 'I' THEN 'INDEX'
        WHEN 'S' THEN 'SEQUENCE'
        WHEN 'V' THEN 'VIEW'
        WHEN 'M' THEN 'MATERIALIZED VIEW'
        WHEN 'F' THEN 'FOREIGN TABLE'
        WHEN 'T' THEN 'TOAST TABLE'
        ELSE UPPER(relkind)
    END AS object_type,
    COUNT(*) AS count
FROM pg_class c
JOIN pg_namespace n ON n.oid = c.relnamespace
WHERE n.nspname NOT IN ('pg_catalog', 'information_schema')
GROUP BY UPPER(relkind)
ORDER BY object_type;
```

object_type	count
INDEX	52
SEQUENCE	6
TABLE	9
VIEW	3
T	45

-- this means some relkind not mapped in CASE

Notice in your output, there is a row T.  
That means your query encountered an unexpected relkind = 'T'.  
In PostgreSQL, 't' = TOAST table (internal storage for large values).  
Since your CASE statement doesn't map 'T', it just shows T.

```
object_type | count
-----+-----
INDEX      |      52
SEQUENCE   |        6
T          |      45
TABLE      |        9
VIEW       |        3
( 5 rows)

postgres=#
```

## Dropping Existing Database

If the target database already exists and you want a full restore, first drop it and create a new one:

```
Copy postgres=#
postgres=#
postgres=#
postgres=# \c dvdrental
You are now connected to database "dvdrental" as user "postgres".

dvdrental=#

dvdrental=# DROP DATABASE postgres ;

DROP DATABASE

dvdrental=#
dvdrental=#
dvdrental=# CREATE DATABASE postgres;
CREATE DATABASE

dvdrental=#

dvdrental=#
```

## Restore using pg\_restore

```
Copy [postgres@node1 ~]$
[postgres@node1 ~]$
[postgres@node1 ~]$ pg_restore -U postgres -d postgres -v -j 4
/pgbackup/

pg_restore: connecting to database for restore
pg_restore: processing item 4362 ENCODING ENCODING
pg_restore: processing item 4363 STDSTRINGS STDSTRINGS
pg_restore: processing item 4364 SEARCHPATH SEARCHPATH
pg_restore: processing item 4365 DATABASE postgres
pg_restore: processing item 2 EXTENSION pg_buffercache
pg_restore: creating EXTENSION "pg_buffercache"
pg_restore: processing item 4366 COMMENT EXTENSION pg_buffercache
pg_restore: creating COMMENT "EXTENSION pg_buffercache"
pg_restore: processing item 3 EXTENSION pg_stat_statements
pg_restore: creating EXTENSION "pg_stat_statements"
```

Count postgres DB object list:

```
SELECT
  CASE UPPER(relkind)
    WHEN 'R' THEN 'TABLE'
    WHEN 'P' THEN 'PARTITION TABLE'
    WHEN 'I' THEN 'INDEX'
    WHEN 'S' THEN 'SEQUENCE'
    WHEN 'V' THEN 'VIEW'
```

```

        WHEN 'M' THEN 'MATERIALIZED VIEW'
        WHEN 'F' THEN 'FOREIGN TABLE'
        WHEN 'T' THEN 'TOAST TABLE'
        ELSE UPPER(relkind)
    END AS object_type,
    COUNT(*) AS count
FROM pg_class c
JOIN pg_namespace n ON n.oid = c.relnamespace
WHERE n.nspname NOT IN ('pg_catalog', 'information_schema')
GROUP BY UPPER(relkind)
ORDER BY object_type;

```

```

postgres=# SELECT
postgres-#     CASE UPPER(relkind)
postgres-#         WHEN 'R' THEN 'TABLE'
postgres-#         WHEN 'P' THEN 'PARTITION TABLE'
postgres-#         WHEN 'I' THEN 'INDEX'
postgres-#         WHEN 'S' THEN 'SEQUENCE'
postgres-#         WHEN 'V' THEN 'VIEW'
postgres-#         WHEN 'M' THEN 'MATERIALIZED VIEW'
postgres-#         WHEN 'F' THEN 'FOREIGN TABLE'
postgres-#         ELSE UPPER(relkind)
postgres-#     END AS object_type,
postgres-#     COUNT(*) AS count
postgres-# FROM pg_class c
postgres-# JOIN pg_namespace n ON n.oid = c.relnamespace
postgres-# WHERE n.nspname NOT IN ('pg_catalog', 'information_schema')
postgres-# GROUP BY UPPER(relkind)
postgres-# ORDER BY object_type;

```

object_type	count
INDEX	52
SEQUENCE	6
T	45
TABLE	9
VIEW	3

5 rows)

## Backup tips

- Directory format with parallel jobs is recommended for large databases.
- Always test your backup by restoring it to a separate database.
- Keep multiple copies of backups in different locations.

Use `pg_dumpall` if you need to backup global objects like roles or tablespaces.

## Conclusion

Using `pg_dump` in directory format with parallel jobs is fast, safe, and flexible. It allows large databases to be backed up and restored efficiently. Always plan your backups, verify them regularly, and take advantage of parallelism to save time.