

Implementing “Always Encrypted” Functionality in PostgreSQL

PostgreSQL doesn't natively have a feature exactly equivalent to SQL Server's **Always Encrypted** option. **Always Encrypted** is a feature in SQL Server that ensures sensitive data is encrypted at the client-side and stored encrypted in the database, without the database server ever having access to the encryption keys.

In PostgreSQL, while there is no direct equivalent, you can achieve similar functionality by using a combination of client-side encryption, third-party extensions, and PostgreSQL's built-in encryption functions.

Methods to Implement "Always Encrypted" in PostgreSQL

1. **Client-Side Encryption** (Best Equivalent to Always Encrypted)
2. **PGP Symmetric and Asymmetric Encryption Functions (Using PostgreSQL Extensions)**
3. **Third-Party Encryption Tools (e.g., pgcrypto, TDE, Vault Integration)**
4. **Transparent Data Encryption (TDE) in Forks like EDB Postgres Advanced Server**

Below is a detailed explanation of how to implement these encryption methods in PostgreSQL.

1. Client-Side Encryption (Most Similar to Always Encrypted)

In PostgreSQL, client-side encryption can be achieved by encrypting sensitive data on the client side (i.e., within the application) before storing it in the database. This ensures that the database never sees the plaintext data or has access to the encryption keys.

Steps to Implement Client-Side Encryption:

1. **Encrypt Data at the Application Layer:**
 - Use an encryption library (like **OpenSSL**, **libsodium**, or a language-specific encryption package like **cryptography** for Python or **crypto** for Node.js) to encrypt data before sending it to PostgreSQL.
 - The encryption key is stored securely within the application, ensuring the database only receives and stores encrypted data.
2. **Store Encrypted Data in PostgreSQL:**
 - Store the encrypted data in a suitable PostgreSQL data type, such as **BYTEA** (for binary data) or **TEXT** (for base64-encoded strings).
3. **Decrypt Data at the Application Layer:**
 - When retrieving sensitive data, the application reads the encrypted data from PostgreSQL and decrypts it using the encryption keys stored within the application.

Example Workflow:

- **Application encrypts sensitive data** (e.g., credit card numbers) using an encryption key and stores it in PostgreSQL.
- PostgreSQL stores the encrypted data, and the decryption key never touches the database.
- When the application retrieves the encrypted data, it decrypts it on the client side.

Python Example with AES Encryption:

Using Python's cryptography library:

Python

```
1 from cryptography.fernet import Fernet
2
3 # Generate a key (store this securely, do not regenerate every time)
4 key = Fernet.generate_key()
5 cipher_suite = Fernet(key)
6
7 # Encrypt the sensitive data
8 plain_text = "SensitiveData123"
9 cipher_text = cipher_suite.encrypt(plain_text.encode())
10
11 # Store cipher_text in PostgreSQL
12 # Later, decrypt when retrieving data
13 decrypted_text = cipher_suite.decrypt(cipher_text).decode()
14 print(decrypted_text)
```

This process ensures that the data is encrypted before being sent to PostgreSQL and decrypted only by the application.

2. PostgreSQL PGP Encryption with pgcrypto

pgcrypto is a popular PostgreSQL extension that provides cryptographic functions, allowing you to encrypt and decrypt data inside the database. However, for "Always Encrypted"-like functionality, encryption keys should not be stored in the database.

Installing pgcrypto:

First, you need to enable the **pgcrypto** extension:

Shell

```
1 CREATE EXTENSION pgcrypto;
```

Symmetric Encryption with pgcrypto:

You can encrypt and decrypt data using **symmetric encryption** functions from **pgcrypto**. This is useful for encrypting small amounts of data like personal information.

- **Encryption Example:**

```
1SELECT pgp_sym_encrypt('SensitiveData123', 'encryption_key');
```

- **Decryption Example:**

```
1SELECT pgp_sym_decrypt(column_name, 'encryption_key');
```

The encryption key should be managed and stored securely outside the database, for example, using an external key management system.

Asymmetric Encryption with pgcrypto:

Asymmetric encryption uses public/private key pairs to encrypt and decrypt data. The public key encrypts the data, and the private key decrypts it.

- **Generate Key Pairs** (outside PostgreSQL using tools like GPG):

```
Shell
```

```
1gpg --gen-key
```

- **Encryption Example:**

```
1SELECT pgp_pub_encrypt('SensitiveData123', dearmor('-----BEGIN PGP PUBLIC KEY BLOCK-----...'));
```

- **Decryption Example:**

```
1SELECT pgp_pub_decrypt(column_name, dearmor('-----BEGIN PGP PRIVATE KEY BLOCK-----...'));
```

3. Third-Party Encryption Tools

HashiCorp Vault Integration:

You can use **HashiCorp Vault** for managing encryption keys and performing encryption/decryption operations outside PostgreSQL. PostgreSQL stores only encrypted data, while Vault securely manages encryption keys.

- **Encryption Workflow:**
 - The application requests Vault to encrypt data before storing it in PostgreSQL.
 - When reading data, the application retrieves the encrypted data from PostgreSQL and requests Vault to decrypt it.

Transparent Data Encryption (TDE):

- **TDE** is not natively supported in PostgreSQL, but some third-party forks, like **EDB Postgres Advanced Server**, offer TDE.
 - **EDB TDE** encrypts data files at rest, ensuring that all data stored on disk is encrypted.
-

4. Transparent Data Encryption (TDE) for PostgreSQL Forks

Some **PostgreSQL forks** or third-party solutions, like **EnterpriseDB Postgres Advanced Server (EPAS)**, provide **Transparent Data Encryption (TDE)**.

- **How TDE Works:**
 - TDE encrypts the entire PostgreSQL database or specific tablespaces at rest.
 - When data is written to disk, it is automatically encrypted, and when read from disk, it is decrypted.
 - Keys are managed separately using an external key management system (KMS).

EnterpriseDB Postgres Advanced Server supports TDE with features like column-level encryption and whole database encryption, but this is only available in their commercial offering.

Summary of Approaches to Achieve Always Encrypted in PostgreSQL:

Method	Description	Pros	Cons
Client-Side Encryption	Encryption and decryption occur at the application layer before data is sent to PostgreSQL.	True client-side encryption.Full control over the process.	Requires application modifications.External management needed.
pgcrypto Extension	Offers symmetric and asymmetric encryption functions within PostgreSQL.	In-database encryption capabilities.	Potential security risk stored in the database.
HashiCorp Vault Integration	External service for key management and encryption, integrated with PostgreSQL.	Robust key management.High-level security.	Increases system complexity.Requires Vault service.
TDE (in PostgreSQL forks)	Encrypts the entire database at rest, typically found in enterprise PostgreSQL variants.	Seamless encryption.Minimal application changes needed.	Not available in standard PostgreSQL.Relies on forks or solutions.

Conclusion:

While PostgreSQL doesn't have a native **Always Encrypted** feature like SQL Server, you can achieve similar results through **client-side encryption**, the **pgcrypto** extension, or third-party encryption tools like **Vault**. Each approach comes with trade-offs in terms of security, complexity, and ease of use, so the right solution depends on your specific needs and requirements for data security.

