

Open in app ↗

Medium

 Search

SQL Explain is your friend: PostgreSQL edition

sweemeng · [Follow](#)

5 min read · Feb 13, 2025

 Listen Share More

Disclosure: I came out with the content, database setup, and queries. But I use ChatGPT and perplexity to fix the grammar, tone, and spelling.

So you discovered an application running slow. You suspect that the application is running a slow database query. How do you find out?

One way is to use `pg_stat_statements`. The `pg_stat_statements` tool collect all queries and their performance metrics. It allows you to identify the slowest query, but it does not show why they are slow. This is where the `EXPLAIN` comes in, complementing `pg_stat_statements` in your investigation.

The SQL `EXPLAIN` query is a way for you view the execution plan of your query. It provides information that will help you make changes to improve performance.

`EXPLAIN` query is not an SQL standard, but many database engines supports it. Since I primarily use PostgreSQL— and different engines output information in various formats — I will focus solely on PostgreSQL's implementation.

There are several forms of the `EXPLAIN` command that I use. While there are additional variations available that you might find useful, the following are the ones I typically rely on:

```

medium_tutorial=# EXPLAIN SELECT id FROM medias WHERE user_id=109;
               QUERY PLAN
-----
Seq Scan on medias  (cost=0.00..40.00 rows=8 width=4)
  Filter: (user_id = 109)
(2 rows)

medium_tutorial=# EXPLAIN ANALYZE SELECT id FROM medias WHERE user_id=109;
               QUERY PLAN
-----
Seq Scan on medias  (cost=0.00..40.00 rows=8 width=4) (actual time=0.180..0.312 rows=8 loops=1)
  Filter: (user_id = 109)
  Rows Removed by Filter: 1992
  Planning Time: 0.071 ms
  Execution Time: 0.330 ms
(5 rows)

medium_tutorial=# EXPLAIN (ANALYZE, BUFFERS) SELECT id FROM medias WHERE user_id=109;
               QUERY PLAN
-----
Seq Scan on medias  (cost=0.00..40.00 rows=8 width=4) (actual time=0.152..0.272 rows=8 loops=1)
  Filter: (user_id = 109)
  Rows Removed by Filter: 1992
  Buffers: shared hit=15
  Planning Time: 0.056 ms
  Execution Time: 0.286 ms
(6 rows)

medium_tutorial=# 

```

Key differences

- **EXPLAIN \$query** : Display the query plan without executing query
- **EXPLAIN ANALYZE \$query** : Executes the query and returns the query plan along with the execution time. Note that if you use this with **UPDATE** , **DELETE** , or **INSERT** statements, it will modify your data.
- **EXPLAIN (ANALYZE, BUFFERS) \$query** : Performs the same as **EXPLAIN ANALYZE** but additionally provides information on shared buffers.

Let's start with an example database I have created:-

```
medium_tutorial=# \d users
```

Column	Type	Collation	Nullable	Default
id	integer		not null	nextval('users_id_seq'::regclass)
name	character varying(255)			
email	character varying(255)			

```
Indexes:
    "users_pkey" PRIMARY KEY, btree (id)
```

```
medium_tutorial=# \d medias
```

Column	Type	Collation	Nullable	Default
id	integer		not null	nextval('medias_id_seq'::regclass)
path	character varying(255)			
user_id	integer			

```
Indexes:
    "medias_pkey" PRIMARY KEY, btree (id)
```

```
medium_tutorial=#
```

Table created for testing

There is around 200 users, and 2000 medias.

My go-to query is the `EXPLAIN (ANALYZE, BUFFERS)` command. Below is an example session:

```
medium_tutorial=# EXPLAIN (ANALYZE, BUFFERS) SELECT id FROM medias WHERE user_id = 109
```

QUERY PLAN

```
Seq Scan on medias (cost=0.00..40.00 rows=8 width=4) (actual time=0.150..0.270)
  Filter: (user_id = 109)
  Rows Removed by Filter: 1992
  Buffers: shared hit=15
Planning Time: 0.063 ms
Execution Time: 0.285 ms
(6 rows)
```

Let's parse the output:

- Seq Scan on medias

This indicates that the query is performing a sequential scan on the `medias` table—meaning it is not using an index. (I will show an example of an index search later.)

- (cost=0.00..40.00 rows=8 width=4)

The query planner estimates an initial cost of 0.00 and a final cost of 40.00 . It also estimates that the query will return 8 rows, with each row being 4 bytes wide. These cost metrics help the planner decide on the most optimal execution plan.

- (actual time=0.150..0.270 rows=8 width=4)

This shows the actual time taken to execute the query. In this case, the execution started at 0.150 ms and ended at 0.270 ms .

- Rows Removed by Filter: 1992

This indicates that 1992 rows were scanned and then discarded based on the filter condition. A high number of rows being filtered out may impact performance.

- Buffers: shared hit=15

This shows that 15 blocks of data were retrieved from the shared buffer cache. (In PostgreSQL, each block is typically 8KB.)

- Execution Time: 0.285 ms

The total time taken to execute the query was 0.285 ms .

Now let's add an index in the `medias` table.

```
CREATE INDEX medias_user_id_idx ON medias(user_id);
```

Now rerun the query:

```
medium_tutorial=# EXPLAIN (ANALYZE, BUFFERS) SELECT id FROM medias WHERE user_id = 109
               QUERY PLAN
-----
Bitmap Heap Scan on medias  (cost=4.34..18.09 rows=8 width=4) (actual time=0.00..0.01 rows=8 width=4)
  Recheck Cond: (user_id = 109)
  Heap Blocks: exact=5
```

```
Buffers: shared hit=5 read=2
-> Bitmap Index Scan on medias_user_id_idx (cost=0.00..4.34 rows=8 width=0)
    Index Cond: (user_id = 109)
    Buffers: shared read=2
Planning:
    Buffers: shared hit=10 read=1
Planning Time: 0.782 ms
Execution Time: 0.082 ms
(11 rows)
```

Now, there are a couple of differences:

- Bitmap Heap Scan on Media

This indicates that the query is now using the index.

- 
- Buffers: shared hit=5 read=2 :

The new metric, `read`, shows the number of blocks read from disk (each block is typically 8KB, so 2 blocks equal 16KB). These blocks are then stored in the cache.

- Time Breakdown:

The output now provides a more detailed breakdown of the execution time, including both planning and execution phases.

- Faster Execution Time:

Finally, note that the overall execution time is faster after adding the index!

Finally, you can also use `EXPLAIN` to test your assumptions. Here are a couple of reasons why it's important to do so:

- Unexpected Query Plans:

The query planner might not produce the query plan you expect, choosing instead a plan that it deems more cost-effective.

- ORM-Generated Queries:

Your ORM may generate queries that are not optimal. For example, it might generate a `SELECT *` query—even when you only need a few columns—which can be inefficient if the table has many columns.

Consider a concrete example. One might assume that a query like selecting `id` from `users` would simply use the index on `medias`. However, the actual plan shows that a join is performed instead:

```
medium_tutorial=# EXPLAIN (ANALYZE, BUFFERS) SELECT id FROM medias WHERE user_id = 1;
               QUERY PLAN

-----
Hash Semi Join (cost=3.25..54.62 rows=1000 width=4) (actual time=0.093..1.102 rows=16)
  Hash Cond: (medias.user_id = users.id)
  Buffers: shared hit=17
    -> Seq Scan on medias (cost=0.00..35.00 rows=2000 width=8) (actual time=0.000..0.000 rows=2000)
        Buffers: shared hit=15
    -> Hash (cost=2.00..2.00 rows=100 width=4) (actual time=0.065..0.066 rows=100)
        Buckets: 1024 Batches: 1 Memory Usage: 12kB
        Buffers: shared hit=2
      -> Limit (cost=0.00..2.00 rows=100 width=4) (actual time=0.009..0.010 rows=100)
          Buffers: shared hit=2
        -> Seq Scan on users (cost=0.00..4.00 rows=200 width=4) (actual time=0.000..0.000 rows=200)
            Buffers: shared hit=2

Planning:
  Buffers: shared hit=14 read=1
Planning Time: 0.397 ms
Execution Time: 1.207 ms
(16 rows)
```

In this case, the plan reveals that the query uses a join rather than simply scanning the index on `medias`. This unexpected join is a perfect example of why it's essential to test your assumptions with `EXPLAIN` before finalizing query optimizations.

In conclusion, if you suspect that a query is running slowly, start by checking it with `pg_stat_statements` and then use `EXPLAIN` to dive deeper into its execution plan. These tools are essential for validating your assumptions and gaining a clear understanding of your database's behavior—especially in complex systems where subtle differences in execution can significantly impact performance.

[Postgresql](#)[Database](#)

[Follow](#)

Written by sweemeng

481 Followers · 789 Following

Codemonkey coding for tea. Startup developer at day, community organizer at night. I also happens to be on patreon <https://www.patreon.com/sweemeng>

No responses yet



Gvadakte

What are your thoughts?

More from sweemeng

al information provided as a name-value (key-value) pair. [Learn more](#)

	Key	Value	
d ▼	Content-Type ▼	application/x-wv ✕	Remove
d ▼	Content-Disposition ▼	Value	Remove
		attachment	
ta			


 sweemeng

Setting content-disposition in S3

Not too long ago, I had this problem. We need to allow users to download json files hosted S3 through a signed URL, but the user on browser...

★ Feb 19, 2023 🖱 15



 sweemeng

Get micropython to work with TTGO T-Display RP2040

Recently I got a TTGO T-Display RP2040. This is a development board based on the Raspberry Pi RP2040 microcontroller. This board features...

M5Stack VIM Escape button

Note: I originally write this on <https://devkami.com/blogs/2021/05/17-vim-escape-key/>

★ May 18, 2021



See all from sweemeng

Recommended from Medium



In Databases by Sergey Egorenkov

Why Uber Moved from Postgres to MySQL

How PostgreSQL's architecture clashed with Uber's scale—and why MySQL offered a better path forward

Mar 29 👏 234 💬 7





Comviva MFS Engineering Tech Blog

The Challenge of Primary Keys in Partitioned Tables in PostgreSQL

— by Parul Vashist

Jan 27 🖱️ 3



In Hack the Stack by Coders Stop

9 Database Optimization Tricks SQL Experts Are Hiding From You

Most developers learn enough SQL to get by—SELECT, INSERT, UPDATE, DELETE, and maybe a few JOINS. They might even know how to create...

★ Mar 27 🖱 186 💬 5

 Tihomir Manushev

Vector Search with pgvector in PostgreSQL

Simple AI-powered similarity search

★ Mar 9

 Ravi Yasas

Multiple database support in NestJS with TypeORM

Supporting multiple databases is necessary for specific scenarios, particularly in migration projects. NestJS offers the flexibility to...

★ Mar 23



 Ajaymaurya

How Often Should You Reindex Your PostgreSQL Database? A Data-Driven Approach

PostgreSQL is one of the most powerful and flexible relational databases available today. However, like any database system, it requires...

★ Mar 24



See more recommendations