

Understanding Checkpoint, LSN, and WAL in PostgreSQL

1. Checkpoint in PostgreSQL

A checkpoint is a crucial event in PostgreSQL that ensures data consistency and reduces recovery time in case of a crash.

It writes all dirty (modified) pages from shared memory (buffer cache) to disk and marks a point in the WAL (Write-Ahead Log) where recovery can begin.

Key Features:

- Ensures Durability: Flushes dirty pages to disk so committed transactions are permanent.
- Speeds Up Crash Recovery: Recovery can start from the last checkpoint rather than replaying the entire WAL.
- Controlled via Parameters:
 - * `checkpoint_timeout`: Maximum time interval between checkpoints.
 - * `checkpoint_completion_target`: Spreads out writes over the interval.
 - * `checkpoint_segments` (pre-9.5) / `max_wal_size` (newer versions): Controls WAL size before triggering a checkpoint.

Checkpoint Process:

1. Flush modified pages to disk.
2. Write a checkpoint record in the WAL.
3. Remove old WAL files based on retention settings.

2. LSN (Log Sequence Number)

A Log Sequence Number (LSN) uniquely identifies each WAL record in PostgreSQL.

It points to an exact location in the WAL, aiding in tracking database changes.

Key Features:

- Used in Replication: Determines how far a replica is synchronized with the primary.
- Aids in WAL Archiving & PITR: Recovery begins from a known LSN.
- Improves Crash Recovery: Only replays WAL records starting from

a specific LSN.

LSN Format & Example:

- Represented in HEX, e.g., 00000002/000000D0
- * First part (00000002): WAL segment number.
- * Second part (000000D0): Offset within the segment.

Common Commands:

- Current WAL LSN: `SELECT pg_current_wal_lsn();`
- Last Checkpoint LSN: `SELECT checkpoint_lsn FROM pg_control_checkpoint();`

3. Relation Between LSN & WAL File

LSN is a pointer to a specific location within the WAL file. PostgreSQL stores WAL files in the `pg_wal` directory (or `pg_xlog` in older versions).

Each WAL file has a fixed size (usually 16MB) and a naming pattern like:

00000001000000020000003C

which consists of:

- Timeline ID
- WAL Segment Number (matching the first part of the LSN)
- WAL Offset

Each WAL file covers a range of LSNs. For example, if a file covers 16MB:

- First LSN: 00000002/00000000
- Last LSN: 00000002/01000000

To determine which WAL file contains a specific LSN:

`SELECT pg_walfile_name('00000002/000000D0');`

This maps the LSN to the corresponding WAL file.

4. Step-by-Step Example: Mapping LSN to WAL File

Step 1: Check the Current LSN

Command:

`SELECT pg_current_wal_lsn();`

Example Output:

0/16B4F78

Step 2: Find the WAL File Containing This LSN

Command:

SELECT pg_walfile_name('0/16B4F78');

Example Output:

00000001000000000000000016

Step 3: Verify the WAL File Path

Command:

ls -lh \$PGDATA/pg_wal/ | grep 00000001000000000000000016

Example Output:

**-rw----- 1 postgres postgres 16M Mar 2 10:30
00000001000000000000000016**

Step 4: Generate a New LSN by Inserting Data

Commands:

CREATE TABLE test_lsn (id SERIAL, name TEXT);

INSERT INTO test_lsn (name) VALUES ('LSN Example');

Then check:

SELECT pg_current_wal_lsn();

Example Output:

0/16B5A40

Step 5: Determine the WAL File for the New LSN

Command:

SELECT pg_walfile_name('0/16B5A40');

Example Output:

00000001000000000000000016

(Indicating the same WAL file is still in use)

Step 6: Check the Last Checkpoint LSN

Command:

SELECT checkpoint_lsn FROM pg_control_checkpoint();

Example Output:

0/16B4A20

Step 7: LSN in Replication (On a Standby Server)

Command:

SELECT sent_lsn, write_lsn, flush_lsn, replay_lsn FROM

pg_stat_replication;

Example Output:

sent_lsn	write_lsn	flush_lsn	replay_lsn
0/16B4F78	0/16B4F78	0/16B4F78	0/16B4F00

Conclusion:

- LSN points to a specific position within the WAL.
- `pg_walfile_name()` maps an LSN to its corresponding WAL file.
- WAL files store database changes sequentially.
- LSNs are essential for replication and recovery.