

## Checkpoint in PostgreSQL

A **checkpoint** is a crucial event in PostgreSQL that ensures data consistency and reduces recovery time in case of a crash. It works by writing all dirty (modified) pages from the shared memory (buffer cache) to disk and marking a point in the WAL (Write-Ahead Log) where recovery can begin.

### Key Features of Checkpoints:

1. **Ensures Durability:** Flushes dirty pages to disk, making sure committed transactions are permanently stored.
2. **Speeds Up Crash Recovery:** The database can start recovery from the last checkpoint instead of replaying the entire WAL.
3. **Controlled via Parameters:**
  1. `checkpoint_timeout`: Maximum time interval between checkpoints.
  2. `checkpoint_completion_target`: Percentage of time between checkpoints used to spread out writes.
  3. `checkpoint_segments` (before PostgreSQL 9.5) / `max_wal_size` (newer versions): Controls the maximum WAL size before a checkpoint triggers.

### Checkpoint Process:

1. Flush all modified data pages to disk.
2. Write a special checkpoint record in the WAL.
3. Remove old WAL files (based on retention settings).

=====

## LSN (Log Sequence Number)

A **Log Sequence Number (LSN)** is a unique identifier for each WAL record in PostgreSQL. It represents the exact position in the WAL and helps track database changes efficiently.

### Key Features of LSN:

1. **Used in Replication:** Determines how far a replica has caught up with the primary.
2. **Helps in WAL Archiving & PITR:** Ensures recovery starts from a known LSN.
3. **Improves Crash Recovery:** PostgreSQL can replay WAL from a specific LSN to recover data.

### LSN Format & Example

- LSN is represented in **HEX format**: 00000002/000000D0
  - **First part (00000002):** WAL segment number.
  - **Second part (000000D0):** Offset within that WAL segment.

Commands to Check LSN:-

Current WAL LSN:

```
SELECT pg_current_wal_lsn();
```

**Last Checkpoint LSN:**

```
SELECT checkpoint_lsn FROM pg_control_checkpoint();
```

**LSN of a Table:**

```
SELECT pg_relation_size('my_table');
```

=====

Relation Between LSN (Log Sequence Number) and WAL (Write-Ahead Log) File in PostgreSQL

LSN (Log Sequence Number) is a **pointer** to a specific location in the WAL (Write-Ahead Log). Each WAL file stores database changes, and every record inside the WAL has an associated LSN.

## 1. WAL File Naming and Structure

PostgreSQL stores WAL files in the `pg_wal` (or `pg_xlog` in older versions) directory. Each WAL file has a **fixed size** (usually 16MB by default) and follows a specific naming pattern:

**WAL File Name Format:**

000000010000000200000003C

This name consists of:

- **Timeline ID (00000001)** → Identifies the current timeline.
- **WAL Segment Number (00000002)** → Corresponds to the first part of the LSN.
- **WAL Offset (00000003C)** → Identifies the WAL file in sequence.

Each WAL file covers a range of LSNs. For example:

- If one WAL file covers 16MB, then:
  - **First LSN in the file:** 00000002/00000000
  - **Last LSN in the file:** 00000002/01000000

## 2. How to Find Which WAL File Contains a Specific LSN

You can determine which WAL file a given LSN belongs to using:

```
SELECT pg_walfile_name('00000002/000000D0');
```

Example Output:

```
000000010000000200000003C
```

This means the WAL file named 000000010000000200000003C contains changes related to the specified LSN.

### 3. How to Find the LSN of the Last WAL Record

You can check the latest LSN using:

```
SELECT pg_current_wal_lsn();
```

Example Output:

```
0/16B4F78
```

This indicates the current WAL insert position.

### 4.LSN to WAL File Mapping Example:

If you get an LSN like 00000002/000000D0, the corresponding WAL file name can be derived using:

```
SELECT pg_walfile_name('00000002/000000D0');
```

which might return something like:

```
000000010000000200000003C
```

This tells you that all changes related to this LSN are recorded inside 000000010000000200000003C.

### 5. LSN to WAL File Mapping Example

If you get an LSN like 00000002/000000D0, the corresponding WAL file name can be derived using:

```
SELECT pg_walfile_name('00000002/000000D0');
```

which might return something like:

```
000000010000000200000003C
```

This tells you that all changes related to this LSN are recorded inside 000000010000000200000003C.

## 6. Check the Last Checkpoint LSN:-

```
SELECT checkpoint_lsn FROM pg_control_checkpoint();
```

This means PostgreSQL will start recovery from **LSN 0/16B4A20** in case of a crash.

## 7. How LSN is Used in Replication

On a standby server (replica), check the last received LSN:

```
SELECT sent_lsn, write_lsn, flush_lsn, replay_lsn FROM pg_stat_replication;
```

sent_lsn	write_lsn	flush_lsn	replay_lsn
0/16B4F78	0/16B4F78	0/16B4F78	0/16B4F00

- sent\_lsn → Last LSN sent to the replica.
- write\_lsn → Last LSN written on the replica.
- flush\_lsn → Last LSN flushed to disk on the replica.
- replay\_lsn → Last LSN replayed on the replica.

If replay\_lsn is behind sent\_lsn, it means replication **lag** exists.

## Conclusion

- LSN **points** to a specific position in WAL.
- pg\_walfile\_name() **maps** an LSN to a WAL file.
- WAL files store database changes **sequentially**.
- LSNs are crucial for **replication** and **recovery**.