

Day 2 Backup and Recovery Strategies

Backup and Restore in PostgreSQL

logical backup

logical backup consist of dumping sql command for whole database or table or schema

sql command consist of create and altering insert query that run on database for it current state when logical backup is taking to take backup for database its pretty straight forward

for this section i have already deploy sample database pagila and will take backup of it

Taking DatabaseBackup

```
pg_dump -d pagila > /backup/pagila-backup.sql
```

```
postgres@postgresql-16-test:/$ pg_dump -d pagila > /backup/pagila-backup.sql
postgres@postgresql-16-test:/$ pg_dump -d pagila -F c -f /backup/pagila-backup.dump
```

-d option to specify database you want the dump

> indication follow by location for storing pg_dump backup

```
pg_dump -d pagila -t actor > /backup/actor-backup.sql
```

```
postgres@postgresql-16-test:/backup$ pg_dump -d pagila -t actor > /backup/actor-backup.sql
postgres@postgresql-16-test:/backup$ ls
actor-backup.sql  dvdrental.tar  pagila-backup.dump  pagila-backup.sql
```

following command will take only backup for table you can mention multilabel tables

-t for specify table or multicable tables to take dump backup

Taking compress backup for pg_dump

pg_dump backup can get huge in size this way its good practices to compress backup

```
pg_dump -d pagila -F c -f /backup/pagila-backup.dump
```

-F used to mentioned format followed by c for compressed

-f used to specify the file name that you want to defined

```
postgres@postgresql-16-test:/backup$ pg_dump -d pagila -t actor > /backup/actor-backup.sql
postgres@postgresql-16-test:/backup$ ls
actor-backup.sql  backup.sql.gz  dvdrental.tar  pagila-backup.dump  pagila-backup.sql
```

pg_dumpall

pg_dumpall is used to take full backup of the entire database cluster including roles and tablespaces

useful if you want to migrate you database to new server or just simply want to migrate it to upgraded version , you can restore pg_dumpall dump to any other PostgreSQL version

```
pg_dumpall > cluster_backup.sql
```

you can use pg_dump and pg_dumpall to take backup from remote postgres you need to specify -h for host -p for port if the remote host use port other than the default one -U for username

```
postgres@postgresql-16-test:/backup$ pg_dumpall > cluster_backup.sql
postgres@postgresql-16-test:/backup$ ls
actor-backup.sql  backup.sql.gz  dvdrental.tar  pagila-backup.dump  pagila-backup.sql
cluster_backup.sql  cluster_backup.sql  pagila-backup.dump
```

compressed backup

```
pg_dumpall | gzip > dumpall.sql.gzip
```

Restore logical backup

restore logical backup for both pg_dump and pg_dumpall is straight forward

for pg_dump i will drop database call pagila and recreate again then restore it with backup i have created
i will restore both the compresses one and normal dump

below command will restore dump compressed file , to restore compressed file best to use pg_restore

```
dropdb pagila
createdb pagila_retsore
pg_restore -d pagila_retsore < pagila-backup.dump
```

the restoration is easy only you need to specify the database using option -d and make pointer < followed by backup path

```
postgres@postgresql-16-test:/backup$ dropdb pagila
postgres@postgresql-16-test:/backup$ psql -e '\l'
psql: error: connection to server on socket "/var/run/postgresql/.s.PGSQL.5432" failed: FATAL: database "\l" does not exist
postgres@postgresql-16-test:/backup$ psql
psql (16.9 (Ubuntu 16.9-0ubuntu0.24.04.1))
Type "help" for help.

postgres=# \l

```

Name	Owner	Encoding	Locale Provider	Collate	Ctype	ICU Locale	ICU Rules	Access privileges
dvdrental	postgres	UTF8	libc	en_US.UTF-8	en_US.UTF-8			
postgres	postgres	UTF8	libc	en_US.UTF-8	en_US.UTF-8			
template0	postgres	UTF8	libc	en_US.UTF-8	en_US.UTF-8			=c/postgres + postgres=CTc/postgres
template1	postgres	UTF8	libc	en_US.UTF-8	en_US.UTF-8			=c/postgres + postgres=CTc/postgres

(4 rows)

```
postgres@postgresql-16-test:/backup$ psql -d pagila_retsore < pagila-backup.dump
The input is a PostgreSQL custom-format dump.
Use the pg_restore command-line client to restore this dump to a database.

postgres@postgresql-16-test:/backup$ psql
psql (16.9 (Ubuntu 16.9-0ubuntu0.24.04.1))
Type "help" for help.

postgres=# \l

```

Name	Owner	Encoding	Locale Provider	Collate	Ctype	ICU Locale	ICU Rules	Access privileges
dvdrental	postgres	UTF8	libc	en_US.UTF-8	en_US.UTF-8			
pagila_retsore	postgres	UTF8	libc	en_US.UTF-8	en_US.UTF-8			
postgres	postgres	UTF8	libc	en_US.UTF-8	en_US.UTF-8			
template0	postgres	UTF8	libc	en_US.UTF-8	en_US.UTF-8			=c/postgres + postgres=CTc/postgres
template1	postgres	UTF8	libc	en_US.UTF-8	en_US.UTF-8			=c/postgres + postgres=CTc/postgres

(5 rows)

```
postgres=# \c pagila_retsore
You are now connected to database "pagila_retsore" as user "postgres".
pagila_retsore=# \dt
Did not find any relations.
pagila_retsore=# \q
postgres@postgresql-16-test:/backup$ pg_restore -d pagila_retsore < pagila-backup.dump
postgres@postgresql-16-test:/backup$ psql
psql (16.9 (Ubuntu 16.9-0ubuntu0.24.04.1))
Type "help" for help.

postgres=# \c pagila_retsore
You are now connected to database "pagila_retsore" as user "postgres".
pagila_retsore=# \dt

```

Schema	Name	Type	Owner
public	actor	table	postgres
public	address	table	postgres
public	category	table	postgres
public	city	table	postgres
public	country	table	postgres
public	customer	table	postgres
public	film	table	postgres
public	film_actor	table	postgres
public	film_category	table	postgres

below command will restore normal dump with .sql extension for that you use psql for restore for that i will create new database and restore from normal uncompressed dump

```
createdb test_restore
```

```
psql -d test_restore < pagila-backup.sql
```

```
postgres@postgresql-16-test:/backup$ psql -d test_restore < pagila-backup.sql
```

```
SET
SET
SET
SET
SET
set_config
-----
(1 row)

SET
SET
SET
SET
ALTER SCHEMA
CREATE DOMAIN
ALTER DOMAIN
CREATE TYPE
ALTER TYPE
CREATE DOMAIN
ALTER DOMAIN
CREATE FUNCTION
ALTER FUNCTION
CREATE FUNCTION
ALTER FUNCTION
CREATE FUNCTION
ALTER FUNCTION
CREATE FUNCTION
ALTER FUNCTION
CREATE FUNCTION
ALTER FUNCTION
CREATE FUNCTION
ALTER FUNCTION
CREATE FUNCTION
ALTER FUNCTION
CREATE SEQUENCE
ALTER SEQUENCE
SET
SET
CREATE TABLE
ALTER TABLE
CREATE FUNCTION
```

```
postgres=# \c test_restore
You are now connected to database "test_restore" as user "postgres".
test_restore=# \dt
```

List of relations			
Schema	Name	Type	Owner
public	actor	table	postgres
public	address	table	postgres
public	category	table	postgres
public	city	table	postgres
public	country	table	postgres
public	customer	table	postgres
public	film	table	postgres
public	film_actor	table	postgres
public	film_category	table	postgres
public	inventory	table	postgres
public	language	table	postgres
public	payment	partitioned table	postgres
public	payment_p2022_01	table	postgres
public	payment_p2022_02	table	postgres
public	payment_p2022_03	table	postgres
public	payment_p2022_04	table	postgres
public	payment_p2022_05	table	postgres
public	payment_p2022_06	table	postgres
public	payment_p2022_07	table	postgres
public	rental	table	postgres
public	staff	table	postgres
public	store	table	postgres

(22 rows)

```
test_restore=# select count(*) from actor;
count
-----
    200
(1 row)
```

```
test_restore=#
```

Restore pg_dumpall

```
psql -f cluster_backup.sql
```

`-f` option to specify the `pg_dumpall` file name

For backups created in custom (`-F c`), directory (`-F d`), or tar (`-F t`) formats, use the `pg_restore` utility

Limitations of Logical Backups

- Performance Impact: Creating logical backups can be resource-intensive and may impact database performance.
- Restoration Time: Restoring from logical backups can be slower than physical backups, especially for large databases, as objects and indexes need to be recreated.
- No Point-in-Time Recovery: Logical backups alone don't support point-in-time recovery; they represent the database state at the time of backup.

Best Practices for Logical Backups

- Regular Scheduling: Schedule regular backups based on your RPO requirements.
- Compression: Use compression to reduce storage requirements, especially for large databases.
- Backup Verification: Regularly verify the integrity of your backups by performing test restores.
- Backup Rotation: Implement a backup rotation strategy to maintain multiple backup copies while managing storage usage.
- Documentation: Document your backup procedures and ensure multiple team members know how to perform backups and restores.

Physical Backups with pg_basebackup

the recommend and most best practices way to take backup for production PostgreSQL is to use physical backup solution such as `pg_basebackup` or `pg_backrest` backup , it consider online backup meaning has very low impact on data base performance during initiation of backup , its support point in time recover and restoring is mush faster

only drawback is for solution such as `pg_basebackup` you cannot take backup for single schema or single database `pg_basebackup` will take full data dir backup

to take backup with point in time recovery we need to enable wal archiving

Normal backup

```
pg_basebackup -D /backup/base/
```

the directory where you specify using `-D` must be empty

```
postgres@postgresql-16-test:/backup$ cd backup.sql
bash: cd: backup.sql: Not a directory
postgres@postgresql-16-test:/backup$ cd base/
postgres@postgresql-16-test:/backup/base$ ls\
backup_label      base              pg_commit_ts      pg_logical         pg_notify          pg_serial          pg_stat            pg_subtrans        pg_twophase        pg_wal            postgresql.auto.conf
backup_manifest    global            pg_dynshmem        pg_multixact        pg_replslot        pg_snapshots       pg_stat_tmp        pg_tblspc          PG_VERSION          pg_xact
```

`pg_basebackup` from you can see take full data dir copy while database operational

compressed and wall stream backup

```
pg_basebackup -Ft -z -D /backup/base/
```

-Ft specify file format
-z the file format will be gzip compressed

```
postgres@postgresql-16-test:/backup/base$ pg_basebackup -Ft -z -D /backup/base/comp/
postgres@postgresql-16-test:/backup/base$ cd comp/
postgres@postgresql-16-test:/backup/base/comp$ ls
backup_manifest base.tar.gz pg_wal.tar.gz
postgres@postgresql-16-test:/backup/base/comp$ rm *
```

compressed with wal stream

```
pg_basebackup -Ft -z -X stream -D /backup/base/comp/
```

```
postgres@postgresql-16-test:/backup/base/comp$ pg_basebackup -Ft -z -X stream -D /backup/base/comp/
postgres@postgresql-16-test:/backup/base/comp$ ls
backup_manifest base.tar.gz pg_wal.tar.gz
```

for this option to work you need to update parameter in postgresql.conf file as follow

```
max_wal_senders = 10 # at least 2 is recommended
wal_level = replica # or logical
```

why you need to stream wal during backup

- pg_basebackup actually does by default -X fetch method, these WAL files pile up on your main server's disk and are only copied at the very end.
- When you take a backup of a large database, the process can take a long time (minutes or even hours). During this entire time, your database is still running and generating transaction logs (WAL files).
- By streaming the WAL files as *they are created*, you prevent this dangerous buildup. The backup process consumes the logs in real-time, putting no extra disk pressure on your primary server.

Restore pg_basebackup backup

first to restore backup you need to disable PostgreSQL services and remove the data directory content and then copy backup content and ensure permission then start the PostgreSQL services

```
systemctl stop postgresql@16-main.service
cd /data # go to data directory

rm -rf * # remove all content

cd /backup/base # go to backup directory for basebackup
mv * /data/ move the files to data directory

#confirm that file permission are correct
systemctl start postgresql@16-main.service
```

<https://www.linkedin.com/in/ahmed-mohamed-423583151>

```

postgres@postgresql-16-test:/backup/base$ ls
backup_label      base      pg_commit_ts  pg_logical  pg_notify  pg_serial  pg_stat  pg_subtrans  pg_twophase  pg_wal  postgresql.auto.conf
backup_manifest   global   pg_dynshmem   pg_multixact pg_replslot pg_snapshots pg_stat_tmp pg_tblspc  PG_VERSION  pg_xact
postgres@postgresql-16-test:/backup/base$ cd ..
postgres@postgresql-16-test:/backup$ cd ..
postgres@postgresql-16-test:/$ su - ahmed
Password:
ahmed@postgresql-16-test:~$ systemctl stop postgresql@16-main.service
==== AUTHENTICATING FOR org.freedesktop.systemd1.manage-units ====
Authentication is required to stop 'postgresql@16-main.service'.
Authenticating as: ahmed
Password:
==== AUTHENTICATION COMPLETE ====
ahmed@postgresql-16-test:~$ sudo su - root
[sudo] password for ahmed:
root@postgresql-16-test:~# cd /
root@postgresql-16-test:/# cd data/
root@postgresql-16-test:/data# rm *
rm: cannot remove 'base': Is a directory
rm: cannot remove 'global': Is a directory
rm: cannot remove 'pg_commit_ts': Is a directory
rm: cannot remove 'pg_dynshmem': Is a directory
rm: cannot remove 'pg_logical': Is a directory
rm: cannot remove 'pg_multixact': Is a directory
rm: cannot remove 'pg_notify': Is a directory
rm: cannot remove 'pg_replslot': Is a directory
rm: cannot remove 'pg_serial': Is a directory
rm: cannot remove 'pg_snapshots': Is a directory
rm: cannot remove 'pg_stat': Is a directory
rm: cannot remove 'pg_stat_tmp': Is a directory
rm: cannot remove 'pg_subtrans': Is a directory
rm: cannot remove 'pg_tblspc': Is a directory
rm: cannot remove 'pg_twophase': Is a directory
rm: cannot remove 'pg_wal': Is a directory
rm: cannot remove 'pg_xact': Is a directory
root@postgresql-16-test:/data# rm -rf *
root@postgresql-16-test:/data# ls
root@postgresql-16-test:/data# cd ..
root@postgresql-16-test:/# cd backup/
root@postgresql-16-test:/backup# cd base
root@postgresql-16-test:/backup/base# mv * /data/
root@postgresql-16-test:/backup/base# cd ..
root@postgresql-16-test:/backup# cd ..
root@postgresql-16-test:/# cd data/
root@postgresql-16-test:/data# ls
backup_label      base      pg_commit_ts  pg_logical  pg_notify  pg_serial  pg_stat  pg_subtrans  pg_twophase  pg_wal  postgresql.auto.conf
backup_manifest   global   pg_dynshmem   pg_multixact pg_replslot pg_snapshots pg_stat_tmp pg_tblspc  PG_VERSION  pg_xact
root@postgresql-16-test:/data# ls -l
total 268
-rw-r----- 1 postgres postgres 225 Jun  7 13:17 backup_label
-rw-r----- 1 postgres postgres 256197 Jun  7 13:17 backup_manifest
drwx----- 7 postgres postgres 59 Jun  7 13:17 base
root@postgresql-16-test:/data# ls
backup_label      base      pg_commit_ts  pg_logical  pg_notify  pg_serial  pg_stat  pg_subtrans  pg_twophase  pg_wal  postgresql.auto.conf
backup_manifest   global   pg_dynshmem   pg_multixact pg_replslot pg_snapshots pg_stat_tmp pg_tblspc  PG_VERSION  pg_xact
root@postgresql-16-test:/data# ls -l
total 268
-rw-r----- 1 postgres postgres 225 Jun  7 13:17 backup_label
-rw-r----- 1 postgres postgres 256197 Jun  7 13:17 backup_manifest
drwx----- 7 postgres postgres 59 Jun  7 13:17 base
drwx----- 2 postgres postgres 4096 Jun  7 13:17 global
drwx----- 2 postgres postgres 6 Jun  7 13:17 pg_commit_ts
drwx----- 2 postgres postgres 6 Jun  7 13:17 pg_dynshmem
drwx----- 4 postgres postgres 68 Jun  7 13:17 pg_logical
drwx----- 4 postgres postgres 36 Jun  7 13:17 pg_multixact
drwx----- 2 postgres postgres 6 Jun  7 13:17 pg_notify
drwx----- 2 postgres postgres 6 Jun  7 13:17 pg_replslot
drwx----- 2 postgres postgres 6 Jun  7 13:17 pg_serial
drwx----- 2 postgres postgres 6 Jun  7 13:17 pg_snapshots
drwx----- 2 postgres postgres 6 Jun  7 13:17 pg_stat
drwx----- 2 postgres postgres 6 Jun  7 13:17 pg_stat_tmp
drwx----- 2 postgres postgres 6 Jun  7 13:17 pg_subtrans
drwx----- 2 postgres postgres 6 Jun  7 13:17 pg_tblspc
drwx----- 2 postgres postgres 6 Jun  7 13:17 pg_twophase
-rw-r----- 1 postgres postgres 3 Jun  7 13:17 PG_VERSION
drwx----- 3 postgres postgres 60 Jun  7 13:17 pg_wal
drwx----- 2 postgres postgres 18 Jun  7 13:17 pg_xact
-rw-r----- 1 postgres postgres 88 Jun  7 13:17 postgresql.auto.conf
root@postgresql-16-test:/data# systemctl start postgresql@16-main.service
postgres@postgresql-16-test:/data# su postgres
postgres@postgresql-16-test:/data$ cd /
postgres@postgresql-16-test:/# cd data/
postgres@postgresql-16-test:/data$ ls
backup_label      base      pg_commit_ts  pg_logical  pg_notify  pg_serial  pg_stat  pg_subtrans  pg_twophase  pg_wal  postgresql.auto.conf
backup_manifest   global   pg_dynshmem   pg_multixact pg_replslot pg_snapshots pg_stat_tmp pg_tblspc  PG_VERSION  pg_xact

```

Restore compressed backup with wall

```

systemctl stop postgresql@16-main.service
cd data # go to data dirtcory

rm -rf * # remove all conent

cd /backup/base-comp # go to backup diretcory for basebackup
sudo tar -xzf base.tar.gz -C /data/
# If WAL files were included, extract them as well
tar -xzf pg_wal.tar.gz -C /data/pg_wal
#confrim that file permisison are correct
systemctl start postgresql@16-main.service

```

```

root@postgresql-16-test:/backup/base-comp# ls
backup_manifest  base.tar.gz  pg_wal.tar.gz
root@postgresql-16-test:/backup/base-comp# tar -xzf base.tar.gz
-C /data/
-C: command not found
root@postgresql-16-test:/backup/base-comp# tar -xzf base.tar.gz -C /data/
root@postgresql-16-test:/backup/base-comp# If WAL files were included, extract them as well
If: command not found
root@postgresql-16-test:/backup/base-comp# tar -xzf pg_wal.tar.gz -C /data/pg_wal
root@postgresql-16-test:/backup/base-comp# cd /data/
root@postgresql-16-test:/data# ls -l
total 20
-rw----- 1 postgres postgres 225 Jun  7 13:22 backup_label
-rw----- 1 postgres postgres 225 Jun  7 13:17 backup_label.old
drwx----- 7 postgres postgres 59 Jun  7 13:22 base
drwx----- 2 postgres postgres 4096 Jun  7 13:26 global
drwx----- 2 postgres postgres  6 Jun  7 13:17 pg_commit_ts
drwx----- 2 postgres postgres  6 Jun  7 13:17 pg_dynshmem
drwx----- 4 postgres postgres 68 Jun  7 13:22 pg_logical
drwx----- 4 postgres postgres 36 Jun  7 13:17 pg_multixact
drwx----- 2 postgres postgres  6 Jun  7 13:17 pg_notify
drwx----- 2 postgres postgres  6 Jun  7 13:22 pg_replslot
drwx----- 2 postgres postgres  6 Jun  7 13:17 pg_serial
drwx----- 2 postgres postgres  6 Jun  7 13:17 pg_snapshots
drwx----- 2 postgres postgres  6 Jun  7 13:17 pg_stat
drwx----- 2 postgres postgres  6 Jun  7 13:17 pg_stat_tmp
drwx----- 2 postgres postgres  6 Jun  7 13:19 pg_subtrans
drwx----- 2 postgres postgres  6 Jun  7 13:17 pg_tblspc
drwx----- 2 postgres postgres  6 Jun  7 13:17 pg_twophase
-rw----- 1 postgres postgres  3 Jun  7 13:17 PG_VERSION
drwx----- 3 postgres postgres 60 Jun  7 13:27 pg_wal
drwx----- 2 postgres postgres 18 Jun  7 13:17 pg_xact
-rw----- 1 postgres postgres 88 Jun  7 13:17 postgresql.auto.conf
-rw----- 1 postgres postgres  0 Jun  7 13:22 tablespace_map
root@postgresql-16-test:/data# systemctl start postgresql@16-main.service
root@postgresql-16-test:/data# █

```

Advantages of Physical Backups

- **Faster Restoration:** Physical backups can be restored more quickly than logical backups, especially for large databases.
- **Complete Backup:** All databases and global objects are backed up together.
- **Point-in-Time Recovery:** When combined with WAL archiving, physical backups enable point-in-time recovery.
- **Lower Impact:** Physical backups can have less impact on database performance than logical backups.

Limitations of Physical Backups

- **Version Dependency:** Physical backups are version-specific and cannot be used for upgrading to newer PostgreSQL versions.
- **All or Nothing:** You cannot selectively backup or restore specific databases or tables.
- **Platform Dependency:** Physical backups may not be portable across different operating systems or architectures.

Advanced Backup Management with pgBackRest

this actually my prefer way of backing up PostgreSQL , it mush advanced with options such as incremental backup , you can defend the retention for backup and its more automated , you can also setup dedicated server for taking backup for multilabel PostgreSQL and you can also move the backup to cloud storage such as S3

Installing pgBackRest

i personally recommended that you install pgBackRest on sperate server and run the backup form there but its not bad idea to run pgBackRest on same PostgreSQL vm

```
sudo apt-get install pgbackrest
```



```

ivdrental=# \q
root@postgresql-16-test:/data# sudo apt-get install pgbackrest
Reading package lists... Done
Building dependency tree... Done
Reading state information... Done
The following additional packages will be installed:
  libssh2-1t64
Suggested packages:
  pgbackrest-doc check-pgbackrest
The following NEW packages will be installed:
  libssh2-1t64 pgbackrest
0 upgraded, 2 newly installed, 0 to remove and 167 not upgraded.
Need to get 646 kB of archives.
After this operation, 1,742 kB of additional disk space will be used.
Do you want to continue? [Y/n] y
Get:1 http://archive.ubuntu.com/ubuntu noble/main amd64 libssh2-1t64 amd64 1.11.0-4.1build2 [120 kB]
Get:2 http://archive.ubuntu.com/ubuntu noble/universe amd64 pgbackrest amd64 2.50-1build2 [527 kB]
Fetched 646 kB in 3s (254 kB/s)
Selecting previously unselected package libssh2-1t64:amd64.
(Reading database ... 85939 files and directories currently installed.)
Preparing to unpack .../libssh2-1t64_1.11.0-4.1build2_amd64.deb ...
Unpacking libssh2-1t64:amd64 (1.11.0-4.1build2) ...
Selecting previously unselected package pgbackrest.
Preparing to unpack .../pgbackrest_2.50-1build2_amd64.deb ...
Unpacking pgbackrest (2.50-1build2) ...
Setting up libssh2-1t64:amd64 (1.11.0-4.1build2) ...
Setting up pgbackrest (2.50-1build2) ...
Processing triggers for libc-bin (2.39-0ubuntu8.4) ...
Processing triggers for man-db (2.12.0-4build2) ...
Scanning processes...
Scanning linux images...

Running kernel seems to be up-to-date.

No services need to be restarted.

No containers need to be restarted.

No user sessions are running outdated binaries.

No VM guests are running outdated hypervisor (qemu) binaries on this host.
root@postgresql-16-test:/data#

```

go to `postgresql.conf` and update the below parameter and restart the services

```

#-----
# WRITE-AHEAD LOG / Archiving
#-----

# This is the main setting to fix the error.
archive_mode = on

# This tells PostgreSQL HOW to archive. pgBackRest provides its own command.
# Make sure 'demo' matches your --stanza name.
archive_command = 'pgbackrest --stanza=demo archive-push %p'

#-----
# WRITE-AHEAD LOG / Settings
#-----

# 'replica' is the minimum level required for pgBackRest backups.
wal_level = replica

#-----
# REPLICATION / Sending Servers
#-----

# pgBackRest needs at least one sender for differential/incremental backups.
# A value like 10 is a safe default.
max_wal_senders = 10

```

open the `/etc/pgbackrest.conf` and update backup location retention for backup

```

[global]
repo1-path=/var/lib/pgbackrest
repo1-retention-full=2

```



```
[demo]
pg1-path=/var/lib/postgresql/14/main
```

In this configuration: - [global] section defines global settings. - repo1-path specifies the location where backups will be stored. - repo1-retention-full specifies how many full backups to retain. - [demo] is a stanza name (a logical grouping of PostgreSQL databases). - pg1-path specifies the PostgreSQL data directory.

Creating a pgBackRest Stanza

Before performing backups, you need to create a stanza:

```
sudo -u postgres pgbackrest --stanza=demo stanza-create
```

Creating a full backup:

```
sudo -u postgres pgbackrest --stanza=demo backup
```

Creating a differential backup:

```
sudo -u postgres pgbackrest --stanza=demo --type=diff backup
```

```
root@postgresql-16-test:/backup# systemctl restart postgresql@16-main.service
root@postgresql-16-test:/backup# sudo -u postgres pgbackrest --stanza=demo backup
WARN: no prior backup exists, incr backup has been changed to full
root@postgresql-16-test:/backup# cd pgbackrest/
root@postgresql-16-test:/backup/pgbackrest# ls
archive backup
root@postgresql-16-test:/backup/pgbackrest#
```

Creating an incremental backup:

```
sudo -u postgres pgbackrest --stanza=demo --type=incr backup
```

Creating a backup with compression:

```
sudo -u postgres pgbackrest --stanza=demo --compress-level=6 backup
```

```
root@postgresql-16-test:/backup/pgbackrest/backup/demo# sudo -u postgres pgbackrest --stanza=demo --compress-level=6 backup
root@postgresql-16-test:/backup/pgbackrest/backup/demo#
```

Restoring with pgBackRest

to restore backup is straight forward

first stop postgresql and then run the command to restore the latest backup

```
systemctl stop postgresql@16-main.service
rm -rf * # remove content of data dir
sudo -u postgres pgbackrest --stanza=demo restore
systemctl start postgresql@16-main.service
```

```
root@postgresql-16-test:/backup/pgbackrest/backup/demo# cd /data/
root@postgresql-16-test:/data# rm -rf *
root@postgresql-16-test:/data# sudo -u postgres pgbackrest --stanza=demo restore
root@postgresql-16-test:/data# ls
backup_label  global      pg_dynshmem  pg_multixact  pg_replslot  pg_snapshots  pg_stat_tmp  pg_tblspc  PG_VERSION  pg_xact      recovery.signal
base          pg_commit_ts  pg_logical    pg_notify     pg_serial    pg_stat       pg_subtrans  pg_twophase  pg_wal      postgresql.auto.conf  tablespace_map.old
root@postgresql-16-test:/data# systemctl start postgresql@16-main.service
root@postgresql-16-test:/data#
```

What it Does: pgBackRest will automatically:

- Find the latest full backup.
- Apply any differential or incremental backups that were taken after it.
- Replay all the archived WAL files to bring the database to the most recent possible state.

restore certain full backup

for that you will first check what are backup there in pgbackrest using the following command

```
sudo -u postgres pgbackrest --stanza=demo info
```

<https://www.linkedin.com/in/ahmed-mohamed-423583151>

```

root@postgresql-16-test:/data# sudo -u postgres pgbackrest --stanza=demo info
stanza: demo
  status: ok (backup/expire running - 99.90% complete)
  cipher: none

  db (current)
    wal archive min/max (16): 00000001000000000000000A/000000020000000000000014

    full backup: 20250607-134346F
      timestamp start/stop: 2025-06-07 13:43:46+00 / 2025-06-07 13:43:59+00
      wal start/stop: 000000010000000000000000C / 000000010000000000000000C
      database size: 52.6MB, database backup size: 52.6MB
      repo1: backup set size: 8.4MB, backup size: 8.4MB

    incr backup: 20250607-134346F_20250607-134916I
      timestamp start/stop: 2025-06-07 13:49:16+00 / 2025-06-07 13:49:17+00
      wal start/stop: 000000010000000000000000E / 000000010000000000000000E
      database size: 52.6MB, database backup size: 8.3KB
      repo1: backup set size: 8.4MB, backup size: 431B
      backup reference list: 20250607-134346F

    incr backup: 20250607-134346F_20250607-135007I
      timestamp start/stop: 2025-06-07 13:50:07+00 / 2025-06-07 13:50:09+00
      wal start/stop: 0000000100000000000000010 / 0000000100000000000000010
      database size: 52.6MB, database backup size: 8.3KB
      repo1: backup set size: 8.4MB, backup size: 430B
      backup reference list: 20250607-134346F, 20250607-134346F_20250607-134916I
root@postgresql-16-test:/data#

```

The output will show you a list of your full and incremental backups with their labels (e.g., 20250607-161005F for a full backup).

we will restore backup of first incremental using backup reference list

```

full backup: 20250607-134346F
  timestamp start/stop: 2025-06-07 13:43:46+00 / 2025-06-07 13:43:59+00
  wal start/stop: 000000010000000000000000C / 000000010000000000000000C
  database size: 52.6MB, database backup size: 52.6MB
  repo1: backup set size: 8.4MB, backup size: 8.4MB

incr backup: 20250607-134346F_20250607-134916I
  timestamp start/stop: 2025-06-07 13:49:16+00 / 2025-06-07 13:49:17+00
  wal start/stop: 000000010000000000000000E / 000000010000000000000000E
  database size: 52.6MB, database backup size: 8.3KB
  repo1: backup set size: 8.4MB, backup size: 431B
  backup reference list: 20250607-134346F

incr backup: 20250607-134346F_20250607-135007I
  timestamp start/stop: 2025-06-07 13:50:07+00 / 2025-06-07 13:50:09+00
  wal start/stop: 0000000100000000000000010 / 0000000100000000000000010
  database size: 52.6MB, database backup size: 8.3KB
  repo1: backup set size: 8.4MB, backup size: 430B
  backup reference list: 20250607-134346F, 20250607-134346F_20250607-134916I

```

```

# Stop PostgreSQL
sudo systemctl stop postgresql@16-main.service
rm -rf * # remove content of data dir
# Perform the restore
sudo -u postgres pgbackrest --stanza=demo --set=20250607-134346F restore

```

<https://www.linkedin.com/in/ahmed-mohamed-423583151>

```

root@postgresql-16-test:/data# sudo -u postgres pgbackrest --stanza=demo info
stanza: demo
status: ok (backup/expire running - 99.90% complete)
cipher: none

db (current)
wal archive min/max (16): 000000010000000000000000A/0000000200000000000000014

full backup: 20250607-134346F
timestamp start/stop: 2025-06-07 13:43:46+00 / 2025-06-07 13:43:59+00
wal start/stop: 000000010000000000000000C / 000000010000000000000000C
database size: 52.6MB, database backup size: 52.6MB
repol: backup set size: 8.4MB, backup size: 8.4MB

incr backup: 20250607-134346F 20250607-134916I
timestamp start/stop: 2025-06-07 13:49:16+00 / 2025-06-07 13:49:17+00
wal start/stop: 000000010000000000000000E / 000000010000000000000000E
database size: 52.6MB, database backup size: 8.3KB
repol: backup set size: 8.4MB, backup size: 431B
backup reference list: 20250607-134346F

incr backup: 20250607-134346F_20250607-135007I
timestamp start/stop: 2025-06-07 13:50:07+00 / 2025-06-07 13:50:09+00
wal start/stop: 0000000100000000000000010 / 0000000100000000000000010
database size: 52.6MB, database backup size: 8.3KB
repol: backup set size: 8.4MB, backup size: 430B
backup reference list: 20250607-134346F, 20250607-134916I
root@postgresql-16-test:/data# sudo systemctl stop postgresql@16-main.service
root@postgresql-16-test:/data# rm -rf *
root@postgresql-16-test:/data# sudo -u postgres pgbackrest --stanza=demo --set=20250607-134346F restore
root@postgresql-16-test:/data# sudo systemctl start postgresql@16-main.service
root@postgresql-16-test:/data# pg_lsclusters
Ver Cluster Port Status Owner    Data directory Log file
16 main    5432 online <unknown> data/      /var/log/postgresql/postgresql-16-main.log
root@postgresql-16-test:/data#

```

Restore to a specific point in time:

```

# Stop PostgreSQL
sudo systemctl stop postgresql@16-main.service
rm -rf * # remove content of data dir
# Perform the restore

sudo -u postgres pgbackrest --stanza=demo --type=time --target="2025-06-07 16:45:00+03" restore
# Start PostgreSQL
sudo systemctl start postgresql@16-main.service

```

You need to use the `--type=time` and `--target` flags. The timestamp must be in a format PostgreSQL understands and should include your timezone.

Managing Backups with pgBackRest

to check the backup available at pgbackrest use the following command

```
sudo -u postgres pgbackrest --stanza=demo info
```

Checking backup integrity:

```
sudo -u postgres pgbackrest --stanza=demo check
```

MySQL Backup Strategies

logical backup using mysqldump

mysqldump is utility for taking dump backup similar to what pg_dump does

it can take all database backup or single database , it can also take specific table backup

taking backup using mysqldump

for this scenario we have sakila database sample already deployed on MySQL we will take dump backup and store in the dedicated backup/ directory

Backing up a single database:

```
mysqldump -h localhost -u root -p --single-transaction --routines --triggers --events mydatabase > mydatabase_backup.sql
```

```
[ahmed@localhost ~]$ ls  
mydatabase_backup.sql  mysql84-community-release-el8-1.noarch.rpm  sakila-db  sakila-db.zip  
[ahmed@localhost ~]$
```

```
mysqldump -h localhost -u root -p --single-transaction --routines --triggers --events --all-databases > all_databases_backup.sql
```

Backing up specific tables:

```
mysqldump -h localhost -u root -p --single-transaction sakila actor film > tables_backup.sql
```

Creating a compressed backup:

```
mysqldump -h localhost -u root -p --single-transaction --routines --triggers --events sakila | gzip > mydatabase_backup.sql.gz
```

```
[ahmed@localhost ~]$ less tables_backup.sql  
[ahmed@localhost ~]$ mysqldump -h localhost -u root -p --single-transaction --routines --triggers --events sakila | gzip > mydatabase_backup.sql.gz  
Enter password:  
[ahmed@localhost ~]$ ls  
all_databases_backup.sql  mydatabase_backup.sql  mydatabase_backup.sql.gz  mysql84-community-release-el8-1.noarch.rpm  sakila-db  sakila-db.zip  tables_backup.sql  
[ahmed@localhost ~]$ ll  
total 9484  
-rw-rw-r-- 1 ahmed ahmed 4620018 Jun  8 02:30 all_databases_backup.sql  
-rw-rw-r-- 1 ahmed ahmed 3398040 Jun  8 02:29 mydatabase_backup.sql  
-rw-rw-r-- 1 ahmed ahmed 715021 Jun  8 02:33 mydatabase_backup.sql.gz  
-rw-rw-r-- 1 ahmed ahmed 15364 Apr 29 2024 mysql84-community-release-el8-1.noarch.rpm  
drwxr-xr-x 2 ahmed ahmed 72 Jun  1 01:05 sakila-db  
-rw-rw-r-- 1 ahmed ahmed 729655 Jun  1 01:05 sakila-db.zip  
-rw-rw-r-- 1 ahmed ahmed 223759 Jun  8 02:31 tables_backup.sql  
[ahmed@localhost ~]$
```

mysqldump options

- `--single-transaction` this recommended to include in mysqldump command it will avoid locking the table while taking database dump
- `--flush-logs` Flushes the MySQL server log files before starting the dump
- `--master-data=2` Includes binary log position in the dump output as a commented CHANGE MASTER statement
- `--routines` Includes stored routines (procedures and functions) in the
- `--triggers` Includes triggers in the dump
- `--events` Includes events in the dump
- `--no-data` Dumps only the database structure, not the contents of the tables

Best Practices for mysqldump

- Use `--single-transaction`: For InnoDB tables, always use `--single-transaction` to create consistent backups without locking.
- Include Routines, Triggers, and Events: Use `--routines`, `--triggers`, and `--events` to ensure all database objects are backed up.
- Compression: Use compression to reduce storage requirements, especially for large databases.

Restoring mysqldump backup

Restoring a single database backup:

```
mysql -h localhost -u root -p < mydatabase_backup.sql
```

Restoring a compressed backup

```
gunzip < mydatabase_backup.sql.gz | mysql -h localhost -u root -p
```

Restoring to a specific database:

```
mysql -h localhost -u root -p mydatabase < mydatabase_backup.sql
```

Physical Backups with Percona XtraBackup

now this is my desire way of taking backup , it better for production has less effect on database during backup , and restoration is much faster , it also support differential backup and the backup and restore process is much more streamline than mysqldump for production environment i highly recommended either use Percona XtraBackup or MySQL enterprise backup if you are using enterprise edition of MySQL

key features of Percona XtraBackup include:

- Non-blocking: Creates consistent backups without locking your database.
- Hot Backup: Backs up your database while it's running.
- Incremental Backup: Supports incremental backups, reducing backup time and storage requirements.
- Compression and Encryption: Supports data compression and encryption.
- Streaming: Can stream backups to another server or to cloud storage.
- Point-in-Time Recovery: Supports point-in-time recovery when used with binary logs.
- Partial Backups: Supports backing up specific databases or tables.

Installing Percona XtraBackup

to install Percona XtraBackup on redhat follow the below steps

```
# Add Percona repository
sudo yum install https://repo.percona.com/yum/percona-release-latest.noarch.rpm
sudo percona-release setup ps80
# Install Percona XtraBackup
sudo yum install percona-xtrabackup-80
```

```
[ahmed@localhost ~]$ sudo yum install https://repo.percona.com/yum/percona-release-latest.noarch.rpm
[sudo] password for ahmed:
Last metadata expiration check: 3:16:10 ago on Sat 07 Jun 2025 11:33:22 PM +03.
percona-release-latest.noarch.rpm
Dependencies resolved.
13 kB/s | 28 kB    00:02
=====
Package                Architecture      Version           Repository        Size
-----
Installing:
percona-release         noarch            1.0-30           @commandline     28 k
Transaction Summary
-----
Install 1 Package
Total size: 28 k
Installed size: 49 k
Is this ok [y/N]: y
Downloading Packages:
Running transaction check
Transaction check succeeded.
Running transaction test
Transaction test succeeded.
Running transaction
  Preparing                : percona-release-1.0-30.noarch
  Installing               : percona-release-1.0-30.noarch
  Running scriptlet        : percona-release-1.0-30.noarch
* Enabling the Percona Release repository
<=> All done!
* Enabling the Percona Telemetry repository
<=> All done!
* Enabling the PMM2 Client repository
<=> All done!
The percona-release package now contains a percona-release script that can enable additional repositories for our newer products.

Note: currently there are no repositories that contain Percona products or distributions enabled. We recommend you to enable Percona Distribution repositories instead of individual product repositories, because with the Distribution you will get not only the database itself but also a set of other components that will help you work with your database.

For example, to enable the Percona Distribution for MySQL 8.0 repository use:

percona-release setup pdps8.0

Note: To avoid conflicts with older product versions, the percona-release setup command may disable our original repository for some products.

For more information, please visit:
https://docs.percona.com/percona-software-repositories/percona-release.html

Verifying                : percona-release-1.0-30.noarch
Installed:
percona-release-1.0-30.noarch
Complete!
```

Creating a full backup

to create full backup use the following syntax , for this instant i will take full backup and store it in the dedicated backup directory with separate mount point

```
xtrabackup --backup --target-dir=/path/to/backup
```

Creating a full backup with compression:

```
xtrabackup --backup --compress --target-dir=/path/to/backup
```

Creating a full backup for specific databases:

```
xtrabackup --backup --databases="db1 db2" --target-dir=/path/to/ backup
```

Creating an incremental backup:

```
xtrabackup --backup --target-dir=/path/to/incremental/backup1 -- incremental-basedir=/path/to/full/backup
```

Preparing Backups for Restoration

Before a backup can be restored, it needs to be prepared. This applies the committed transactions and rolls back the uncommitted ones:

Preparing a full backup:

```
xtrabackup --prepare --target-dir=/path/to/backup
```

Preparing a full backup with incremental backups:

Prepare the full backup with --apply-log-only

```
xtrabackup --prepare --apply-log-only --target-dir=/path/to/ full/backup
```

Apply the first incremental backup with --apply-log-only

```
xtrabackup --prepare --apply-log-only --target-dir=/path/to/ full/backup --incremental-dir=/path/to/incremental/backup1
```

Apply the second incremental backup (last one without --apply-log-only)

```
xtrabackup --prepare --target-dir=/path/to/full/backup -- incremental-dir=/path/to/incremental/backup2
```

Stop MySQL:

```
sudo systemctl stop mysqld
```

Move or rename the current data directory:

```
sudo mv /var/lib/mysql /var/lib/mysql.bak
```

Create a new data directory:

```
sudo mkdir /var/lib/mysql
```

Copy the prepared backup to the data directory:

```
sudo xtrabackup --copy-back --target-dir=/path/to/backup
```

Set the correct ownership:

```
sudo chown -R mysql:mysql /var/lib/mysql
```

Start MySQL:

```
sudo systemctl start mysqld
```

<https://www.linkedin.com/in/ahmed-mohamed-423583151>