

# Understanding the Difference Between dblink and postgres\_fdw in PostgreSQL

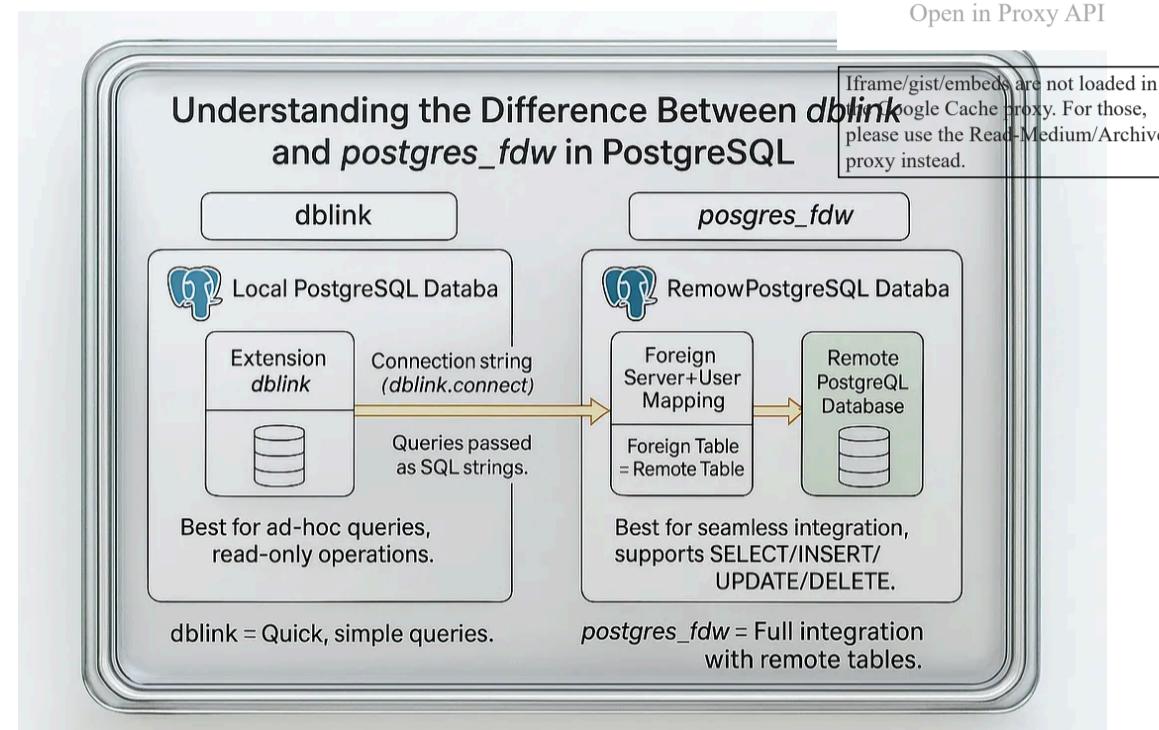
[Open in Google Cache](#)[Open in Read-Medium](#)[Open in Freedium](#)

Jeyaram Ayyalusamy

Following



9 min read · 9 hours ago

[Open in Archive.today](#)[Open in Archive.is](#)[Open in Proxy API](#)

PostgreSQL is not just a relational database; it is a highly extensible system. Two of its most commonly used extensions for working across multiple PostgreSQL databases are **dblink** and **postgres\_fdw** (Foreign Data Wrapper).

At first glance, they might look similar because both allow you to connect to another PostgreSQL database. But in reality, they serve different purposes and have different strengths.

## ◆ 1. What is dblink ?

The **dblink** extension is one of the oldest ways to connect one PostgreSQL database to another. It allows you to run **queries as text strings** against a remote database.

- Think of it as a **bridge**: you connect to the other database, send a query as text, and get the results back.
- You must manually define the expected column structure when pulling data.
- It's lightweight but also limited — best suited for **ad-hoc queries**.

### Example workflow with dblink :

1. Install and enable the extension (`CREATE EXTENSION dblink`).
2. Configure a remote connection.
3. Run a query like:

```
SELECT *
FROM dblink('conn_db_link', 'SELECT item_id, quantity, price FROM items')
AS result(item_id INT, quantity INT, price NUMERIC(10,2));
```

```
postgres=# SELECT *
postgres-# FROM dblink('conn_db_link', 'SELECT item_id, quantity, price FROM items')
postgres-# AS result(item_id INT, quantity INT, price NUMERIC(10,2));
      item_id   |   quantity   |   price
-----+-----+-----+
    9093121 |     297 | 395.80
    9093122 |     640 | 226.10
    9093123 |     238 | 173.32
    9093124 |     200 | 80.91
    9093125 |      18 | 222.13
    9093126 |     945 | 101.12
    9093127 |     419 | 192.06
    9093128 |     739 | 493.85
    9093129 |     755 | 163.90
    9093130 |     374 | 69.76
    9093131 |     613 | 339.56
    9093132 |     923 | 75.71
    9093133 |     917 | 461.97
    9093134 |     916 | 156.12
    9093135 |     719 | 433.66
```

Here, the SQL query is written as a string and executed on the remote server.

## ◆ 2. What is postgres\_fdw?

The **postgres\_fdw extension** is a modern, more powerful replacement for many use cases of `dblink`. FDW stands for **Foreign Data Wrapper**, which is part of PostgreSQL's implementation of the SQL/MED standard (Management of External Data).

- Instead of sending queries as strings, you can create **foreign tables** in your local database that directly map to remote tables.
- These foreign tables look and behave like normal local tables.
- You can **SELECT, INSERT, UPDATE, and DELETE** records across databases.
- PostgreSQL's **query planner** can optimize queries across local and remote data sources.

### Example workflow with `postgres_fdw`:

1. Install and enable the extension (`CREATE EXTENSION postgres_fdw`).
2. Create a foreign server that points to the remote database.
3. Create a user mapping with credentials.
4. Define a foreign table that mirrors the remote table.
5. Run queries against it directly, without writing raw SQL strings.

```
SELECT * FROM items_remote WHERE price > 100;
```

The query above looks like a local query, but it is executed on the remote database behind the scenes.

### ◆ 3. Key Differences at a Glance

Feature	<code>dblink</code>	<code>postgres_fdw</code>	Query style	SQL passed as a string	Normal SQL
on foreign tables			Integration level	Lightweight, ad-hoc	Full integration with local schema
<b>Operations supported</b>			Mostly SELECT (writes are clunky)		
SELECT, INSERT, UPDATE, DELETE			Query optimization	No	Yes (Postgres planner pushes conditions to remote server)
			Ease of use	Simple for quick queries	
			More setup, but easier for ongoing use	Use case	One-time lookups, quick joins
			Long-term integration, production workloads		

👉 In short:

- Use `dblink` if you just need a quick lookup across databases.
- Use `postgres_fdw` if you need continuous, reliable, and optimized access to remote tables.

## Step-by-Step Guide: Implementing `postgres_fdw` in PostgreSQL

### 15

Now that we understand the difference, let's dive into the step-by-step process of setting up and using `postgres_fdw` in PostgreSQL 15.

#### Step 1: Install the contrib package

If you haven't already installed contrib modules for PostgreSQL 15, do so:

```
sudo yum install postgresql15-contrib
```

```
[root@ggnode1 dbs]# sudo yum install postgresql15-contrib      # or: sudo yum inst
ol7_UEKR6
ol7_addons
ol7_developer
ol7_developer_EPEL
ol7_latest
ol7_optional_latest
ol7_software_collections
pgdg-common/7Server/x86_64/signature
pgdg-common/7Server/x86_64/signature
pgdg12/7Server/x86_64/signature
pgdg12/7Server/x86_64/signature
pgdg13/7Server/x86_64/signature
pgdg13/7Server/x86_64/signature
pgdg14/7Server/x86_64/signature
pgdg14/7Server/x86_64/signature
pgdg15/7Server/x86_64/signature
pgdg15/7Server/x86_64/signature
Resolving Dependencies
--> Running transaction check
--> Package postgresql15-contrib.x86_64 0:15.14-1PGDG.rhel7 will be installed
--> Processing Dependency: libpython3.6m.so.1.0()(64bit) for package: postgresql
--> Running transaction check
--> Package python3-libs.x86_64 0:3.6.8-21.0.1.el7_9 will be installed
--> Processing Dependency: python(abi) = 3.6 for package: python3-libs-3.6.8-21.
--> Running transaction check
--> Package python3.x86_64 0:3.6.8-21.0.1.el7_9 will be installed
```

```
--> Processing Dependency: python3-pip for package: python3-3.6.8-21.0.1.el7_9.x
--> Processing Dependency: python3-setuptools for package: python3-3.6.8-21.0.1.
--> Running transaction check
--> Package python3-pip.noarch 0:9.0.3-8.0.3.el7 will be installed
--> Package python3-setuptools.noarch 0:39.2.0-10.0.3.el7 will be installed
--> Finished Dependency Resolution
```

Dependencies Resolved

```
=====
Package                               Arch
=====
Installing:
postgresql15-contrib                  x86_64
Installing for dependencies:
python3                                x86_64
python3-libs                            x86_64
```

≡ Medium

Search

Write

2



Home

Library

Profile

Stories

Stats

Following

Level Up Coding

Jeyaram Ayyalus...

Murat Bilal

Sheikh Wasiu Al H...

Rizqi Mulki

DbVisualizer

Yasemin Büşra Kar...

The Medium Blog

Monowar Mukul

Dario Radečić

More

#### Transaction Summary

Install 1 Package (+4 Dependent packages)

```
Total download size: 10 M
Installed size: 50 M
Is this ok [y/d/N]: y
Downloading packages:
(1/5): python3-3.6.8-21.0.1.el7_9.x86_64.rpm
(2/5): postgresql15-contrib-15.14-1PGDG.rhel7.x86_64.rpm
(3/5): python3-pip-9.0.3-8.0.3.el7.noarch.rpm
(4/5): python3-setuptools-39.2.0-10.0.3.el7.noarch.rpm
(5/5): python3-libs-3.6.8-21.0.1.el7_9.x86_64.rpm
```

Total

This package contains both `dblink` and `postgres_fdw`.

## Step 2: Enable the extension in your database

Log in to PostgreSQL and enable the FDW extension:

```
CREATE EXTENSION postgres_fdw;
```

```
postgres=# CREATE EXTENSION postgres_fdw;
CREATE EXTENSION
postgres=#

```

Verify that it's installed:

```
\dx postgres_fdw
```

```
postgres=# \dx postgres_fdw
              List of installed extensions
   Name    | Version | Schema | Description
   postgres_fdw | 1.1     | public  | foreign-data wrapper for remote PostgreSQL se
(1 row)
```

```
postgres=#
```

### Step 3: Create a foreign server

Define a connection to the remote PostgreSQL server:

```
CREATE SERVER server_remote FOREIGN DATA WRAPPER postgres_fdw
postgres=# OPTIONS (host '192.168.50.11', dbname 'postgres', port '5432');
```

```
postgres=# CREATE SERVER server_remote FOREIGN DATA WRAPPER postgres_fdw
postgres=# OPTIONS (host '192.168.50.11', dbname 'postgres', port '5432');
CREATE SERVER
postgres=#

```

Here:

- **host** = IP address of the remote server
- **dbname** = remote database name
- **port** = default PostgreSQL port (5432)

### Step 4: Create a user mapping

Map the current user to the remote user by providing login credentials:

```
CREATE USER MAPPING FOR current_user SERVER server_remote
OPTIONS (user 'postgres', password 'oracle123');
```

```
postgres=# CREATE USER MAPPING FOR current_user SERVER server_remote
postgres=# OPTIONS (user 'postgres', password 'oracle123');
CREATE USER MAPPING
postgres=#

```

This ensures that when the local user connects to the foreign server, it uses the specified username and password.

### Step 5: Import schema or create foreign tables

At this point, you can either import an entire schema from the remote database or manually define a foreign table.

## Option A: Import all tables from a schema

```
IMPORT FOREIGN SCHEMA public FROM SERVER server_remote
INTO public;
```

```
postgres=# IMPORT FOREIGN SCHEMA public FROM SERVER server_remote
postgres=# INTO public;
IMPORT FOREIGN SCHEMA
postgres=#

```

This automatically creates local foreign tables for everything in the remote public schema.

## Viewing Local and Foreign Tables

After setting up `postgres_fdw` and importing or creating foreign tables, you can list all relations (tables, views, sequences, etc.) in your current schema. In PostgreSQL, the command for this is:

```
\d+
```

Example output:

```
postgres=# \d+
                                         List of relations
 Schema |        Name        |   Type   | Owner | Persistence | Acc
-----+----------------+-----+-----+-----+-----+
 public | items          | foreign table | postgres | permanent | 
 public | orders         | foreign table | postgres | permanent | 
 public | products        | table      | postgres | permanent | 
 public | products_product_id_seq | sequence | postgres | permanent | 
(4 rows)

postgres=#

```

## What This Means

- **items (foreign table):**

This is a table that resides on the remote PostgreSQL server but is accessible locally through `postgres_fdw`.

- **orders (foreign table):**

Another remote table, seamlessly integrated into the local database. Queries on this table are executed against the remote database.

- **products (local table):**

A regular table that exists locally in the `public` schema.

- **products\_product\_id\_seq (sequence):**

A sequence object used to generate IDs for the `products` table.

## Key Takeaway

With `postgres_fdw`, your local PostgreSQL instance can contain a mix of local and remote objects.

- To your application, these look almost the same.
- The big difference is that foreign tables delegate operations to the remote database.

This setup makes it possible to query and join data across multiple databases without complex ETL processes.

## Option B: Define a specific foreign table

If you want to add new tables after importing a foreign schema, you cannot re-import the foreign tables. Instead, you can use this option to add additional tables.

```
postgres=# IMPORT FOREIGN SCHEMA public FROM SERVER server_remote
postgres=# INTO public;
ERROR: relation "items" already exists
CONTEXT: importing foreign table "items"
postgres=#

```

```
CREATE FOREIGN TABLE items_remote (
    item_id    INT,
    quantity   INT,
    price      NUMERIC(10,2)
)
SERVER server_remote
OPTIONS (table_name 'items');
```

```
postgres=# CREATE FOREIGN TABLE items_remote (
postgres(#     item_id    INT,
postgres(#     quantity   INT,
postgres(#     price      NUMERIC(10,2)
postgres(# )
postgres-# SERVER server_remote
postgres-# OPTIONS (table_name 'items');
CREATE FOREIGN TABLE
postgres=#

```

Here, `items_remote` is the local representation of the remote `items` table or any new tables.

List of relations					
Schema	Name	Type	Owner	Persistence	Acc
public	items	foreign table	postgres	permanent	
public	items_remote	foreign table	postgres	permanent	

```

public | orders           | foreign table | postgres | permanent | 
public | products         | table        | postgres | permanent | hea
public | products_product_id_seq | sequence    | postgres | permanent | 
(5 rows)

postgres=#

```

## Step 6: Query the foreign table

Now you can query the remote table as if it were local:

```
SELECT * FROM items WHERE price > 200;
```

```
postgres=# SELECT * FROM items WHERE price > 200;
```

item_id	quantity	price
9093121	297	395.80
9093122	640	226.10
9093125	18	222.13
9093128	739	493.85
9093131	613	339.56
9093133	917	461.97
9093135	719	433.66
9093137	211	212.29
9093138	732	212.20
9093140	319	295.23

## Step 7: Perform DML operations

Unlike `dblink`, `postgres_fdw` allows you to modify data on the remote server.

- Delete data:

```
DELETE FROM items WHERE item_id = 9093129;
```

```

postgres=# DELETE FROM items WHERE item_id = 9093129;
DELETE 1
postgres=#

```

- Insert data:

```
INSERT INTO items VALUES (9093129, 500, 250.00);
```

```
postgres=# INSERT INTO items VALUES (9093129, 500, 250.00);
INSERT 0 1
postgres=#

```

- Update data:

```
UPDATE items SET price = 199.99 WHERE item_id = 9093121;

```

```
postgres=# UPDATE items SET price = 199.99 WHERE item_id = 9093121;
UPDATE 1
postgres=#

```

All changes are applied directly to the remote database.

## ◆ Practical Considerations

### 1. Performance:

- `postgres_fdw` pushes down WHERE conditions to the remote server, reducing network traffic.
- Joins between local and foreign tables may be less efficient, so use filters.

### 2. Security:

- Avoid storing plain-text passwords in user mappings. Use role-based authentication or password files.

### 3. Use cases:

- Ideal for distributed reporting, microservices accessing shared data, or consolidating analytics.
- Not a replacement for replication — use logical or streaming replication for high availability.

## ◆ Final Thoughts

- `dblink` is the older, simpler way: best for occasional queries.
- `postgres_fdw` is the modern, integrated way: best for production systems where remote tables should behave like local ones.

With the steps above, you can seamlessly connect PostgreSQL databases using `postgres_fdw` and perform both read and write operations across

## 💡 Stay Updated with Daily PostgreSQL & Cloud Tips!

If you've been finding my blog posts helpful and want to stay ahead with daily insights on PostgreSQL, Cloud Infrastructure, Performance Tuning, and DBA Best Practices — I invite you to subscribe to my Medium account.

🔔 [Subscribe here ↗ https://medium.com/@jramcloud1/subscribe](https://medium.com/@jramcloud1/subscribe)

Your support means a lot — and you'll never miss a practical guide again!

## 🔗 Let's Connect!

If you enjoyed this post or would like to connect professionally, feel free to reach out to me on LinkedIn:

👉 [Jeyaram Ayyalusamy](#)

I regularly share content on **PostgreSQL, database administration, cloud technologies, and data engineering**. Always happy to connect, collaborate, and discuss ideas!

[Postgresql](#) [Oracle](#) [Open Source](#) [MySQL](#) [AWS](#)

 **Written by Jeyaram Ayyalusamy** 
171 followers · 2 following
Following ▾

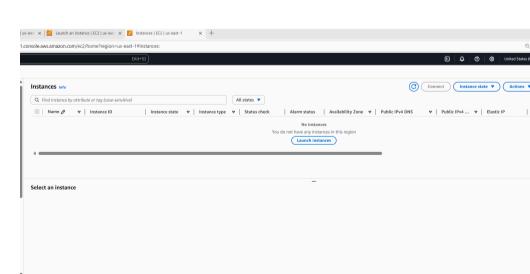
Oracle DBA | AWS Certified | 18+ yrs in DB Admin, RAC, GoldenGate, RDS, MySQL, PostgreSQL | Author of MySQL & RDS books | Cloud & Performance Expert

## No responses yet

 Gvadakte

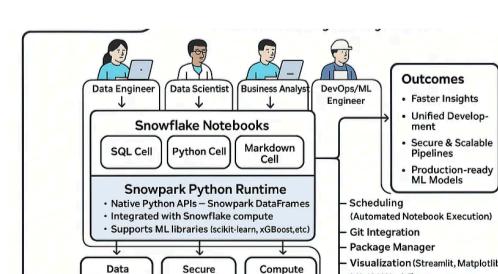
What are your thoughts?

## More from Jeyaram Ayyalusamy



 Jeyaram Ayyalusamy 

**Upgrading PostgreSQL from Version 16 to Version 17 Using...**



 Jeyaram Ayyalusamy 

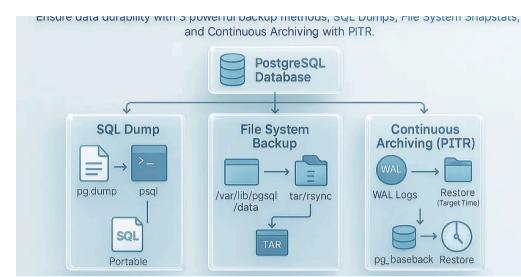
**16—Experience Snowflake with Notebooks and Snowpark Python...**

 A Complete Step-by-Step Guide to Installing PostgreSQL 16 on a Linux Server...

Aug 4  40

In the fast-moving world of data, organizations are no longer just collecting...

 ...



J Jeyaram Ayyalusamy 

## 💡 Essential Techniques Every DBA Should Know: PostgreSQL...

In the world of databases, data is everything—and losing it can cost your business time,...

Aug 20  26



J Jeyaram Ayyalusamy 

## PostgreSQL Users and Roles Explained: A Complete Guide for...

Managing access and permissions is at the heart of any database system—and...

 ...

Jun 10  53

See all from Jeyaram Ayyalusamy

## Recommended from Medium



R Rizqi Mulki

## SQL Query Optimization: Tips That Actually Work

Stop writing slow queries that make your users (and your database) cry

 Aug 10  2



The CS Engineer

## Forget SQL Joins—These 4 Patterns Made My Queries 10x...

One four-table join turned a 30-millisecond request into a two-second disaster. I ripped...

 ...

 343

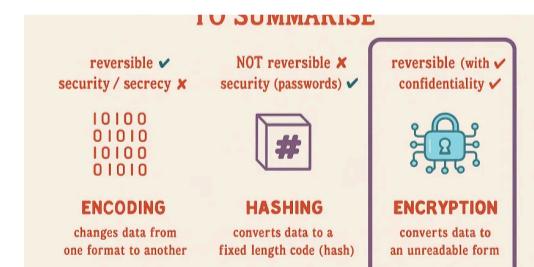


# Hash Block

## 10 Postgres Logical Replication Recipes for Zero-Downtime...

Practical, cut-pasteable playbooks that move data between clusters without pausing the...

 4d ago  3



Dev Cookies

## Encoding vs Hashing vs Encryption—Explained with...

When working with data security, three concepts often create confusion: encoding,...

 ...

Aug 22  8



Saumya Bhatt

Chronicles

## Inside PostgreSQL Replication: WAL, Logical Slots, and CDC

PostgreSQL is an incredibly powerful database—and at the heart of its reliability...

Jun 27 36

+ ...

Aug 9 90

+ ...

[See more recommendations](#)

[Help](#) [Status](#) [About](#) [Careers](#) [Press](#) [Blog](#) [Privacy](#) [Rules](#) [Terms](#) [Text to speech](#)