

Finding & Fixing Missing Indexes in Under 10 Minutes

"Most slow queries are just an index away from being fast." — every seasoned DBA

Unindexed columns turn your sleek Postgres into a table-scanning slug. This crash course shows how to detect, prioritize, and create the right indexes before your next coffee refill. No extensions needed — just core catalog views (`pg_stat_*`, `pg_index`, `pg_constraint`) and a dash of SQL.

1 • Spot Sequential Scans at Scale (1 min)

```
SELECT relname      AS table,
       seq_scan      AS seq_scans,
       idx_scan      AS idx_scans,
       round(100*seq_scan/NULLIF(seq_scan+idx_scan,0),2) AS seq_pct,
       n_live_tup     AS rowsFROM   pg_stat_user_tablesWHERE  n_live_tup > 10000    -- big
enough to matterORDER BY seq_scan DESC
LIMIT 15;
```

- `seq_pct > 10 %` on tables > 10 k rows → likely missing or unused indexes.
- Small tables can live with seq scans; indexes add overhead.

2 • Pinpoint Offending Columns with EXPLAIN (2 min)

1. Get top queries for the table:

```
SELECT queryFROM   pg_stat_statementsWHERE  query LIKE '%big_table%'ORDER BY
total_exec_time DESC
LIMIT 5;
```

1. Run `EXPLAIN (ANALYZE, BUFFERS)` on the slowest query.

Look for:

- Seq Scan with a filter → **add normal / partial index**
- Bitmap Heap Scan removing many rows → **covering index**
- Joins lacking Index Scan on FK side → **index the FK column**

Need help reading plans? See **EXPLAIN ANALYZE Demystified**.

3 • Automatic Foreign-Key Audit (3 min)

Find FK columns missing indexes:

Copy

```
WITH fks AS (  
  SELECT conrelid, conname, conkey  
  FROM pg_constraint  
  WHERE contype = 'f'  
) , ix AS (  
  SELECT indrelid, indkey  
  FROM pg_index  
  WHERE indisvalid AND indpred IS NULL  
) SELECT n.nspname || '.' || c.relname AS table,  
  f.conname AS fk_name,  
  array_to_string(ARRAY(  
    SELECT a.attname  
    FROM pg_attribute a  
    WHERE a.attrelid = f.conrelid  
    AND a.attnum = ANY(f.conkey)  
    ORDER BY array_position(f.conkey, a.attnum)  
  ), ',') AS fk_cols FROM fks f JOIN pg_class c ON c.oid = f.conrelid JOIN pg_namespace n ON  
n.oid = c.relnamespace  
LEFT JOIN ix ON ix.indrelid = f.conrelid  
AND f.conkey = ix.indkey[0:array_length(f.conkey,1)-1] WHERE ix.indrelid IS NULL -- no  
supporting index ORDER BY table;
```

Result lists FK constraints that need indexes.

4 • Generate Index DDL (1 min)

Auto-craft index statements:

Copy

```
SELECT format(  
  'CREATE INDEX CONCURRENTLY IF NOT EXISTS idx_%I_%s ON %s USING btree (%s);',  
  relname, -- table name  
  string_agg(col, '_'), -- suffix  
  relid::regclass, -- schema.table  
  string_agg(col, ',') -- column list  
) AS ddl FROM your_fk_missing_index_query GROUP BY relid, relname;
```

Run the generated `CREATE INDEX CONCURRENTLY` outside a transaction block.

5 • Partial & Covering Index Patterns (2 min)

Scenario	Index Recipe	Why
Soft deletes (WHERE deleted_at IS NULL)	CREATE INDEX ... WHERE deleted_at IS NULL	Smaller, faster scans
Recent rows (created_at > NOW()-30d)	CREATE INDEX ... ON ... (created_at) WHERE created_at > ...	Keeps old data out of index
Filter + sort (status='paid' ORDER BY date)	CREATE INDEX ... (status, created_at DESC)	Supports filter *and* order by
Covering lookup (select few cols)	CREATE INDEX ... (id) INCLUDE (col1, col2)	Enables index-only scan

Use `pg_size_pretty(pg_relation_size('index_name'))` to verify size savings.

6 • Validate Impact (1 min)

Copy

```
EXPLAIN (ANALYZE, BUFFERS)<your query again>;
```

`Seq Scan` should disappear and time should drop. Monitor:

Copy

```
SELECT idx_scan, seq_scan FROM pg_stat_user_tables WHERE relname = 'big_table';
```

Expect `idx_scan` to climb after deployment.

7 • Gotchas & Best Practices

1. **Don't over-index** — writes pay the price; review `idx_scan = 0` quarterly.
2. **CREATE INDEX CONCURRENTLY** in prod to avoid write locks.
3. Drop unused indexes:

Copy

```
SELECT relname FROM pg_stat_user_indexes WHERE idx_scan = 0  
AND pg_relation_size(indexrelid) > 10*8192;
```

1. Avoid duplicates (compare `pg_index.indkey`).
2. Use `fillfactor` for hot rows (`ALTER INDEX ... SET (fillfactor=90)`).

TL;DR

1. Use `pg_stat_user_tables` to locate heavy seq scans.
2. `EXPLAIN ANALYZE` reveals columns to index.
3. Audit FKs for missing indexes.
4. Create indexes **concurrently**; favour partial and covering patterns.
5. Re-run stats and queries to confirm wins.