

The Shocking Truth About PostgreSQL Locks (PostgreSQL documentation won't tell you)

- **Picture this scenario:** A developer pushes seemingly innocent code to production. Within hours, the application grinds to a halt. Users are locked out. The database becomes unresponsive. Sound familiar?
- This isn't a rare occurrence — it happens thousands of times daily across companies worldwide. The culprit? PostgreSQL locks that nobody truly understands.

The Hidden Reality of PostgreSQL Locking

- While most tutorials cover basic SELECT and INSERT operations, they completely ignore the intricate dance of locks happening beneath the surface. PostgreSQL employs over 15 different lock types, each with distinct behaviors that can make or break application performance.
- The most dangerous part? These locks often appear harmless in development environments but become catastrophic under production load.

Lock Type #1: The Silent Killer — Advisory Locks

- Advisory locks are PostgreSQL's best-kept secret and biggest trap. Unlike traditional locks, these are entirely application-controlled, meaning PostgreSQL won't automatically release them if something goes wrong.
- Real-world disaster scenario:

```
-- This innocent-looking function can crash your entire system
BEGIN;
SELECT pg_advisory_lock(12345);
-- Application crashes before COMMIT
-- Lock remains FOREVER
```

- **The shocking truth:** Advisory locks survive connection drops, server restarts, and even application crashes. Production systems have been found with advisory locks held for months, completely blocking critical operations.

The fix most developers miss:

```
-- Always use TRY variants with timeouts
SELECT pg_try_advisory_lock(12345);

-- Or use session-level locks that auto-release
```

```
SELECT pg_advisory_lock_shared(12345);
```

Lock Type #2: The Performance Assassin — Row Share Locks

- Here's what nobody tells developers about SELECT FOR UPDATE:

```
-- This query looks harmless
SELECT * FROM users WHERE id = 123 FOR UPDATE;
```

- **The hidden reality:** This single line can cascade into thousands of blocked queries. PostgreSQL creates a row share lock that blocks ANY UPDATE or DELETE on that specific row, even from unrelated transactions.
- Production nightmare example: A major e-commerce platform discovered that their "view user profile" feature was using SELECT FOR UPDATE unnecessarily. During Black Friday traffic, this single query pattern created a lock chain affecting 47,000 concurrent users.

The solution that saved millions:

```
-- Use FOR SHARE for read operations
SELECT * FROM users WHERE id = 123 FOR SHARE;

-- Or eliminate locks entirely
SELECT * FROM users WHERE id = 123;

-- Handle concurrency at application level
```

Lock Type #3: The Database Destroyer — Exclusive Locks

- ALTER TABLE operations in PostgreSQL are atomic bombs in disguise. They acquire ACCESS EXCLUSIVE locks that block everything — reads, writes, even simple SELECT statements.

The production horror story:

```
-- This "quick" column addition
ALTER TABLE users ADD COLUMN phone VARCHAR(20);
```

- On a table with 10 million rows, this operation can hold an exclusive lock for 8+ hours. During this time, the entire application becomes unusable.

The advanced solution:

```
-- Create column without lock
ALTER TABLE users ADD COLUMN phone VARCHAR(20) DEFAULT "";
```

```
-- Backfill data in batches
UPDATE users SET phone = calculate_phone()
WHERE id BETWEEN 1 AND 10000;

-- Add constraints after data population
ALTER TABLE users ALTER COLUMN phone SET NOT NULL;
```

The Lock Detection Query Nobody Shares

- Most developers have no visibility into lock conflicts. Here's the diagnostic query that can save careers:

```
SELECT
  blocked_locks.pid AS blocked_pid,
  blocked_activity.username AS blocked_user,
  blocking_locks.pid AS blocking_pid,
  blocking_activity.username AS blocking_user,
  blocked_activity.query AS blocked_statement,
  blocking_activity.query AS current_statement_in_blocking_process,
  blocked_activity.application_name AS blocked_application,
  blocking_activity.application_name AS blocking_application
FROM pg_catalog.pg_locks blocked_locks
JOIN pg_catalog.pg_stat_activity blocked_activity
  ON blocked_activity.pid = blocked_locks.pid
JOIN pg_catalog.pg_locks blocking_locks
  ON blocking_locks.locktype = blocked_locks.locktype
  AND blocking_locks.database IS NOT DISTINCT FROM blocked_locks.database
  AND blocking_locks.relation IS NOT DISTINCT FROM blocked_locks.relation
  AND blocking_locks.page IS NOT DISTINCT FROM blocked_locks.page
  AND blocking_locks.tuple IS NOT DISTINCT FROM blocked_locks.tuple
  AND blocking_locks.virtualxid IS NOT DISTINCT FROM blocked_locks.virtualxid
  AND blocking_locks.transactionid IS NOT DISTINCT FROM blocked_locks.transactionid
  AND blocking_locks.classid IS NOT DISTINCT FROM blocked_locks.classid
  AND blocking_locks.objid IS NOT DISTINCT FROM blocked_locks.objid
  AND blocking_locks.objsubid IS NOT DISTINCT FROM blocked_locks.objsubid
  AND blocking_locks.pid != blocked_locks.pid
JOIN pg_catalog.pg_stat_activity blocking_activity
  ON blocking_activity.pid = blocking_locks.pid
WHERE NOT blocked_locks.granted;
```

The Transaction Timeout Trap

- PostgreSQL's default configuration allows transactions to run indefinitely. This seemingly innocent setting has caused more production outages than any other database configuration.

The fix:

```
-- Set in postgresql.conf
statement_timeout = '30s'
lock_timeout = '10s'
```

```
idle_in_transaction_session_timeout = '60s'
```

Lock Escalation: The Elephant in the Room

- Unlike other databases, PostgreSQL doesn't automatically escalate row locks to table locks. This sounds good until realizing it can lead to deadlock cascades that bring down entire systems.

Deadlock prevention strategy:

```
-- Always acquire locks in consistent order
BEGIN;
LOCK TABLE table_a IN SHARE MODE;
LOCK TABLE table_b IN SHARE MODE;
-- Perform operations
COMMIT;
```

The Monitoring Solution That Actually Works

- Set up this alert to catch lock problems before they destroy production:

```
-- Monitor lock wait times
SELECT
    schemaname,
    tablename,
    attname,
    n_distinct,
    correlation
FROM pg_stats
WHERE schemaname NOT IN ('information_schema', 'pg_catalog')
AND n_distinct < 100 -- Potential lock contention
ORDER BY correlation DESC;
```

The Bottom Line

- PostgreSQL locks are not just a technical detail — they're the difference between a scalable application and a production nightmare. The examples above represent real scenarios that have cost companies millions in downtime and lost revenue.
- Understanding these lock patterns isn't optional for serious PostgreSQL developers. It's the foundation of building systems that can handle real-world traffic without collapsing under their own weight.
- The next time someone mentions PostgreSQL performance, remember: it's not about query optimization or indexing strategies. It's about mastering the invisible world of locks that controls everything happening in the database.

- Start monitoring locks today. The alternative is waiting for the 3 AM emergency call.