

# PostgreSQL Performance: Real-Life Scenarios & Solutions

## Table Of Contents

- [Intro](#)
- [1. Handling High Read & Write Throughput](#)
  - [Scenario](#)
  - [Solution](#)
- [2. Slow Query Performance](#)
  - [Scenario](#)
  - [Solution](#)
- [3. Managing Large Data Volumes](#)
  - [Scenario](#)
  - [Solution](#)
- [4. Addressing Lock Contention](#)
  - [Scenario](#)
  - [Solution](#)
- [5. High Latency in Replication](#)
  - [Scenario](#)
  - [Solution](#)
- [References](#)

## Intro

We all know theory is different than practice, this article provides real-life situations that you can face with PostgreSQL performance and how to solve them.

## 1. Handling High Read & Write Throughput

### Scenario

An e-commerce website experiences high read and write throughput, especially during peak shopping seasons like Christmas. The database should be able to handle thousands of transactions per second without degrading performance.

### Solution

- **Connection Pooling:** use *PgBouncer* to manage and reuse database connections efficiently.

```
[databases]
ecommerce = host=localhost dbname=ecommerce

[pgbouncer]
listen_addr = 127.0.0.1
listen_port = 6432
auth_type = md5
auth_file = /etc/pgbouncer/userlist.txt
pool_mode = transaction
max_client_conn = 500
default_pool_size = 50
```

- **Partitioning:** implement range partitioning on transaction date to optimize read and write operations.

```
CREATE TABLE transactions (
  id serial PRIMARY KEY,
  user_id int,
  product_id int,
  transaction_date date,
  amount numeric
) PARTITION BY RANGE (transaction_date);
```

```
CREATE TABLE transactions_2024 PARTITION OF transactions FOR VALUES FROM ('2024-08-28') TO ('2025-08-28');
```

- **Indexing:** create appropriate indexes to speed up frequent queries.

```
CREATE INDEX idx_user_product ON transactions (user_id, product_id);
```

## 2. Slow Query Performance

### Scenario

A SaaS application reports slow performance for several key analytical queries affecting the user experience.

### Solution

- **Query Optimization:** analyze and rewrite inefficient queries using *EXPLAIN ANALYZE* to understand execution plans.

```
EXPLAIN ANALYZE
SELECT user_id, COUNT(*)
FROM user_activity
WHERE activity_date > CURRENT_DATE - INTERVAL '1 month'
GROUP BY user_id;
```

- **Hypothetical Indexing:** use the HypoPG extension to test the impact of new indexes without actually creating them.

```
CREATE EXTENSION hypopg;
SELECT * FROM hypopg create index('CREATE INDEX ON user_activity (activity_date)');
EXPLAIN SELECT user_id, COUNT(*) FROM user_activity WHERE activity_date > CURRENT_DATE -
INTERVAL '1 month' GROUP BY user_id;
```

- **Materialized Views:** create materialized views to pre-compute expensive joins and aggregations.

```
CREATE MATERIALIZED VIEW monthly_user_activity AS
SELECT user_id, COUNT(*) AS activity_count
FROM user_activity
WHERE activity_date > CURRENT_DATE - INTERVAL '1 month'
GROUP BY user_id;
```

### 3. Managing Large Data Volumes

#### Scenario

A financial services company needs to manage and analyze terabytes of historical transaction data to close the fiscal year while maintaining performance.

#### Solution

- **Data Archiving:** regularly archive old data to separate storage systems to keep the primary database lean.

```
CREATE TABLE transactions_archive (LIKE transactions INCLUDING ALL);

INSERT INTO transactions_archive
SELECT * FROM transactions
WHERE transaction_date < CURRENT_DATE - INTERVAL '1 year';
DELETE FROM transactions
WHERE transaction_date < CURRENT_DATE - INTERVAL '1 year';
```

- **Advanced Partitioning:** use list partitioning to segment data by year.

```
CREATE TABLE transactions (
  id serial PRIMARY KEY,
  user_id int,
  product_id int,
  transaction_date date,
  amount numeric
) PARTITION BY LIST (EXTRACT(YEAR FROM transaction_date));

CREATE TABLE transactions_2024 PARTITION OF transactions FOR VALUES IN (2024);
```

- **Parallel Query Execution:** enable and configure parallel query execution for large-scale data processing.

```
SET max_parallel_workers_per_gather = 4;
SET max_parallel_workers = 8;
```

## 4. Addressing Lock Contention

### Scenario

A content management system or simply CMS, experiences frequent lock contention causing delays and timeouts during peak usage.

### Solution

- **Monitoring Locks:** use *pg\_stat\_activity* and *pg\_locks* to identify and analyze lock contention.

```
SELECT pid, locktype, relation::regclass, mode, granted
FROM pg_locks
JOIN pg_stat_activity ON pg_locks.pid = pg_stat_activity.pid
WHERE NOT granted;
```

- **Optimizing Transactions:** break long transactions into smaller units to reduce lock durations.

```
BEGIN;
UPDATE articles SET content = '...' WHERE id IN (1, 2, 3);
COMMIT;

BEGIN;
UPDATE articles SET content = '...' WHERE id IN (4, 5, 6);
COMMIT;
```

- **Index Maintenance:** regularly reindex to ensure indexes are efficient and reduce the likelihood of lock contention.

```
REINDEX TABLE articles;
```

## 5. High Latency in Replication

### Scenario

A universal application with multiple data centers requires low-latency replication between primary and standby servers.

### Solution

- **Synchronous Replication:** configure synchronous replication for critical data to ensure data consistency across regions.

```
ALTER SYSTEM SET synchronous_standby_names = 'node1, node2';  
SELECT pg_reload_conf();
```

- **Replication Slots:** use replication slots to ensure data is not lost due to replication lag.

```
SELECT * FROM pg_create_physical_replication_slot('replication_slot1');
```

- **Network Optimization:** optimize TCP settings to reduce replication latency.

```
tcp_keepalives_idle = 60  
tcp_keepalives_interval = 10  
tcp_keepalives_count = 5
```