

[Open in app ↗](#)**Medium**

Search



Get unlimited access to the best of Medium for less than \$1/week. [Become a member](#)



How to Enable and Secure Remote PostgreSQL 17 Connections on Red Hat Linux: Complete Step-by-Step Guide

7 min read · Jun 2, 2025



Jeyaram Ayyalusamy

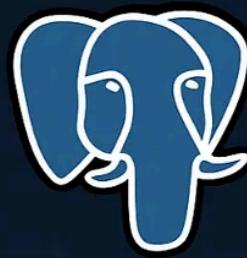
Following

Listen

Share

More

How to Enable and Secure Remote PostgreSQL 17 Connections on Red Hat Linux, CentOS, Rocky Linux, AlmaLinux, Oracle Linux: Complete Step-by-Step Guide



Red Hat



Rocky LI



AlmaLinux

PostgreSQL is one of the most popular open-source relational database management systems (RDBMS) used across industries. It powers mission-critical systems, SaaS platforms, business intelligence tools, and data lakes globally. But while PostgreSQL is extremely flexible, its **default security posture disables remote access** to ensure databases are protected out-of-the-box.

In real-world production, however, you often need to allow external connections:

- Application servers connecting from other machines
- Developers accessing databases remotely
- BI tools querying PostgreSQL over the network
- Multi-tier or microservices architecture where database servers are decoupled

In this comprehensive blog, we'll explore **how to enable remote connections to PostgreSQL 17 running on Red Hat Enterprise Linux (RHEL)** and its derivatives (CentOS, Rocky Linux, AlmaLinux, Oracle Linux), while ensuring best security and performance practices.

Why Remote Access Is Not Enabled by Default?

PostgreSQL disables remote access by default for a very good reason: **security**. Open database ports exposed to the internet are one of the most common ways databases get compromised.

That's why **enabling remote access requires deliberate configuration, proper network security, and access control** to protect your data.

- ✓ **Good practice:** Open PostgreSQL to trusted internal networks or VPN-connected machines, but avoid exposing directly to public internet unless absolutely necessary (and secured with SSL + proper firewalling).

PostgreSQL Configuration on Red Hat Linux — File Locations

Before diving into the steps, let's review where PostgreSQL stores its configuration files on Red Hat Linux:

File Default Path

<code>postgresql.conf</code>	<code>/var/lib/pgsql/17/data/postgresql.conf</code>
<code>pg_hba.conf</code>	<code>/var/lib/pgsql/17/data/pg_hba.conf</code>
<code>PostgreSQL logs (default)</code>	
<code>/var/lib/pgsql/17/data/log/</code>	

 If unsure where PostgreSQL installed, verify using:

```
psql -U postgres -c 'SHOW config_file;'
```

```
[postgres@node1~]$ psql -U postgres -c 'SHOW config_file;'  
config_file  
-----  
/var/lib/pgsql/17/data/postgresql.conf  
(1 row)  
  
[postgres@node1~]$
```

 **Tip:** Always back up these files before making configuration changes.

Step 1 — Modify PostgreSQL Configuration for Remote Listening

By default, PostgreSQL only listens on `localhost`. We need to change this behavior.

1.1 Edit `postgresql.conf`

```
sudo vi /var/lib/pgsql/17/data/postgresql.conf
```

```
[root@node1 ~]# cat /var/lib/pgsql/17/data/postgresql.conf | grep listen_address  
listen_addresses = 'localhost'          # what IP address(es) to listen on;  
[root@node1 ~]#
```

Locate:

```
listen_addresses = 'localhost'

[root@node1 ~]# cat /var/lib/pgsql/17/data/postgresql.conf | grep listen_addresses
listen_addresses = '*'          # what IP address(es) to listen on;
[root@node1 ~]#
```

Change to:

```
listen_addresses = '*'
```

- * means PostgreSQL will listen on all available network interfaces.
- You may also restrict to specific IPs:

```
listen_addresses = '192.168.100.50'
```

 **Security Tip:** If your server has multiple interfaces (private/public), it's often safer to bind PostgreSQL only to private IPs.

Of course! Here's a **shortened, clean, and Medium-ready version** of your blog post:

1.2 Understanding PostgreSQL 17 port Parameter in postgresql.conf

PostgreSQL 17's `postgresql.conf` file allows administrators to control how the database server operates. One of the key parameters in this file is `port`, which defines where the server listens for client connections.

What is the port Parameter?

The `port` parameter specifies the TCP port PostgreSQL uses to accept connections. By default, PostgreSQL listens on port 5432:

```
port = 5432 # (change requires restart)
```

If you need PostgreSQL to listen on a different port (for example, to avoid conflicts or run multiple instances), you can change this value.

Why Change the Port?

- Run multiple PostgreSQL instances on the same server.
- Avoid conflicts with other services.
- Customize configurations for different environments.
- As part of network security hardening.

How to Change the Port

- 1 Edit `postgresql.conf` (location may vary depending on OS; e.g. `/var/lib/pgsql/17/data/postgresql.conf` on Red Hat):

```
sudo vi /var/lib/pgsql/17/data/postgresql.conf
```

- 2 Update the port number:

```
port = 5433
```

3 Restart PostgreSQL to apply:

```
sudo systemctl restart postgresql-17
```

Important Notes

- Changing the port requires restarting the PostgreSQL service.
- Make sure the new port is open in your firewall and cloud security groups.
- Clients must update connection strings:

```
psql -h <server_ip> -p 5433 -U <username> -d <database>
```

1.3 Edit pg_hba.conf — Host-Based Access Control

This file defines *who* can connect, from *where*, and *how* they authenticate.

Open the file:

```
sudo vi /var/lib/pgsql/17/data/pg_hba.conf
```

```
cat /var/lib/pgsql/17/data/pg_hba.conf
```

#	TYPE	DATABASE	USER	ADDRESS	METHOD
---	------	----------	------	---------	--------

```

# "local" is for Unix domain socket connections only
local  all      all                                     peer
# IPv4 local connections:
host   all      all          127.0.0.1/32           scram-sha-256
# IPv6 local connections:
host   all      all          ::1/128                scram-sha-256
# Allow replication connections from localhost, by a user with the
# replication privilege.
local  replication all                                     peer
host   replication all          127.0.0.1/32           scram-sha-256
host   replication all          ::1/128                scram-sha-256
[root@node1 ~]#

```

Add entries to allow remote clients.

Example 1 — Allow one specific client IP:

```
host   all      all      0.0.0.0/0      md5
```

```

cat /var/lib/pgsql/17/data/pg_hba.conf
# TYPE  DATABASE        USER        ADDRESS            METHOD

# "local" is for Unix domain socket connections only
local  all      all                                     peer
# IPv4 local connections:
host   all      all          0.0.0.0/0           scram-sha-256
# IPv6 local connections:
host   all      all          ::1/128               scram-sha-256
# Allow replication connections from localhost, by a user with the
# replication privilege.
local  replication all                                     peer
host   replication all          127.0.0.1/32         scram-sha-256
host   replication all          ::1/128               scram-sha-256
[root@node1 ~]#
[root@node1 ~]#

```

⚠ Important Note

- While `md5` provides basic password security, exposing your database to `0.0.0.0/0` is **not safe for production** because it allows global internet access.

Best practice:

- Use more restrictive IP ranges (e.g. `192.168.1.0/24` for your internal network).
- Use `scram-sha-256` authentication for stronger security.
- Always secure PostgreSQL with a firewall or VPN.

Example 2 — Allow an entire subnet:

```
host      all      all      192.168.32.131/24      md5
```

Explanation:

- `host` → allows TCP/IP connections.
- `all` → allows all databases.
- `all` → allows all users.
- `IP/subnet` → defines allowed client IP addresses.
- `md5` → requires password authentication.

 **Warning:** Avoid using `trust` in production as it allows passwordless access.

Step 2 — Open Firewall Ports on Red Hat Linux

PostgreSQL listens on TCP port `5432` by default. Red Hat Linux uses `firewalld` as its firewall manager.

2.1 Verify that firewalld is active:

```
sudo systemctl status firewalld
```

2.2 Open port 5432 for PostgreSQL:

```
sudo firewall-cmd --permanent --add-port=5432/tcp
```

2.3 Apply firewall changes:

```
sudo firewall-cmd --reload
```

2.4 Verify:

```
sudo firewall-cmd --list-ports
```

Best Practice:

- Only open port 5432 to trusted networks.
- For public clouds, also configure AWS Security Groups, Azure NSGs, or GCP firewall rules appropriately.

Step 3 — Restart PostgreSQL Service

PostgreSQL must be restarted to apply configuration changes.

```
sudo systemctl restart postgresql-17
```

Verify the service status:

```
sudo systemctl status postgresql-17
```

 **Tip:** Always check logs in case PostgreSQL fails to restart:

```
sudo journalctl -xeu postgresql-17
```

Step 4 — Test Remote Connection

Once configured, test the connection from a remote machine using:

```
psql -h <PostgreSQL_server_IP> -p 5432 -U <username> -d <database>
```

Example:

```
psql -h 192.168.32.131 -p 5432 -U myuser -d mydatabase
```

- You'll be prompted for your PostgreSQL user's password.
- On success, you'll get the psql prompt:

```
[postgres@node1 ~]$ psql -h 192.168.32.131 -p 5432 -U postgres -d postgres
Password for user postgres:
psql (17.5)
Type "help" for help.
```

```
postgres=#  
postgres=#
```

Step 5 — Secure Your Remote Connections (Highly Recommended)

Simply opening PostgreSQL to remote clients can expose you to serious security risks unless properly secured.

Best Practice Description Use SSL/TLS Encrypt all client-server traffic using certificates. Restrict IP Access Allow only known trusted IPs/subnets in both firewalls and `pg_hba.conf`. Use Strong Passwords Enforce complex, unique passwords for all database roles. Enable Logging Enable detailed connection logging for auditability. Use VPN Access PostgreSQL via a private VPN whenever possible. Rotate Credentials Periodically rotate user credentials and revoke unused roles.

-  PostgreSQL fully supports SSL encryption. You can configure `ssl = on` in `postgresql.conf` and manage certificates using self-signed or enterprise PKI solutions.

Common Problems and Fixes

Problem Solution `psql: could not connect to server` Verify `listen_addresses`, firewalls, and PostgreSQL status. Authentication failed Ensure correct username/password and matching `pg_hba.conf` entry. Timeout Verify network connectivity and firewall rules. Permission denied on firewall Check cloud security group configurations. Port still blocked Use `'ss -plnt`

Full Quick Checklist

- ✓ Update `postgresql.conf` — `listen_addresses`
- ✓ Update `pg_hba.conf` — allow specific IPs
- ✓ Open firewall port 5432 via `firewalld`
- ✓ Restart PostgreSQL service
- ✓ Test remote connection
- ✓ Secure connections with firewall + SSL

Bonus: Automate Everything with Ansible

For repeatable deployments across multiple servers, you can automate remote PostgreSQL configuration using tools like:

- Ansible
- Terraform + Ansible combo (for cloud-native automation)
- Bash provisioning scripts

 This allows you to deploy hardened, fully configured PostgreSQL servers at scale with minimal human error.

Real-World Use Case Scenario

Scenario:

A mid-sized SaaS company is running PostgreSQL 17 on Red Hat Linux virtual machines inside AWS private subnets. Application servers reside in separate private subnets. They:

- Open port 5432 internally using AWS Security Groups.
- Restrict `pg_hba.conf` to only allow connections from application server subnets.
- Use SSL for encrypted traffic.
- Automate PostgreSQL configuration with Ansible.
- Monitor logs for unauthorized connection attempts.

Result:

 Secure, scalable, remote PostgreSQL connections with full control over compliance, performance, and auditability.

Conclusion: PostgreSQL Remote Access on Red Hat — Powerful, Flexible & Secure

With careful configuration, PostgreSQL 17 can be securely accessed remotely — enabling modern cloud-native, multi-tier, and microservices architectures to

flourish.

- Red Hat Linux gives you full control over PostgreSQL configuration.
- `postgresql.conf` + `pg_hba.conf` provide fine-grained access control.
- `firewalld` ensures proper network segmentation.
- SSL encryption secures data-in-transit.

By following this guide, you not only enable remote access — you do it **safely** and **professionally**.

👉 Follow me on Medium to continue the full PostgreSQL deep dive!

🔗 Let's Connect!

If you enjoyed this post or would like to connect professionally, feel free to reach out to me on LinkedIn:

👉 [Jeyaram Ayyalusamy](#)

I regularly share content on **PostgreSQL, database administration, cloud technologies, and data engineering**. Always happy to connect, collaborate, and discuss ideas!

Database

Postgresql

Sql

Linux



Following ▾

Written by Jeyaram Ayyalusamy

76 followers · 2 following

Oracle DBA | AWS Certified | 18+ yrs in DB Admin, RAC, GoldenGate, RDS, MySQL, PostgreSQL | Author of MySQL & RDS books | Cloud & Performance Expert

No responses yet



Gvadakte

What are your thoughts?

More from Jeyaram Ayyalusamy

The diagram illustrates the progression of database operations and their impact on performance:

- NORMAL OPERATION:** Shows a sequence of operations:
 - WRITES (represented by green cubes) lead to a VACUUM operation.
 - VACUUM leads to MODERATE BLOAT (represented by yellow cubes).
 - MODERATE BLOAT leads to REINDEX.
 - DISK (represented by red disks) leads to HIGH BLOAT (represented by red cubes).
 - HIGH BLOAT leads to a user icon followed by a crossed-out wrench icon, indicating a critical state.

PostgreSQL logo is displayed on the right.

J Jeyaram Ayyalusamy

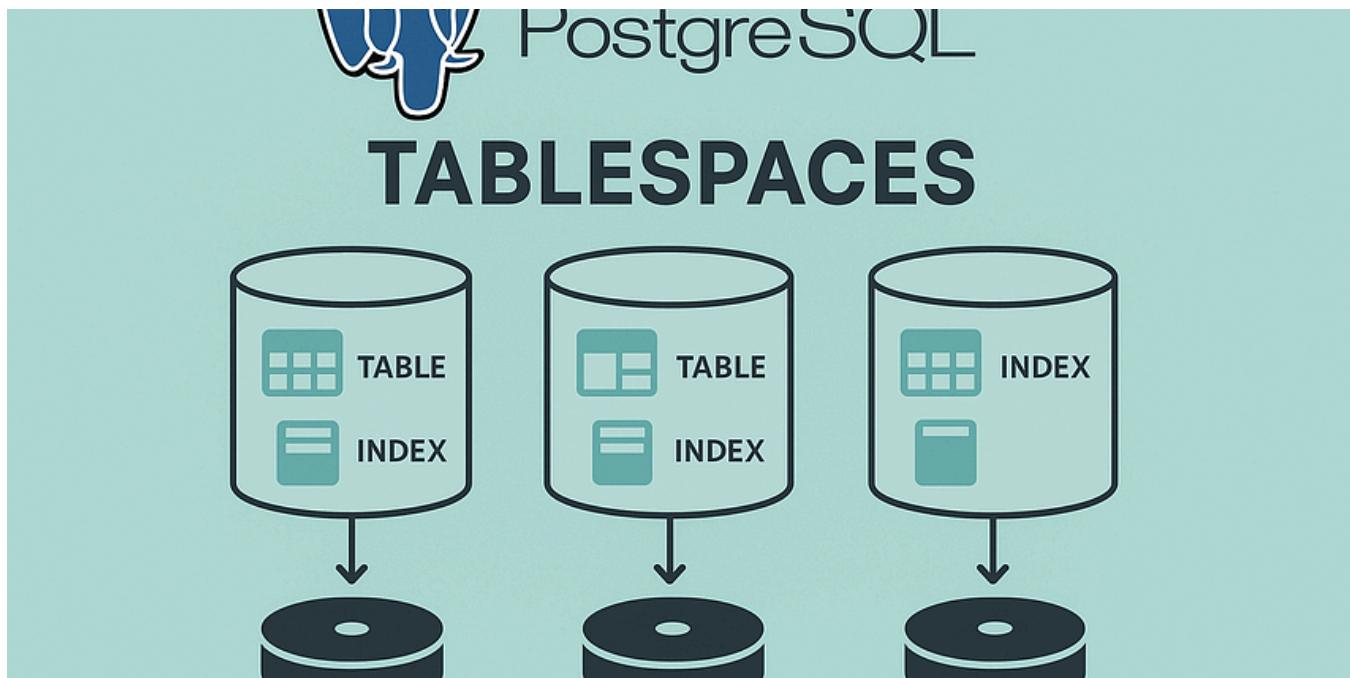
Understanding and Managing Bloating in PostgreSQL: A Complete Guide for DBAs

PostgreSQL is a powerful, reliable, and feature-rich open-source relational database system. It's praised for its extensibility, ACID...

Jun 25 52



...



J Jeyaram Ayyalusamy

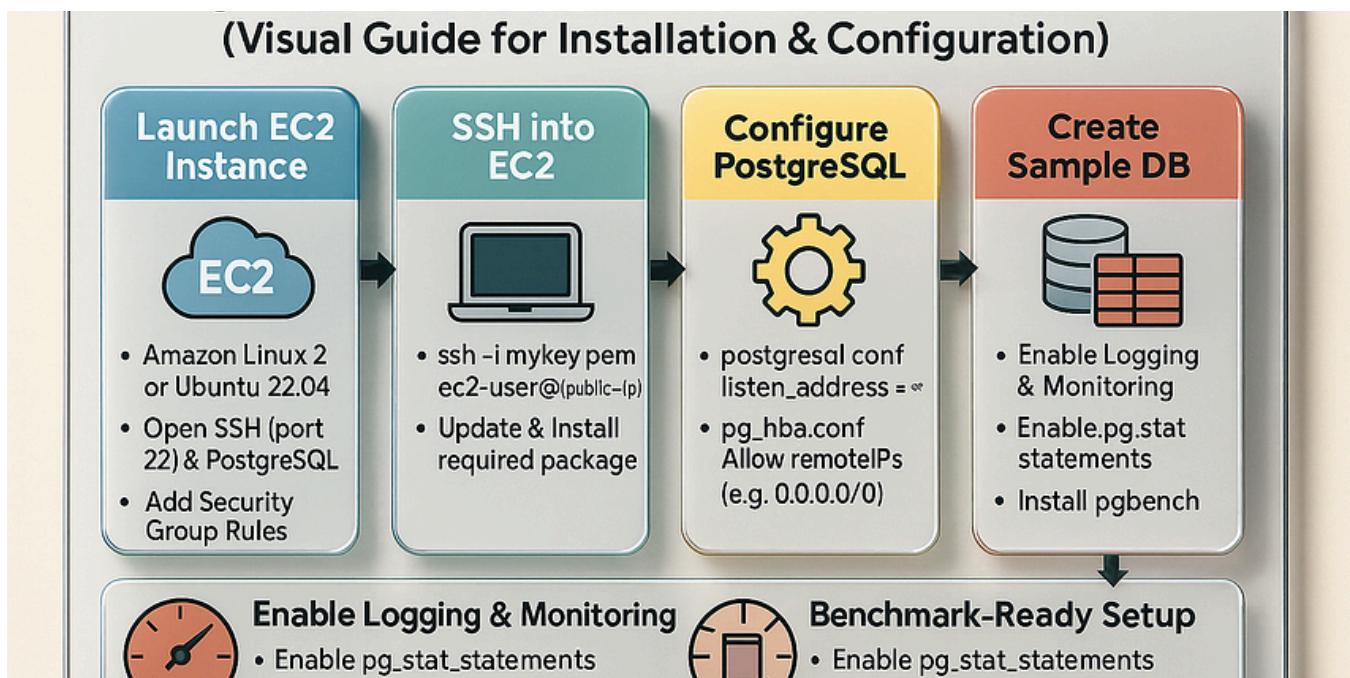
PostgreSQL Tablespaces Explained: Complete Guide for PostgreSQL 17 DBAs

PostgreSQL offers a powerful feature called tablespaces, allowing database administrators to take control of how and where data is...

Jun 12 8



...



J Jeyaram Ayyalusamy

PostgreSQL 17 on AWS EC2—Full Installation & Configuration Walkthrough

Setting up PostgreSQL 17 on AWS EC2 might seem complex at first—but once you break down each component, it becomes an efficient and...

1d ago 👏 50



...



J Jeyaram Ayyalusamy

How to Install PostgreSQL 17 on Red Hat, Rocky, AlmaLinux, and Oracle Linux (Step-by-Step Guide)

PostgreSQL 17 is the latest release of one of the world's most advanced open-source relational databases. If you're using a Red Hat-based...

Jun 3



...

See all from Jeyaram Ayyalusamy

Recommended from Medium

 Azlan Jamal

Stop Using SERIAL in PostgreSQL: Here's What You Should Do Instead

In PostgreSQL, there are several ways to create a PRIMARY KEY for an id column, and they usually involve different ways of generating the...

 Jul 12  33

...

 Rizqi Mulki

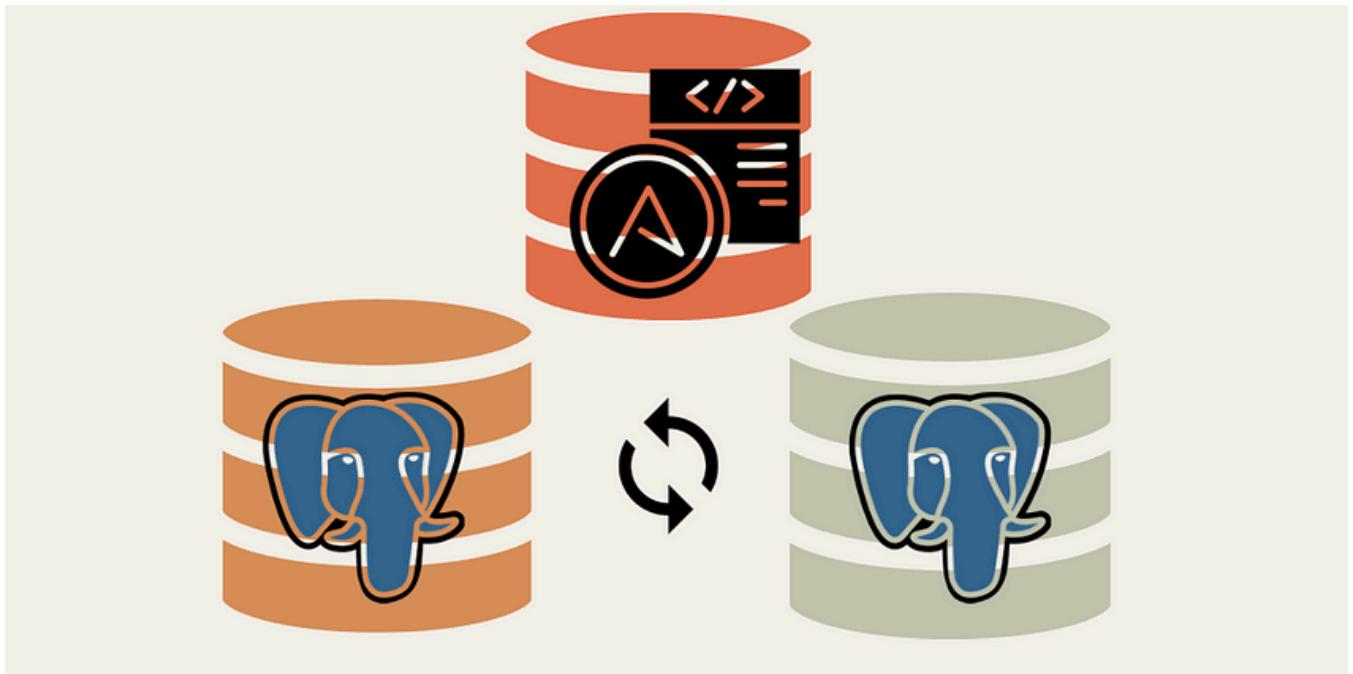
Full-Text Search in PostgreSQL: Better Than You Think

Why developers are abandoning Elasticsearch for PostgreSQL's built-in search—and how it handles 10M records without breaking a sweat

5d ago 2

Upvote icon

...



Oz

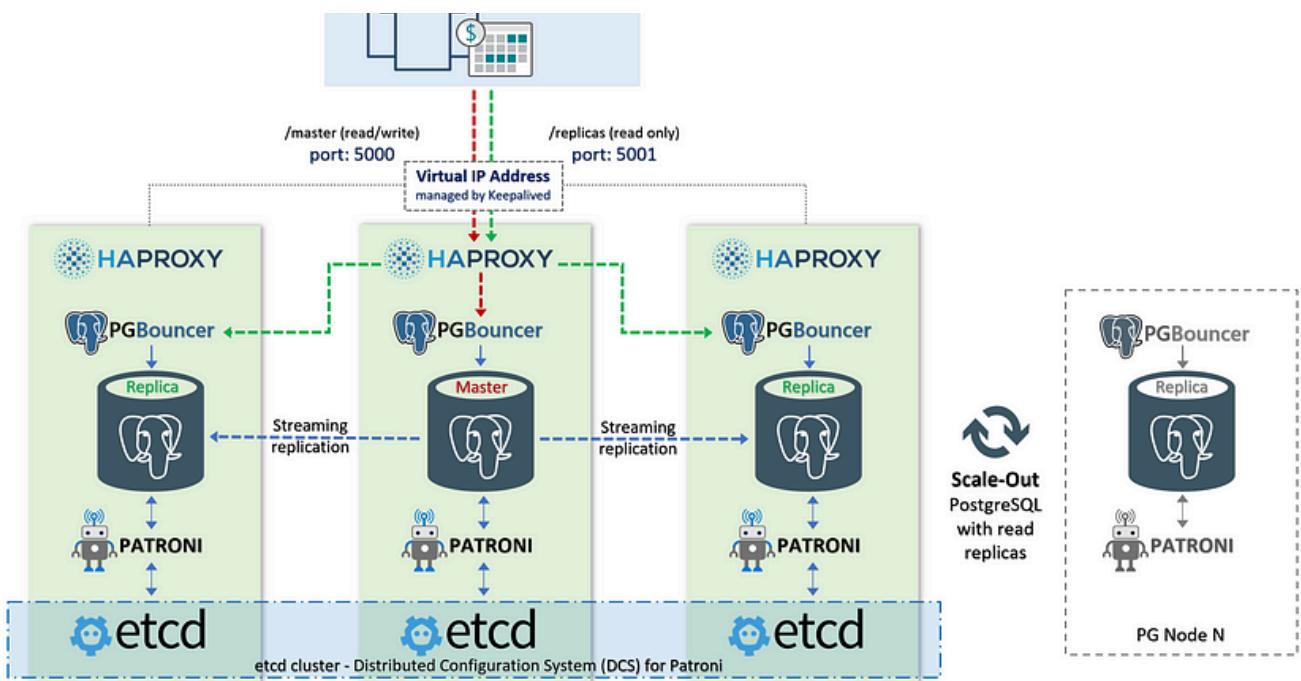
Automating PostgreSQL Streaming Replication Setup with Ansible

Setting up PostgreSQL streaming replication manually can be a tedious and error-prone task —involving installing PostgreSQL, configuring...

Jul 8 58

Upvote icon

...

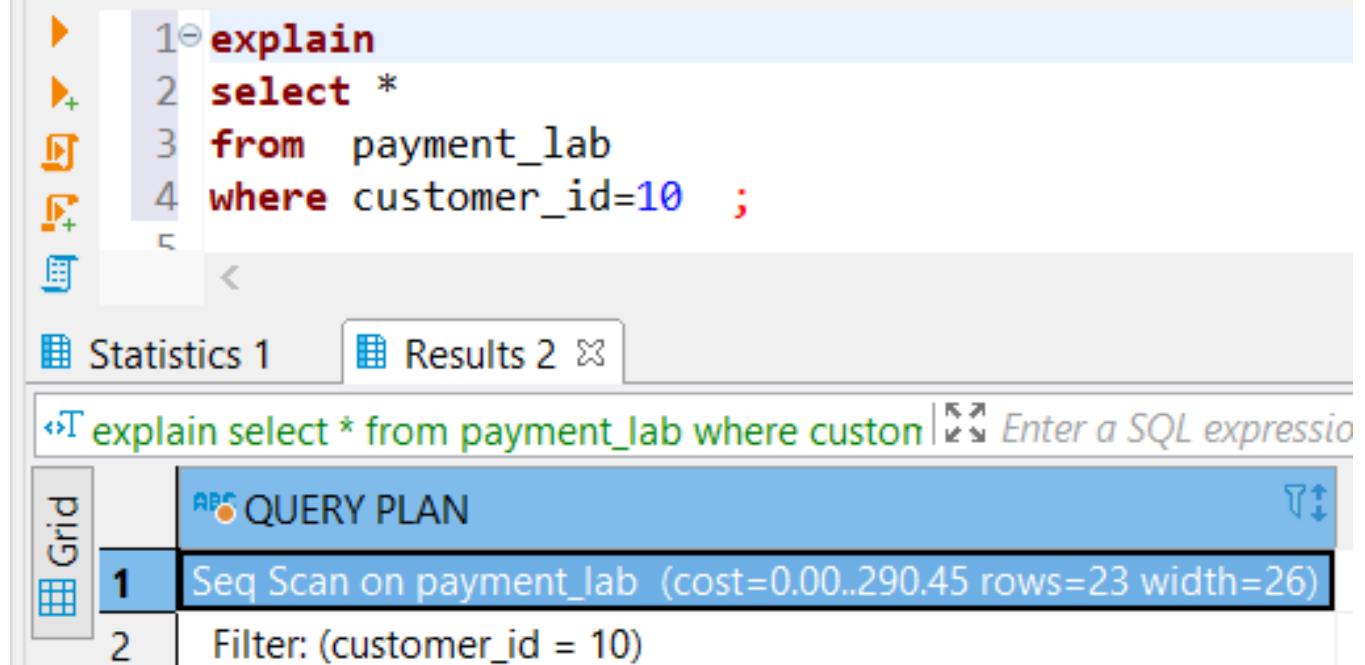


Kanat Akylson

How to Deploy a High-Availability PostgreSQL Cluster in 5 Minutes Using Ansible and Patroni

This tutorial shows how to spin up a production-grade HA PostgreSQL cluster in just 5 minutes using m2y ready-made GitHub repository with...

Jun 9 65 1



The screenshot shows a PostgreSQL query editor interface. At the top, there is a code editor with the following SQL query:

```
1 explain
2 select *
3 from payment_lab
4 where customer_id=10 ;
```

Below the code editor are two tabs: "Statistics 1" and "Results 2". The "Results 2" tab is active, displaying the query plan:

Grid	QUERY PLAN
1	Seq Scan on payment_lab (cost=0.00..290.45 rows=23 width=26)
2	Filter: (customer_id = 10)

The first row of the plan table is highlighted.

Muhammet Kurtoglu

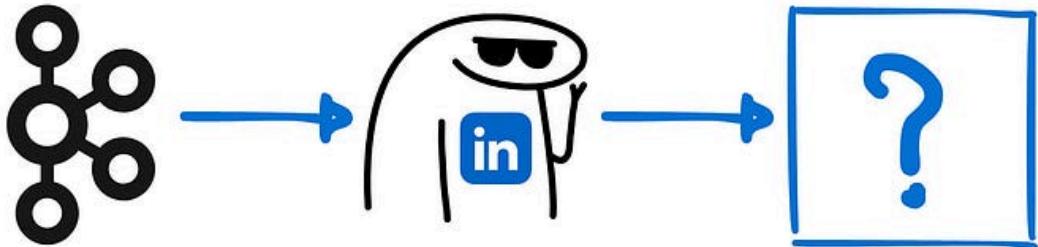
Postgresql Query Performance Analysis

SQL tuning is critically important for ensuring database performance, reliability, and scalability. In most cases, performance issues in a...

6d ago 10



Linkedin is moving from Kafka to this



The company that created Kafka is replacing it with a new solution

In Data Engineer Things by Vu Trinh

The company that created Kafka is replacing it with a new solution

How did LinkedIn build Northguard, the new scalable log storage

Jul 17 330 6

...

See more recommendations