

[Open in app](#)

Medium



Get unlimited access to the best of Medium for less than \$1/week. [Become a member](#)



20 - PostgreSQL 17 Performance Tuning: Full-Text Search Index (TSVECTOR)

4 min read · Sep 5, 2025



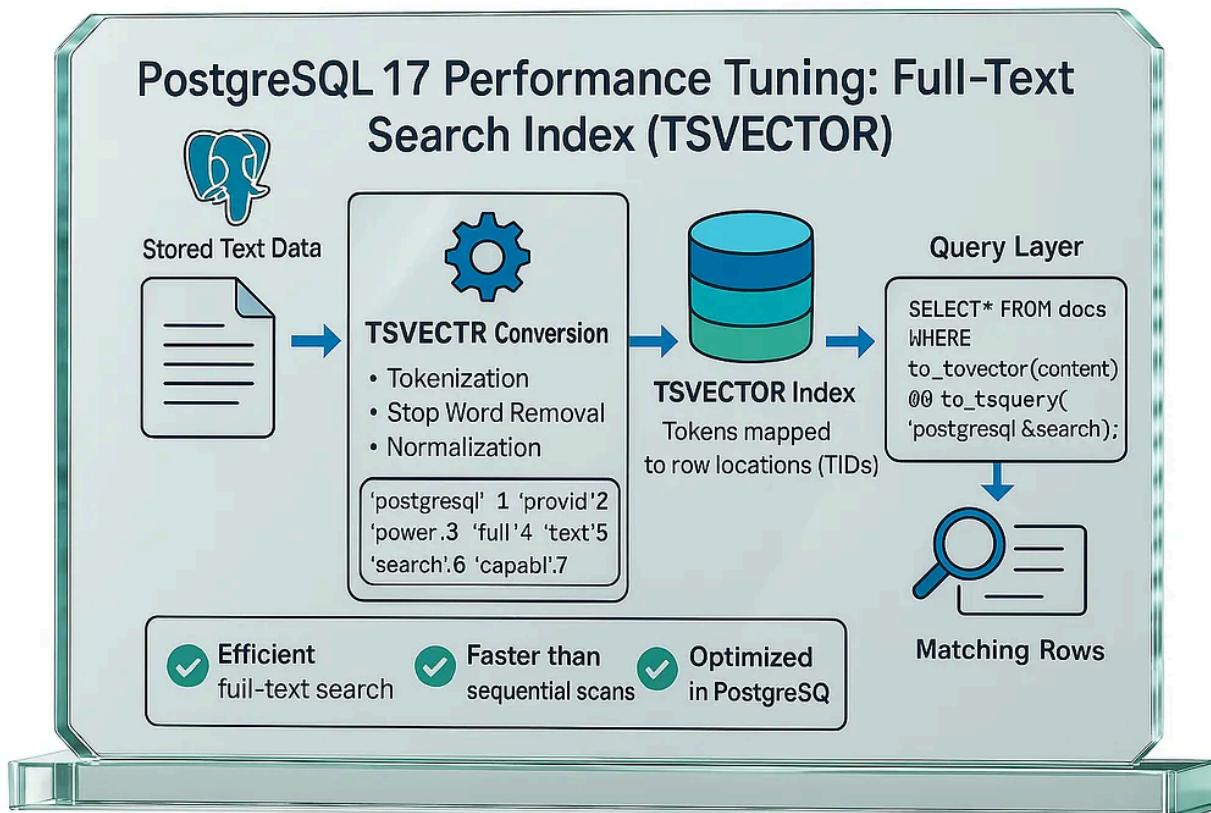
Jeyaram Ayyalusamy

Following

Listen

Share

More



A **Full-Text Search Index (TSVECTOR)** in PostgreSQL is a special way of indexing text data so that it can be searched quickly and efficiently. Instead of storing raw text, PostgreSQL converts the text into a `tsvector`, which is a preprocessed form

that breaks the text into tokens (words), removes stop words (like “a”, “the”), and normalizes words to their root form (e.g., “running”, “ran” → “run”).

This `tsvector` is then indexed, usually with a **GIN (Generalized Inverted Index)** or **GiST (Generalized Search Tree)** index, which allows PostgreSQL to match search queries (`tsquery`) against documents at high speed.

By using `tsvector`, PostgreSQL supports full-text search features like:

- Finding documents that contain specific words or phrases.
- Ranking results based on relevance.
- Highlighting matched words in search results.

In short, a Full-Text Search Index with `tsvector` transforms text into a searchable format and, when combined with an index, provides fast and accurate full-text search inside PostgreSQL.

Step 1 — Create the table

```
CREATE TABLE articles (
    id          BIGSERIAL PRIMARY KEY,
    title       TEXT,
    body        TEXT,
    published_at TIMESTAMPTZ DEFAULT now()
);
```

```
postgres=# CREATE TABLE articles (
    id          BIGSERIAL PRIMARY KEY,
    title       TEXT,
    body        TEXT,
```

```
published_at TIMESTAMPTZ DEFAULT now()
);
CREATE TABLE
postgres=#
```

Step 2 — Load 1,000,000 rows (synthetic corpus)

We'll sprinkle common tech keywords so the search actually matches.

```
INSERT INTO articles (title, body)
SELECT
'Post ' || g || ' about ' ||
(ARRAY['database','performance','indexing','postgresql','tuning',
      'sql','json','cache','query','optimizer'])
[1 + (random()*9)::int] AS title,
'This article discusses ' ||
(ARRAY['database','performance tuning','indexing strategies','PostgreSQL 17',
      'GIN indexes','full-text search','planning and costing','caching',
      [1 + (random()*9)::int] ||
' with practical tips on ' ||
(ARRAY['tsvector','tsquery','plainto_tsquery','to_tsquery','ranking','highl
[1 + (random()*5)::int] || '.' AS body
FROM generate_series(1, 1000000) g;
```

```
postgres=# INSERT INTO articles (title, body)
SELECT
'Post ' || g || ' about ' ||
(ARRAY['database','performance','indexing','postgresql','tuning',
      'sql','json','cache','query','optimizer'])
[1 + (random()*9)::int] AS title,
'This article discusses ' ||
(ARRAY['database','performance tuning','indexing strategies','PostgreSQL 17',
      'GIN indexes','full-text search','planning and costing','caching',
      [1 + (random()*9)::int] ||
' with practical tips on ' ||
(ARRAY['tsvector','tsquery','plainto_tsquery','to_tsquery','ranking','highl
```

```
[1 + (random()*5)::int] || '.' AS body  
FROM generate_series(1, 1000000) g;
```

Step 3 — Analyze statistics

```
ANALYZE articles;
```

```
postgres=# ANALYZE articles;  
ANALYZE  
postgres=#
```

Step 4 — Baseline query without an index

We'll search for “database performance”. Computing the vector on the fly forces a full scan.

```
EXPLAIN ANALYZE  
SELECT count(*)  
FROM articles  
WHERE to_tsvector('english', coalesce(title,'') || ' ' || coalesce(body,'')) @@  
plainto_tsquery('english','database performance');
```

```
postgres=# EXPLAIN ANALYZE
SELECT count(*)
FROM articles
WHERE to_tsvector('english', coalesce(title,'') || ' ' || coalesce(body,'')) @@
plainto_tsquery('english','database performance');
```

```
Finalize Aggregate  (cost=130449.57..130449.58 rows=1 width=8) (actual time=13184.60..13184.60 rows=1 width=8)
  ->  Gather  (cost=130449.36..130449.57 rows=2 width=8) (actual time=13184.60..13184.60 rows=2 width=8)
      Workers Planned: 2
      Workers Launched: 2
        ->  Partial Aggregate  (cost=129449.36..129449.37 rows=1 width=8) (actual time=13184.60..13184.60 rows=1 width=8)
            ->  Parallel Seq Scan on articles  (cost=0.00..129449.33 rows=1000000 width=8)
                Filter: (to_tsvector('english')::regconfig, ((COALESCE(title, '')) @@ plainto_tsquery('english','database performance')) = true)
                Rows Removed by Filter: 329211
Planning Time: 0.090 ms
Execution Time: 13191.826 ms
(10 rows)
```

postgres=#

```
Time: 13191.826 ms (00:13.826)
postgres=#
```

With no index, PostgreSQL scans all 1M rows (~13.826 s).

Step 5 — Add a generated tsvector column (stored)

This keeps document vectors always up-to-date without triggers.

```
ALTER TABLE articles
ADD COLUMN tsv tsvector
GENERATED ALWAYS AS (
```

```
to_tsvector('english', coalesce(title,'')) || ' ' || coalesce(body,''))
) STORED;
```

```
postgres=# ALTER TABLE articles
ADD COLUMN tsv tsvector
GENERATED ALWAYS AS (
    to_tsvector('english', coalesce(title,'')) || ' ' || coalesce(body,''))
) STORED;
ALTER TABLE
postgres=#

```

Step 6 — Create a GIN index on the tsvector

```
CREATE INDEX idx_articles_tsv_gin ON articles USING gin (tsv);
```

```
postgres=# CREATE INDEX idx_articles_tsv_gin ON articles USING gin (tsv);
CREATE INDEX
postgres=#

```

(If building online in production, use `CREATE INDEX CONCURRENTLY`.)

Analyze statistics

```
ANALYZE articles;
```

```
postgres=# ANALYZE articles;
ANALYZE
postgres=#
```

Step 7 — Search again (now using the index)

```
postgres=# EXPLAIN ANALYZE
SELECT count(*)
FROM articles
WHERE tsv @@ plainto_tsquery('english','database performance');

QUERY PLAN
-----
Aggregate  (cost=34497.68..34497.69 rows=1 width=8) (actual time=39.425..39.42
-> Bitmap Heap Scan on articles  (cost=190.83..34440.90 rows=22714 width=0)
    Recheck Cond: (tsv @@ '''databas'' & 'perform'::tsquery)
    Heap Blocks: exact=10602
-> Bitmap Index Scan on idx_articles_tsv_gin  (cost=0.00..185.15 rows
    Index Cond: (tsv @@ '''databas'' & 'perform'::tsquery)
Planning Time: 0.101 ms
Execution Time: 39.453 ms
(8 rows)
```

postgres=#

Time: 39.453 ms
postgres=#

GIN flips the plan to Bitmap Index Scan, dropping runtime from *Time reduction:*

$(13,191.826 - 39.453) = 13,152.373 \text{ ms saved} (\approx 99.701\% \text{ reduction})$

🔔 Stay Updated with Daily PostgreSQL & Cloud Tips!

If you've been finding my blog posts helpful and want to stay ahead with daily insights on PostgreSQL, Cloud Infrastructure, Performance Tuning, and DBA Best Practices — I invite you to subscribe to my Medium account.

🔔 [Subscribe here](https://medium.com/@jramcloud1/subscribe) 👉 <https://medium.com/@jramcloud1/subscribe>

Your support means a lot — and you'll never miss a practical guide again!

🔗 Let's Connect!

If you enjoyed this post or would like to connect professionally, feel free to reach out to me on LinkedIn:

👉 [Jeyaram Ayyalusamy](#)

I regularly share content on **PostgreSQL, database administration, cloud technologies, and data engineering**. Always happy to connect, collaborate, and discuss ideas!

Postgresql

AWS

Open Source

Mongodb

MySQL

J

Following ▾

Written by Jeyaram Ayyalusamy ✎

158 followers · 2 following

Oracle DBA | AWS Certified | 18+ yrs in DB Admin, RAC, GoldenGate, RDS, MySQL, PostgreSQL | Author of MySQL & RDS books | Cloud & Performance Expert

No responses yet



Gvadakte

What are your thoughts?



More from Jeyaram Ayyalusamy

No instances
You do not have any instances in this region
[Launch instances](#)

© 2025, Amazon Web Services, Inc. or its affiliates.

J Jeyaram Ayyalusamy

Upgrading PostgreSQL from Version 16 to Version 17 Using pg_upgrade on a Linux Server AWS EC2...

A Complete Step-by-Step Guide to Installing PostgreSQL 16 on a Linux Server (With Explanations)

Aug 4 40



...



J Jeyaram Ayyalusamy

24 - PostgreSQL 17 Performance Tuning: Monitoring Table-Level Statistics with pg_stat_user_tables

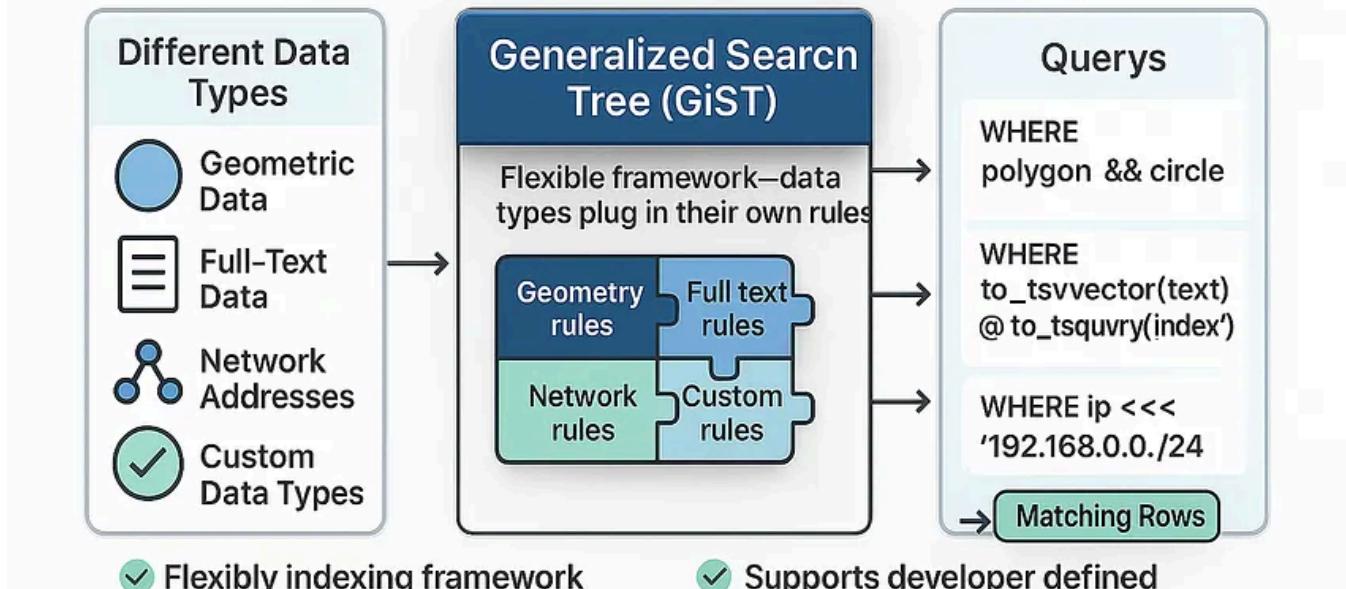
When tuning PostgreSQL, one of the most important steps is to observe table-level statistics. You cannot optimize what you cannot measure...

Sep 7 19 1



...

GiST (Generalized Search Tree)



J Jeyaram Ayyalusamy

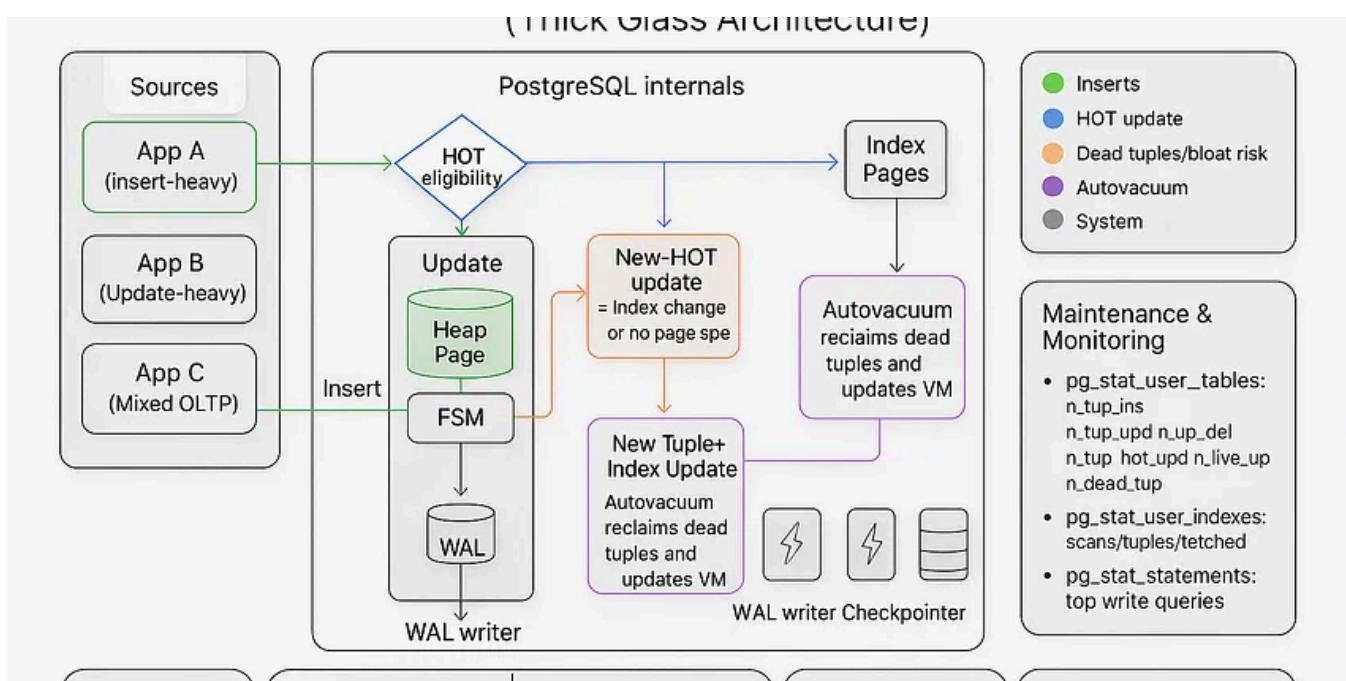
21—PostgreSQL 17 Performance Tuning: GiST (Generalized Search Tree)

GiST (Generalized Search Tree) in PostgreSQL is a flexible indexing framework that allows developers to build indexes for many different...

Sep 5 1



...



 Jeyaram Ayyalusamy 

23 - PostgreSQL 17 Performance Tuning: Monitoring Inserts, Updates, and HOT Updates

When tuning PostgreSQL, it is very important to understand the INSERT, UPDATE, and DELETE patterns of your tables. Different workloads...

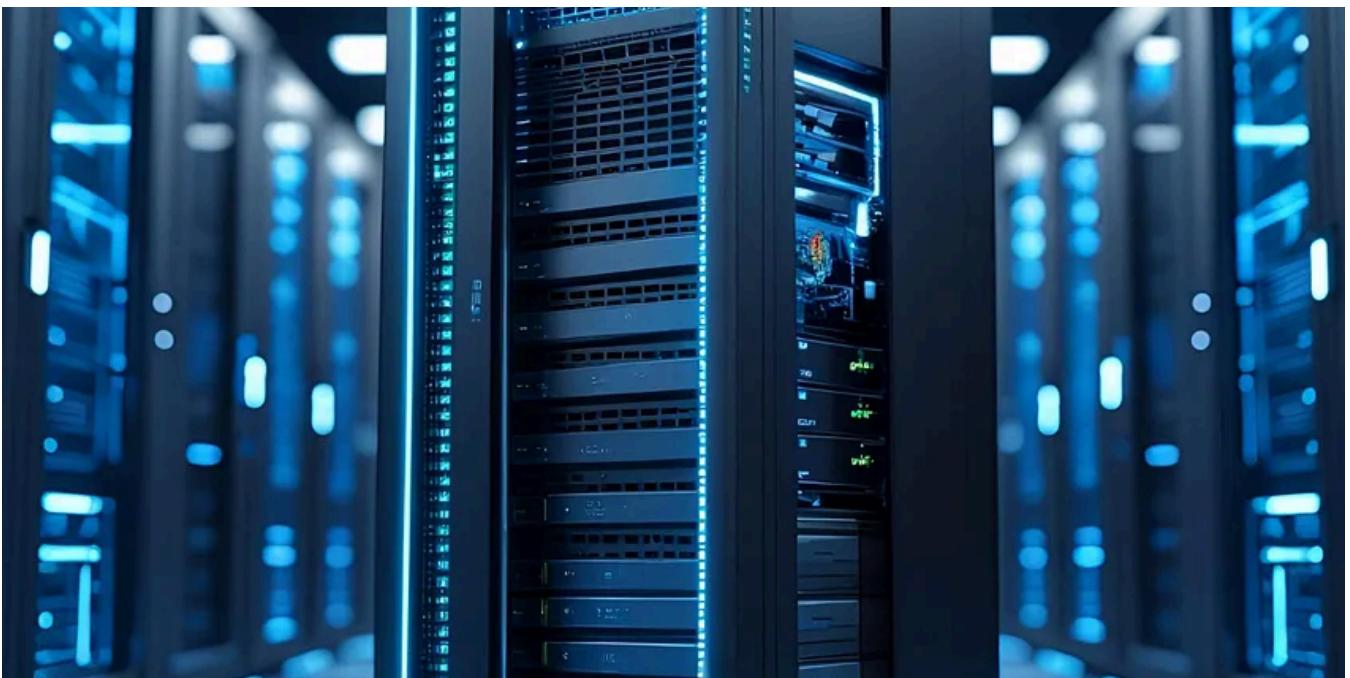
Sep 7  11



...

See all from Jeyaram Ayyalusamy

Recommended from Medium



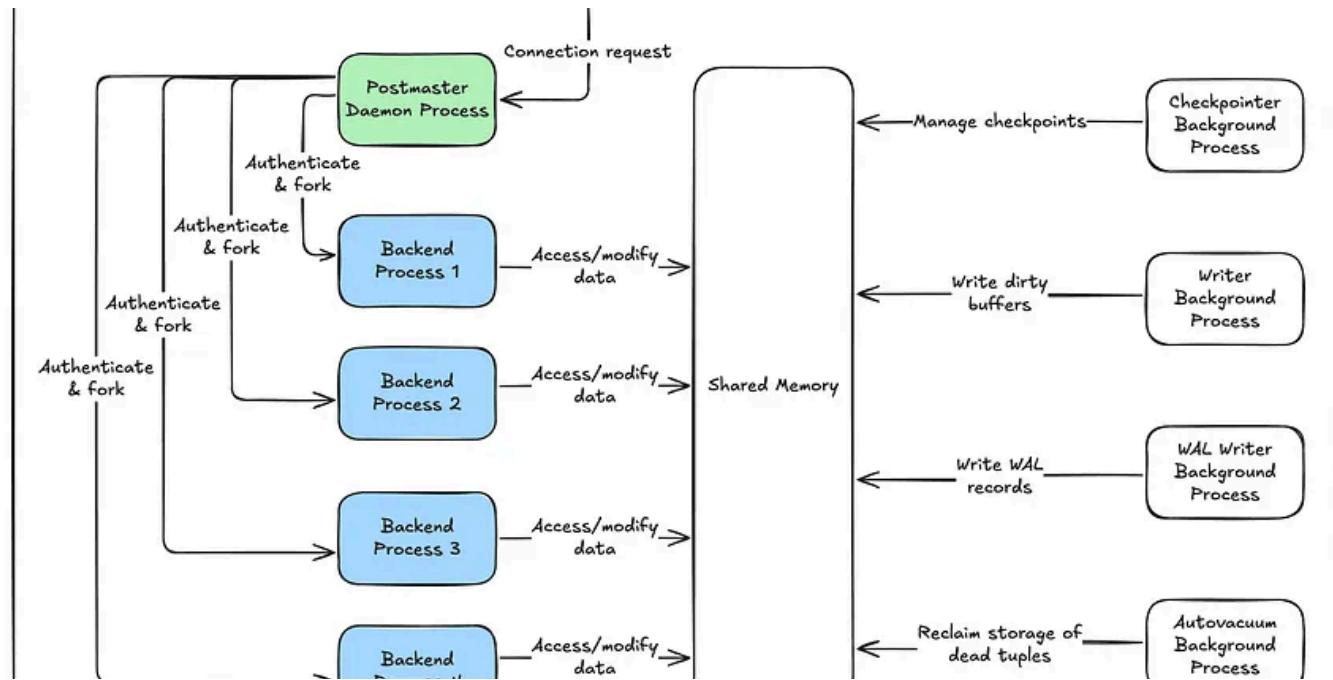
 Rizqi Mulki

PostgreSQL Index Bloat: The Silent Killer of Performance

How a 10TB database became 3x faster by fixing the invisible problem that plagues 90% of production PostgreSQL deployments

★ Sep 15 11 1

[+]

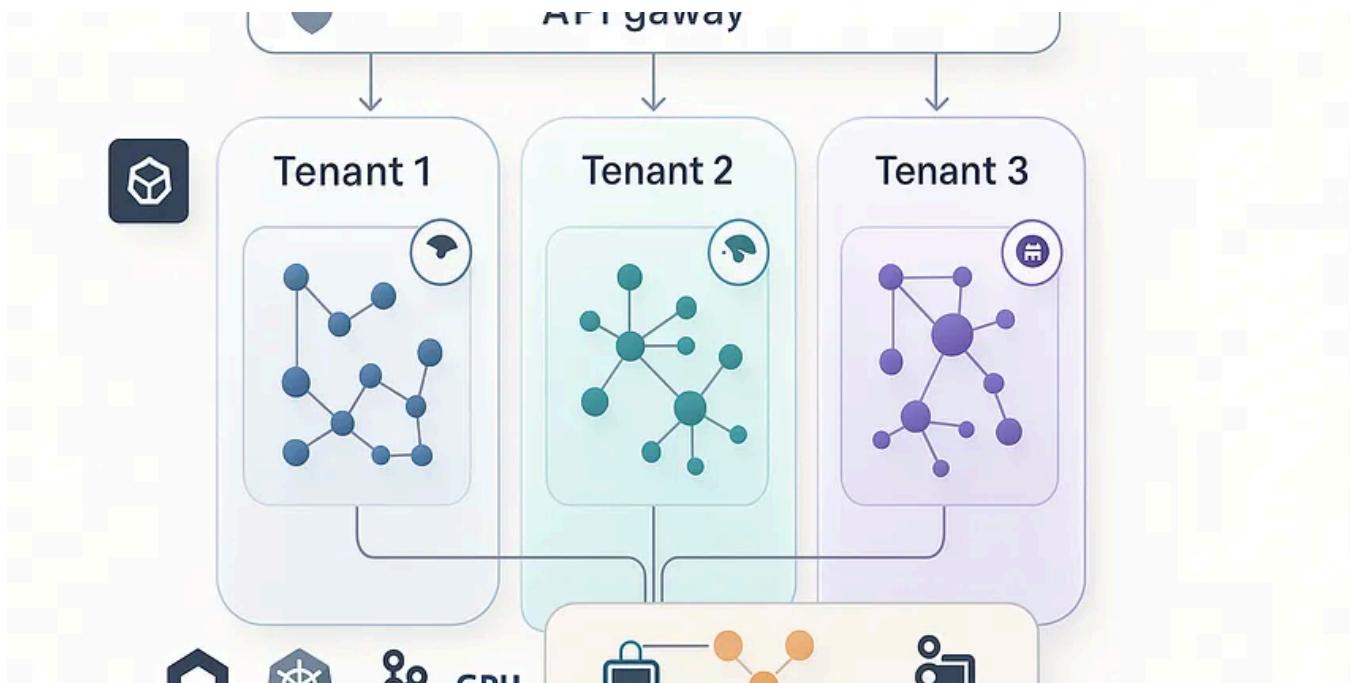


 Kareem Mohllal

Database Connection Pooling with PgBouncer

Database connections are expensive, let's explore why and how PgBouncer can save your database from drowning in connections.

Aug 5 ⌘ 10

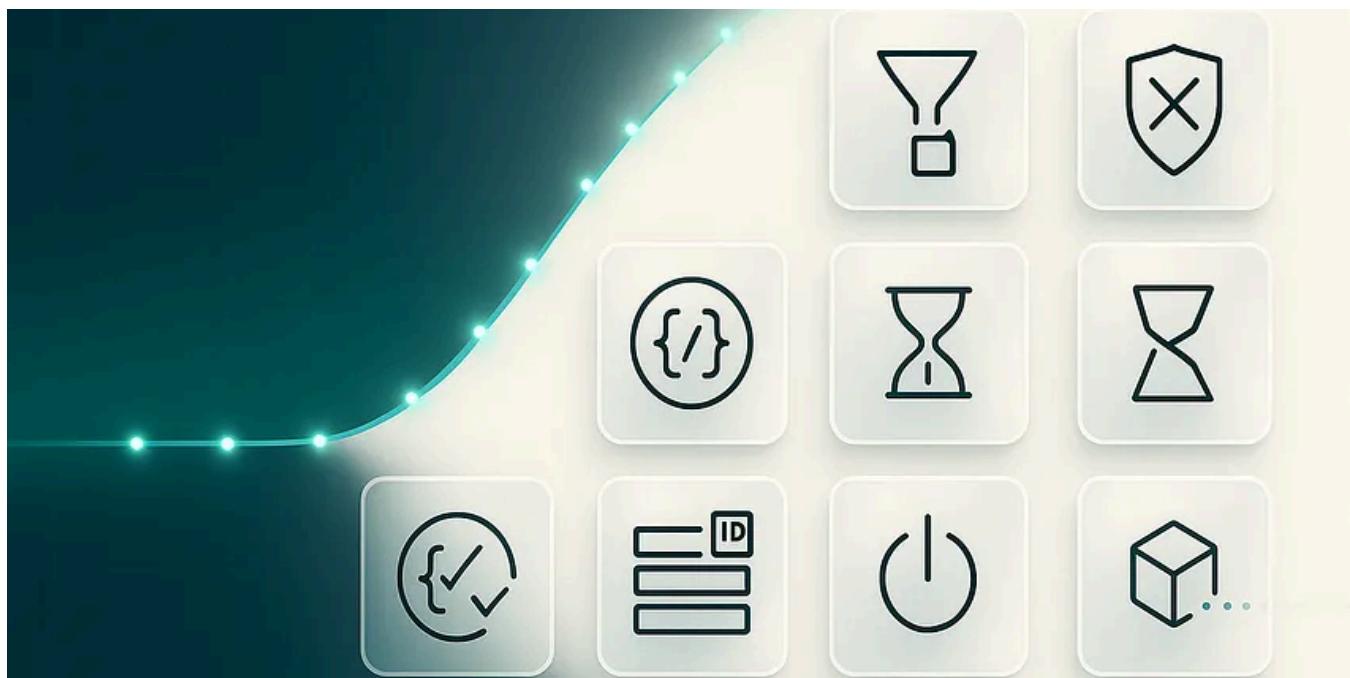


Thinking Loop

5 Ironclad Rules for Multi-Tenant Vector Isolation

Keep noisy neighbors out and SLAs intact—without overprovisioning your vector cluster.

★ Sep 15 ⌘ 17

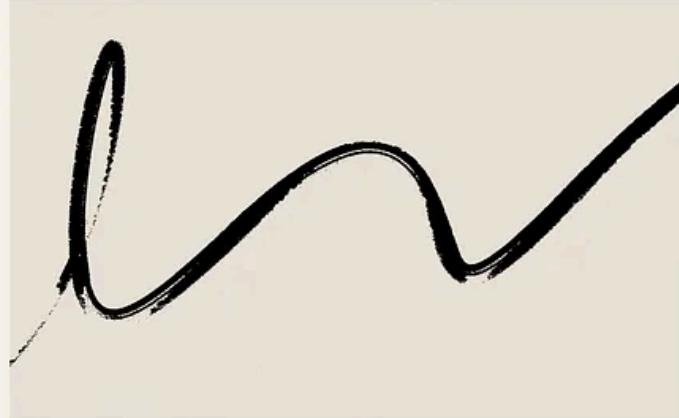


Hash Block

10 Node.js Error-Handling Tricks That Saved Me

Practical patterns to catch bugs early, keep services alive, and turn chaos into readable logs.

4d ago 17



R Rohan

JSONB in PostgreSQL: The Fast Lane to Flexible Data Modeling 🚀

“Why spin up yet another database just to store semi-structured data?”—Every engineer who discovered JSONB

Jul 18 12 1





Tomasz Gintowt

Security in PostgreSQL

PostgreSQL is a powerful open-source database. Security is very important when working with sensitive data like passwords, customer...

6d ago 5



...

[See more recommendations](#)