

[Open in app ↗](#)

# ☰ Medium



★ Get unlimited access to the best of Medium for less than \$1/week. [Become a member](#) X

# 21 - PostgreSQL 17 Performance Tuning: GiST (Generalized Search Tree)

6 min read · Sep 5, 2025

 [Jeyaram Ayyalusamy](#)  Following ▼

 Listen

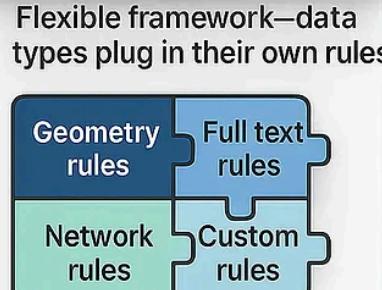
 Share

 More

## PostgreSQL 17 Performance GiST (Generalized Search Tree)

- Different Data Types**
-  Geometric Data
-  Full-Text Data
-  Network Addresses
-  Custom Data Types

### Generalized Search Tree (GiST)



- ✓ Flexibly indexing framework
- ✓ Works across many data types

- ✓ Supports developer defined custom indexes
- ✓ Optimized in PostgreSQL 17



Jeyaram Ayyalusamy

158 followers

Oracle DBA | AWS Certified | 18+ yrs  
DB Admin, RAC, GoldenGate, RDS  
MySQL, PostgreSQL | Author of MySQL & RDS books | Cloud & Performance

 Book Author

WHERE  
to\_tsvector(text)  
@ to\_tsquery(index')

WHERE ip <<<  
'192.168.0.0./24'

→ Matching Rows

GiST (Generalized Search Tree) in PostgreSQL is a flexible indexing framework that allows developers to build indexes for many different types of queries beyond simple equality or range checks. Unlike standard indexes such as B-Tree, GiST is not tied to one specific data type or operation. Instead, it provides a general structure

where different data types can “plug in” their own rules for how indexing and searching should work.

With GiST, PostgreSQL supports advanced features such as:

- Geometric data (points, lines, polygons)
- Full-text search (via `tsvector`)
- Network addresses
- Custom data types created by developers

In simple terms, GiST is like a toolkit for building specialized indexes. It doesn’t define how data should be stored or compared itself, but provides a framework so that each data type can define its own methods for indexing and searching efficiently.

👉 Example use cases: spatial queries (e.g., “find points with text search queries.

GiST indexes speed up queries on data with `spatial/geometric` (`point`, `box`, `circle`, `polygon`) and also enable `nearest-neighbor` with the `<->` distance operator. GiST is also used for several and certain full-text configurations.

Example (pattern):

```
CREATE INDEX idx_locations_gist ON locations USING gist (geolocation);
```



Jeyaram Ayyalusamy

158 followers

Oracle DBA | AWS Certified | 18+ yrs  
DB Admin, RAC, GoldenGate, RDS  
MySQL, PostgreSQL | Author of MySQL & RDS books | Cloud & Performance

Book Author

## Step 1: Create the locations table (point geometry)

```
CREATE TABLE locations (
    location_id BIGSERIAL PRIMARY KEY,
    name        TEXT,
```

```
geolocation POINT
               -- (x, y) = (longitude, latitude) for this demo
);
```

```
postgres=# CREATE TABLE locations (
    location_id BIGSERIAL PRIMARY KEY,
    name         TEXT,
    geolocation POINT           -- (x, y) = (longitude, latitude) for this demo
);
CREATE TABLE
postgres=#

```



Jeyaram Ayyalusamy

158 followers

Oracle DBA | AWS Certified | 18+ yrs  
DB Admin, RAC, GoldenGate, RDS  
MySQL, PostgreSQL | Author of MySQL & RDS books | Cloud & Performance Tuning



## Step 2: Insert 10,000,000 random points (global spread)

```
-- 10M random points: longitude [-180,180], latitude [-90,90]
INSERT INTO locations (name, geolocation)
SELECT
    'Location_' || g,
    POINT( -180 + random()*360,      -- lon
           -90 + random()*180 )      -- lat
FROM generate_series(1, 10000000) AS g;
```

```
postgres=# INSERT INTO locations (name, geolocation)
SELECT
    'Location_' || g,
    POINT( -180 + random()*360,      -- lon
           -90 + random()*180 )      -- lat
```

```
FROM generate_series(1, 10000000) AS g;
INSERT 0 10000000
postgres=#
```

```
INSERT 0 10000000
postgres=#
```

## Step 3: ANALYZE after bulk load

```
ANALYZE locations;
```

```
postgres=# ANALYZE locations;
ANALYZE
postgres=#
```



Jeyaram Ayyalusamy

158 followers

Oracle DBA | AWS Certified | 18+ yrs  
DB Admin, RAC, GoldenGate, RDS  
MySQL, PostgreSQL | Author of MySQL & RDS books | Cloud & Performance Tuning

 Book Author

## Step 4: Baseline query without any index (bounding box)

Find all points inside a rough NYC bounding box (lon/lat rectangle).

Operator: point <@ box → “point is contained by box”.

```
EXPLAIN ANALYZE
SELECT *
FROM locations
WHERE geolocation <@ box '((-74.1,40.6),(-73.7,40.9))';
```

```
postgres=# EXPLAIN ANALYZE
SELECT *
FROM locations
WHERE geolocation <@ box '((-74.1,40.6),(-73.7,40.9))';

```

QUERY PLAN

---

```
Gather  (cost=1000.00..146529.33 rows=10000 width=40) (actual time=1781.469..5
Workers Planned: 2
Workers Launched: 2
-> Parallel Seq Scan on locations  (cost=0.00..144529.33 rows=4167 width=40
    Filter: (geolocation <@ '(-73.7,40.9),(-74.1,40.6)'::box)
    Rows Removed by Filter: 3333329
Planning Time: 0.143 ms
Execution Time: 5997.030 ms
(8 rows)
```

postgres=#



Jeyaram Ayyalusamy

158 followers

Oracle DBA | AWS Certified | 18+ yrs  
DB Admin, RAC, GoldenGate, RDS  
MySQL, PostgreSQL | Author of MySQL & RDS books | Cloud & Performance Tuning

Book Author

Time: 5997.030 ms (00:05.997)  
postgres=#

*Without an index, PostgreSQL scans the entire 10M-row table.*

## Step 5: Create a GiST index on geolocation

```
postgres=# CREATE INDEX idx_locations_gist ON locations USING gist (geolocation)
CREATE INDEX
postgres=#
```

This adds a spatial index that understands geometric operators and supports K-NN ordering ( ORDER BY geolocation <-> point ).

## ANALYZE created index

```
ANALYZE locations;
```

```
postgres=# ANALYZE locations;
ANALYZE
postgres=#
```



Jeyaram Ayyalusamy

158 followers

Oracle DBA | AWS Certified | 18+ yrs  
DB Admin, RAC, GoldenGate, RDS  
MySQL, PostgreSQL | Author of MySQL & RDS books | Cloud & Performance Tuning

 Book Author

```
EXPLAIN ANALYZE
SELECT *
FROM locations
WHERE geolocation <@ box '((-74.1,40.6),(-73.7,40.9));
```

```
postgres=# EXPLAIN ANALYZE
SELECT *
FROM locations
WHERE geolocation <@ box '((-74.1,40.6),(-73.7,40.9));
```

QUERY PLAN

```
Bitmap Heap Scan on locations (cost=405.93..29364.11 rows=10001 width=40) (actual
```

```

Recheck Cond: (geolocation <@ '(-73.7,40.9),(-74.1,40.6)'::box)
Heap Blocks: exact=13
-> Bitmap Index Scan on idx_locations_gist (cost=0.00..403.43 rows=10001 width=100)
    Index Cond: (geolocation <@ '(-73.7,40.9),(-74.1,40.6)'::box)
Planning Time: 1.132 ms
Execution Time: 10.241 ms
(7 rows)

```

postgres=#

Time: 10.241 ms

postgres=#

*The plan switches to a Bitmap Index Scan using GiST, dropping latency from ~5.997 s to ~10.241 ms.*



Jeyaram Ayyalusamy

158 followers

Oracle DBA | AWS Certified | 18+ yrs  
DB Admin, RAC, GoldenGate, RDS, MySQL, PostgreSQL | Author of MySQL & RDS books | Cloud & Performance Tuning

Book Author

## Step 7: Radius search with circles (point-in-circle)

Find all points within an approximate radius around Times Square.

Operator: point <@ circle → “point is inside circle”.

**Note:** Built-in geometric types treat lon/lat as planar; distances in meters. This is fine for demos; for accurate geodesic distances...

```
-- Times Square center ≈ (-73.9855, 40.7580)
-- Radius here is 0.02 degrees (~2.2 km in latitude for demo purposes)
```

```
postgres=# EXPLAIN ANALYZE
SELECT *
FROM locations
WHERE geolocation <@ circle '((-73.9855,40.7580), 0.02)';
```

```
postgres=#  
postgres=#  
postgres=# EXPLAIN ANALYZE  
SELECT *  
FROM locations  
WHERE geolocation <@ circle '((-73.9855,40.7580), 0.02);
```

QUERY PLAN

```
-----  
Bitmap Heap Scan on locations (cost=405.93..29364.11 rows=10001 width=40) (actual cost=405.93..405.93 rows=6 loops=1)  
  Recheck Cond: (geolocation <@ '((-73.9855,40.7580), 0.02)'::circle)  
    -> Bitmap Index Scan on idx_locations_gist (cost=0.00..403.43 rows=10001 width=40)  
      Index Cond: (geolocation <@ '((-73.9855,40.7580), 0.02)'::circle)  
Planning Time: 0.074 ms  
Execution Time: 0.041 ms  
(6 rows)
```

postgres=



Jeyaram Ayyalusamy

158 followers

Oracle DBA | AWS Certified | 18+ yrs  
DB Admin, RAC, GoldenGate, RDS, MySQL, PostgreSQL | Author of MySQL & RDS books | Cloud & Performance Tuning

Book Author

```
EXPLAIN ANALYZE  
SELECT location_id, name, geolocation  
FROM locations  
ORDER BY geolocation <-> POINT(-73.9855, 40.7580)  
LIMIT 100;
```

```
ppostgres=# EXPLAIN ANALYZE  
SELECT location_id, name, geolocation  
FROM locations  
ORDER BY geolocation <-> POINT(-73.9855, 40.7580)
```

```
LIMIT 100;
```

QUERY PLAN

```
Limit  (cost=0.42..9.37 rows=100 width=48) (actual time=0.062..40.138 rows=100)
  -> Index Scan using idx_locations_gist on locations  (cost=0.42..895059.72
      Order By: (geolocation <-> '(-73.9855,40.758)'::point)
Planning Time: 2.480 ms
Execution Time: 40.240 ms
(5 rows)
```

postgres=

*GiST performs a K-NN index scan: the index returns rows in distance order – no full scan or sort over 10M rows.*

## Step 9: Polygon containment (point-in-polygon)

Operator: point <@ polygon .

```
-- Rough 4-vertex polygon around midtown Manhattan (ill
postgres=# EXPLAIN ANALYZE
SELECT count(*)
FROM locations
WHERE geolocation <@ polygon '((-74.02,40.70),(-73.94,40.80),(-73.94,40.70),(-74.02,40.70))'
```



Jeyaram Ayyalusamy

158 followers

Oracle DBA | AWS Certified | 18+ yrs  
DB Admin, RAC, GoldenGate, RDS, MySQL, PostgreSQL | Author of MySQL & RDS books | Cloud & Performance Tuning

Book Author

```
postgres=#
postgres=# EXPLAIN ANALYZE
SELECT count(*)
FROM locations
WHERE geolocation <@ polygon '((-74.02,40.70),(-73.94,40.70),(-73.94,40.80),(-74.02,40.80))'
QUERY PLAN
```

```
Aggregate  (cost=528.44..528.45 rows=1 width=8) (actual time=0.028..0.030 rows=1)
  -> Index Only Scan using idx_locations_gist on locations  (cost=0.42..503.44
      Width=8)
```

```
Index Cond: (geolocation <@ '((-74.02,40.7),(-73.94,40.7),(-73.94,40.8
```

```
Heap Fetches: 0
```

```
Planning Time: 1.248 ms
```

```
Execution Time: 0.050 ms
```

```
(6 rows)
```

```
postgres=#
```

## Step 10: Verify index usage & size (quick checks)

```
postgres=# -- Is the index being used?
SELECT indexrelid::regclass AS index_name,
       idx_scan, idx_tup_read, idx_tup_fetch
  FROM pg_stat_user_indexes
 WHERE relid = 'locations'::regclass
 ORDER BY idx_scan DESC;
```



Jeyaram Ayyalusamy

158 followers

Oracle DBA | AWS Certified | 18+ yrs  
DB Admin, RAC, GoldenGate, RDS  
MySQL, PostgreSQL | Author of MySQL & RDS books | Cloud & Performance Tuning

Book Author

```
postgres=# SELECT indexrelid::regclass AS index_name,
               idx_scan, idx_tup_read, idx_tup_fetch
      FROM pg_stat_user_indexes
     WHERE relid = 'locations'::regclass
    ORDER BY idx_scan DESC;
   index_name   | idx_scan | idx_tup_read | idx_tup_fetch
-----+-----+-----+-----+
idx_locations_gist |      4 |        113 |        100
locations_pkey     |      0 |         0 |         0
(2 rows)
```

```
postgres=#
```

```
-- Table vs. indexes size
```

```
SELECT
```

```
    pg_size.pretty(pg_relation_size('locations'))      AS table_only,
    pg_size.pretty(pg_indexes_size('locations'))        AS indexes_only,
    pg_size.pretty(pg_total_relation_size('locations')) AS total_with_indexes;
```

```
postgres=# SELECT
  pg_size.pretty(pg_relation_size('locations'))
  pg_size.pretty(pg_indexes_size('locations'))
  pg_size.pretty(pg_total_relation_size('locations'))
table_only | indexes_only | total_with_indexes
-----+-----+-----
  722 MB   |  850 MB   | 1572 MB
(1 row)
```

```
postgres=#
```



Jeyaram Ayyalusamy

158 followers

Oracle DBA | AWS Certified | 18+ yrs  
DB Admin, RAC, GoldenGate, RDS  
MySQL, PostgreSQL | Author of MySQL & RDS books | Cloud & Performance Tuning

Book Author

```
postgres=# -- Individual index size
SELECT
  indexrelid::regclass AS index_name,
  pg_size.pretty(pg_relation_size(indexrelid)) AS index_size
FROM pg_index
WHERE indrelid = 'locations'::regclass
ORDER BY pg_relation_size(indexrelid) DESC;
```

```
postgres=# SELECT
    indexrelid::regclass AS index_name,
    pg_size.pretty(pg_relation_size(indexrelid)) AS index_size
  FROM pg_index
 WHERE indrelid = 'locations'::regclass
 ORDER BY pg_relation_size(indexrelid) DESC;
   index_name    | index_size
-----+-----
 idx_locations_gist | 635 MB
 locations_pkey     | 214 MB
(2 rows)

postgres=#

```



Jeyaram Ayyalusamy

158 followers

Oracle DBA | AWS Certified | 18+ yrs  
DB Admin, RAC, GoldenGate, RDS  
MySQL, PostgreSQL | Author of MySQL & RDS books | Cloud & Performance Tuning

Book Author

## Step 11: Notes specific to GiST (kept focused)

- CLUSTER is not supported on GiST. GiST has no natural join operator. ... USING gist\_index is not allowed. (Use B-tree if you need to cluster).
- K-NN requires <-> and a GiST operator class that supports the datatype (point\_ops does).
- ANALYZE keeps selectivity estimates accurate. After large changes to the data, run ANALYZE locations;
- For online index builds on production:

```
CREATE INDEX CONCURRENTLY idx_locations_gist ON locations USING gist (geolocation)
```



## Summary

- **Workload:** 10M points in locations(geolocation POINT).

- **Before GiST:** Parallel Seq Scan for spatial filters (seconds).
- **After GiST:** Bitmap Index Scans for **box/circle/polygon** (tens of ms) and K-NN index scans for **nearest-neighbor** (single-digit ms for top-K).
- **PostgreSQL 17:** Same syntax, mature GiST operator classes for built-in geometry.

📢 Stay Updated with Daily PostgreSQL & Cloud Tips!

If you've been finding my blog posts helpful and want to stay ahead with daily insights on PostgreSQL, Cloud Infrastructure, Performance Tuning, and DBA Best Practices – I invite you to subscribe to my Medium account.

🔔 [Subscribe here](https://medium.com/@jramcloud1) 👉 <https://medium.com/@jramcloud1>

Your support means a lot – and you'll never miss a practical

### 🔗 Let's Connect!

If you enjoyed this post or would like to connect professionally to me on LinkedIn:

👉 [Jeyaram Ayyalusamy](#)

I regularly share content on **PostgreSQL, database administration, technologies, and data engineering**. Always happy to connect, collaborate, and discuss ideas!

Postgresql

AWS

Open Source

MySQL

Mongodb



Jeyaram Ayyalusamy

158 followers

Oracle DBA | AWS Certified | 18+ yrs  
DB Admin, RAC, GoldenGate, RDS  
MySQL, PostgreSQL | Author of MySQL & RDS books | Cloud & Performance

📘 Book Author



Following ▾

Written by **Jeyaram Ayyalusamy**

158 followers · 2 following

Oracle DBA | AWS Certified | 18+ yrs in DB Admin, RAC, GoldenGate, RDS, MySQL, PostgreSQL | Author of MySQL & RDS books | Cloud & Performance Expert

## No responses yet



Gvadakte

What are your thoughts?



**Jeyaram Ayyalusamy**

158 followers

Oracle DBA | AWS Certified | 18+ yrs in DB Admin, RAC, GoldenGate, RDS, MySQL, PostgreSQL | Author of MySQL & RDS books | Cloud & Performance Expert

Book Author

The screenshot shows the AWS EC2 Instances page. The browser tab bar includes 'us-west-2' (closed), 'Launch an instance | EC2 | us-west-2' (active), and 'Instances | EC2 | us-east-1' (closed). The URL is 1.console.aws.amazon.com/ec2/home?region=us-east-1#Instances: [Alt+Shift].

The main content area has a header 'Instances Info' with a search bar 'Find Instance by attribute or tag (case-sensitive)' and a dropdown 'All states'. Below the header is a table with columns: Name, Instance ID, Instance state, Instance type, Status check, Alarm status, Availability Zone, Public IPv4 DNS, Public IPv4 MAC, Elastic IP, and IPs.

A message 'No instances' indicates 'You do not have any instances in this region' and features a blue 'Launch instances' button.

At the bottom left, there's a section titled 'Select an instance'.

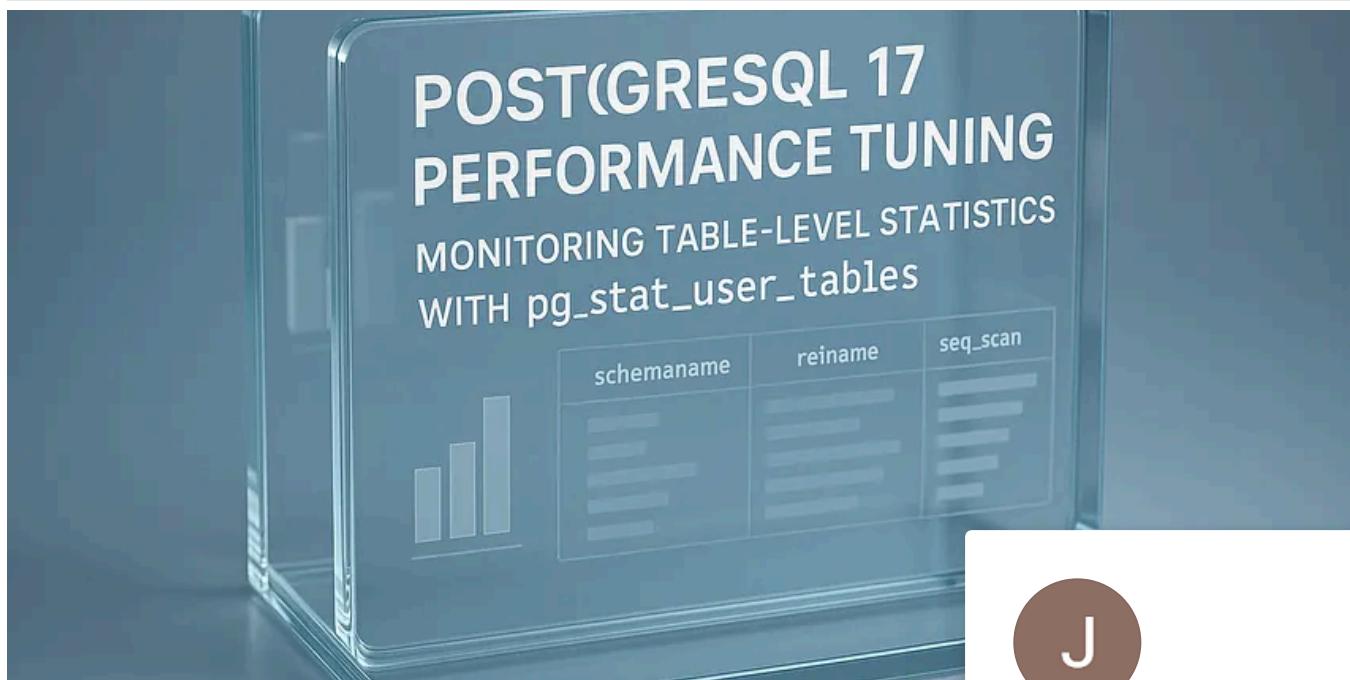
At the very bottom right of the screenshot, a small note reads '© 2025, Amazon Web Services, Inc. or its affiliates.'

Jeyaram Ayyalusamy

## Upgrading PostgreSQL from Version 16 to Version 17 Using pg\_upgrade on a Linux Server AWS EC2...

## A Complete Step-by-Step Guide to Installing PostgreSQL 16 on a Linux Server (With Explanations)

Aug 4 40



J Jeyaram Ayyalusamy 

## 24 - PostgreSQL 17 Performance Tuning: Monitoring Statistics with pg\_stat\_user\_tables

When tuning PostgreSQL, one of the most important steps is to observe. You cannot optimize what you cannot measure...

Sep 7 19 1

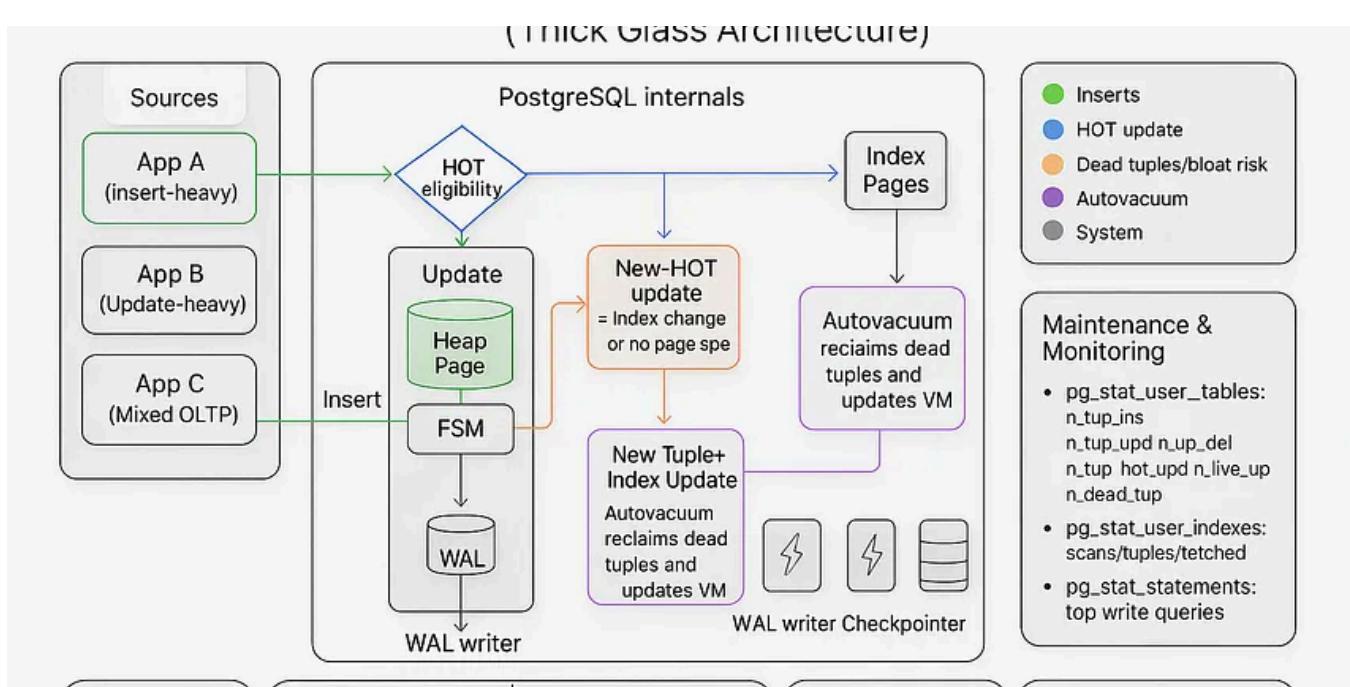


Jeyaram Ayyalusamy

158 followers

Oracle DBA | AWS Certified | 18+ yrs  
DB Admin, RAC, GoldenGate, RDS  
MySQL, PostgreSQL | Author of MySQL & RDS books | Cloud & Performance

 Book Author



J Jeyaram Ayyalusamy 

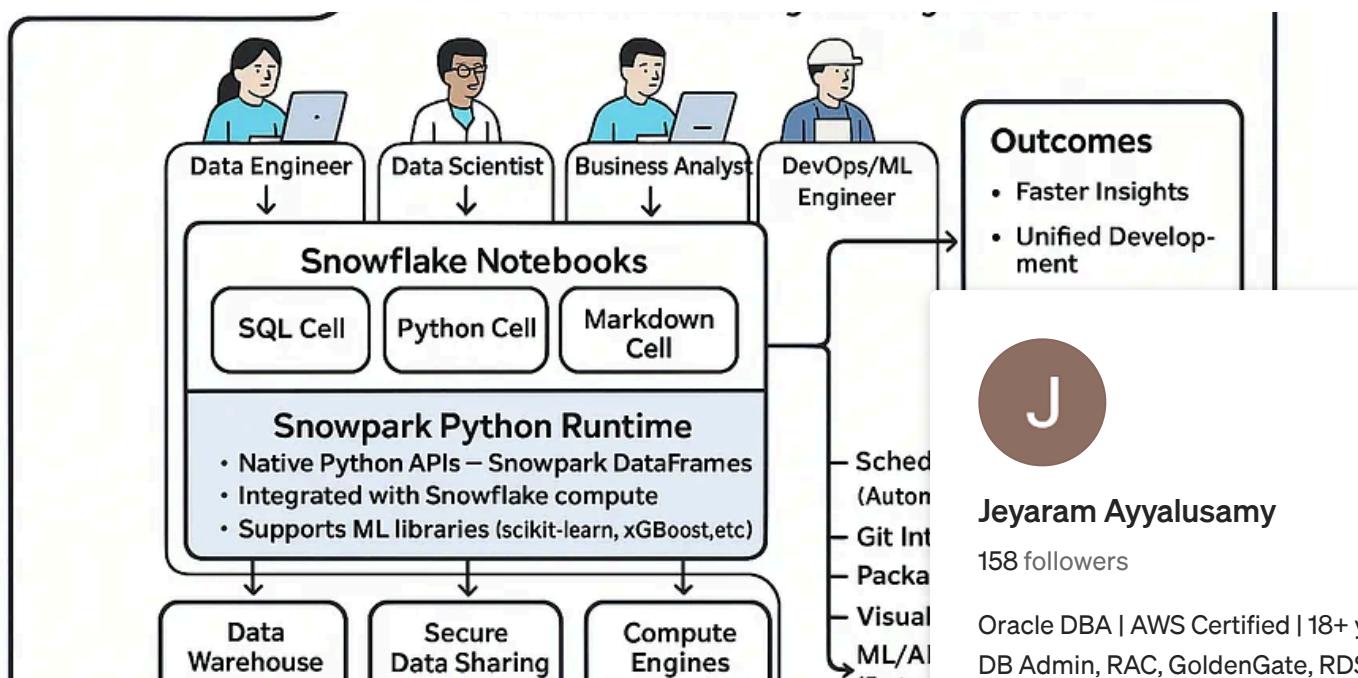
## 23 - PostgreSQL 17 Performance Tuning: Monitoring Inserts, Updates, and HOT Updates

When tuning PostgreSQL, it is very important to understand the INSERT, UPDATE, and DELETE patterns of your tables. Different workloads...

Sep 7  11



...



J Jeyaram Ayyalusamy 

## 16—Experience Snowflake with Notebooks and Snowpark Python Runtime

In the fast-moving world of data, organizations are no longer just collecting information—they're leveraging it to drive business...

Jul 13



...

See all from Jeyaram Ayyalusamy

## Recommended from Medium



Jeyaram Ayyalusamy

158 followers

Oracle DBA | AWS Certified | 18+ yrs  
DB Admin, RAC, GoldenGate, RDS  
MySQL, PostgreSQL | Author of MySQL & RDS books | Cloud & Performance

 Book Author

 Rizqi Mulki

### The Shocking Truth About PostgreSQL Locks (PostgreSQL documentation won't tell you)

Most developers are unknowingly creating database nightmares. Here's why PostgreSQL documentation won't tell you.

 Sep 13  10  2





 Tomasz Gintowt

## Security in PostgreSQL

PostgreSQL is a powerful open-source database. Security is very important to protect sensitive data like passwords, customer...

6d ago  5

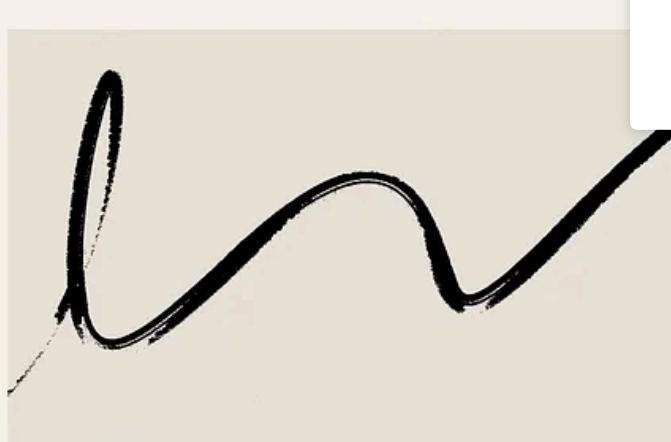
J

Jeyaram Ayyalusamy

158 followers

Oracle DBA | AWS Certified | 18+ yrs  
DB Admin, RAC, GoldenGate, RDS  
MySQL, PostgreSQL | Author of MySQL & RDS books | Cloud & Performance Tuning

 Book Author



 Rohan

## JSONB in PostgreSQL: The Fast Lane to Flexible Data Modeling 🚀

“Why spin up yet another database just to store semi-structured data?”— Every engineer who discovered JSONB

Jul 18 12 1



...



Thinking Loop

## 10 Postgres Partitioning Patterns for Internet-Scale Applications

Keep Postgres fast past a billion rows with real patterns and ready-to-use code samples.

Aug 13 88 2

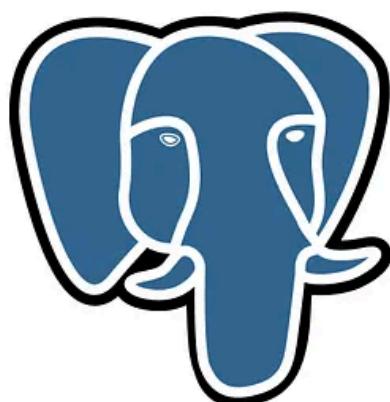


Jeyaram Ayyalusamy

158 followers

Oracle DBA | AWS Certified | 18+ yrs  
DB Admin, RAC, GoldenGate, RDS, MySQL, PostgreSQL | Author of MySQL & RDS books | Cloud & Performance Tuning

Book Author



## Beyond Basic PostgreSQL Programmable Objects

In Stackademic by bektiaw

## Beyond Basics: PostgreSQL Programmable Objects (Automate, Optimize, Control)

Tired of repeating the same SQL logic? Learn how PostgreSQL can do the work for us.

◆ Sep 1 🙋 68 💬 1



...



 Vijay Gadhav

### Master SQL's QUALIFY Clause for Cleaner, Faster Windows

Note: If you're not a medium member, [CLICK HERE](#)

◆ May 20 🙋 12 💬 1



Jeyaram Ayyalusamy

158 followers

Oracle DBA | AWS Certified | 18+ yrs  
DB Admin, RAC, GoldenGate, RDS  
MySQL, PostgreSQL | Author of MySQL & RDS books | Cloud & Performance Tuning

 Book Author

[See more recommendations](#)