

[Open in app ↗](#)

Medium



Get unlimited access to the best of Medium for less than \$1/week. [Become a member](#) X

18 - PostgreSQL 17 Performance Tuning: Hash Index

5 min read · Sep 4, 2025



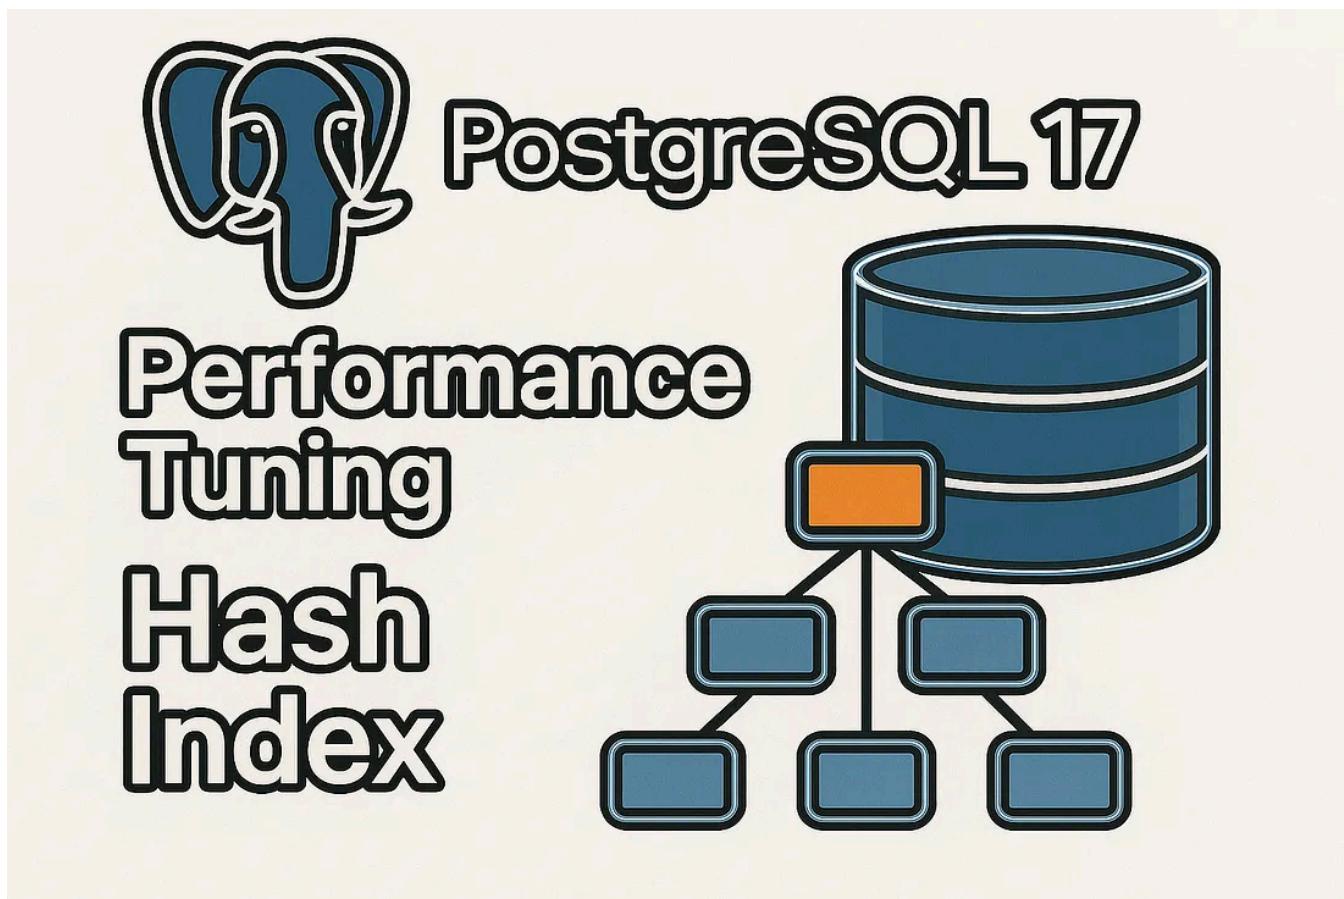
Jeyaram Ayyalusamy

Following ▼

Listen

Share

More



Indexes are the backbone of database performance tuning. In PostgreSQL, the default index type is B-Tree, which handles a wide variety of queries very efficiently. But sometimes, a different type of index is a better fit.

What is a Hash Index in PostgreSQL?

A **Hash Index** is a special type of index in PostgreSQL that is optimized for **equality lookups** — queries where you check if a column equals a specific value (e.g., WHERE email = 'user123@example.com').

Instead of storing values in a sorted order like a **B-Tree index**, a hash index uses a **hashing function**. The value you search for is hashed, and PostgreSQL jumps directly to the matching location in the index.

Key Characteristics

- **Best for = operator**

Works only for equality conditions. You cannot use it for range queries (< , > , BETWEEN) or sorting.

- **Fast lookups**

For exact matches, hash indexes can be extremely fast, especially with large datasets.

- **Compact structure**

They don't need to maintain order, so storage is usually smaller compared to B-Tree indexes.

- **Production-ready in PostgreSQL 17**

Older PostgreSQL versions had limitations: hash indexes weren't WAL-logged (not crash safe). Starting from PostgreSQL 10 and now in **PostgreSQL 17**, they are fully durable and safe to use.

When to Use Hash Indexes

- ✓ Use when:

- You often query with equality checks (=).
- The column has many unique values (e.g., email, username, ID).

✖ Avoid when:

- You need range lookups (`>`, `<`, `BETWEEN`).
- You want the index to support sorting or ordering.
- You already benefit from a **B-Tree index**, which supports equality and ranges.

👉 In short:

A **Hash Index** in PostgreSQL is a **special-purpose index** that makes equality lookups **blazing fast**, but it's not a general replacement for the default **B-Tree index**.

Step 1: Creating the `employees` table

We'll start by creating a simple `employees` table. It contains employee details, including `email`, which will be the focus of our equality lookups.

```
CREATE TABLE employees (
    emp_id      BIGINT,
    emp_name    TEXT,
    email       TEXT,
    dept        TEXT,
    salary      NUMERIC
);
```

```
postgres=# CREATE TABLE employees (
    emp_id      BIGINT,
    emp_name    TEXT,
    email       TEXT,
    dept        TEXT,
    salary      NUMERIC
```

```
)  
CREATE TABLE  
postgres=#
```

📌 Here:

- `emp_id` is a unique identifier.
- `email` will be queried often, and we will later add a **hash index** on it.
- Other columns (`name`, `dept`, `salary`) are for realism and testing.

Step 2: Inserting 10 million rows

To test performance properly, we need a large dataset. We'll generate 10 million employee records with distinct email addresses.

```
-- Insert 10 million rows with random emails and departments
INSERT INTO employees (emp_name, email, dept, salary)
SELECT
  'Employee_' || g,
  'user' || g || '@example.com',
  'Dept_' || (g % 20),           -- 20 departments
  (random()*100000)::NUMERIC(10,2)  -- salary between 0 and 100k
FROM generate_series(1, 10000000) g;
```

```
postgres=# INSERT INTO employees (emp_name, email, dept, salary)
SELECT
  'Employee_' || g,
  'user' || g || '@example.com',
  'Dept_' || (g % 20),           -- 20 departments
  (random()*100000)::NUMERIC(10,2)  -- salary between 0 and 100k
FROM generate_series(1, 10000000) g;
```

```
INSERT 0 100000000
```

```
postgres=#
```

- ✓ After this step, the table has 10 million rows, each with a unique email (`user1@example.com`, `user2@example.com`, ..., `user10000000@example.com`).

Step 3: Analyzing the table

Whenever you bulk-load data, it's important to run `ANALYZE`. This updates the PostgreSQL query planner with accurate statistics.

```
ANALYZE employees;
```

```
postgres=# ANALYZE employees;
ANALYZE
postgres=#
```

Without this step, the planner might guess row counts incorrectly and choose suboptimal execution plans.

Step 4: Running a query without an index

Let's run a simple lookup: find the row where `email = 'user5000000@example.com'`.

```
EXPLAIN ANALYZE
```

```
SELECT * FROM employees WHERE email = 'user5000000@example.com';
```

```
postgres=# EXPLAIN ANALYZE
SELECT * FROM employees WHERE email = 'user5000000@example.com';
                                         QUERY PLAN
-----
Gather  (cost=1000.00..164395.43 rows=1 width=62) (actual time=4931.459..4936.
  Workers Planned: 2
  Workers Launched: 2
    -> Parallel Seq Scan on employees  (cost=0.00..163395.33 rows=1 width=62) (
        Filter: (email = 'user5000000@example.com'::text)
        Rows Removed by Filter: 3333333
Planning Time: 0.469 ms
Execution Time: 4936.386 ms
(8 rows)

postgres#
```

Time: 4936.386 ms (00:04.936)
postgres#

📊 What happened?

- PostgreSQL performed a **sequential scan**, reading all 10M rows.
- Only **one row** matched, but 3,333,333 rows had to be checked and discarded.
- The execution time was ~4.9 seconds — too slow for real-world workloads.

Step 5: Creating a hash index

Now, let's create a **Hash Index** on the `email` column.

```
CREATE INDEX idx_employee_email_hash ON employees USING hash (email);
```

```
postgres=# CREATE INDEX idx_employee_email_hash ON employees USING hash (email)
CREATE INDEX
postgres=#

```

📌 Explanation:

- `USING hash` tells PostgreSQL to build a **hash-based index** instead of the default B-Tree.
- Hash indexes are optimized for **equality checks**, not range scans.

Analyzing the table

Whenever you bulk-load data or creating new indexes, it's important to run `ANALYZE`. This updates the PostgreSQL query planner with accurate statistics.

```
ANALYZE employees;
```

```
postgres=# ANALYZE employees;
ANALYZE
```

```
postgres=#
```

Step 6: Query plan after using the hash index

Run the same query again:

```
postgres=# EXPLAIN ANALYZE
SELECT * FROM employees WHERE email = 'user5000000@example.com';
                                         QUERY PLAN
-----
Index Scan using idx_employee_email_hash on employees (cost=0.00..8.02 rows=1
    Index Cond: (email = 'user5000000@example.com'::text)
Planning Time: 0.130 ms
Execution Time: 0.030 ms
(4 rows)

postgres=#

```

Time: 0.030 ms
postgres=#

🔥 Huge improvement:

- The sequential scan is replaced with an **Index Scan**.
- The database jumps directly to the matching row using the hash index.
- Execution time dropped from ~5.2 seconds to less than 1 millisecond.

Step 7: Understanding when to use Hash Indexes

Hash indexes are **specialized**. They are not a replacement for B-Tree indexes, but they excel in certain cases:

✓ When Hash Indexes Shine:

- Queries with equality conditions (`=`).
- Lookups on columns with high cardinality (many unique values), like `email`, `username`, or `UUID`.

✖ When Not to Use:

- Range queries (`<`, `>`, `BETWEEN`) → Use B-Tree instead.
- Sorting or ordering → B-Tree indexes are required.

Step 8: Improvements in PostgreSQL 17

Historically, hash indexes had limitations (not WAL-logged, crash unsafe). But now:

- **WAL-logged:** Hash indexes survive crashes safely.
- **VACUUM/ANALYZE aware:** They integrate smoothly with normal maintenance.
- **Production-ready:** Fully safe and recommended for workloads needing equality lookups.

Step 9: Verifying index size

Sometimes indexes can be large. You can check how much space they take compared to the table:

```
SELECT
    pg_size.pretty(pg_relation_size('employees')) AS table_size,
    pg_size.pretty(pg_indexes_size('employees')) AS indexes_size,
    pg_size.pretty(pg_total_relation_size('employees')) AS total_size;
```

```
postgres=# SELECT
    pg_size.pretty(pg_relation_size('employees')) AS table_size,
    pg_size.pretty(pg_indexes_size('employees')) AS indexes_size,
    pg_size.pretty(pg_total_relation_size('employees')) AS total_size;
table_size | indexes_size | total_size
-----+-----+
  870 MB   |  256 MB     |  1126 MB
(1 row)

postgres=#
```

This helps you understand the **storage overhead** of adding hash indexes.

Final Thoughts

This experiment shows how **Hash Indexes in PostgreSQL 17** can make equality lookups **thousands of times faster**.

- Without indexes, queries degrade to slow sequential scans.
- With a hash index, lookups become nearly instant.
- The choice of index type matters — **match the index to the query pattern**.

👉 If your workload involves frequent equality lookups on columns like email, username, or unique IDs, hash indexes are a powerful option.

📢 Stay Updated with Daily PostgreSQL & Cloud Tips!

If you've been finding my blog posts helpful and want to stay ahead with daily insights on PostgreSQL, Cloud Infrastructure, Performance Tuning, and DBA Best Practices — I invite you to subscribe to my Medium account.

🔔 Subscribe here 👉 <https://medium.com/@jramcloud1/subscribe>

Your support means a lot — and you'll never miss a practical guide again!

🔗 Let's Connect!

If you enjoyed this post or would like to connect professionally, feel free to reach out to me on LinkedIn:

👉 [Jeyaram Ayyalusamy](#)

I regularly share content on **PostgreSQL, database administration, cloud technologies, and data engineering**. Always happy to connect, collaborate, and discuss ideas!

Postgresql

MySQL

AWS

Open Source

Oracle

A circular profile picture of a man with dark hair and a beard, wearing a dark shirt.

Following ▾

Written by Jeyaram Ayyalusamy

158 followers · 2 following

Oracle DBA | AWS Certified | 18+ yrs in DB Admin, RAC, GoldenGate, RDS, MySQL, PostgreSQL | Author of MySQL & RDS books | Cloud & Performance Expert

No responses yet



Gvadakte

What are your thoughts?



More from Jeyaram Ayyalusamy

The screenshot shows the AWS Management Console interface for the EC2 service. The top navigation bar includes tabs for 'Launch an instance' and 'Instances'. The main content area is titled 'Instances Info' and displays a search bar and filter options. A message states 'No instances' and 'You do not have any instances in this region'. A prominent blue 'Launch instances' button is centered below the message. The bottom right corner of the page footer contains the text '© 2025, Amazon Web Services, Inc. or its affiliates.'

J Jeyaram Ayyalusamy

Upgrading PostgreSQL from Version 16 to Version 17 Using pg_upgrade on a Linux Server AWS EC2...

A Complete Step-by-Step Guide to Installing PostgreSQL 16 on a Linux Server (With Explanations)

Aug 4 40



J Jeyaram Ayyalusamy 

24 - PostgreSQL 17 Performance Tuning: Monitoring Table-Level Statistics with pg_stat_user_tables

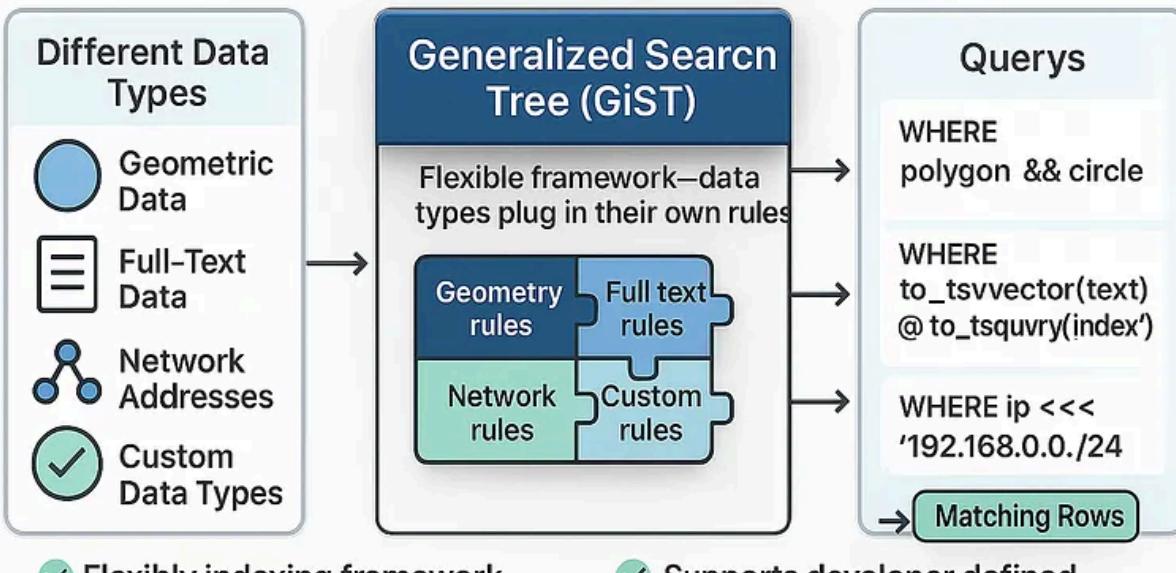
When tuning PostgreSQL, one of the most important steps is to observe table-level statistics. You cannot optimize what you cannot measure...

Sep 7  19  1



...

GiST (Generalized Search Tree)



J Jeyaram Ayyalusamy 

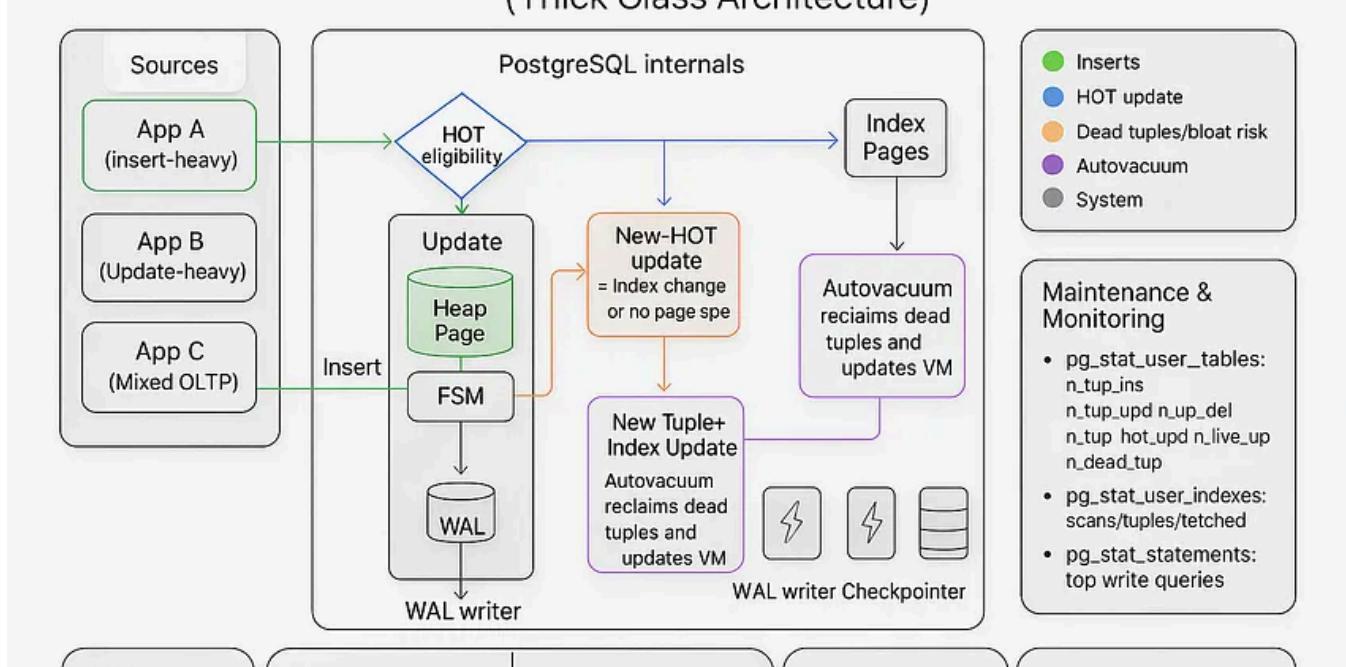
21—PostgreSQL 17 Performance Tuning: GiST (Generalized Search Tree)

GiST (Generalized Search Tree) in PostgreSQL is a flexible indexing framework that allows developers to build indexes for many different...

Sep 5  1



...



J Jeyaram Ayyalusamy

23 - PostgreSQL 17 Performance Tuning: Monitoring Inserts, Updates, and HOT Updates

When tuning PostgreSQL, it is very important to understand the INSERT, UPDATE, and DELETE patterns of your tables. Different workloads...

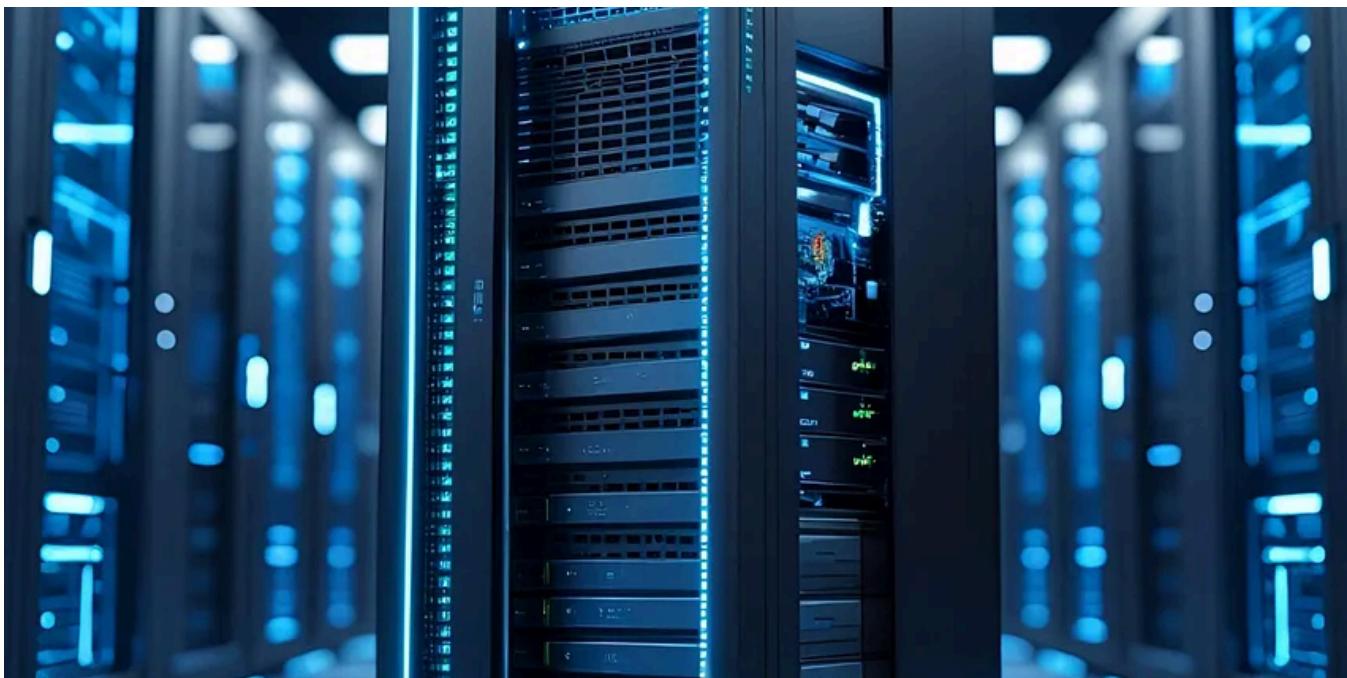
Sep 7 11



...

See all from Jeyaram Ayyalusamy

Recommended from Medium



Rizqi Mulki

PostgreSQL Index Bloat: The Silent Killer of Performance

How a 10TB database became 3x faster by fixing the invisible problem that plagues 90% of production PostgreSQL deployments

★ Sep 15

11

1



...

#PostgreSQL

security

TOMASZ GINTOWT



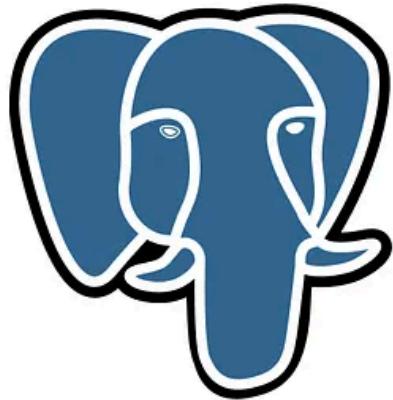
Tomasz Gintowt

Security in PostgreSQL

PostgreSQL is a powerful open-source database. Security is very important when working with sensitive data like passwords, customer...

6d ago 5

...



Beyond Basic PostgreSQL Programmable Objects



In Stackademic by bektiaw

Beyond Basics: PostgreSQL Programmable Objects (Automate, Optimize, Control)

Tired of repeating the same SQL logic? Learn how PostgreSQL can do the work for us.



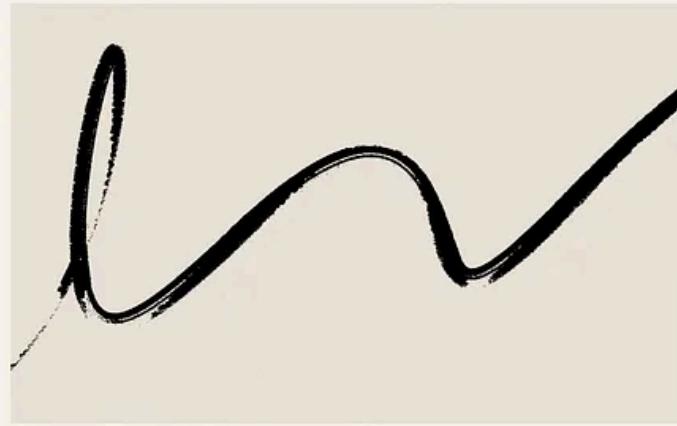
Sep 1

56

1



...



Rohan

JSONB in PostgreSQL: The Fast Lane to Flexible Data Modeling 🚀

“Why spin up yet another database just to store semi-structured data?”—Every engineer who discovered JSONB

Jul 18 12 1

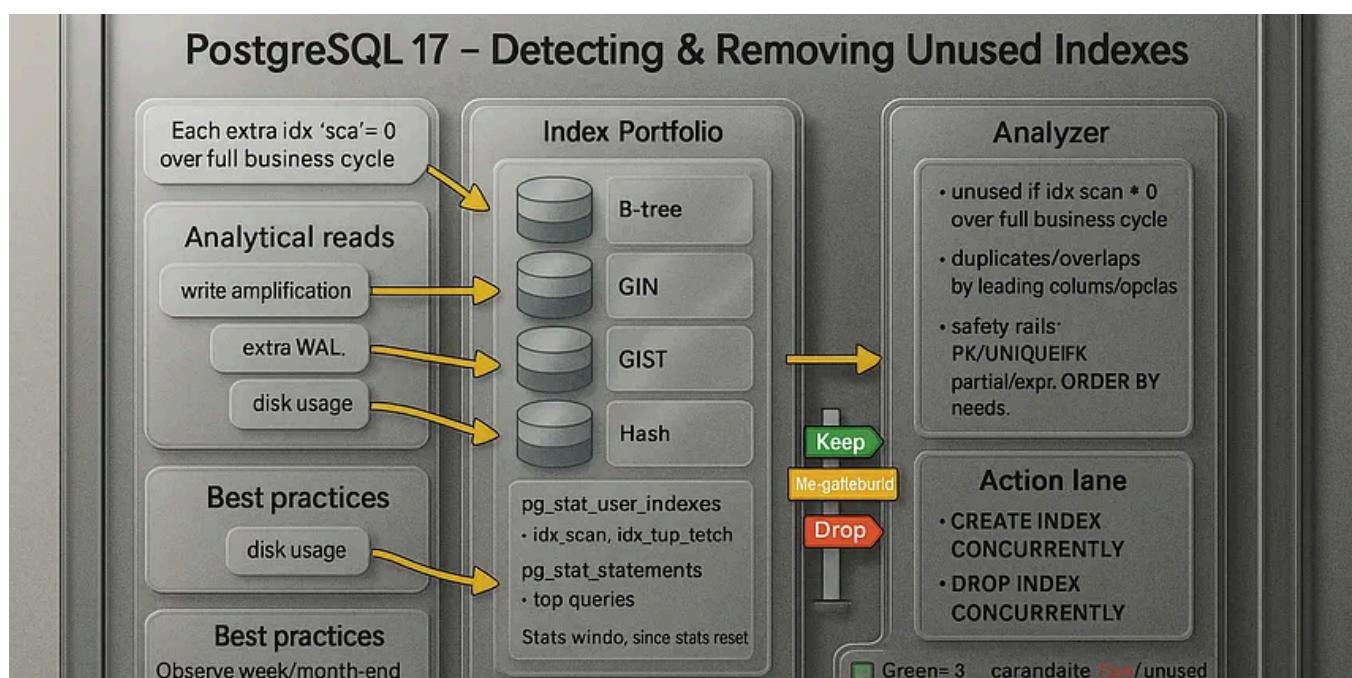


Ajaymaurya

PostgreSQL 18 Features That Change Everything

When a new PostgreSQL release drops, the developer world always pauses. PostgreSQL 18 is no exception—it's bold, innovative, and packed...

Sep 14 19 2



 Jeyaram Ayyalusamy 

25 - PostgreSQL 17 Performance Tuning: Detecting and Removing Unused Indexes

Indexes are one of the most important tools for query performance in PostgreSQL. But while a missing index slows down queries, an unused or...

Sep 10  40



...

See more recommendations