

[Open in app ↗](#)

Medium



Get unlimited access to the best of Medium for less than \$1/week. [Become a member](#)



13- PostgreSQL 17 Performance Tuning: Indexing and Its Impact on Performance

8 min read · Sep 2, 2025



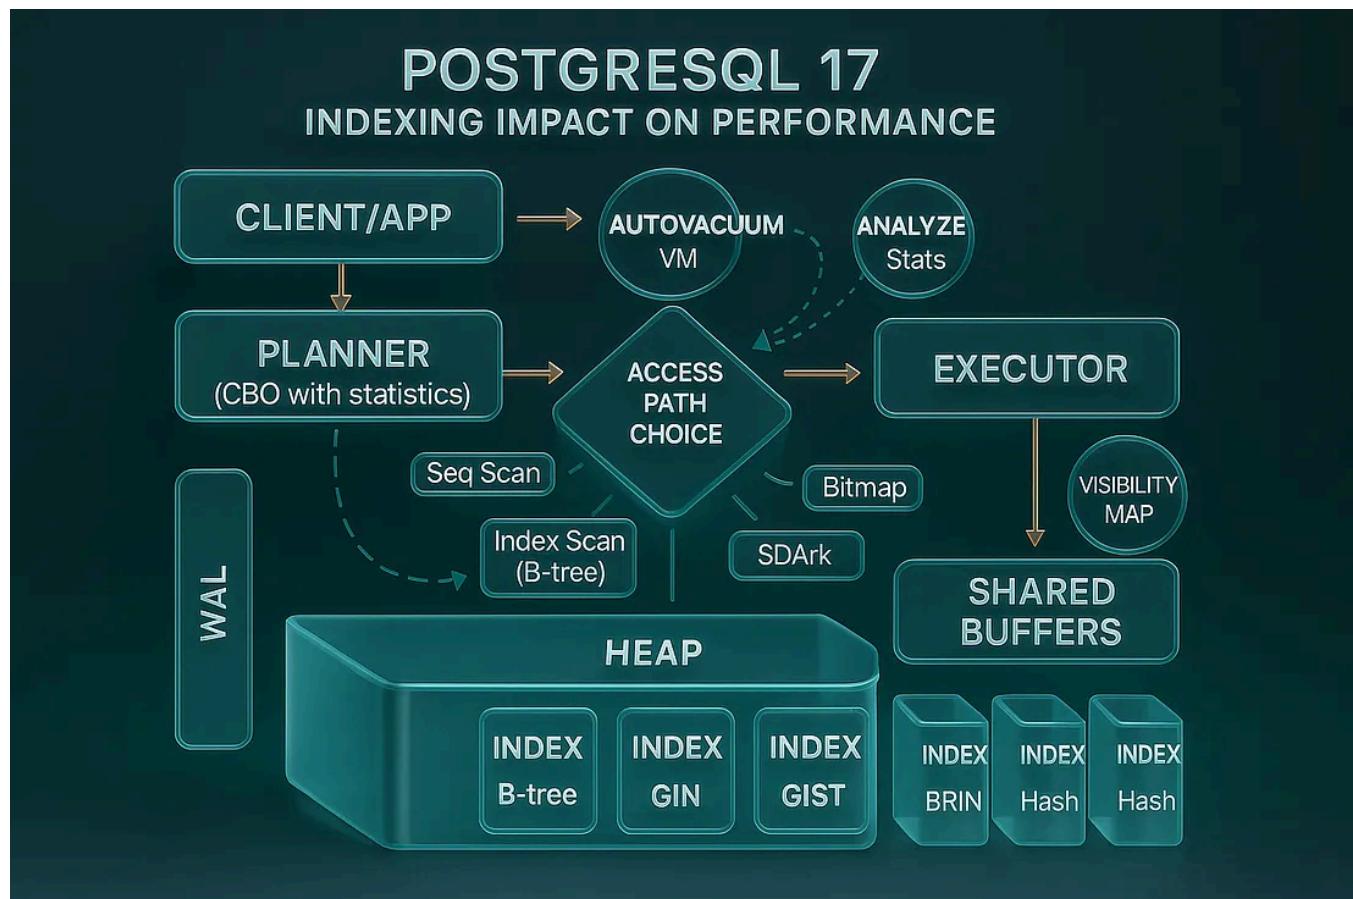
Jeyaram Ayyalusamy

Following

Listen

Share

More



The importance of indexing in PostgreSQL cannot be stressed enough. While memory and autovacuum tuning are important, they cannot compensate for missing indexes. There is simply no replacement for a missing index.

To achieve good performance, proper indexing is essential. In this demo, we will create a large table, run queries without an index, examine performance with `EXPLAIN ANALYZE`, and then add an index to see how it changes query speed and resource usage.

Step 1: Create a Large Table

We'll create a table called `products` with multiple columns to simulate a realistic workload.

```
CREATE TABLE products (
    product_id    BIGINT,
    product_name  TEXT,
    category      TEXT,
    price         NUMERIC,
    stock_qty     INT
);
```

```
postgres=# CREATE TABLE products (
    product_id    BIGINT,
    product_name  TEXT,
    category      TEXT,
    price         NUMERIC,
    stock_qty     INT
);
CREATE TABLE
postgres=#
```

Now we load 10 million rows of test data:

```
-- Insert 5 million rows for category 'Electronics'
INSERT INTO products
```

```
SELECT generate_series(1, 5000000),
       'Product_' || generate_series(1, 5000000),
       'Electronics',
       (random() * 500)::NUMERIC,
       (random() * 100)::INT;
```

```
postgres=# -- Insert 5 million rows for category 'Electronics'
INSERT INTO products
SELECT generate_series(1, 5000000),
       'Product_' || generate_series(1, 5000000),
       'Electronics',
       (random() * 500)::NUMERIC,
       (random() * 100)::INT;
INSERT 0 5000000
postgres=#
```

```
-- Insert 5 million rows for category 'Clothing'
INSERT INTO products
SELECT generate_series(5000001, 10000000),
       'Product_' || generate_series(5000001, 10000000),
       'Clothing',
       (random() * 200)::NUMERIC,
       (random() * 500)::INT;
```

```
postgres=# -- Insert 5 million rows for category 'Clothing'
INSERT INTO products
SELECT generate_series(5000001, 10000000),
       'Product_' || generate_series(5000001, 10000000),
```

```
'Clothing',
(random() * 200)::NUMERIC,
(random() * 500)::INT;
INSERT 0 5000000
postgres=#
```

👉 At this point, the table `products` has **10 million rows** with realistic columns and only two distinct categories (`Electronics`, `Clothing`).

Step 2: Update Statistics

Before running queries, refresh optimizer statistics so PostgreSQL knows about the data distribution:

```
ANALYZE products;
```

```
postgres=# ANALYZE products;
ANALYZE
postgres=#
```

The query planner uses these statistics to decide whether an index is useful.

Step 3: Query Without an Index

Now let's try to fetch a single product:

Reading psql \timing and a “slow single-row lookup”

You turned on psql's timing, ran a point lookup, and saw ~3.1 seconds:

```
\timing  
SELECT * FROM products WHERE product_id = 424242;
```

```
postgres=# \timing  
Timing is on.  
postgres=#  
postgres=# SELECT * FROM products WHERE product_id = 424242;  
 product_id | product_name | category | price | stock_qty  
-----+-----+-----+-----+-----  
 424242 | Product_424242 | Category_2 | 40.16 | 8  
(1 row)  
  
Time: 7.590 ms  
postgres=#
```

Run the query a few times to avoid caching artifacts.

👉 You'll notice the query is **very slow**.

Why? Because PostgreSQL has no index, so it must scan the entire `products` table.

What `\timing` actually measures

- `\timing` is a psql client feature. The Time printed is wall-clock elapsed time from the client's perspective.
- It includes:
 - Server planning + execution time,
 - Network round-trip latency,
 - Client rendering time (printing rows).
- Because this query returned **one row**, the 3.124s is almost certainly **server work and/or waiting**, not printing overhead.

Check the execution plan:

```
EXPLAIN ANALYZE SELECT * FROM products WHERE product_id = 424242;
```

```
postgres=# EXPLAIN ANALYZE SELECT * FROM products WHERE product_id = 424242;
                                         QUERY PLAN
-----
Gather  (cost=1000.00..153967.66 rows=1 width=48) (actual time=2784.976..2789.
    Workers Planned: 2
    Workers Launched: 2
        -> Parallel Seq Scan on products  (cost=0.00..152967.56 rows=1 width=48) (a
            Filter: (product_id = 424242)
            Rows Removed by Filter: 3333333
Planning Time: 0.063 ms
Execution Time: 2789.897 ms
(8 rows)

Time: 2790.251 ms (00:02.790)
postgres=#

```

What this plan says at a glance

- A **parallel sequential scan** read essentially the **entire `products` table** looking for one row.
- **Three participants** did the work: 2 parallel workers plus the leader (coordinated by a `Gather` node).
- About **~10 million rows** were scanned and discarded ($3,333,333 \text{ per participant} \times 3 \approx 9,999,999$), producing **exactly 1 row**.
- End-to-end latency was **~2.79 seconds**.

Decoding each part

```
Gather (cost=1000.00..153967.66 rows=1 width=48) (actual time=2784.976..2789.880 rows=1
loops=1)
```

- **Node role:** Gather orchestrates parallel workers and merges their outputs.
- **cost=...**: planner's *estimated* effort (abstract units, not time). `1000.00` is startup cost; `153967.66` is total estimated cost to produce all rows.
- **rows=1**: the planner expected **one** result row.
- **width=48**: estimated average row size (payload only).
- **actual_time=...**: measured wall-clock times in ms to first and last row emitted by this node: ~2785 ms to first output, ~2790 ms to finish.
- **rows=1 loops=1**: produced 1 row, executed once.

Workers Planned: 2 / Workers Launched: 2

- The planner requested two parallel workers; both started successfully. The **leader** often participates, so **3 participants total** processed table chunks in parallel.

```
-> Parallel Seq Scan on products (cost=0.00..152967.56 rows=1 width=48) (actual
time=1868.044..2780.034 rows=0 loops=3)
```

- **Node role:** each participant performs a **sequential scan** over a disjoint slice of products .
- **Timing:** workers began returning tuples around **1.868 s** and the last finished around **2.780 s**.
- **rows=0 loops=3**: this is per-participant reporting. Because **only one single row** matched among ~10M scanned and it was found by just one participant, the averaged per-worker output rounds to 0. Don't be misled—**the plan as a whole returned 1 row** (as shown at the `Gather`).

Filter: (product_id = 424242)

- The predicate was applied **row-by-row** during the scan. With a sequential scan, every visited tuple is checked against this filter.

Rows Removed by Filter: 3333333

- Roughly **3.33M rows per participant** failed the predicate. With 3 participants, that implies ~10M rows scanned in total to find the single match.

Planning Time: 0.063 ms

- The planner chose this plan almost instantly. The time was spent in **execution**, not planning.

Execution Time: 2789.897 ms **and final Time:** 2790.251 ms

- About 2.79 s from start to finish. The final Time: line is the psql-reported wall time and aligns with the executor's Execution Time.

Why did PostgreSQL 17 choose a Parallel Seq Scan?

Given the predicate `product_id = 424242` (a selective equality), we'd *expect* an index-based plan. This plan indicates that, for this query at this moment, the executor believed a **parallel full-table read** was the cheapest way to get the row. The most common interpretations of this plan output are:

- No usable B-tree index on `product_id` exists, so a scan of all pages is required and parallelism helps reduce elapsed time.
- A usable index exists but **wasn't chosen** because the cost model for this query favored a parallel scan (for example, the planner estimated many heap fetches anyway, or statistics suggested poor selectivity). The plan itself doesn't show whether an index exists; it only shows the path taken.

Step 4: Add an Index

Let's fix the issue by creating an index on `product_id`:

```
CREATE INDEX idx_products_id ON products(product_id);
```

```
postgres=# CREATE INDEX idx_products_id ON products(product_id);
CREATE INDEX
Time: 11811.905 ms (00:11.812)
```

```
postgres=#  
postgres=#
```

- **CREATE INDEX idx_products_id ON products(product_id);**

This statement builds a B-tree index (the default method) named `idx_products_id` on the `product_id` column of the `products` table. B-tree is optimal for equality lookups like `product_id = 424242`.

- **CREATE INDEX**

Confirms the index was created successfully.

Step 5: Query With an Index

Now run the same query again:

```
EXPLAIN ANALYZE SELECT * FROM products WHERE product_id = 424242;
```

```
postgres=# EXPLAIN ANALYZE SELECT * FROM products WHERE product_id = 424242;  
QUERY PLAN  
-----  
Index Scan using idx_products_id on products (cost=0.43..8.45 rows=1 width=48)  
  Index Cond: (product_id = 424242)  
Planning Time: 0.175 ms  
Execution Time: 2.069 ms  
(4 rows)
```

```
Time: 3.616 ms
```

```
postgres=#
```

Plan header

- **Index Scan using idx_products_id on products**

The query now uses the newly created index to locate the target row. An **Index Scan** walks the index to find the matching key(s) and fetches the corresponding row(s) from the table.

- **(cost=0.43..8.45 rows=1 width=48)**
 - **cost** values are **planner estimates** in abstract units (not time).
 - 0.43 : estimated startup cost before producing the first row.
 - 8.45 : estimated total cost to produce all rows for this node.
 - **rows=1** : the planner expects to return **one** row.
 - **width=48** : estimated average payload size (bytes) per row at this node.
 - **(actual_time=2.046..2.048 rows=1 loops=1)**
 - **actual_time** shows **measured** milliseconds for the first and last output row from this node.
 - The first row arrived at ~2.046 ms, the last at ~2.048 ms.
 - **rows=1** confirms that exactly **one** row was produced.
 - **loops=1** means the node executed once.

Predicate

- **Index Cond: (product_id = 424242)**

This is the **sargable** condition the index uses. The engine navigates the B-tree directly to the key 424242 instead of scanning many unrelated rows.

Planning and execution timings

- **Planning Time: 0.175 ms**

The optimizer produced this plan extremely quickly (well under a millisecond).

- **Execution Time: 2.069 ms**

The executor spent ~2.07 ms running the query end-to-end (as measured by PostgreSQL).

psql wall-clock

- **Time: 3.616 ms**

psql's wall-clock for the command, which includes client-side overhead in addition to the server's execution time.

👉 PostgreSQL now uses an **index scan**.

- Instead of reading 10 million rows, it jumps directly to the correct row using the index.
- Query execution time improves dramatically.

Step 6: The Cost of Indexes

Indexes are not free. Let's compare the size of the table and the index:

```
SELECT
    pg_size.pretty(pg_relation_size('products')) AS table_size,
    pg_size.pretty(pg_relation_size('idx_products_id')) AS index_size;
```

```
postgres=# SELECT
    pg_size.pretty(pg_relation_size('products')) AS table_size,
```

```
pg_size.pretty(pg_relation_size('idx_products_id')) AS index_size;
table_size | index_size
-----+-----
 785 MB    | 214 MB
(1 row)

postgres=#
```

👉 You will see that the index consumes a **significant amount of disk space** — sometimes more than half the size of the table itself.

Trade-offs of indexes:

- **Disk usage:** Indexes consume additional storage.
- **Write overhead:** Inserts, updates, and deletes must also update the index.
- **Maintenance cost:** Each index must remain synchronized with the table, slowing down write-heavy workloads.

For example, inserting new rows into `products` now requires not only writing data to the table but also updating the index. If a table has **multiple indexes**, each insert/update must update all of them, which can drastically slow down writes.

Key Takeaways

- Proper indexing is **essential** for PostgreSQL performance.
 - Without indexes, PostgreSQL may perform full table scans, even for queries fetching a single row.
 - Indexes greatly improve read performance but come at a cost in disk usage and write speed.
 - Always check the size and necessity of an index before adding it — not all indexes are worth their overhead.
- ✓ In PostgreSQL 17, just as in earlier versions, proper indexing is the difference between queries taking **seconds vs. milliseconds**.

📢 Stay Updated with Daily PostgreSQL & Cloud Tips!

If you've been finding my blog posts helpful and want to stay ahead with daily insights on PostgreSQL, Cloud Infrastructure, Performance Tuning, and DBA Best Practices — I invite you to subscribe to my Medium account.

🔔 [Subscribe here](https://medium.com/@jramcloud1/subscribe) 👉 <https://medium.com/@jramcloud1/subscribe>

Your support means a lot — and you'll never miss a practical guide again!

🔗 Let's Connect!

If you enjoyed this post or would like to connect professionally, feel free to reach out to me on LinkedIn:

👉 [Jeyaram Ayyalusamy](#)

I regularly share content on PostgreSQL, database administration, cloud technologies, and data engineering. Always happy to connect, collaborate, and discuss ideas!

Postgresql

AWS

Open Source

Oracle

MySQL



Following ▾

Written by Jeyaram Ayyalusamy

158 followers · 2 following

Oracle DBA | AWS Certified | 18+ yrs in DB Admin, RAC, GoldenGate, RDS, MySQL, PostgreSQL | Author of MySQL & RDS books | Cloud & Performance Expert

No responses yet





Gvadakte



What are your thoughts?

More from Jeyaram Ayyalusamy

The screenshot shows the AWS EC2 Instances page. At the top, there are three tabs: 'us-west-2' (selected), 'Launch an instance | EC2', and 'Instances | EC2'. Below the tabs, the URL is 1.console.aws.amazon.com/ec2/home?region=us-east-1#instances:. The main content area is titled 'Instances Info' and contains a search bar with placeholder 'Find Instance by attribute or tag (case-sensitive)' and a dropdown menu set to 'All states'. There are several filter buttons for 'Name', 'Instance ID', 'Instance state', 'Instance type', 'Status check', 'Alarm status', 'Availability Zone', 'Public IPv4 DNS', 'Public IPv4 ...', 'Elastic IP', and 'IPv6'. A message 'No instances' is displayed, followed by the sub-message 'You do not have any instances in this region' and a blue 'Launch instances' button. On the left side, there is a sidebar with the heading 'Select an instance'.

J Jeyaram Ayyalusamy

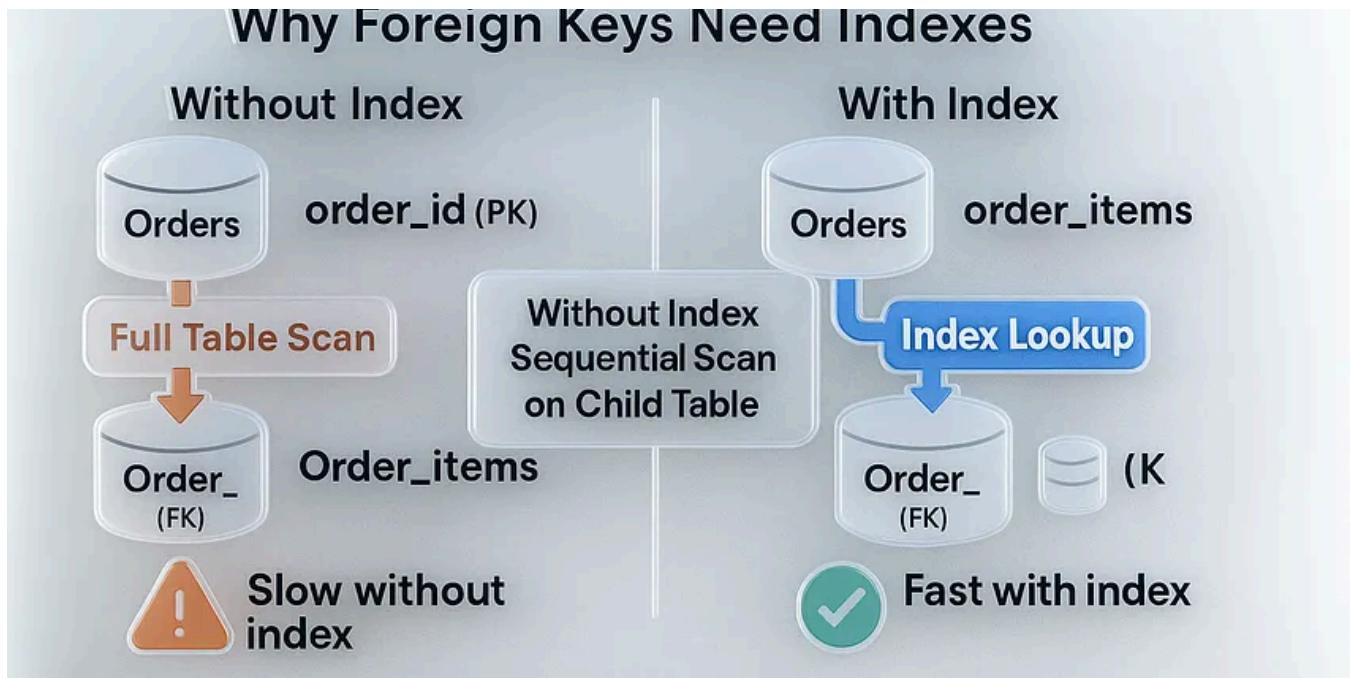
Upgrading PostgreSQL from Version 16 to Version 17 Using pg_upgrade on a Linux Server AWS EC2...

A Complete Step-by-Step Guide to Installing PostgreSQL 16 on a Linux Server (With Explanations)

Aug 4 40



...

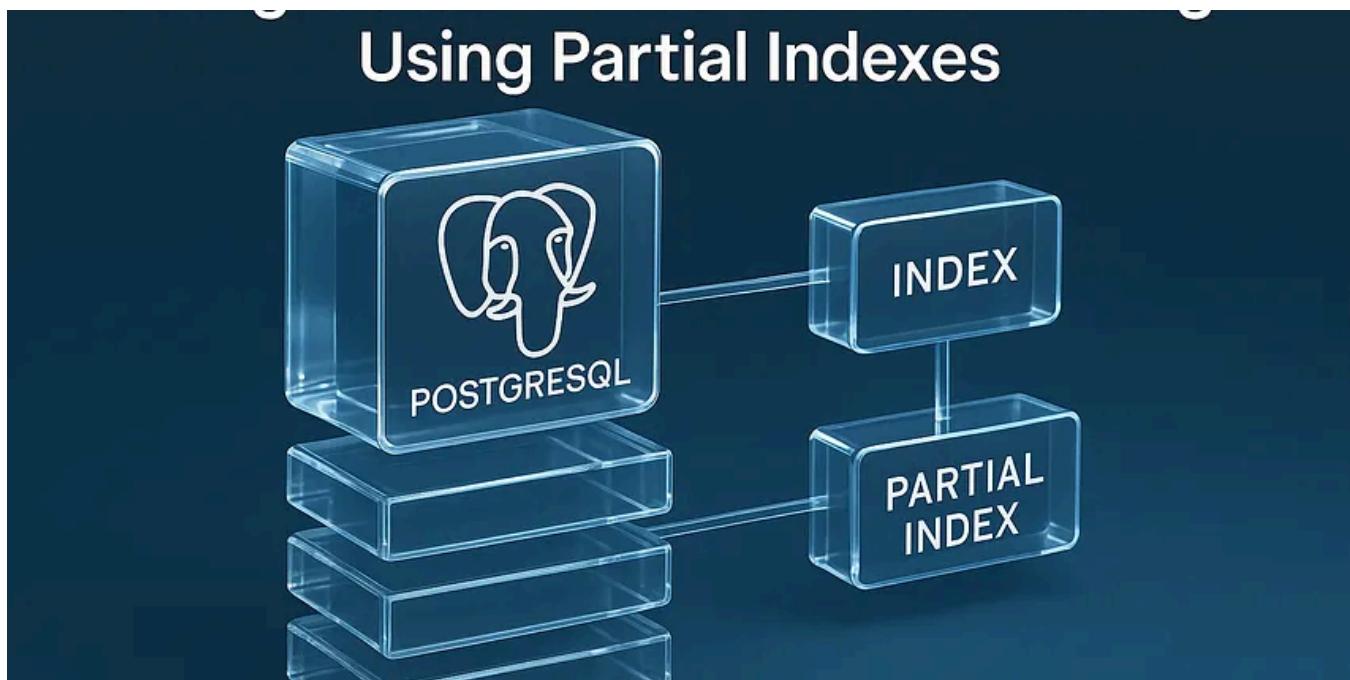


J Jeyaram Ayyalusamy

16 -PostgreSQL 17 Performance Tuning: Why Foreign Keys Need Indexes

When designing relational databases, parent-child relationships are common. A classic case is the Orders table (parent) and the Items table...

Sep 3 3 2



J Jeyaram Ayyalusamy

17 - PostgreSQL 17 Performance Tuning: Using Partial Indexes

Even though we usually create indexes on all rows of a table, PostgreSQL also allows you to create an index on just a subset of records...

Sep 4

3



...



Jeyaram Ayyalusamy

24 - PostgreSQL 17 Performance Tuning: Monitoring Table-Level Statistics with `pg_stat_user_tables`

When tuning PostgreSQL, one of the most important steps is to observe table-level statistics. You cannot optimize what you cannot measure...

Sep 7

19

1



...

See all from Jeyaram Ayyalusamy

Recommended from Medium

security

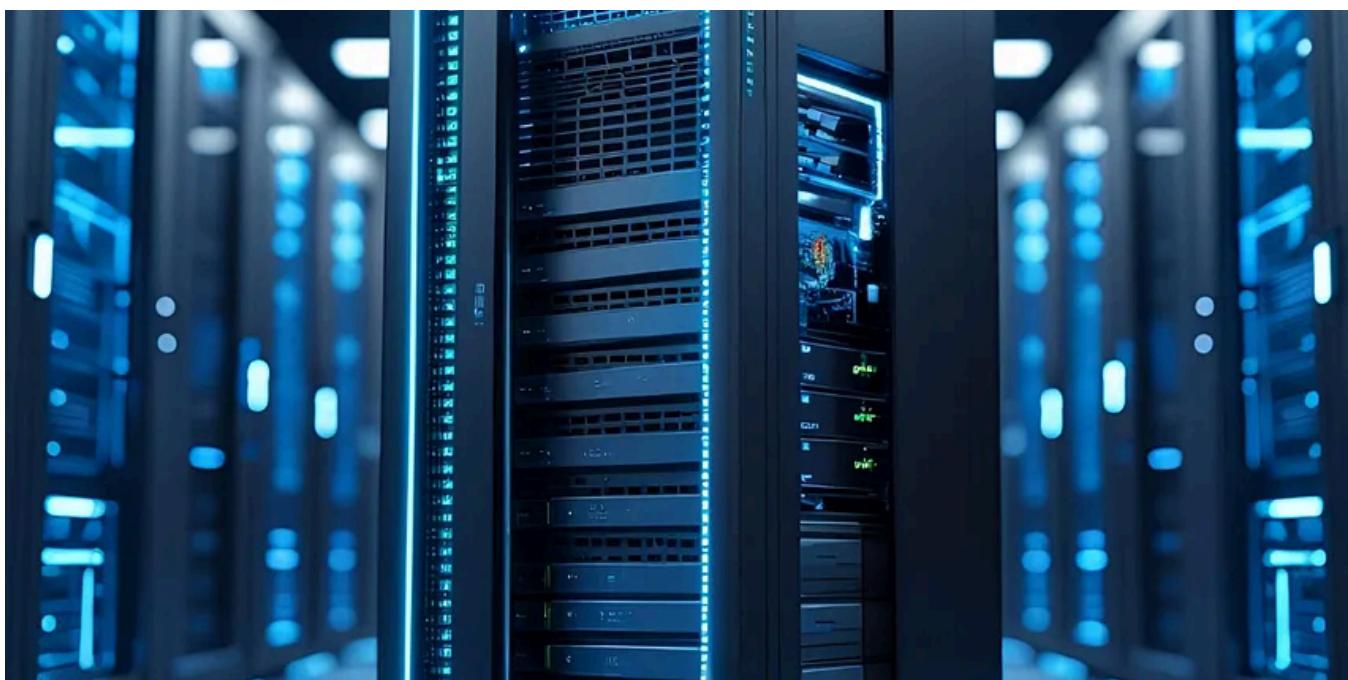
TOMASZ GINTOWT

 Tomasz Gintowt

Security in PostgreSQL

PostgreSQL is a powerful open-source database. Security is very important when working with sensitive data like passwords, customer...

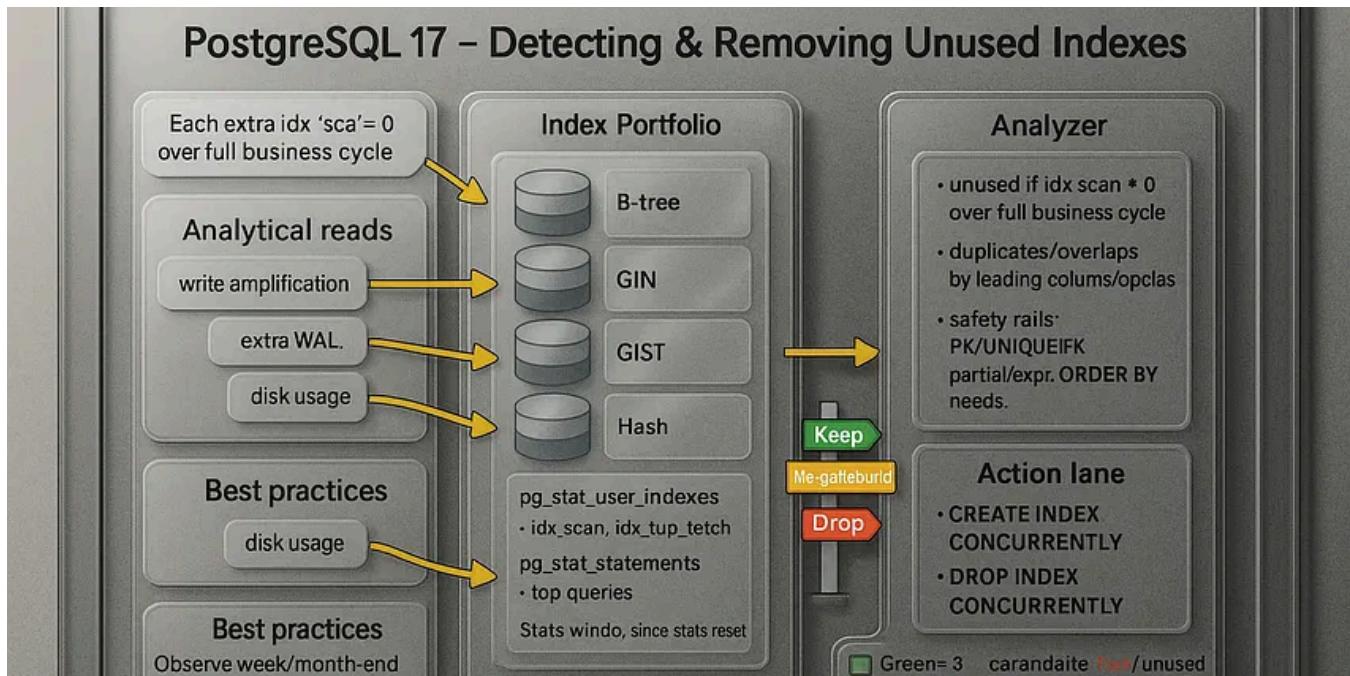
6d ago  5



 Rizqi Mulki

PostgreSQL Index Bloat: The Silent Killer of Performance

How a 10TB database became 3x faster by fixing the invisible problem that plagues 90% of production PostgreSQL deployments



J Jeyaram Ayyalusamy

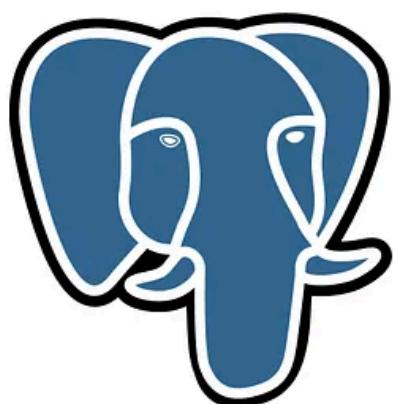
25 - PostgreSQL 17 Performance Tuning: Detecting and Removing Unused Indexes

Indexes are one of the most important tools for query performance in PostgreSQL. But while a missing index slows down queries, an unused or...

Sep 10 40



...



**Beyond Basic
PostgreSQL
Programmable
Objects**

 In Stackademic by bektiaw

Beyond Basics: PostgreSQL Programmable Objects (Automate, Optimize, Control)

Tired of repeating the same SQL logic? Learn how PostgreSQL can do the work for us.

★ Sep 1 ⌘ 68 🗣 1



...

 Rohan

JSONB in PostgreSQL: The Fast Lane to Flexible Data Modeling 🚀

“Why spin up yet another database just to store semi-structured data?”—Every engineer who discovered JSONB

★ Jul 18 ⌘ 12 🗣 1



...



 Thinking Loop

10 Postgres Partitioning Patterns for Internet-Scale Apps

Keep Postgres fast past a billion rows with real patterns and ready-to-run SQL.

 Aug 13  88  2



See more recommendations