

[Open in app ↗](#)

Medium



Get unlimited access to the best of Medium for less than \$1/week. [Become a member](#) X

14 - PostgreSQL 17 Performance Tuning: When Bitmap Scans Beat Index & Seq Scans

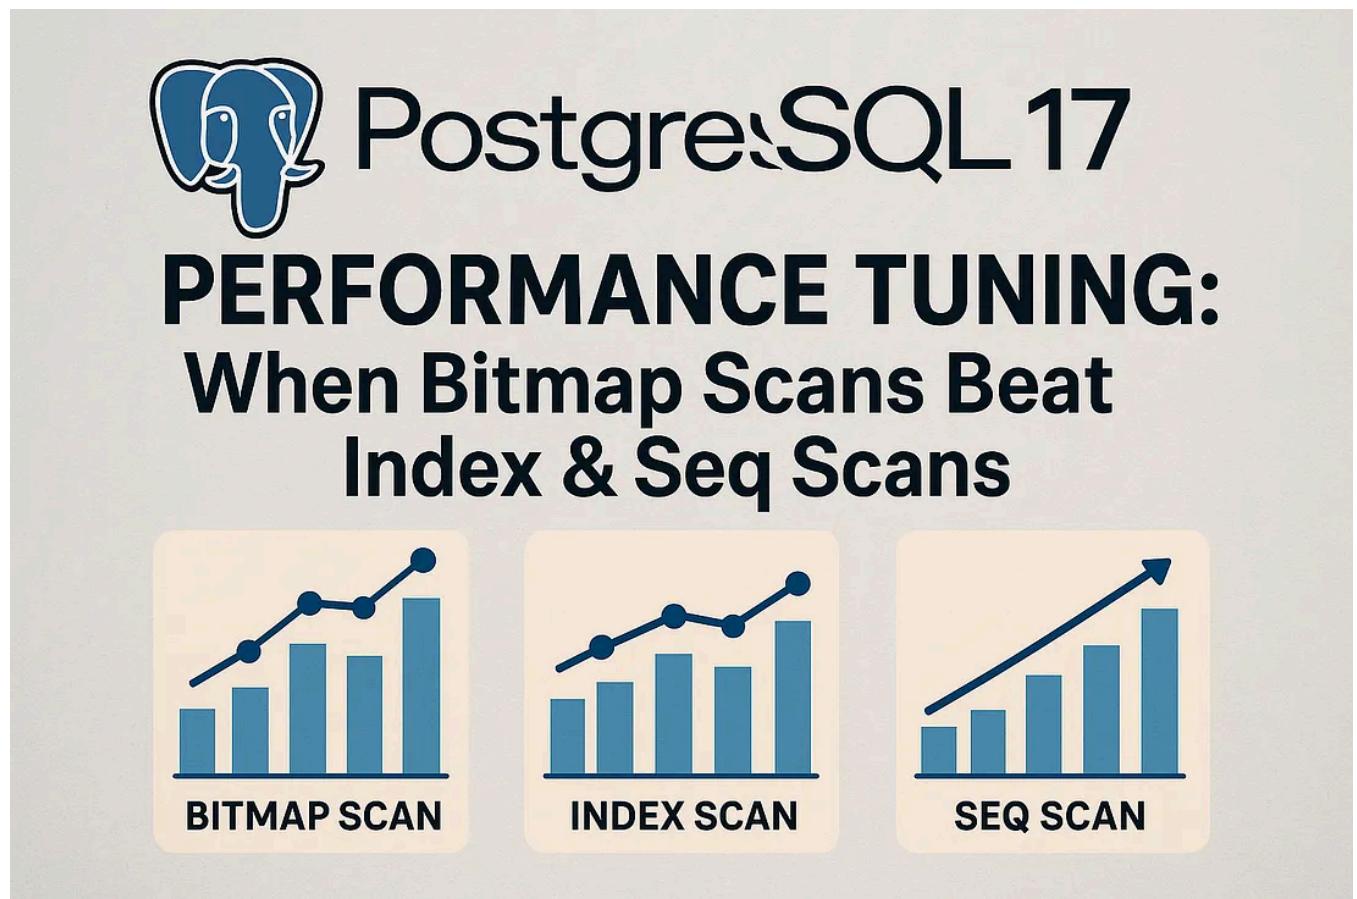
7 min read · Sep 2, 2025

 Jeyaram Ayyalusamy  Following ▼

 Listen

 Share

 More



PostgreSQL chooses among three broad strategies for filtering rows:

- Index Scan when the result set is tiny (needle in a haystack).
- Sequential Scan when most rows match (just read the whole table).

- **Bitmap (Index + Heap) Scan** for the *in-between*: too many rows for a fast Index Scan, too few to justify a full Seq Scan, or when **combining multiple indexes**.

Bitmap scans build an in-memory **bitmap of row locations (TIDs)** from one or more indexes, then visit each heap page **once**, pulling all needed rows per page. This dramatically reduces random I/O compared to a plain Index Scan over many matching rows.

Below is a reproducible demo you can paste into `psql`.

PostgreSQL 17 Performance Tuning: Why Bitmap Heap Scan Can Be Slower Than Seq Scan — and How CLUSTER Fixes It

When tuning queries in PostgreSQL, indexes don't always guarantee faster execution. In fact, sometimes a **Bitmap Heap Scan with an index** can be significantly slower than a simple **Sequential Scan**. In this section, we'll walk through a hands-on test with **10 million rows** and explain the query plans in detail.

Step 1: Creating the `products` table

```
CREATE TABLE products (
    product_id      BIGINT,
    product_name    TEXT,
    category        TEXT,
    price           NUMERIC,
    stock_qty       INT
);
```

This table will hold product information such as name, category, price, and stock quantity.

Step 2: Insert 10M rows

```
INSERT INTO products (product_name, category, price, stock_qty)
SELECT
    'Product_' || g,
    'Category_' || (g % 10),
    (random()*500)::NUMERIC(10,2),
    (random()*100)::INT
FROM generate_series(1, 10000000) AS g;
```

- 10 evenly distributed categories (Category_0 ... Category_9).
- Random prices between 0.00 and 500.00 .
- Stock quantities between 0 and 100 .

Step 3: Analyze the table

```
ANALYZE products;
```

This updates statistics for the planner.

Step 4: Query plan without index (Seq Scan)

```
EXPLAIN ANALYZE SELECT * FROM products WHERE stock_qty = 76;
```

PostgreSQL 17 Performance Tuning: Understanding the Parallel Seq Scan

When analyzing query performance in PostgreSQL, looking at the execution plan is crucial. Let's break down the following plan in detail:

```
postgres=# EXPLAIN ANALYZE SELECT * FROM products WHERE stock_qty = 76;
                                         QUERY PLAN
-----
 Gather  (cost=1000.00..146183.08 rows=97665 width=44) (actual time=3.077..3272
   Workers Planned: 2
   Workers Launched: 2
     -> Parallel Seq Scan on products  (cost=0.00..135416.58 rows=40694 width=44
          Filter: (stock_qty = 76)
          Rows Removed by Filter: 3299940
 Planning Time: 0.046 ms
 Execution Time: 3363.344 ms
(8 rows)
```

- **What it means:**

The `Gather` node collects results from multiple parallel workers. PostgreSQL planned for 2 workers, and both were successfully launched.

- **Why it matters:**

This is the coordination point where PostgreSQL merges all rows found by parallel workers and returns them as a single result set.

2. Parallel Seq Scan on `products`

- **What it means:**

Each worker performed a **Sequential Scan** across a different portion of the `products` table. A sequential scan means PostgreSQL reads rows in physical order directly from the heap.

- **Rows processed:**

- Each worker scanned about **3.3 million rows** (`rows=33393` actually returned per worker).
- The filter `stock_qty = 76` discarded the vast majority of rows, as shown in `Rows Removed by Filter: 3,299,940`.
- **Why it matters:**
Even though every row in the table was touched, parallel workers allowed the job to be divided, making it faster than if a single process had done all the scanning.

3. Costs vs. Actual Times

- **Cost estimates:**
- `cost=1000.00..146183.08` means the planner estimated this query would have a base startup cost of 1000 and a total cost of ~146k.
- **Actual execution:**
- The query started returning rows after ~3 ms.
- The entire scan finished in about **3.27 seconds** (`actual_time=3.077..3272.108`).

👉 The actual time aligns reasonably with planner expectations, which means PostgreSQL chose an appropriate strategy.

4. Planning and Execution Time

- **Planning Time:** `0.046 ms` — PostgreSQL needed almost no time to decide on this plan.
- **Execution Time:** `3363.344 ms` — The real cost of scanning 10 million rows and filtering them.

Why the Sequential Scan Was Efficient

Even though scanning the whole table sounds expensive, it performed well here because:

1. **Parallelism:** Multiple workers shared the workload.
2. **Sequential I/O:** Reading pages in order is much cheaper than jumping around randomly.
3. **Data selectivity:** The condition matched about 1% of rows, but scanning was still competitive because sequential reads are highly optimized.

Key Takeaway:

In PostgreSQL, a sequential scan (especially parallelized) can be faster than an index-based plan when the query touches a large percentage of the table or when rows are spread across many pages.

Step 5: Create index on stock_qty

```
CREATE INDEX idx_products_stock ON products(stock_qty);
```

Step 6: Query plan with index (Bitmap Heap Scan)

```
EXPLAIN ANALYZE SELECT * FROM products WHERE stock_qty = 76;
```

Query Plan Output

```
postgres=# EXPLAIN ANALYZE SELECT * FROM products WHERE stock_qty = 76;
                                         QUERY PLAN
-----
Bitmap Heap Scan on products  (cost=1089.35..89823.31 rows=97667 width=44) (ac
    Recheck Cond: (stock_qty = 76)
    Heap Blocks: exact=58440
    -> Bitmap Index Scan on idx_products_stock  (cost=0.00..1064.94 rows=97667
        Index Cond: (stock_qty = 76)
Planning Time: 4.174 ms
Execution Time: 42559.194 ms
(7 rows)

postgres=#

```

Explanation:

- **Bitmap Index Scan:** Quickly finds ~100k row pointers (TIDs) matching `stock_qty = 76`.
- **Bitmap Heap Scan:** Fetches those rows from the heap but must touch **58,440 pages** scattered across the table.
- **Execution Time (~42.6s):** Much slower than the Seq Scan because of **random I/O**.

👉 Lesson: An index doesn't help if matching rows are scattered everywhere.

Step 7: Cluster the Table

```
CLUSTER products USING idx_products_stock;
```

When you run the **CLUSTER** command in PostgreSQL, the database physically **rewrites the table** so that its rows are stored on disk in the same order as the specified index (in this case, `idx_products_stock`).

In our example, rows are reordered so that all entries with the same `stock_qty` value are stored **contiguously** on disk.

Why Does This Improve Performance?

1. Data Locality

- Normally, even if you have an index, the corresponding rows may be scattered across different disk pages.
- After clustering, rows with the same `stock_qty` are stored together.
- This means fewer page reads are needed when scanning rows that share the same value, reducing I/O.

2. Faster Sequential Access

- PostgreSQL can fetch grouped rows in **sequential blocks** instead of random disk access.
- Sequential reads are much faster than random reads, especially on large tables.

3. Improved Index Efficiency

- The index `idx_products_stock` now directly maps to data that is already in physical order.
- This reduces the number of heap fetches, since PostgreSQL doesn't have to jump around to find matching tuples.

4. Better Cache Utilization

- When rows are stored together, PostgreSQL can load an entire block of related rows into memory.
- This improves buffer cache hit ratio, making repeated queries faster.

5. Query Patterns Benefit Most

- Queries like:

```
SELECT * FROM products WHERE stock_qty = 76;
```

- or

```
SELECT * FROM products WHERE stock_qty BETWEEN 70 AND 80;
```

- run significantly faster because the database engine reads fewer scattered pages.

Key Considerations

- **One-time Reorganization:**

CLUSTER is not automatic. If new rows are inserted or updated, the physical order can get out of sync again.

- **Use Periodically:**

For frequently queried tables where data distribution changes over time, consider re-clustering during maintenance windows.

- **Table Locking:**

CLUSTER requires an **exclusive lock**, meaning the table is not available for reads/writes during the operation.

- ✓ In summary, clustering improves performance by **reducing I/O, improving sequential access, and aligning physical storage with logical query patterns**. This leads to **faster lookups and more efficient scans** for queries that filter or group by the clustered column.

Step 8: Query plan after clustering

```
EXPLAIN ANALYZE SELECT * FROM products WHERE stock_qty = 76;
```

Query Plan Output

```
postgres=# EXPLAIN ANALYZE SELECT * FROM products WHERE stock_qty = 76;
                                         QUERY PLAN
-----
Bitmap Heap Scan on products  (cost=1089.35..89823.31 rows=97667 width=44) (ac
Recheck Cond: (stock_qty = 76)
Heap Blocks: exact=835
-> Bitmap Index Scan on idx_products_stock  (cost=0.00..1064.94 rows=97667
      Index Cond: (stock_qty = 76)
Planning Time: 2.360 ms
Execution Time: 430.219 ms
(7 rows)

postgres=#

```

Explanation:

- The Bitmap Index Scan still identifies 100k rows.
- But this time, only 835 heap pages need to be fetched (instead of 58k).
- **Execution Time (~0.43s):** Dramatic improvement because rows are now physically contiguous.

Step 9: Analyze after clustering

```
ANALYZE products;
```

Step 10: Query plan after Analyze (Index Scan)

```
EXPLAIN ANALYZE SELECT * FROM products WHERE stock_qty = 76;
```

Query Plan Output

```
postgres=# EXPLAIN ANALYZE SELECT * FROM products WHERE stock_qty = 76;
                                         QUERY PLAN
-----
 Index Scan using idx_products_stock on products  (cost=0.43..3170.28 rows=1083
   Index Cond: (stock_qty = 76)
 Planning Time: 0.148 ms
 Execution Time: 128.913 ms
 (4 rows)

postgres=#

```

Explanation:

- PostgreSQL switches to a **direct Index Scan**.
- Since rows are tightly packed and statistics are updated, the planner knows the index scan is optimal.
- **Execution Time (~0.13s)**: Fastest plan so far.

Why Bitmap Heap Scan Was Slower Than Seq Scan

- **Without clustering**: The Bitmap Heap Scan had to jump to 58,440 scattered pages → costly random I/O.
- **Seq Scan** read the table linearly with parallel workers → faster despite reading more data.

Why CLUSTER Improved Execution Time

- Physically orders data according to the index.
- Rows with `stock_qty = 76` became adjacent → only 835 pages to fetch.
- Execution time dropped from **42.6s** → **0.43s** → **0.13s** after Analyze.

Final Takeaways

- An index doesn't always guarantee better performance.
- Bitmap Heap Scans can be very slow if data is **scattered across the heap**.
- Sequential scans may outperform them because of efficient linear I/O and parallelism.
- `CLUSTER` dramatically improves performance by grouping related rows together.
- After clustering and analyzing, PostgreSQL can switch to the fastest **Index Scan**.

👉 **Lesson for tuning:** Always look at the query plan, not just the presence of indexes. Data organization matters just as much as indexes.

📢 Stay Updated with Daily PostgreSQL & Cloud Tips!

If you've been finding my blog posts helpful and want to stay ahead with daily insights on PostgreSQL, Cloud Infrastructure, Performance Tuning, and DBA Best Practices — I invite you to subscribe to my Medium account.

🔔 Subscribe here 👉 <https://medium.com/@jramcloud1/subscribe>

Your support means a lot — and you'll never miss a practical guide again!

🔗 Let's Connect!

If you enjoyed this post or would like to connect professionally, feel free to reach out to me on LinkedIn:

👉 [Jeyaram Ayyalusamy](#)

I regularly share content on **PostgreSQL, database administration, cloud technologies, and data engineering**. Always happy to connect, collaborate, and discuss ideas!

Postgresql

Oracle

AWS

Open Source

MySQL

A circular profile picture of a man with dark hair and a beard, wearing a dark shirt.

Following ▾

Written by Jeyaram Ayyalusamy

158 followers · 2 following

Oracle DBA | AWS Certified | 18+ yrs in DB Admin, RAC, GoldenGate, RDS, MySQL, PostgreSQL | Author of MySQL & RDS books | Cloud & Performance Expert

No responses yet



Gvadakte

What are your thoughts?



More from Jeyaram Ayyalusamy

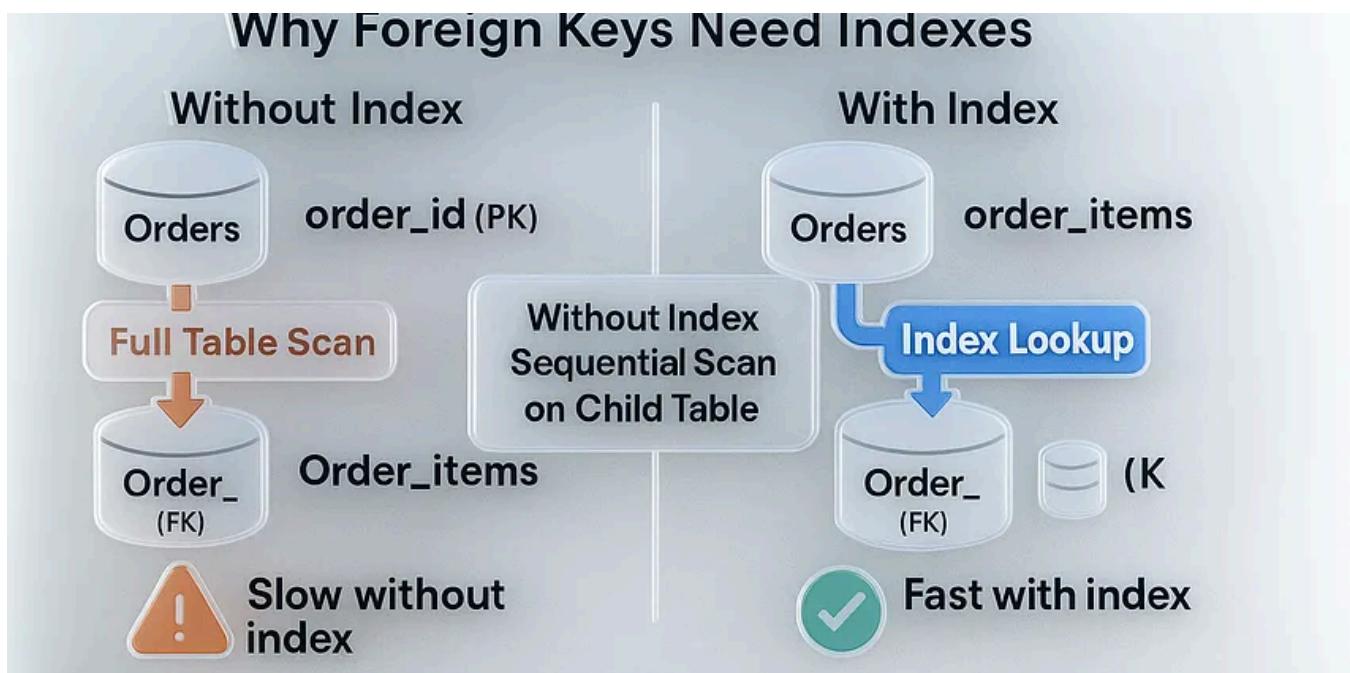
The screenshot shows the AWS EC2 Instances page. The browser tab bar includes 'us-west-2' (active), 'Launch an instance | EC2 | us-west-2', 'Instances | EC2 | us-east-1', and a '+' button. The URL is '1.console.aws.amazon.com/ec2/home?region=us-east-1#Instances'. The main content area is titled 'Instances Info' with a search bar and filters for 'Name', 'Instance ID', 'Instance state', 'Instance type', 'Status check', 'Alarm status', 'Availability Zone', 'Public IPv4 DNS', 'Public IPv4 IP', 'Elastic IP', and 'IPs'. A message states 'No instances' and 'You do not have any instances in this region'. A blue 'Launch instances' button is visible. On the left, a sidebar says 'Select an instance'. At the bottom right, it says '© 2025, Amazon Web Services, Inc. or its affiliates.'

J Jeyaram Ayyalusamy

Upgrading PostgreSQL from Version 16 to Version 17 Using pg_upgrade on a Linux Server AWS EC2...

A Complete Step-by-Step Guide to Installing PostgreSQL 16 on a Linux Server (With Explanations)

Aug 4 40



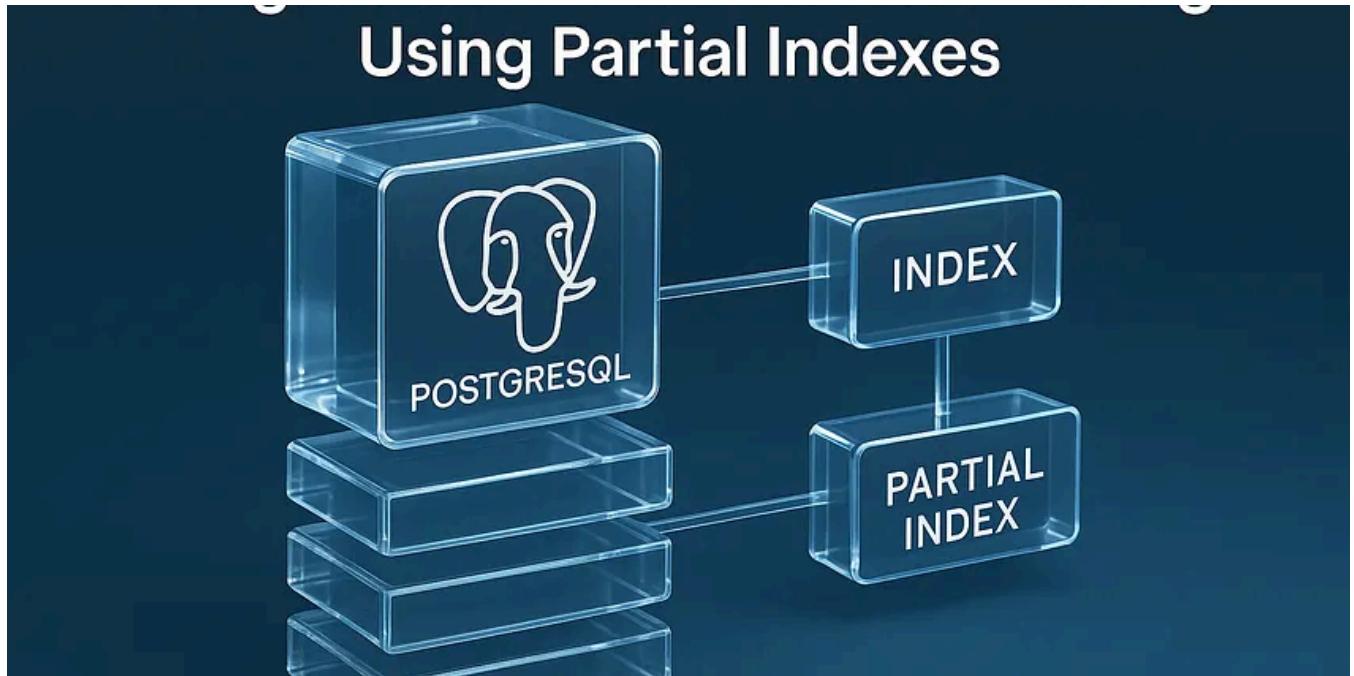
J Jeyaram Ayyalusamy 

16 -PostgreSQL 17 Performance Tuning: Why Foreign Keys Need Indexes

When designing relational databases, parent-child relationships are common. A classic case is the Orders table (parent) and the Items table...

Sep 3  3  2



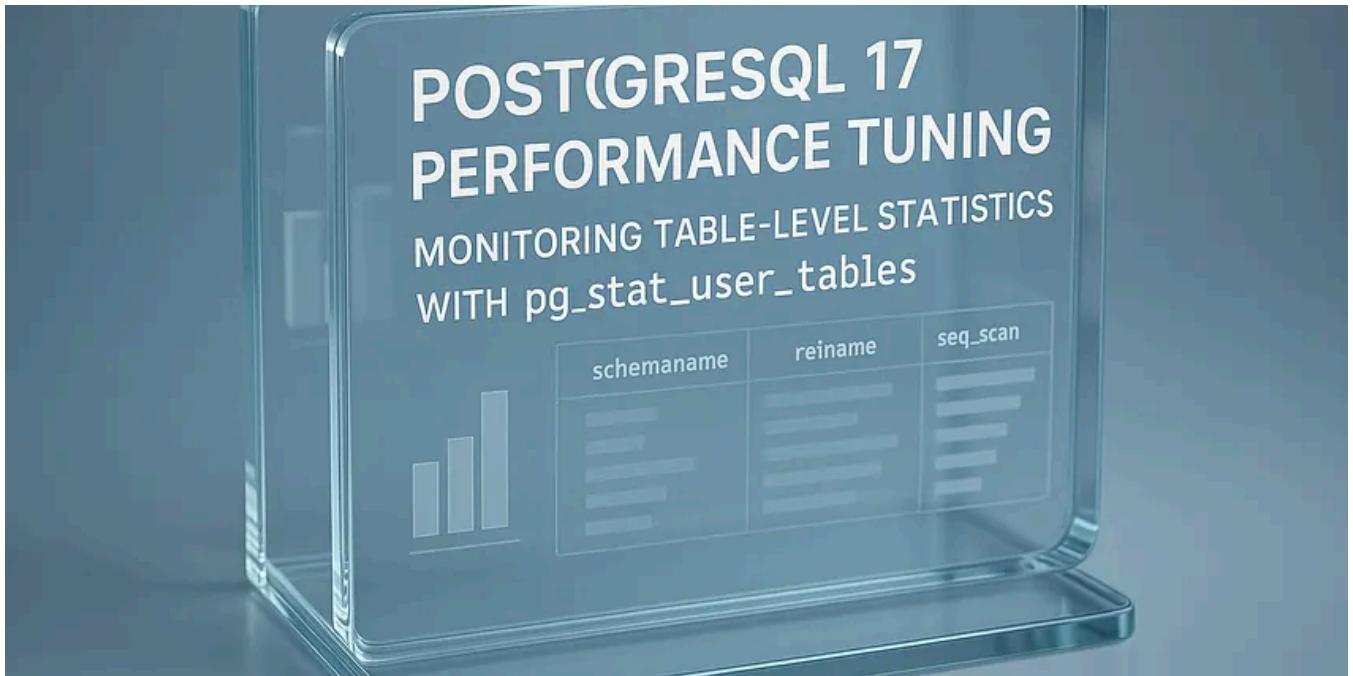
J Jeyaram Ayyalusamy 

17 - PostgreSQl 17 Performance Tuning: Using Partial Indexes

Even though we usually create indexes on all rows of a table, PostgreSQL also allows you to create an index on just a subset of records...

Sep 4  3



J Jeyaram Ayyalusamy

24 - PostgreSQL 17 Performance Tuning: Monitoring Table-Level Statistics with pg_stat_user_tables

When tuning PostgreSQL, one of the most important steps is to observe table-level statistics. You cannot optimize what you cannot measure...

Sep 7 19 1



See all from Jeyaram Ayyalusamy

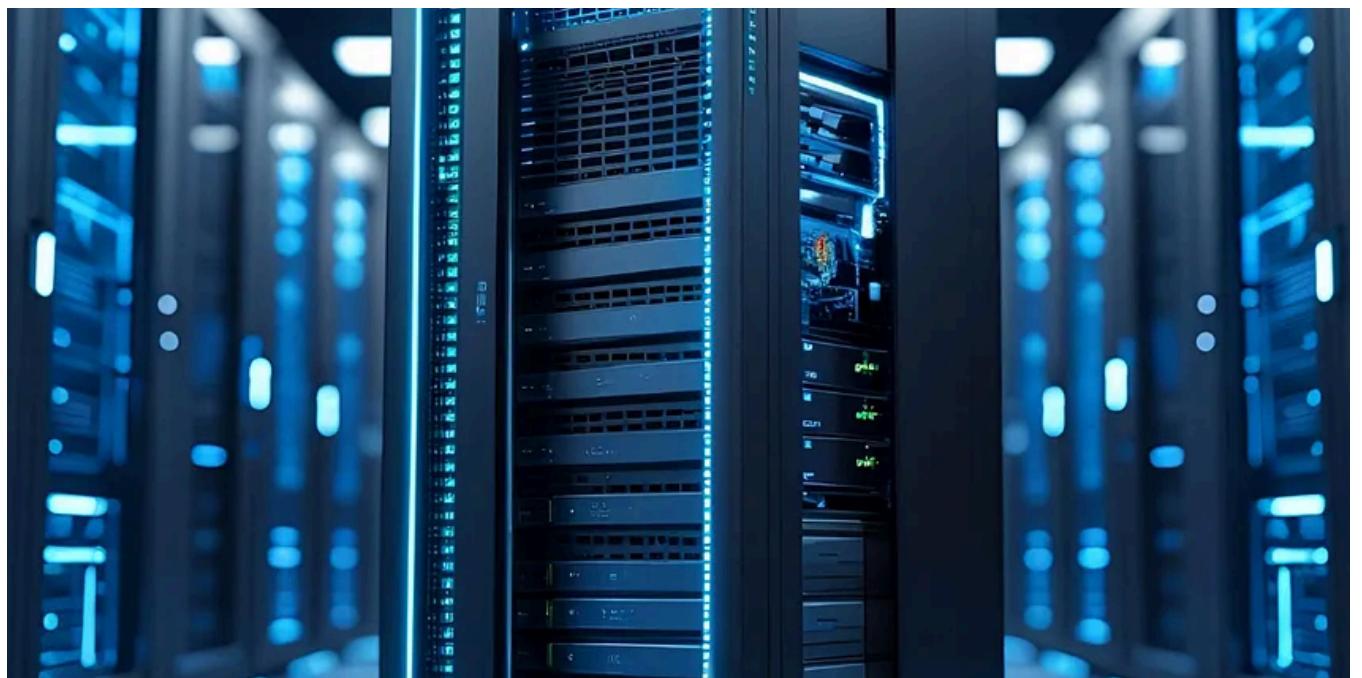
Recommended from Medium

 Tomasz Gintowt

Security in PostgreSQL

PostgreSQL is a powerful open-source database. Security is very important when working with sensitive data like passwords, customer...

6d ago  5



 Rizqi Mulki

PostgreSQL Index Bloat: The Silent Killer of Performance

How a 10TB database became 3x faster by fixing the invisible problem that plagues 90% of production PostgreSQL deployments

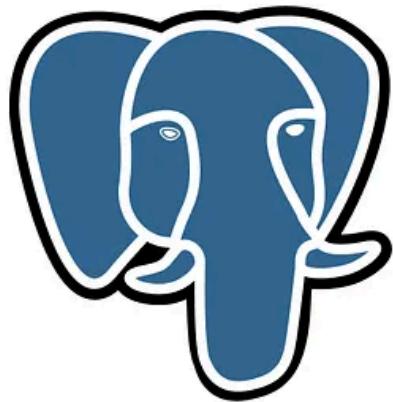
★ Sep 15

👏 11

💬 1



...



Beyond Basic PostgreSQL Programmable Objects



In Stackademic by bektiaw

Beyond Basics: PostgreSQL Programmable Objects (Automate, Optimize, Control)

Tired of repeating the same SQL logic? Learn how PostgreSQL can do the work for us.



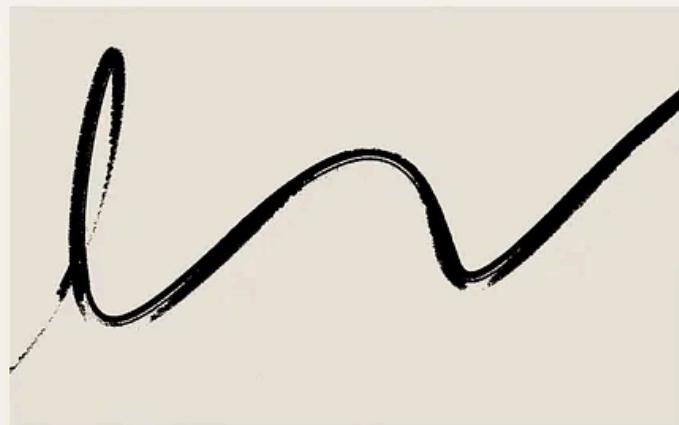
Sep 1

👏 68

💬 1



...



Rohan

JSONB in PostgreSQL: The Fast Lane to Flexible Data Modeling 🚀

“Why spin up yet another database just to store semi-structured data?”—Every engineer who discovered JSONB

◆ Jul 18 ⌘ 12 ⏰ 1

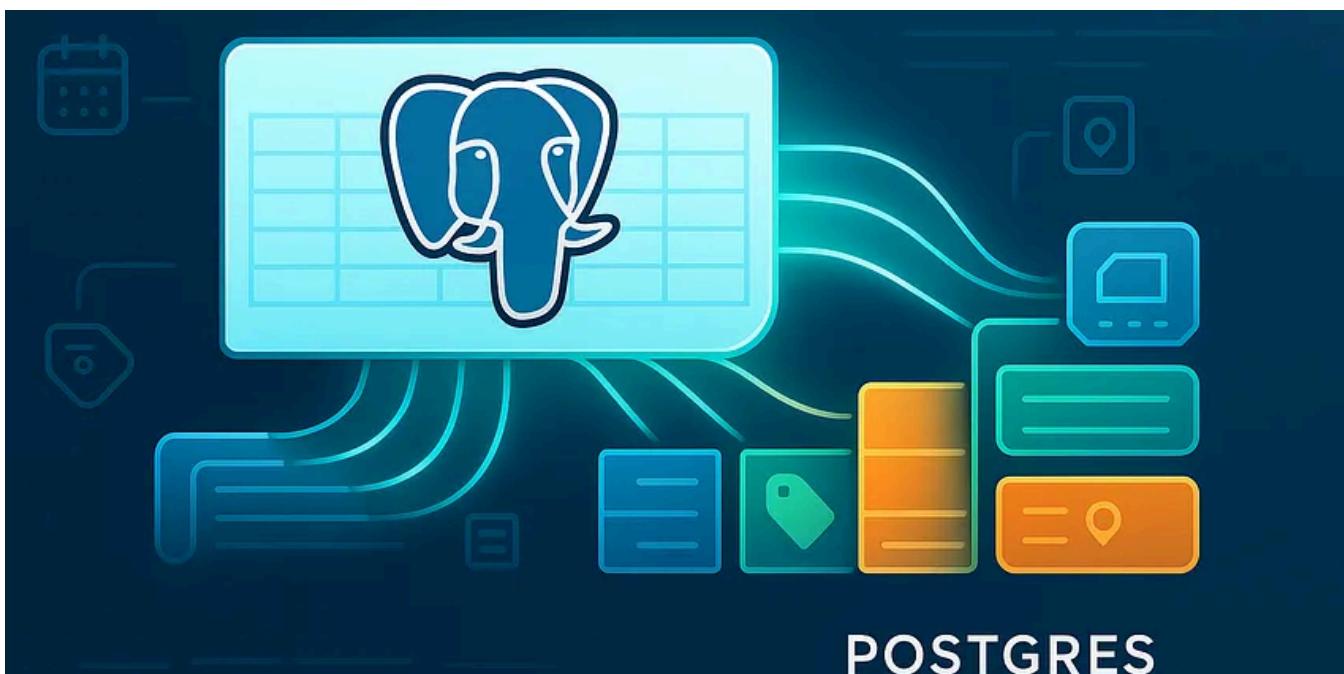


 Thread Whisperer

Postgres 18 Arrives: Async I/O You Should Turn On First

Turn disk waits into throughput with a few safe switches

◆ Sep 15 ⌘ 77





Thinking Loop

10 Postgres Partitioning Patterns for Internet-Scale Apps

Keep Postgres fast past a billion rows with real patterns and ready-to-run SQL.

Aug 13

88

2



...

See more recommendations