

[Open in app](#)

Medium



Get unlimited access to the best of Medium for less than \$1/week. [Become a member](#)



16 -PostgreSQL 17 Performance Tuning: Why Foreign Keys Need Indexes

5 min read · Sep 3, 2025

J Jeyaram Ayyalusamy Following

Listen

Share

More

PostgreSQL 17 Performance Tuning: Why Foreign Keys Need Indexes

Without Index



order_id (PK)

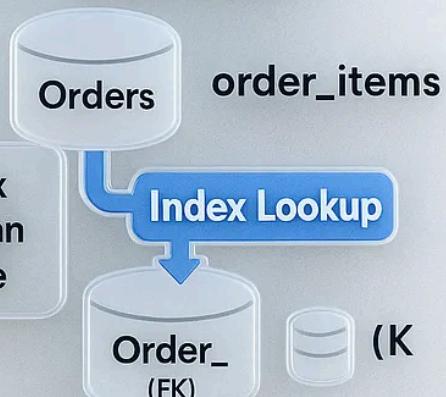
Without Index
Sequential Scan
on Child Table

Order_items



Slow without
index

With Index



order_items

Index Lookup

Order_items
(FK)



Fast with index



PostgreSQL Always index foreign keys for efficient deletes, updates, and

When designing relational databases, parent-child relationships are common. A classic case is the **Orders** table (parent) and the **Items** table (child). Each order can have many items, so the child table typically has far more rows than the parent.

👉 In such cases, it's always a good practice to create an index on the **foreign key column** in the child table.

There are two main reasons:

1. **Faster joins** → Queries joining parent and child tables almost always join on the foreign key. An index dramatically speeds this up.
2. **Faster parent changes** → When deleting or updating parent rows, PostgreSQL must check child rows for constraint enforcement. Without an index, these checks can be painfully slow.

Let's build a demo to see the impact.

Step 1: Create Parent and Child Tables

```
-- Parent table: Orders
CREATE TABLE orders (
    order_no    SERIAL PRIMARY KEY,
    order_date  DATE DEFAULT CURRENT_DATE
);
```

```
postgres=# -- Parent table: Orders
CREATE TABLE orders (
    order_no    SERIAL PRIMARY KEY,
    order_date  DATE DEFAULT CURRENT_DATE
);
CREATE TABLE
postgres=#
```

```
-- Child table: Items
CREATE TABLE items (
```

```
item_no      SERIAL PRIMARY KEY,  
order_no     INT,  
product_name TEXT,  
description  TEXT,  
created_at   TIMESTAMP DEFAULT now(),  
CONSTRAINT fk_order FOREIGN KEY (order_no)  
    REFERENCES orders(order_no)  
    ON UPDATE CASCADE  
    ON DELETE CASCADE  
);
```

```
postgres=# -- Child table: Items  
CREATE TABLE items (  
    item_no      SERIAL PRIMARY KEY,  
    order_no     INT,  
    product_name TEXT,  
    description  TEXT,  
    created_at   TIMESTAMP DEFAULT now(),  
    CONSTRAINT fk_order FOREIGN KEY (order_no)  
        REFERENCES orders(order_no)  
        ON UPDATE CASCADE  
        ON DELETE CASCADE  
);  
CREATE TABLE  
postgres=#
```

Step 2: Populate Tables with Data

We'll insert 1 million rows into `orders` and 3 million rows into `items` (4 items per order).

```
-- Insert 1M orders  
INSERT INTO orders (order_no)  
SELECT g FROM generate_series(1, 1000000) g;
```

```
postgres=# -- Insert 1M orders
INSERT INTO orders (order_no)
SELECT g FROM generate_series(1, 1000000) g;
INSERT 0 1000000
postgres=#
```

```
-- Insert 3M items (4 per order)
INSERT INTO items (order_no, product_name, description)
SELECT (g/4)+1, 'Product_' || g, 'Desc_' || g
FROM generate_series(1, 3000000) g;
```

```
postgres=# -- Insert 3M items (4 per order)
INSERT INTO items (order_no, product_name, description)
SELECT (g/4)+1, 'Product_' || g, 'Desc_' || g
FROM generate_series(1, 3000000) g;
INSERT 0 3000000
postgres=#
```

👉 Now:

- `orders` has 1M rows.
- `items` has 3M rows linked by the foreign key.

Step 3: Query Without Index on Foreign Key

Let's try a join query:

```
\timing
SELECT *
FROM orders o
JOIN items i ON o.order_no = i.order_no
WHERE o.order_no = 120;
```

```
postgres=# \timing
Timing is on.
postgres=#
postgres=# SELECT *
FROM orders o
JOIN items i ON o.order_no = i.order_no
WHERE o.order_no = 120;
order_no | order_date | item_no | order_no | product_name | description | quantity
-----+-----+-----+-----+-----+-----+-----+
 120 | 2025-09-02 | 12000476 |      120 | Product_476 | Desc_476 | 202
 120 | 2025-09-02 | 12000477 |      120 | Product_477 | Desc_477 | 202
 120 | 2025-09-02 | 12000478 |      120 | Product_478 | Desc_478 | 202
 120 | 2025-09-02 | 12000479 |      120 | Product_479 | Desc_479 | 202
(4 rows)

Time: 194.021 ms
postgres=#
```

Result:

- PostgreSQL performs a **sequential scan** on the `items` table (4 million rows) to find just 4 matches.
- Execution is slow.

Check the plan:

```
EXPLAIN ANALYZE
SELECT *
FROM orders o
JOIN items i ON o.order_no = i.order_no
WHERE o.order_no = 120;
```

```
postgres=# EXPLAIN ANALYZE
SELECT *
FROM orders o
JOIN items i ON o.order_no = i.order_no
WHERE o.order_no = 120;
```

QUERY PLAN

```
Nested Loop  (cost=1000.42..44661.88 rows=4 width=51) (actual time=0.249..198.
->  Index Scan using orders_pkey on orders o  (cost=0.42..8.44 rows=1 width=
      Index Cond: (order_no = 120)
->  Gather  (cost=1000.00..44653.40 rows=4 width=43) (actual time=0.230..198
      Workers Planned: 2
      Workers Launched: 2
      ->  Parallel Seq Scan on items i  (cost=0.00..43653.00 rows=2 width=43
          Filter: (order_no = 120)
          Rows Removed by Filter: 999999
Planning Time: 0.105 ms
Execution Time: 198.853 ms
(11 rows)

Time: 200.196 ms
postgres=#

```

👉 You'll see a **Seq Scan** on **items**, which is inefficient.

Step 4: Add Index on the Foreign Key Column

Now add an index:

```
CREATE INDEX idx_items_order_no ON items(order_no);
```

```
postgres=# CREATE INDEX idx_items_order_no ON items(order_no);
CREATE INDEX
postgres=#
```

Step 5: Query With Index

Run the same query again:

```
EXPLAIN ANALYZE
SELECT *
FROM orders o
JOIN items i ON o.order_no = i.order_no
WHERE o.order_no = 120;
```

```
postgres=# EXPLAIN ANALYZE
SELECT *
FROM orders o
JOIN items i ON o.order_no = i.order_no
WHERE o.order_no = 120;
```

QUERY PLAN

```
Nested Loop  (cost=0.85..16.98 rows=4 width=51) (actual time=0.030..0.041 rows
->  Index Scan using orders_pkey on orders o  (cost=0.42..8.44 rows=1 width=
    Index Cond: (order_no = 120)
->  Index Scan using idx_items_order_no on items i  (cost=0.43..8.50 rows=4
```

```
Index Cond: (order_no = 120)
```

```
Planning Time: 0.091 ms
```

```
Execution Time: 0.060 ms
```

```
(7 rows)
```

```
Time: 0.442 ms
```

```
postgres=#
```

Result:

- PostgreSQL switches to an **Index Scan on items**.
- The query time drops dramatically.

👉 Instead of scanning 3M rows, it jumps directly to the 4 matching rows.



Step 6: Parent Delete Performance

Indexes also help when deleting from the parent:

```
DELETE FROM orders WHERE order_no = 500;
```

```
postgres=# DELETE FROM orders WHERE order_no = 500;
```

```
DELETE 1
```

```
Time: 4.355 ms
```

```
postgres=#
```

- Without an index on `items.order_no`, PostgreSQL must scan the entire `items` table to check for child rows.
- With the index, it can instantly find and remove matching rows.

Key Takeaways

- Always index foreign key columns in child tables.
 - Benefits:
 - **Joins are faster** (avoid sequential scans).
 - **Parent updates/deletes are efficient** (quick referential checks).
 - Without indexes, queries and deletes on large parent-child datasets can degrade to **full-table scans**, crippling performance.
- In PostgreSQL 17, just as in earlier versions, indexing foreign keys is one of the simplest yet most impactful tuning steps for real-world workloads.

📢 Stay Updated with Daily PostgreSQL & Cloud Tips!

If you've been finding my blog posts helpful and want to stay ahead with daily insights on PostgreSQL, Cloud Infrastructure, Performance Tuning, and DBA Best Practices — I invite you to subscribe to my Medium account.

🔔 [Subscribe here](https://medium.com/@jramcloud1/subscribe) 👉 <https://medium.com/@jramcloud1/subscribe>

Your support means a lot — and you'll never miss a practical guide again!

🔗 **Let's Connect!**

If you enjoyed this post or would like to connect professionally, feel free to reach out to me on LinkedIn:

👉 [Jeyaram Ayyalusamy](#)

I regularly share content on **PostgreSQL, database administration, cloud technologies, and data engineering**. Always happy to connect, collaborate, and discuss ideas!

Postgresql

MySQL

AWS

Oracle

Open Source



Following

Written by Jeyaram Ayyalusamy

158 followers · 2 following

Oracle DBA | AWS Certified | 18+ yrs in DB Admin, RAC, GoldenGate, RDS, MySQL, PostgreSQL | Author of MySQL & RDS books | Cloud & Performance Expert

Responses (2)



Gvadakte



What are your thoughts?



Mishra Ram

Sep 7

...

Insightful



1

[Reply](#)



Awsrmzn

Sep 3 (edited)

...

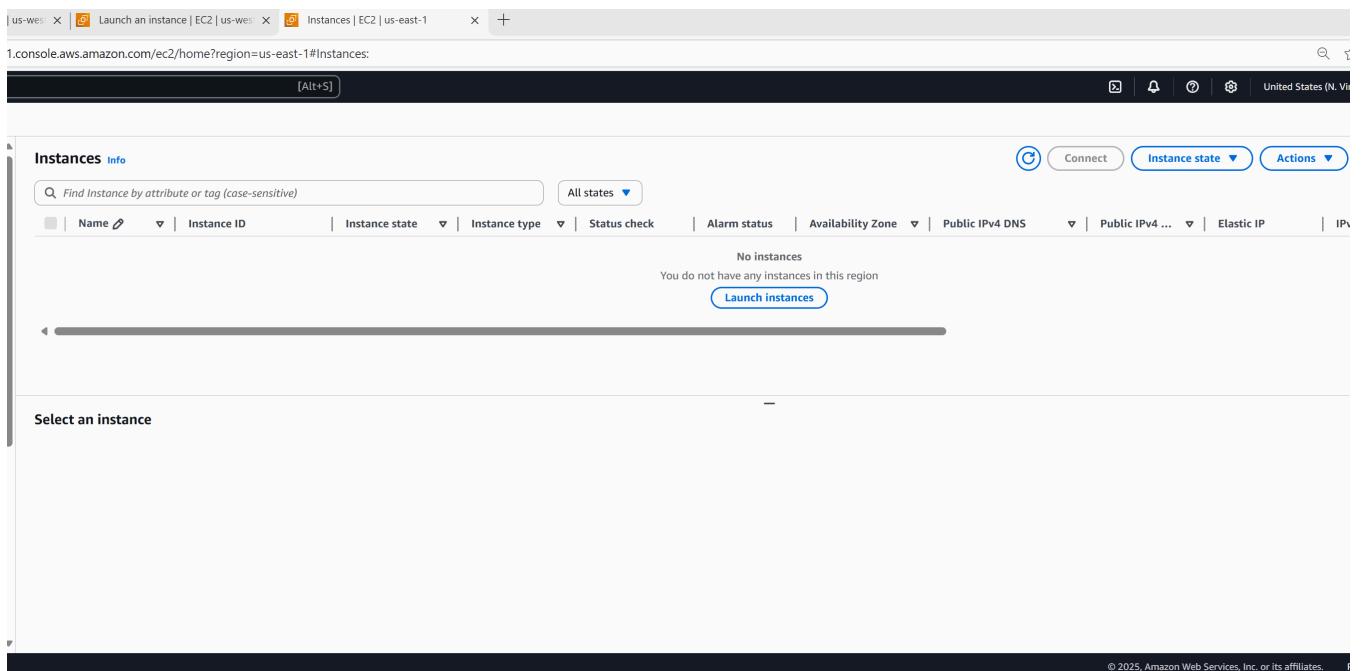
very insightful



1

[Reply](#)

More from Jeyaram Ayyalusamy



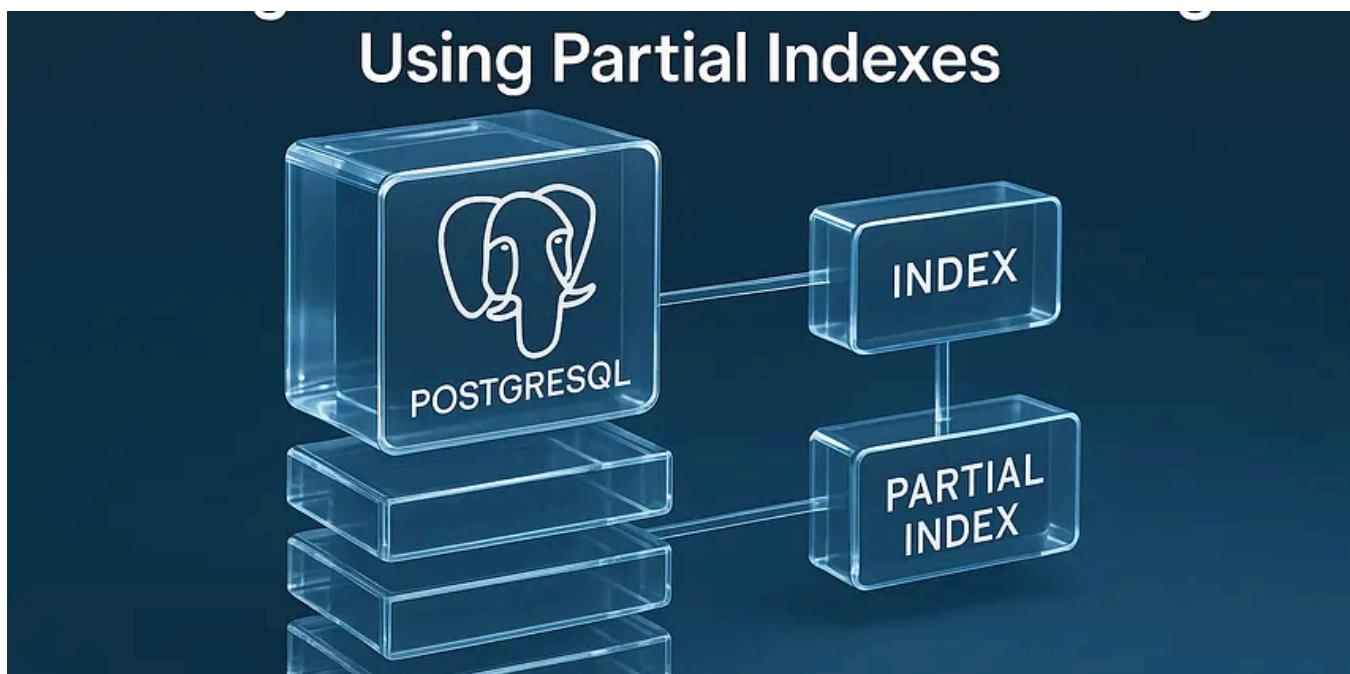
The screenshot shows the AWS EC2 Instances page. The browser tab bar includes 'us-west-2' (active), 'Launch an instance | EC2', 'us-west-2', 'Instances | EC2', and 'us-east-1'. The main content area has a header 'Instances Info' with a search bar and filters for 'Name', 'Instance ID', 'Instance state', 'Instance type', 'Status check', 'Alarm status', 'Availability Zone', 'Public IPv4 DNS', 'Public IPv4 IP', 'Elastic IP', and 'IPs'. Below this, a message says 'No instances' and 'You do not have any instances in this region'. A blue 'Launch instances' button is visible. On the left, a sidebar says 'Select an instance'. At the bottom right, there's a copyright notice: '© 2025, Amazon Web Services, Inc. or its affiliates.'

J Jeyaram Ayyalusamy 

Upgrading PostgreSQL from Version 16 to Version 17 Using pg_upgrade on a Linux Server AWS EC2...

 A Complete Step-by-Step Guide to Installing PostgreSQL 16 on a Linux Server (With Explanations)

Aug 4  40



J Jeyaram Ayyalusamy 

17 - PostgreSQL 17 Performance Tuning: Using Partial Indexes

Even though we usually create indexes on all rows of a table, PostgreSQL also allows you to create an index on just a subset of records...

Sep 4  3



...



J Jeyaram Ayyalusamy 

24 - PostgreSQL 17 Performance Tuning: Monitoring Table-Level Statistics with pg_stat_user_tables

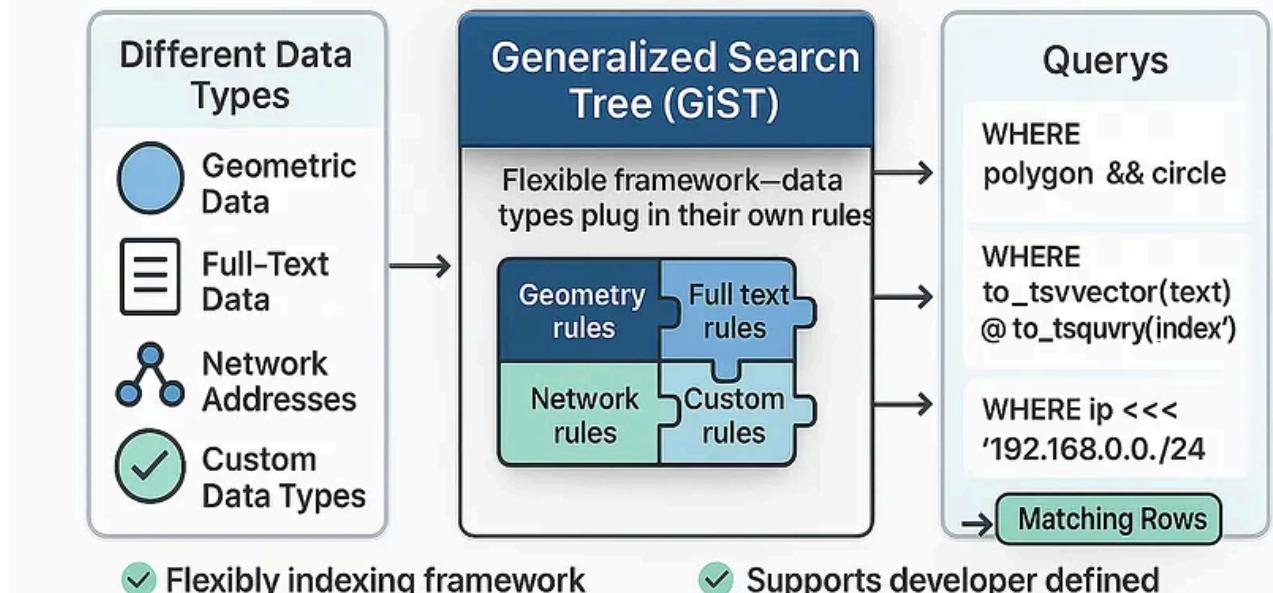
When tuning PostgreSQL, one of the most important steps is to observe table-level statistics. You cannot optimize what you cannot measure...

Sep 7  19  1



...

GiST (Generalized Search Tree)



J Jeyaram Ayyalusamy 

21—PostgreSQL 17 Performance Tuning: GiST (Generalized Search Tree)

GiST (Generalized Search Tree) in PostgreSQL is a flexible indexing framework that allows developers to build indexes for many different...

Sep 5  1



...

See all from Jeyaram Ayyalusamy

Recommended from Medium

 Rizqi Mulki

PostgreSQL Index Bloat: The Silent Killer of Performance

How a 10TB database became 3x faster by fixing the invisible problem that plagues 90% of production PostgreSQL deployments

★ Sep 15 11 1



#PostgreSQL

security

TOMASZ GINTOWT

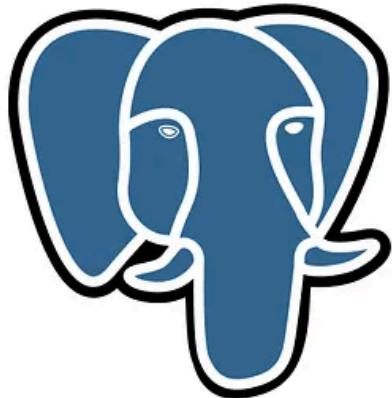
 Tomasz Gintowt

Security in PostgreSQL

PostgreSQL is a powerful open-source database. Security is very important when working with sensitive data like passwords, customer...

6d ago 5

...



Beyond Basic PostgreSQL Programmable Objects

In Stackademic by bektiaw

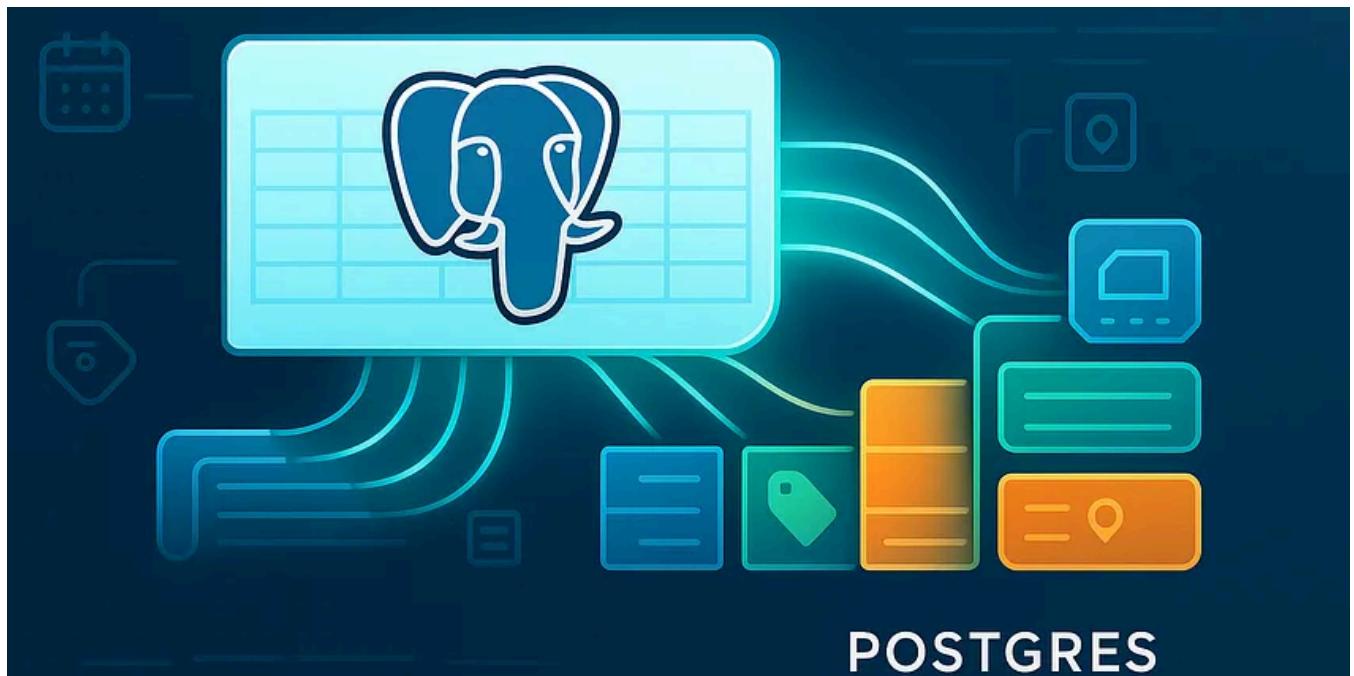
Beyond Basics: PostgreSQL Programmable Objects (Automate, Optimize, Control)

Tired of repeating the same SQL logic? Learn how PostgreSQL can do the work for us.

★ Sep 1 68 1



...

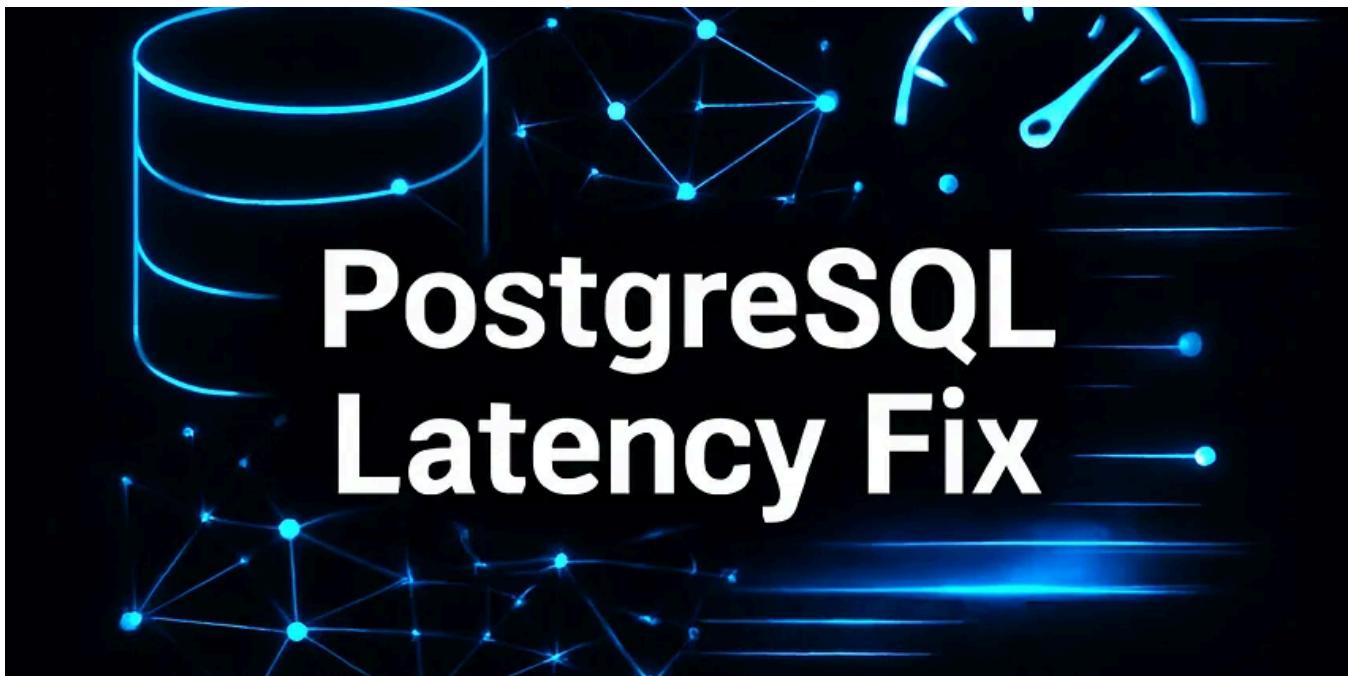


Thinking Loop

10 Postgres Partitioning Patterns for Internet-Scale Apps

Keep Postgres fast past a billion rows with real patterns and ready-to-run SQL.

Aug 13 88 2



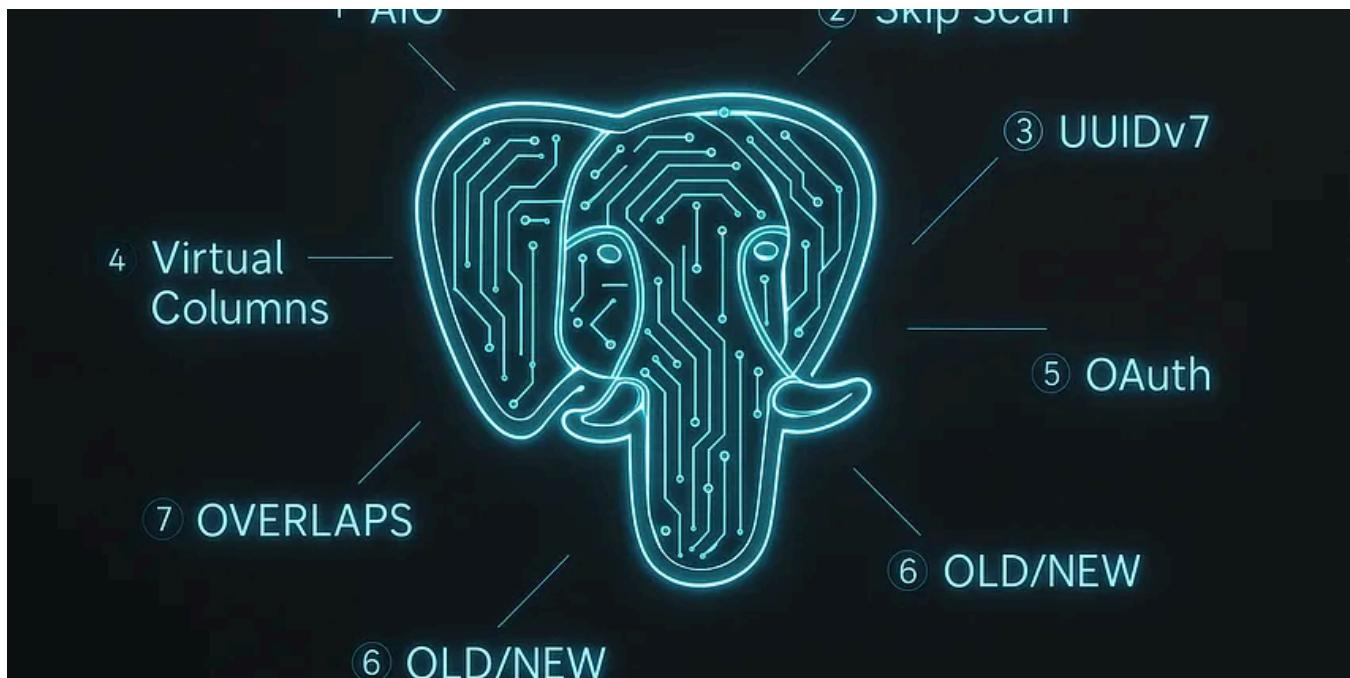
Chronicles

We Lost Users Because of PostgreSQL Latency—The Fix Wasn't More Hardware

We thought our PostgreSQL setup was rock solid. Queries were running, dashboards were loading, and everything “looked fine” in staging.

Aug 9 18





 Thread Whisperer

Postgres 18 Arrives: 7 Features That Will Change Your Stack

A field-tested teardown of the Postgres 18 release for busy builders

★ Sep 13 ⌐ 307 💬 8



See more recommendations