

[Open in app ↗](#)

Medium



Get unlimited access to the best of Medium for less than \$1/week. [Become a member](#) X

25 - PostgreSQL 17 Performance Tuning: Detecting and Removing Unused Indexes

5 min read · Sep 10, 2025



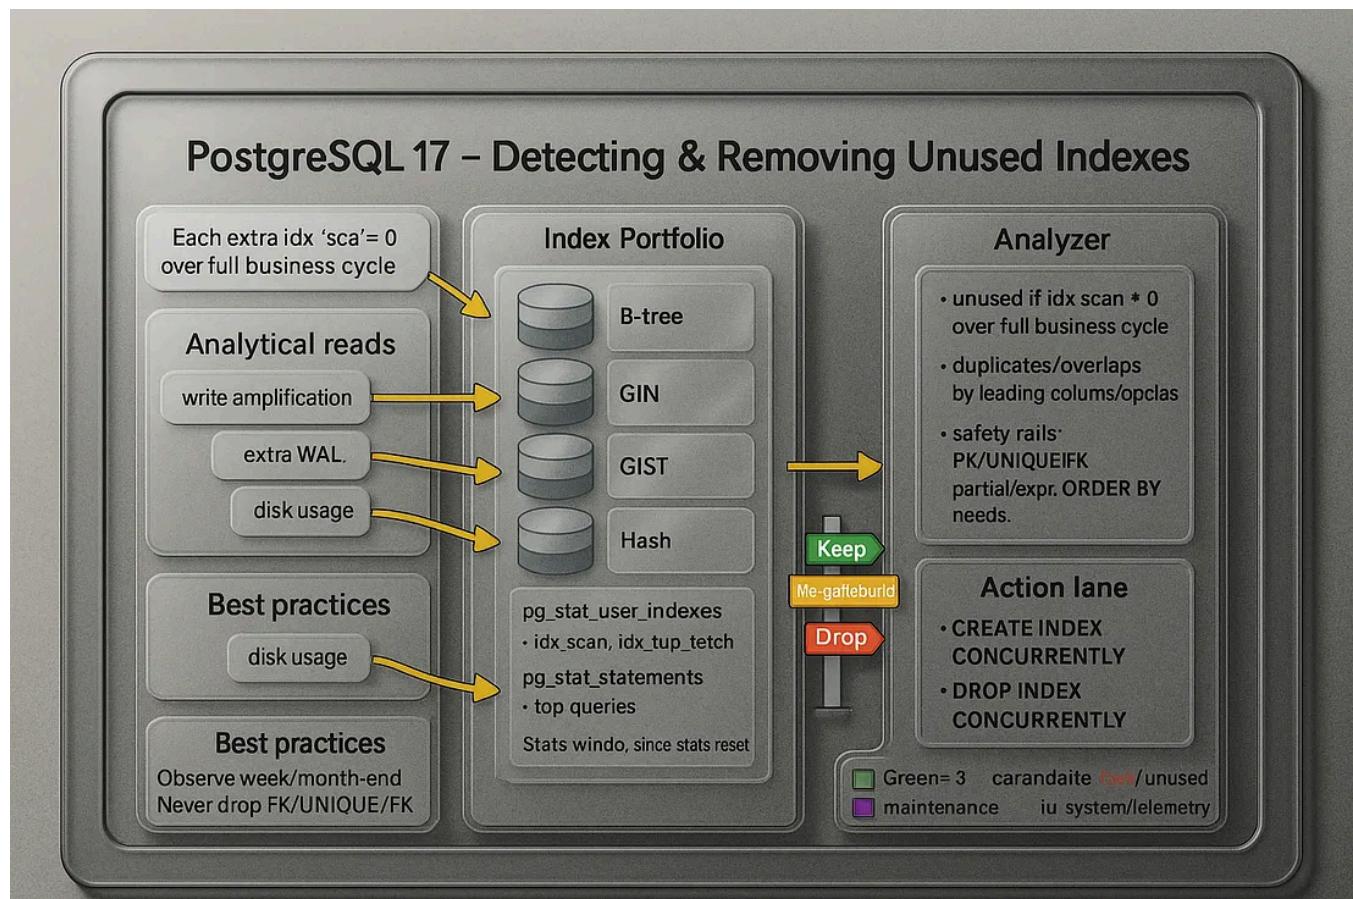
Jeyaram Ayyalusamy

Following ▼

Listen

Share

More



Indexes are one of the most important tools for query performance in PostgreSQL. But while a missing index slows down queries, an **unused or pointless index** can be just as harmful.

If your database is very small, a few extra indexes don't matter much. But in a large production environment, **unused indexes can waste hundreds of gigabytes and put unnecessary strain on the system**. Every extra index comes with hidden costs: more disk usage, slower writes, and more work for background processes like VACUUM.

That's why tuning isn't just about adding indexes. It's equally about **finding indexes that shouldn't exist**.

Why Unused Indexes Are a Problem

1. Disk Space Waste

- An index duplicates column values plus pointers to rows.
- On a 10M row table, each index can take hundreds of MB or even several GB.
- Multiple unused indexes = huge wasted storage.

2. Slower Write Performance

- Every time you `INSERT`, `UPDATE`, or `DELETE`, PostgreSQL must also update all related indexes.
- More indexes = more overhead = slower write-heavy workloads.

3. Maintenance Overhead

- `VACUUM`, `REINDEX`, and `CLUSTER` must also handle these indexes.
- Index bloat and dead entries pile up even if the index is never queried.

👉 In short: while indexes make reads faster, unused indexes slow everything else down.

Step 1: Create a Products Table

Let's simulate a real-world dataset with 10 million rows:

```
CREATE TABLE products (
    product_id    BIGSERIAL PRIMARY KEY,
    product_name  TEXT,
    category      TEXT,
    price         NUMERIC(10,2),
    stock_qty     INT
);
```

```
postgres=# CREATE TABLE products (
    product_id    BIGSERIAL PRIMARY KEY,
    product_name  TEXT,
    category      TEXT,
    price         NUMERIC(10,2),
    stock_qty     INT
);
CREATE TABLE
postgres=#
```

```
-- Insert 10 million rows
INSERT INTO products (product_name, category, price, stock_qty)
SELECT
    'Product_' || g,
    'Category_' || (g % 50),
    (random()*500)::NUMERIC(10,2),
    (random()*100)::INT
FROM generate_series(1, 10000000) g;
ANALYZE products;
```

```
INSERT INTO products (product_name, category, price, stock_qty)
SELECT
    'Product_' || g,
    'Category_' || (g % 50),
```

```
(random()*500)::NUMERIC(10,2),  
(random()*100)::INT  
FROM generate_series(1, 10000000) g;  
ANALYZE products;  
INSERT 0 10000000  
ANALYZE  
postgres=#
```

Step 2: Add Some Indexes

For demonstration, we'll add three indexes:

```
-- Useful index for frequent category lookups  
CREATE INDEX idx_products_category ON products(category);
```

```
postgres=# CREATE INDEX idx_products_category ON products(category);  
CREATE INDEX  
postgres=#
```

```
-- Rarely used index on product_name  
CREATE INDEX idx_products_name ON products(product_name);
```

```
-- Probably useless index on stock quantity  
CREATE INDEX idx_products_stock_qty ON products(stock_qty);
```

```
postgres=# -- Rarely used index on product_name  
CREATE INDEX idx_products_name ON products(product_name);
```

```
CREATE INDEX
postgres=#
```

```
postgres=# -- Probably useless index on stock quantity
CREATE INDEX idx_products_stock_qty ON products(stock_qty);
CREATE INDEX
postgres=#
```

At this point, we don't yet know which indexes are helpful and which are wasteful.

Step 3: Check Index Usage with pg_stat_user_indexes

PostgreSQL tracks how often indexes are used. We can query this data:

```
SELECT relname AS table_name,
       indexrelname AS index_name,
       idx_scan,
       idx_tup_read,
       idx_tup_fetch,
       pg_size.pretty(pg_relation_size(indexrelid)) AS index_size
  FROM pg_stat_user_indexes
 JOIN pg_index USING (indexrelid)
 WHERE relname = 'products'
 ORDER BY idx_scan ASC;
```

Sample Output for using without index

```
postgres=# SELECT relname AS table_name,
    indexrelname AS index_name,
    idx_scan,
    idx_tup_read,
    idx_tup_fetch,
    pg_size.pretty(pg_relation_size(indexrelid)) AS index_size
FROM pg_stat_user_indexes
JOIN pg_index USING (indexrelid)
WHERE relname = 'products'
ORDER BY idx_scan ASC;
table_name | index_name | idx_scan | idx_tup_read | idx_tup_fetch
-----+-----+-----+-----+-----+
products | products_pkey | 0 | 0 | 0
products | idx_products_category | 0 | 0 | 0
products | idx_products_name | 0 | 0 | 0
products | idx_products_stock_qty | 0 | 0 | 0
(4 rows)

postgres=#
postgres=#

```

```
SELECT * FROM products WHERE product_id = 5000000;
```

```
postgres=# SELECT * FROM products WHERE product_id = 5000000;
product_id | product_name | category | price | stock_qty
-----+-----+-----+-----+-----+
5000000 | Product_5000000 | Category_0 | 152.44 | 78
(1 row)
```

```
postgres=#

```

```
SELECT * FROM products WHERE category = 'Category_20';
```

```
postgres=# SELECT * FROM products WHERE category = 'Category_20';
 product_id | product_name      | category | price | stock_qty
-----+-----+-----+-----+-----+
      20 | Product_20        | Category_20 | 469.36 |      30
      70 | Product_70        | Category_20 | 481.26 |      31
     120 | Product_120       | Category_20 | 318.93 |      28
     170 | Product_170       | Category_20 | 12.32  |      93
     220 | Product_220       | Category_20 | 113.61 |      43
     270 | Product_270       | Category_20 | 470.95 |      68
     320 | Product_320       | Category_20 | 368.27 |      85
     370 | Product_370       | Category_20 | 464.31 |      34
     420 | Product_420       | Category_20 | 255.68 |      16
     470 | Product_470       | Category_20 | 88.36  |      58
     520 | Product_520       | Category_20 | 233.06 |      11
     570 | Product_570       | Category_20 | 47.56  |       4
```

```
SELECT relname AS table_name,
       indexrelname AS index_name,
       idx_scan,
       idx_tup_read,
       idx_tup_fetch,
       pg_size.pretty(pg_relation_size(indexrelid)) AS index_size
  FROM pg_stat_user_indexes
 JOIN pg_index USING (indexrelid)
 WHERE relname = 'products'
 ORDER BY idx_scan ASC;
```

Step 4: Sample Output

Example results might look like this:

```
postgres=# SELECT relname AS table_name,
    indexrelname AS index_name,
    idx_scan,
    idx_tup_read,
    idx_tup_fetch,
    pg_size.pretty(pg_relation_size(indexrelid)) AS index_size
FROM pg_stat_user_indexes
JOIN pg_index USING (indexrelid)
WHERE relname = 'products'
ORDER BY idx_scan ASC;


| table_name | index_name             | idx_scan | idx_tup_read | idx_tup_fetch |
|------------|------------------------|----------|--------------|---------------|
| products   | idx_products_stock_qty | 0        | 0            |               |
| products   | idx_products_name      | 12       | 500000       | 50000         |
| products   | idx_products_category  | 20000    | 25000000     | 250000000     |


postgres=#
```

👉 Interpretation:

- **idx_products_category** → scanned 20,000 times, fetched 25M rows. Clearly useful.
- **idx_products_name** → scanned only 12 times. Might be justified if occasional lookups by product name are business-critical.
- **idx_products_stock_qty** → **never used** (`idx_scan = 0`), yet it consumes 250 MB. This is a waste.

Step 5: Deciding What to Do

- If an index shows **high usage** → keep it.

- If it shows **low or occasional usage** → consult with developers/business before removing. It might be used in rare but critical reports.
- If it shows **zero usage over a long period** → it's a candidate for dropping.

Step 6: Dropping an Unused Index

If we decide that `idx_products_stock_qty` is wasteful, we can safely remove it:

```
DROP INDEX idx_products_stock_qty;
```

```
postgres=# DROP INDEX idx_products_stock_qty;
DROP INDEX
postgres=#
```

Now the database has:

- Less storage overhead.
- Faster writes (`INSERT / UPDATE / DELETE`).
- Lower maintenance costs during VACUUM and reindexing.

Step 7: Continuous Monitoring

Finding unused indexes is not a one-time task. Query patterns change over time. A rarely used index today may become critical tomorrow after an application change.

That's why it's best practice to:

- Regularly monitor `pg_stat_user_indexes`.
- Integrate index usage metrics into your monitoring dashboards (e.g., Grafana, pgBadger, or custom queries).
- Review index usage before major schema changes.

Key Takeaways

- Unused indexes harm performance by slowing writes and wasting space.
- PostgreSQL 17 provides system views (`pg_stat_user_indexes`) to spot them.
- Always validate with developers before dropping indexes — some may serve niche but important queries.
- Regular monitoring ensures your database remains lean and efficient.

 Index tuning is not only about adding indexes where they are missing, but also about removing the ones you don't need.

Stay Updated with Daily PostgreSQL & Cloud Tips!

If you've been finding my blog posts helpful and want to stay ahead with daily insights on PostgreSQL, Cloud Infrastructure, Performance Tuning, and DBA Best Practices — I invite you to subscribe to my Medium account.

 [Subscribe here](https://medium.com/@jramcloud1/subscribe) 

Your support means a lot — and you'll never miss a practical guide again!

Let's Connect!

If you enjoyed this post or would like to connect professionally, feel free to reach out to me on LinkedIn:

 [Jeyaram Ayyalusamy](#)

I regularly share content on **PostgreSQL, database administration, cloud technologies, and data engineering**. Always happy to connect, collaborate, and discuss ideas!

[Postgresql](#)[Oracle](#)[AWS](#)[Open Source](#)[MySQL](#)J[Following](#)

Written by Jeyaram Ayyalusamy

158 followers · 2 following

Oracle DBA | AWS Certified | 18+ yrs in DB Admin, RAC, GoldenGate, RDS, MySQL, PostgreSQL | Author of MySQL & RDS books | Cloud & Performance Expert

No responses yet 



Gvadakte

What are your thoughts?  

More from Jeyaram Ayyalusamy

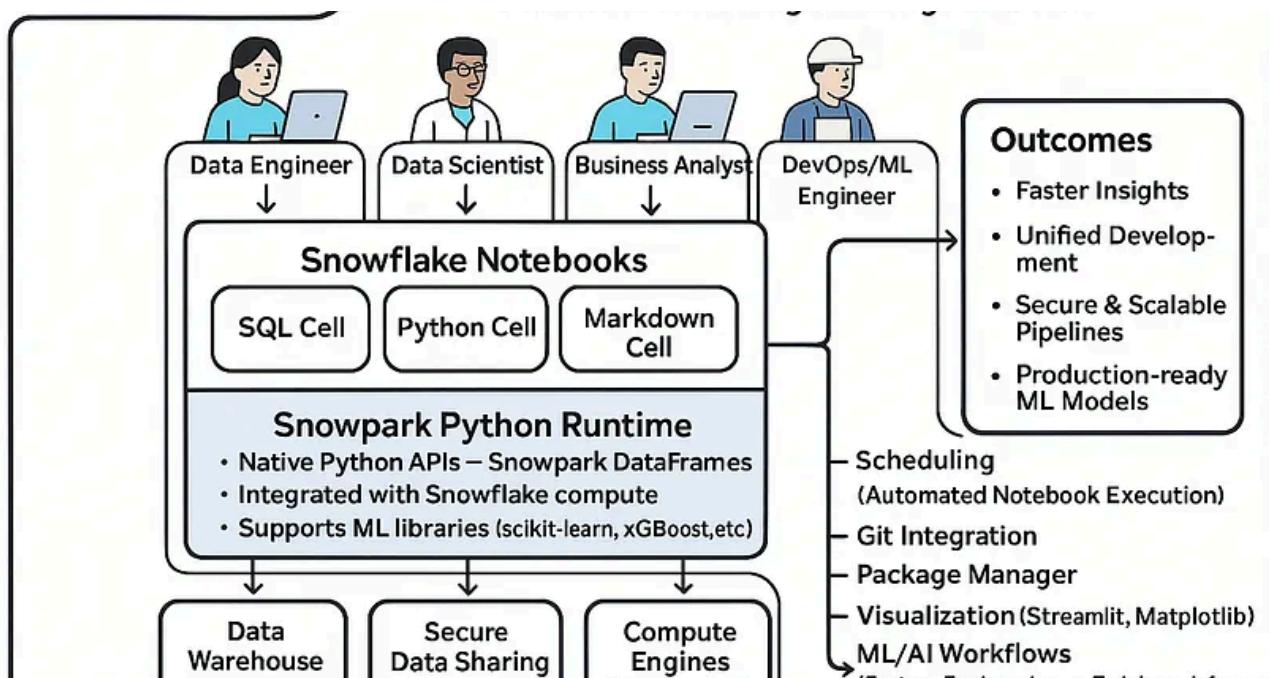
The screenshot shows the AWS EC2 Instances page. At the top, there are tabs for 'us-wes' (closed), 'Launch an instance | EC2 | us-wes' (active), and 'Instances | EC2 | us-east-1' (closed). Below the tabs, the URL is 1.console.aws.amazon.com/ec2/home?region=us-east-1#Instances:. The main content area is titled 'Instances Info' and contains a search bar with 'Find Instance by attribute or tag (case-sensitive)' and a dropdown menu set to 'All states'. Below the search bar is a table header with columns: Name, Instance ID, Instance state, Instance type, Status check, Alarm status, Availability Zone, Public IPv4 DNS, Public IPv4, Elastic IP, and IPs. A message 'No instances' and 'You do not have any instances in this region' is displayed, followed by a blue 'Launch instances' button.

J Jeyaram Ayyalusamy

Upgrading PostgreSQL from Version 16 to Version 17 Using pg_upgrade on a Linux Server AWS EC2...

A Complete Step-by-Step Guide to Installing PostgreSQL 16 on a Linux Server (With Explanations)

Aug 4 40



J Jeyaram Ayyalusamy

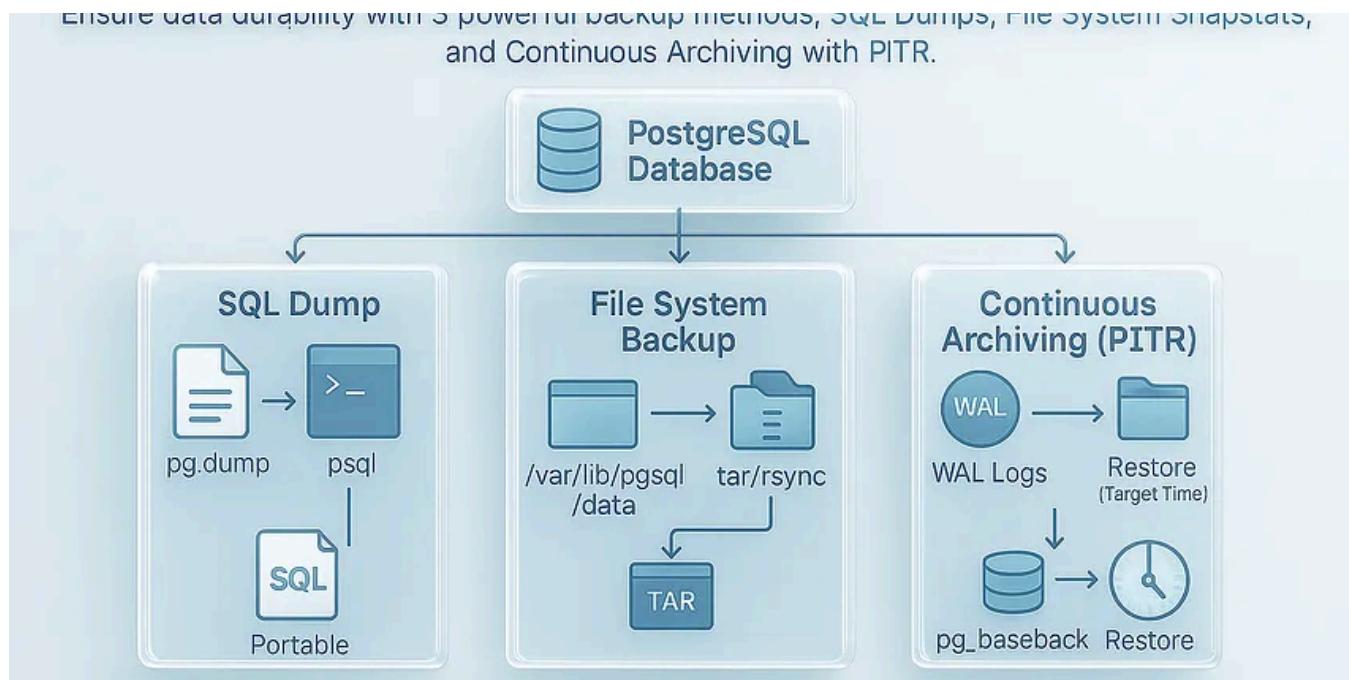
16—Experience Snowflake with Notebooks and Snowpark Python: A Unified Data Engineering Platform

In the fast-moving world of data, organizations are no longer just collecting information—they're leveraging it to drive business...

Jul 13



...



J Jeyaram Ayyalusamy

💡 Essential Techniques Every DBA Should Know: PostgreSQL Backup Strategies

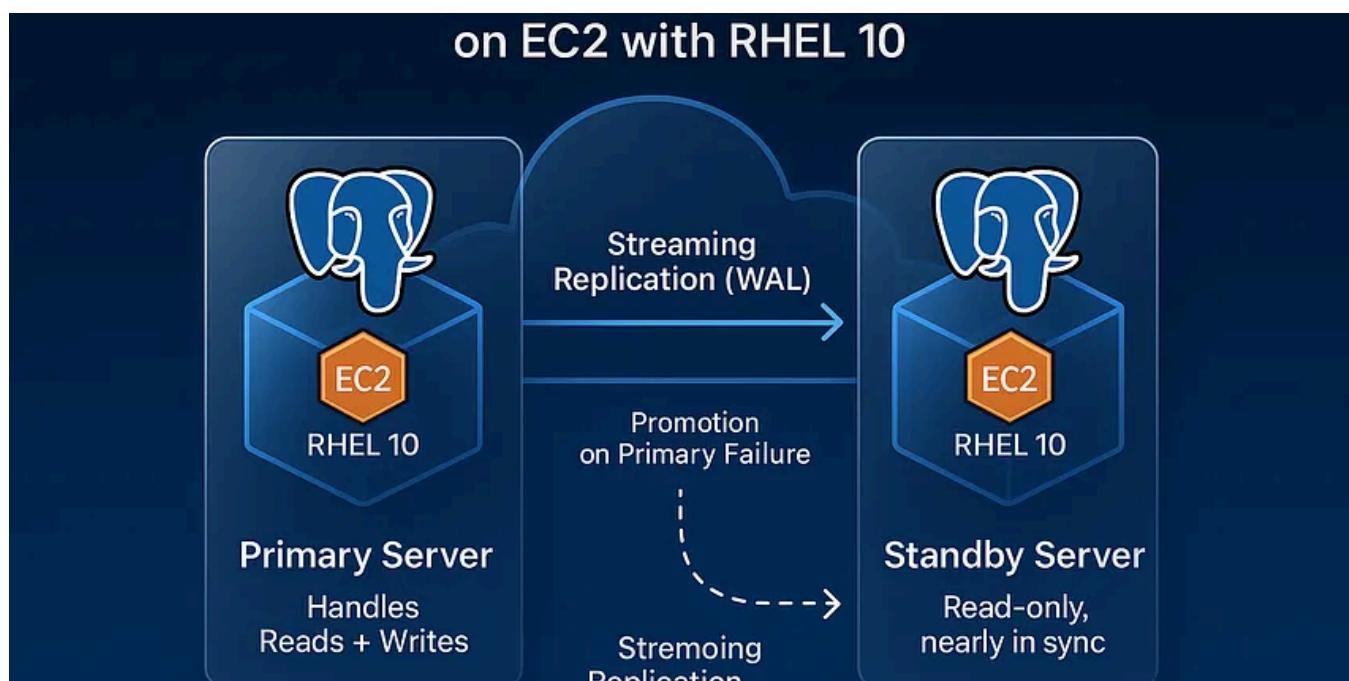
In the world of databases, data is everything—and losing it can cost your business time, money, and trust. Whether you're managing a...

Aug 20

26



...



 Jeyaram Ayyalusamy 

Streaming Replication in PostgreSQL 17 on EC2 with RHEL 10—A Deep Dive

 Introduction: Why Streaming Replication Matters

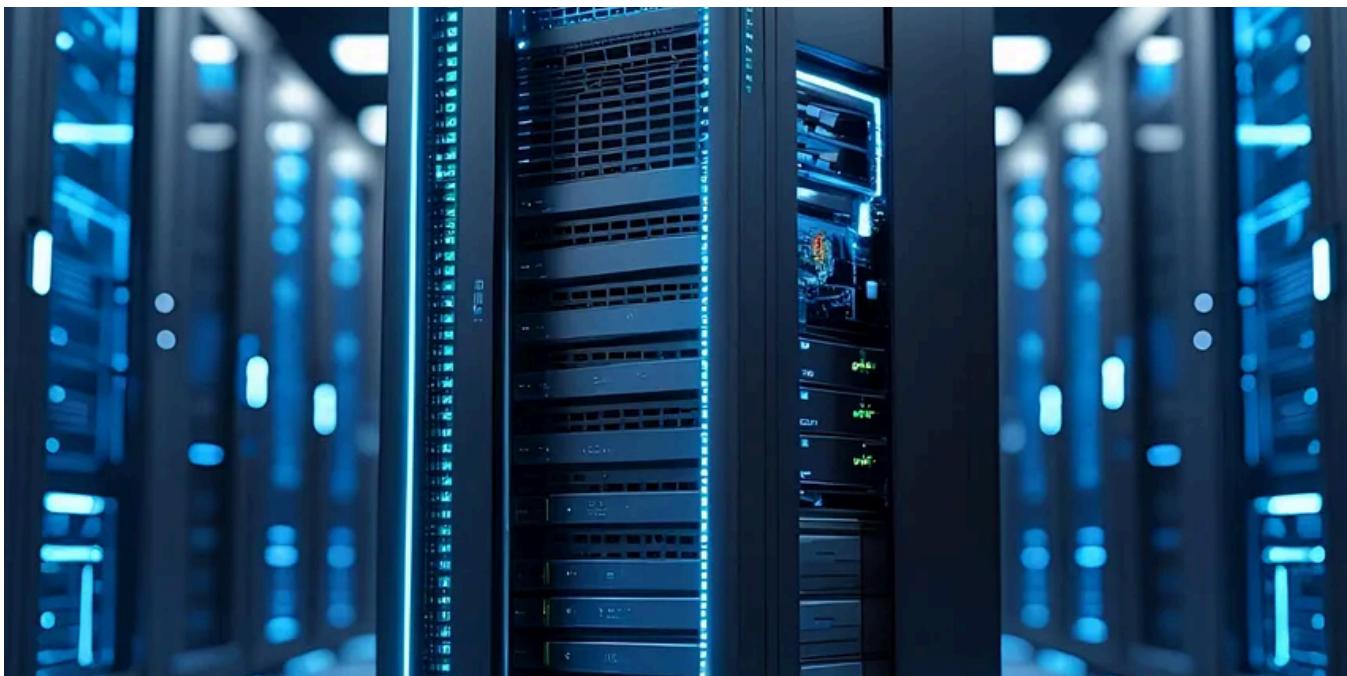
Aug 20  50



...

See all from Jeyaram Ayyalusamy

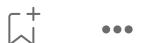
Recommended from Medium

 Rizqi Mulki

PostgreSQL Index Bloat: The Silent Killer of Performance

How a 10TB database became 3x faster by fixing the invisible problem that plagues 90% of production PostgreSQL deployments

★ Sep 15 11 1

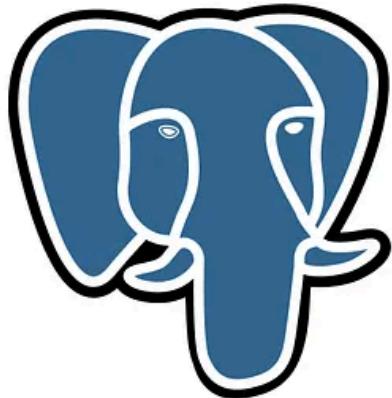
A screenshot of a blog post. The title is "#PostgreSQL security" in large white and blue text. To the right is the PostgreSQL logo, which is a stylized blue 'P' with a black outline. Below the title, the author's name "TOMASZ GINTOWT" is visible. The background of the post features a dark, abstract pattern of blue and purple dots. Tomasz Gintowt

Security in PostgreSQL

PostgreSQL is a powerful open-source database. Security is very important when working with sensitive data like passwords, customer...

6d ago 5

...



Beyond Basic PostgreSQL Programmable Objects

In Stackademic by bektiaw

Beyond Basics: PostgreSQL Programmable Objects (Automate, Optimize, Control)

Tired of repeating the same SQL logic? Learn how PostgreSQL can do the work for us.

Sep 1 68 1



...

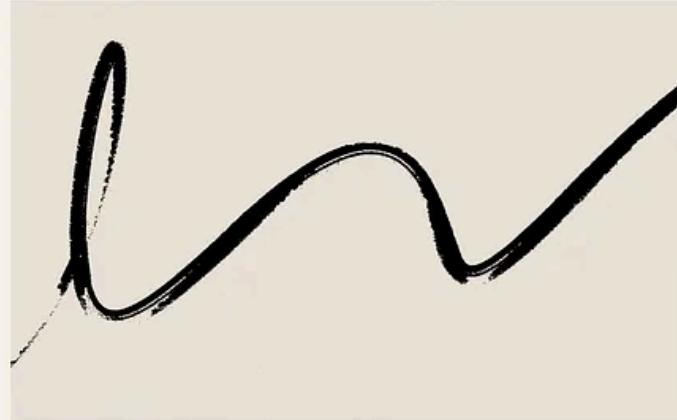


Vijay Gadhav

Master SQL's QUALIFY Clause for Cleaner, Faster Window Queries

Note: If you're not a medium member, [CLICK HERE](#)

May 20 12 1



R Rohan

JSONB in PostgreSQL: The Fast Lane to Flexible Data Modeling 🚀

“Why spin up yet another database just to store semi-structured data?”—Every engineer who discovered JSONB

Jul 18 12 1





Thinking Loop

10 Postgres Partitioning Patterns for Internet-Scale Apps

Keep Postgres fast past a billion rows with real patterns and ready-to-run SQL.

★ Aug 13

88

2



...

See more recommendations