

[Open in app](#)

Medium



Get unlimited access to the best of Medium for less than \$1/week. [Become a member](#)



27 - PostgreSQL 17 Performance Tuning: Monitoring Queries with pg_stat_statements

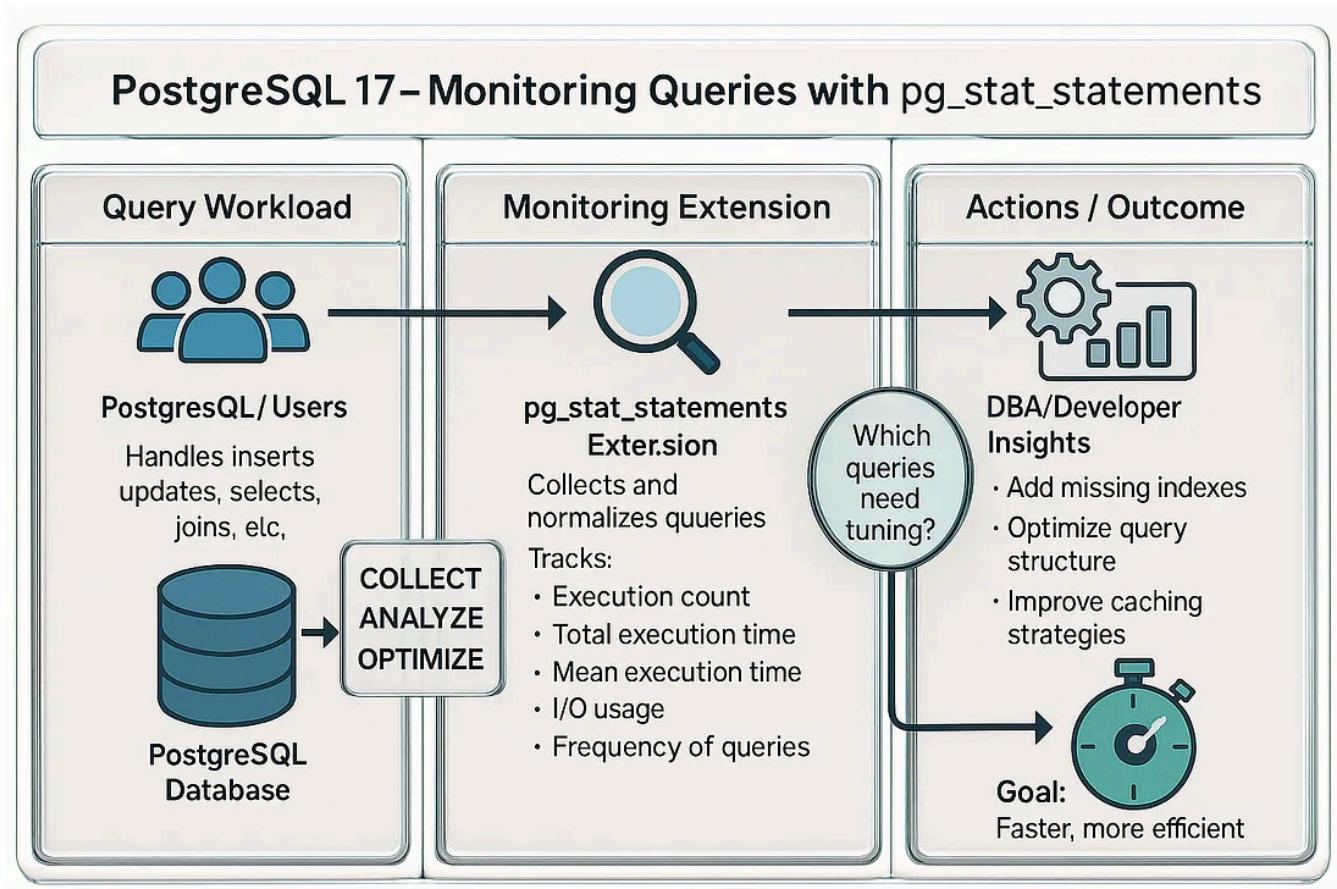
11 min read · 3 days ago

J Jeyaram Ayyalusamy Following

Listen

Share

More



You cannot improve database performance without first collecting the right data. PostgreSQL provides a wide range of **system views** that let administrators and developers understand what's really happening inside their database.

One of the most powerful tools for this purpose is the `pg_stat_statements` extension.

It helps you identify:

- Which queries are slow.
- How often queries are executed.
- Whether queries have stable or fluctuating runtimes.

Instead of guessing or parsing log files, `pg_stat_statements` provides a clear, aggregated view of query behavior.

Step 1: Installing PostgreSQL Extension Modules

To enable advanced features like `pg_stat_statements`, you need to install the PostgreSQL contrib package. This package includes a collection of officially supported extensions that enhance PostgreSQL's functionality.

On Red Hat-based systems (RHEL, CentOS, Amazon Linux), run:

```
sudo yum install postgresql17-contrib
```

```
[ec2-user@ip-172-31-29-78 ~]$ sudo yum install postgresql17-contrib
Updating Subscription Management repositories.
Unable to read consumer identity
```

This system is not registered with an entitlement server. You can use "rhc" or

Last metadata expiration check: 2:44:15 ago on Fri Sep 19 18:47:07 2025.
Dependencies resolved.

```
=====
 Package                                Architecture      Version
 =====
 Installing:
 postgresql17-contrib                  x86_64          17.0.0
```

Installing dependencies:

libxslt	x86_64	1.
---------	--------	----

Transaction Summary

Install 2 Packages

Total download size: 920 k

Installed size: 3.2 M

Is this ok [y/N]: y

Downloading Packages:

```
(1/2): libxslt-1.1.39-8.el10_0.x86_64.rpm
(2/2): postgresql17-contrib-17.6-1PGDG.rhel10.x86_64.rpm
```

Total

Running transaction check

Transaction check succeeded.

Running transaction test

Transaction test succeeded.

Running transaction

Preparing :

Installing : libxslt-1.1.39-8.el10_0.x86_64

Installing : postgresql17-contrib-17.6-1PGDG.rhel10.x86_64

Running scriptlet: postgresql17-contrib-17.6-1PGDG.rhel10.x86_64

Installed products updated.

Installed:

libxslt-1.1.39-8.el10_0.x86_64

Complete!

[ec2-user@ip-172-31-29-78 ~]\$

This command:

- Installs the **contrib module** for PostgreSQL 17
- Adds support for extensions like `pg_stat_statements`, `tablefunc`, `uuid-ossp`, and more
- Is required before enabling certain features in `postgresql.conf` (e.g.,
`shared_preload_libraries = 'pg_stat_statements'`)



After installation, don't forget to restart the PostgreSQL service and run `CREATE EXTENSION pg_stat_statements;` inside your database to activate it.

Step 2: Enable pg_stat_statements

By default, this extension is not active. You need to enable it in `postgresql.conf`.

1. Open your PostgreSQL config file (e.g., `postgresql.conf`).

2. Find the line:

```
#shared_preload_libraries = ''
```

```
postgres=# show shared_preload_libraries;
shared_preload_libraries
-----
```

```
(1 row)
```

```
postgres=#
```

3. Uncomment it and set the value to include `pg_stat_statements`:

```
shared_preload_libraries = 'pg_stat_statements'
```

```
[postgres@ip-172-31-29-78 data]$ cat postgresql.conf | grep -i shared_preload_l
shared_preload_libraries = 'pg_stat_statements'          # (change requires rest
[postgres@ip-172-31-29-78 data]$
```

```
postgres=# show shared_preload_libraries;
shared_preload_libraries
-----
pg_stat_statements
(1 row)

postgres=#
```

4. Restart PostgreSQL for the change to take effect.

```
[ec2-user@ip-172-31-29-78 ~]$ sudo systemctl restart postgresql-17
[ec2-user@ip-172-31-29-78 ~]$
```

Finally, install the extension in the database you want to monitor:

```
CREATE EXTENSION pg_stat_statements;
```

This will create the system view pg_stat_statements .

```
postgres=# CREATE EXTENSION pg_stat_statements;
CREATE EXTENSION
postgres=#
```

Step 3: Generate Some Demo Load

For demonstration, let's create a `products` table with 10 million rows and run a few queries against it.

```
CREATE TABLE products (
    product_id    BIGSERIAL PRIMARY KEY,
    product_name  TEXT,
    category      TEXT,
    price         NUMERIC,
    stock_qty     INT
);
```

```
postgres=# CREATE EXTENSION pg_stat_statements;
CREATE EXTENSION
postgres=#
postgres=# drop table products;
DROP TABLE
postgres=# CREATE TABLE products (
    product_id    BIGSERIAL PRIMARY KEY,
    product_name  TEXT,
    category      TEXT,
    price         NUMERIC,
    stock_qty     INT
);
CREATE TABLE
postgres=#
```

```
INSERT INTO products (product_name, category, price, stock_qty)
SELECT
    'Product_' || g,
    'Category_' || (g % 10),
    (random()*500)::NUMERIC,
    (random()*100)::INT
FROM generate_series(1, 10000000) g;
ANALYZE products;
```

```
postgres=# INSERT INTO products (product_name, category, price, stock_qty)
SELECT
    'Product_' || g,
    'Category_' || (g % 10),
    (random()*500)::NUMERIC,
    (random()*100)::INT
FROM generate_series(1, 10000000) g;
ANALYZE products;
INSERT 0 10000000
ANALYZE
postgres=#
```

Now run queries such as:

```
SELECT * FROM products WHERE product_id = 100;
SELECT * FROM products WHERE category = 'Category_3';
UPDATE products SET stock_qty = stock_qty + 1 WHERE category = 'Category_5';
```

```
postgres=# SELECT * FROM products WHERE product_id = 100;
 product_id | product_name | category |      price       | stock_qty
-----+-----+-----+-----+-----+
 100 | Product_100 | Category_0 | 228.865334176435 |      36
(1 row)

postgres=#
```

```
postgres=# SELECT * FROM products WHERE category = 'Category_3';
 product_id | product_name | category |      price       | stock_qty
-----+-----+-----+-----+-----+
   3 | Product_3   | Category_3 | 216.496292091521 |      9
  13 | Product_13  | Category_3 | 158.801528884299 |     50
  23 | Product_23  | Category_3 | 402.177798888444 |     60
  33 | Product_33  | Category_3 | 188.638645203066 |     26
  43 | Product_43  | Category_3 | 217.520085735964 |     95
  53 | Product_53  | Category_3 | 204.260688438616 |     43
  63 | Product_63  | Category_3 | 4.38792827933232 |     36
  73 | Product_73  | Category_3 | 179.192418212635 |     13
  83 | Product_83  | Category_3 | 177.750125505008 |     33
  93 | Product_93  | Category_3 | 406.374932408794 |     32
 103 | Product_103 | Category_3 | 47.2825437678309 |      5
 113 | Product_113 | Category_3 | 347.047753193479 |     34
 123 | Product_123 | Category_3 | 243.453650026469 |     46
 133 | Product_133 | Category_3 | 400.878856999775 |     99
 143 | Product_143 | Category_3 | 110.710054872779 |     84
 153 | Product_153 | Category_3 | 422.519471685829 |     43
 163 | Product_163 | Category_3 | 418.99101277994 |     10
 173 | Product_173 | Category_3 | 330.138006726809 |     66
 183 | Product_183 | Category_3 | 266.622715912463 |     60
```

```
postgres=# UPDATE products SET stock_qty = stock_qty + 1 WHERE category = 'Category A'
UPDATE 1000000
postgres=#

```

Step 4: Query pg_stat_statements

Now check what PostgreSQL recorded:

```
SELECT query, calls, total_exec_time, mean_exec_time, rows, stddev_exec_time
FROM pg_stat_statements
ORDER BY total_exec_time DESC
LIMIT 10;
```

👉 Each row in this view represents one **normalized query**. Queries with different parameter values are aggregated together. For example:

- `SELECT * FROM products WHERE product_id = 100;`
- `SELECT * FROM products WHERE product_id = 200;`

Both are shown as:

```
postgres=# SELECT query, calls, total_exec_time, mean_exec_time, rows, stddev_exec_time
FROM pg_stat_statements
ORDER BY total_exec_time DESC
LIMIT 10;
```

query	calls
<code>INSERT INTO products (product_name, category, price, stock_qty)</code>	1
<code>SELECT \$1 g,</code>	3
<code>\$2 (g % \$3),</code>	
<code>(random()*\$4)::NUMERIC,</code>	

```

(random()*$5)::INT
FROM generate_series($6, $7) g
UPDATE products SET stock_qty = stock_qty + $1 WHERE category = $2
ANALYZE products
SELECT * FROM products WHERE category = $1
CREATE EXTENSION pg_stat_statements
CREATE TABLE products (
    product_id BIGSERIAL PRIMARY KEY,
    product_name TEXT,
    category TEXT,
    price NUMERIC,
    stock_qty INT
)
drop table products
SELECT * FROM products WHERE product_id = $1
SELECT query, calls, total_exec_time, mean_exec_time, rows
FROM pg_stat_statements
ORDER BY total_exec_time DESC
LIMIT $1
show shared_preload_libraries
(10 rows)

```

postgres=#

This makes it easier to measure query patterns, not just individual runs.

Step 5: Useful Columns in pg_stat_statements

- **calls** → how often the query ran.
- **total_exec_time** → total time spent across all executions.
- **mean_exec_time** → average execution time.
- **rows** → average number of rows returned.
- **stddev_exec_time** → (standard deviation) shows if query times are stable or fluctuate heavily.

Why is this important?

- If runtimes are unstable, it could mean the query sometimes hits the cache and sometimes reads from disk.

- Shared blocks vs read blocks columns help identify whether data came from PostgreSQL's buffer cache (`shared_blk_hit`) or from disk (`shared_blk_read`).

```
SELECT
    query,
    calls,
    shared_blk_hit,
    shared_blk_read,
    ROUND(shared_blk_hit * 100.0 / GREATEST(shared_blk_hit + shared_blk_read, 1)) AS hit_ratio
FROM
    pg_stat_statements
ORDER BY
    hit_ratio ASC
LIMIT 20;
```

```
query,
calls,
shared_blk_hit,
shared_blk_read,
ROUND(shared_blk_hit * 100.0 / GREATEST(shared_blk_hit + shared_blk_read, 1)) AS hit_ratio
FROM
    pg_stat_statements
ORDER BY
    hit_ratio ASC
LIMIT 20;
```

query

```
show shared_preload_libraries
SELECT query, calls, total_exec_time, mean_exec_time, rows, stddev_exec_time+
FROM pg_stat_statements
ORDER BY total_exec_time DESC
LIMIT $1
SELECT * FROM products WHERE category = $1
ANALYZE products
SELECT * FROM products WHERE product_id = $1
CREATE TABLE products (
    product_id BIGSERIAL PRIMARY KEY,
    product_name TEXT,
    category TEXT,
    price NUMERIC,
```

```

    stock_qty      INT
)
CREATE EXTENSION pg_stat_statements
drop table products
UPDATE products SET stock_qty = stock_qty + $1 WHERE category = $2
SELECT query, calls, total_exec_time, mean_exec_time, rows
FROM pg_stat_statements
ORDER BY total_exec_time DESC
LIMIT $1
INSERT INTO products (product_name, category, price, stock_qty)
SELECT
    $1 || g,
    $2 || (g % $3),
    (random()*$4)::NUMERIC,
    (random()*$5)::INT
FROM generate_series($6, $7) g
(11 rows)

postgres=#

```

Step 6: Tracking Temporary File Usage

Queries that use too much memory can spill to disk, creating **temporary files**. In OLTP systems (transaction-heavy workloads), this is usually a red flag because it slows everything down.

`pg_stat_statements` tracks this too. High temporary file I/O may indicate:

- Work memory (`work_mem`) is too low.
- Large sorts, hash joins, or index builds are happening.

```

SELECT
    query,
    calls,
    temp_blk_written,
    temp_blk_read,
    ROUND(temp_blk_written * 8 / 1024.0, 2) AS temp_written_MB,
    ROUND(temp_blk_read * 8 / 1024.0, 2) AS temp_read_MB
FROM
    pg_stat_statements
WHERE

```

```

temp_blk_written > 0
ORDER BY
temp_blk_written DESC
LIMIT 20;

```

```

postgres=# SELECT
query,
calls,
temp_blk_written,
temp_blk_read,
ROUND(temp_blk_written * 8 / 1024.0, 2) AS temp_written_MB,
ROUND(temp_blk_read * 8 / 1024.0, 2) AS temp_read_MB

```

```

FROM
pg_stat_statements
WHERE
temp_blk_written > 0
ORDER BY
temp_blk_written DESC
LIMIT 20;

```

query	calls	temp
<code>INSERT INTO products (product_name, category, price, stock_qty)</code>	1	
<code>SELECT</code>		
<code>\$1 g,</code>		
<code>\$2 (g % \$3),</code>		
<code>(random()*\$4)::NUMERIC,</code>		
<code>(random()*\$5)::INT</code>		
<code>FROM generate_series(\$6, \$7) g</code>		
<code>(1 row)</code>		

```
postgres=#
```

Step 7: Enable I/O Timing

By default, I/O timing is not recorded because it adds some overhead. To track it, enable:

```
track_io_timing = on
```

```
postgres=# show track_io_timing;
track_io_timing
-----
off
(1 row)

postgres=#
```

```
[postgres@ip-172-31-29-78 data]$ cat postgresql.conf | grep -i track_io_timing
track_io_timing = on
[postgres@ip-172-31-29-78 data]$
```



```
[ec2-user@ip-172-31-29-78 ~]$ sudo systemctl restart postgresql-17
[ec2-user@ip-172-31-29-78 ~]$
```

```
postgres=# show track_io_timing;
track_io_timing
-----
on
(1 row)
```

```
postgres=#
```

This will let you see how much query time was spent on actual disk I/O.

🔍 Using EXPLAIN (ANALYZE, BUFFERS) to Diagnose Query Performance

PostgreSQL's EXPLAIN (ANALYZE, BUFFERS) command provides deep visibility into how a query interacts with memory and disk. Here's a real-world example:

```
EXPLAIN (ANALYZE, BUFFERS)
SELECT * FROM products WHERE category = 'Category_5';
```

Sample Output:

```
postgres=# EXPLAIN (ANALYZE, BUFFERS)
SELECT * FROM products WHERE category = 'electronics';
                                         QUERY PLAN
-----
Gather  (cost=1000.00..187116.59 rows=1 width=49) (actual time=5987.833..5998.
    Workers Planned: 2
    Workers Launched: 2
    Buffers: shared hit=12025 read=111592 dirtied=19987 written=13943
    I/O Timings: shared read=13342.256 write=1658.516
    -> Parallel Seq Scan on products  (cost=0.00..186116.49 rows=1 width=49) (a
        Filter: (category = 'electronics'::text)
        Rows Removed by Filter: 333333
        Buffers: shared hit=12025 read=111592 dirtied=19987 written=13943
        I/O Timings: shared read=13342.256 write=1658.516
Planning:
    Buffers: shared hit=6 read=3
    I/O Timings: shared read=1.824
Planning Time: 1.928 ms
Execution Time: 6000.232 ms
(15 rows)

postgres=#

```

💡 What This Tells Us

- **Sequential Scan:** PostgreSQL scanned the entire `products` table—this is often a red flag in OLTP workloads.
- **Rows Removed by Filter:** 9 million rows were scanned but didn't match the filter, indicating poor selectivity.
- **Buffers:**
 - `shared hit=14584` : These blocks were found in memory (PostgreSQL's buffer cache).
 - `shared read=109033` : These blocks were read from disk—this is expensive and slows down performance.
- **I/O Timings:**
 - `shared read=4033.907 ms` : Over 4 seconds were spent just reading from disk.
- **Execution Time:** The query took over 8 seconds to complete.

⚠ Optimization Opportunity

This query would benefit from:

- A B-tree index on `category`:

```
CREATE INDEX idx_products_category ON products (category);
```

```
postgres=# CREATE INDEX idx_products_category ON products (category);
CREATE INDEX
```

```
EXPLAIN (ANALYZE, BUFFERS)
SELECT * FROM products WHERE category = 'Category_5';
```

```
postgres=# EXPLAIN (ANALYZE, BUFFERS)
SELECT * FROM products WHERE category = 'Category_5';
```

QUERY PLAN

```
Bitmap Heap Scan on products  (cost=11229.68..230035.48 rows=1003000 width=49)
  Recheck Cond: (category = 'Category_5'::text)
  Heap Blocks: exact=10321
  Buffers: shared hit=468 read=10711 written=59
  I/O Timings: shared read=4543.616 write=0.454
    -> Bitmap Index Scan on idx_products_category  (cost=0.00..10978.93 rows=1003000 width=49)
      Index Cond: (category = 'Category_5'::text)
      Buffers: shared read=858 written=25
      I/O Timings: shared read=6.589 write=0.144
Planning:
  Buffers: shared hit=43 read=2
  I/O Timings: shared read=0.945
Planning Time: 1.156 ms
Execution Time: 5955.813 ms
(14 rows)
```

```
postgres=#
```

- Increasing `work_mem` if sorts or joins are involved.
- Rewriting the query to reduce the result set or improve filtering.

 Always pair `EXPLAIN (ANALYZE, BUFFERS)` with `track_io_timing = on` in `postgresql.conf` to capture disk latency.

Key Takeaways

- `pg_stat_statements` is the go-to extension for query monitoring in PostgreSQL 17.
- It aggregates queries with parameters, helping spot patterns instead of noise.
- Key insights:
- **Most expensive queries** (by time).
- **Most frequent queries** (by calls).
- **Queries with unstable runtimes** (cache vs disk).
- **Queries causing temp file I/O** (tune `work_mem`).
- Without monitoring, tuning is guesswork — with `pg_stat_statements`, it's **data-driven optimization**.

 Before adjusting indexes, memory, or autovacuum, always check what queries are actually causing the load.

Stay Updated with Daily PostgreSQL & Cloud Tips!

If you've been finding my blog posts helpful and want to stay ahead with daily insights on PostgreSQL, Cloud Infrastructure, Performance Tuning, and DBA Best Practices — I invite you to subscribe to my Medium account.

 [Subscribe here](https://medium.com/@jramcloud1/subscribe) 

Your support means a lot — and you'll never miss a practical guide again!

Let's Connect!

If you enjoyed this post or would like to connect professionally, feel free to reach out to me on LinkedIn:

 [Jeyaram Ayyalusamy](#)

I regularly share content on **PostgreSQL, database administration, cloud technologies, and data engineering**. Always happy to connect, collaborate, and

discuss ideas!

Postgresql

Oracle

Open Source

Mongodb

Sql

J

Following ▾

Written by Jeyaram Ayyalusamy

158 followers · 2 following

Oracle DBA | AWS Certified | 18+ yrs in DB Admin, RAC, GoldenGate, RDS, MySQL, PostgreSQL | Author of MySQL & RDS books | Cloud & Performance Expert

No responses yet



Gvadakte

What are your thoughts?



More from Jeyaram Ayyalusamy

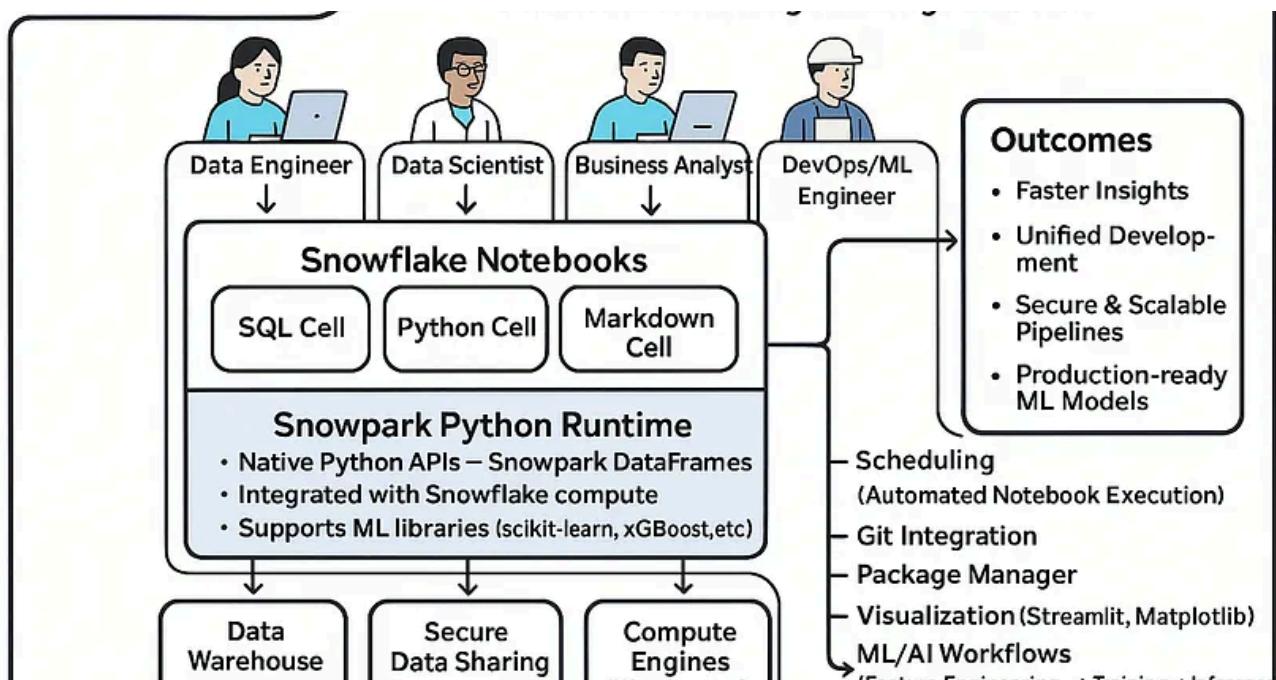
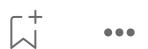
The screenshot shows the AWS EC2 Instances page. At the top, there are three tabs: 'us-wes' (closed), 'Launch an instance | EC2 | us-wes' (active), and 'Instances | EC2 | us-east-1' (closed). Below the tabs, the URL is 1.console.aws.amazon.com/ec2/home?region=us-east-1#Instances:. The main content area is titled 'Instances Info' and contains a search bar with 'Find Instance by attribute or tag (case-sensitive)' and a dropdown menu set to 'All states'. Below the search bar is a table header with columns: Name, Instance ID, Instance state, Instance type, Status check, Alarm status, Availability Zone, Public IPv4 DNS, Public IPv4, Elastic IP, and IPs. A message 'No instances' and 'You do not have any instances in this region' is displayed, followed by a blue 'Launch instances' button.

J Jeyaram Ayyalusamy

Upgrading PostgreSQL from Version 16 to Version 17 Using pg_upgrade on a Linux Server AWS EC2...

A Complete Step-by-Step Guide to Installing PostgreSQL 16 on a Linux Server (With Explanations)

Aug 4 40

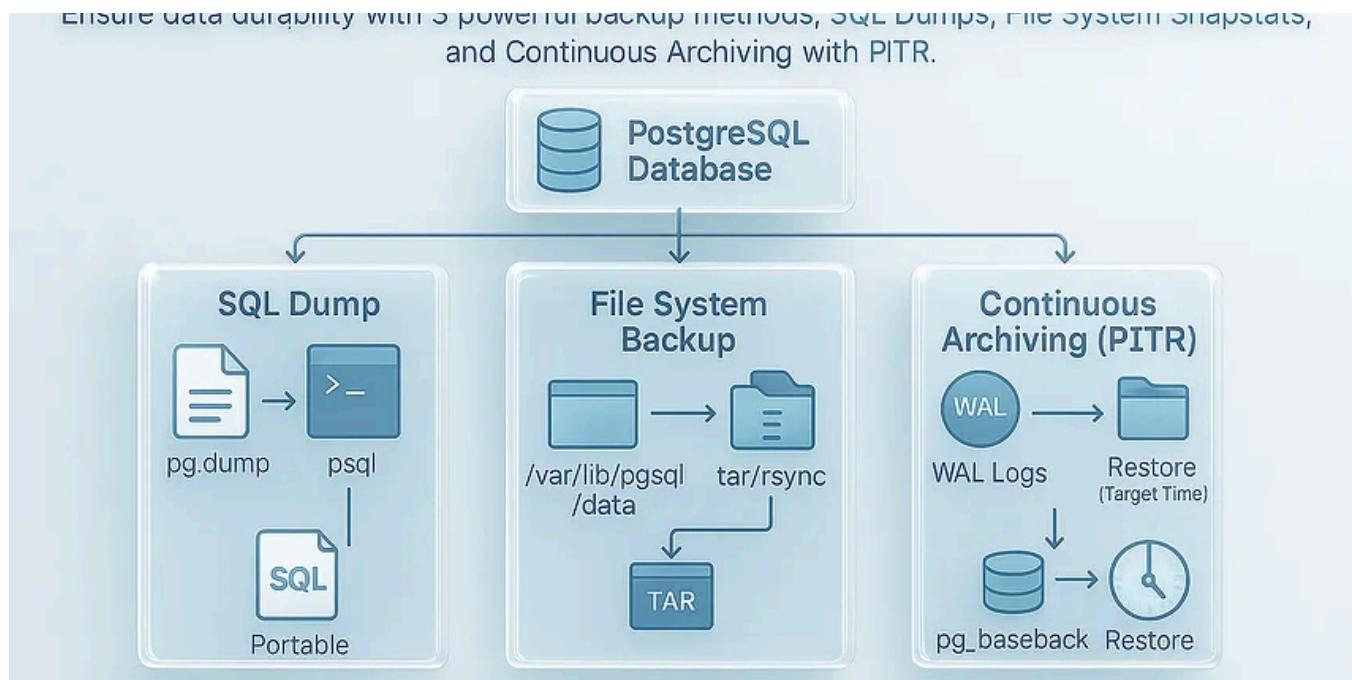


J Jeyaram Ayyalusamy

16—Experience Snowflake with Notebooks and Snowpark Python: A Unified Data Engineering Platform

In the fast-moving world of data, organizations are no longer just collecting information—they're leveraging it to drive business...

Jul 13

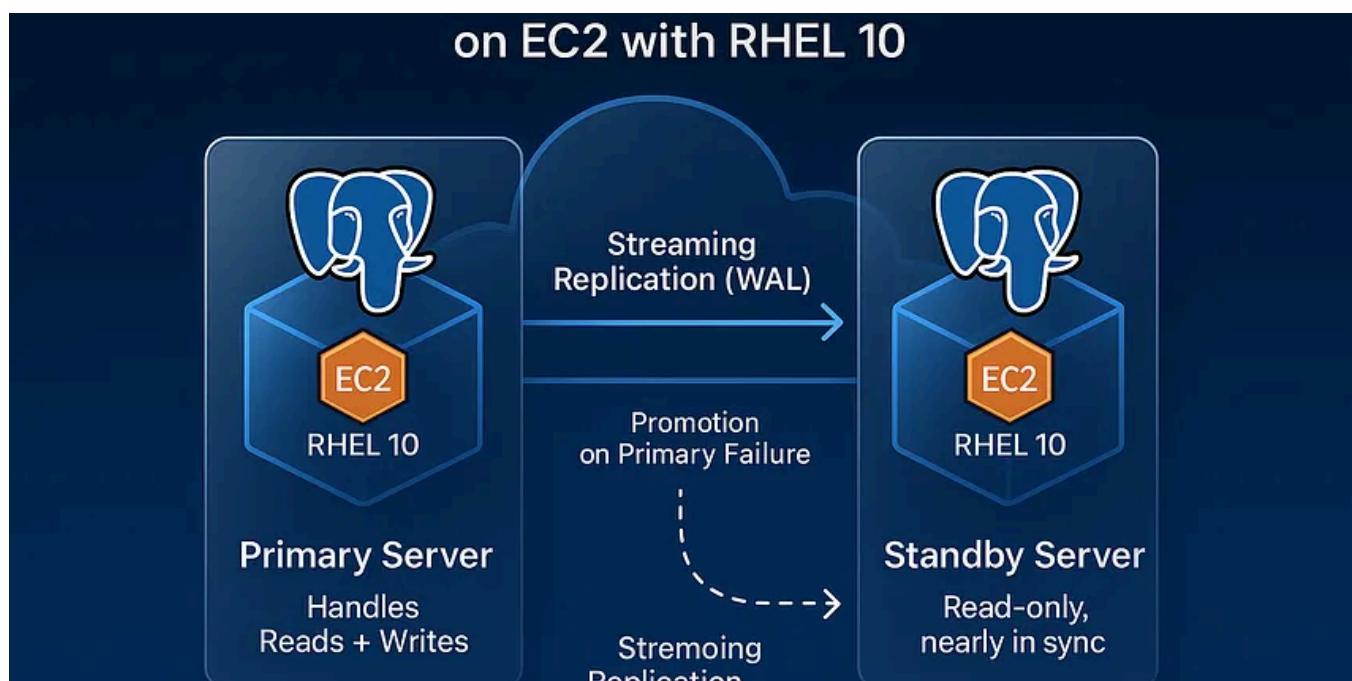


J Jeyaram Ayyalusamy

💡 Essential Techniques Every DBA Should Know: PostgreSQL Backup Strategies

In the world of databases, data is everything—and losing it can cost your business time, money, and trust. Whether you're managing a...

Aug 20 26



 Jeyaram Ayyalusamy 

Streaming Replication in PostgreSQL 17 on EC2 with RHEL 10—A Deep Dive

 Introduction: Why Streaming Replication Matters

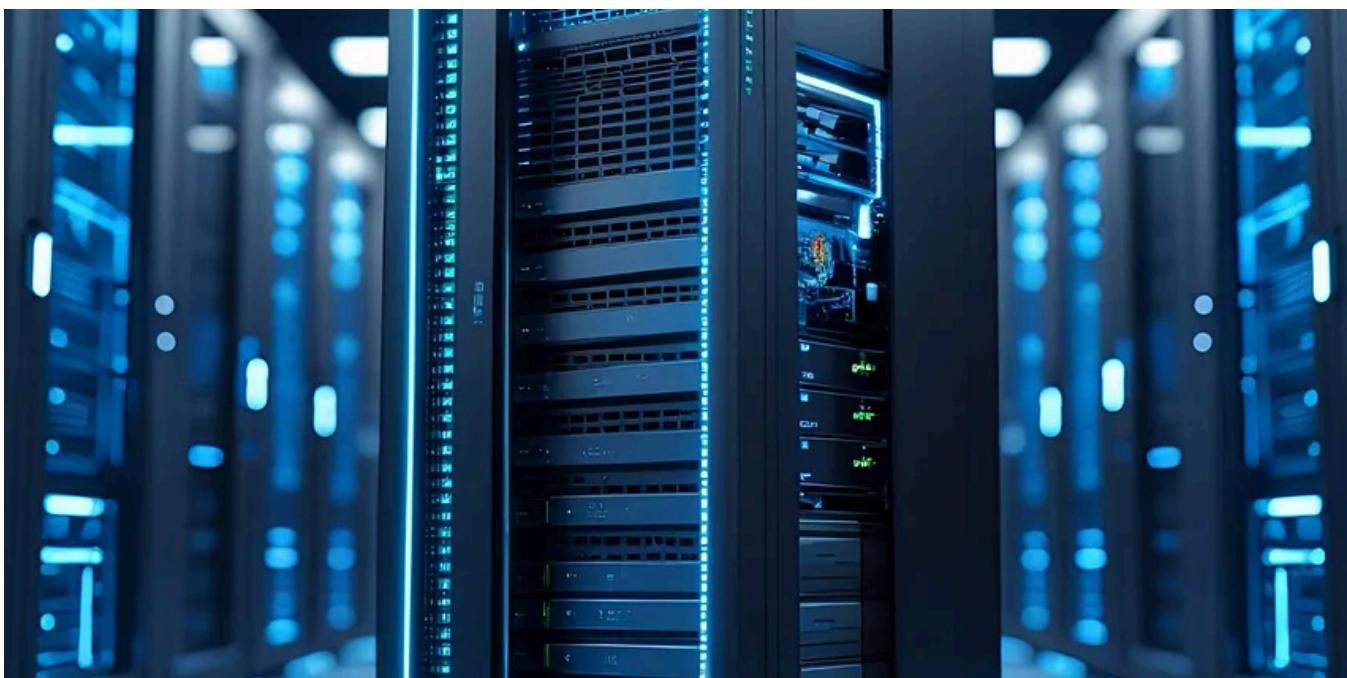
Aug 20  50



...

See all from Jeyaram Ayyalusamy

Recommended from Medium

 Rizqi Mulki

PostgreSQL Index Bloat: The Silent Killer of Performance

How a 10TB database became 3x faster by fixing the invisible problem that plagues 90% of production PostgreSQL deployments

★ Sep 15 11 1

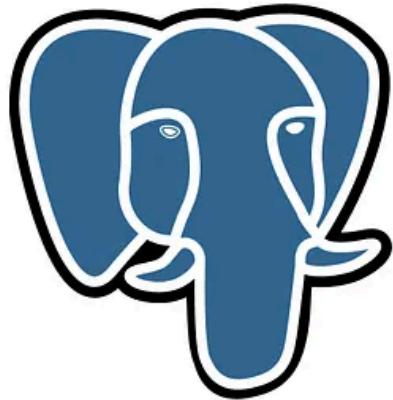
A screenshot of a blog post cover. The title "#PostgreSQL" is at the top left in white, and the subtitle "security" is in blue at the bottom left. The main image is a blue-toned graphic of the PostgreSQL logo. The author's name, "TOMASZ GINTOWT", is at the bottom left of the image. Tomasz Gintowt

Security in PostgreSQL

PostgreSQL is a powerful open-source database. Security is very important when working with sensitive data like passwords, customer...

6d ago 5

...



Beyond Basic PostgreSQL Programmable Objects

In Stackademic by bektiaw

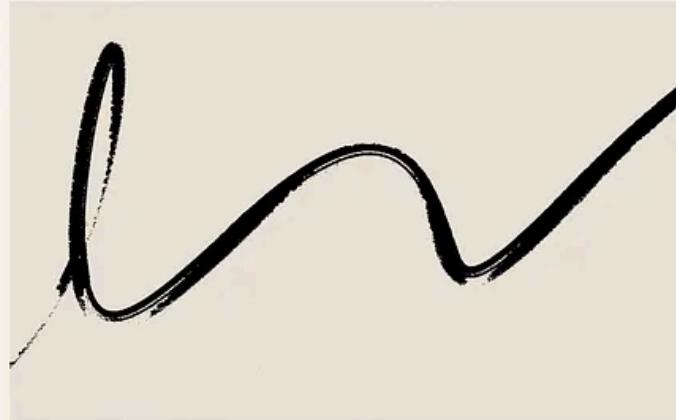
Beyond Basics: PostgreSQL Programmable Objects (Automate, Optimize, Control)

Tired of repeating the same SQL logic? Learn how PostgreSQL can do the work for us.

Sep 1 68 1



...



R Rohan

JSONB in PostgreSQL: The Fast Lane to Flexible Data Modeling 🚀

“Why spin up yet another database just to store semi-structured data?”—Every engineer who discovered JSONB

◆ Jul 18 ⌘ 12 💬 1

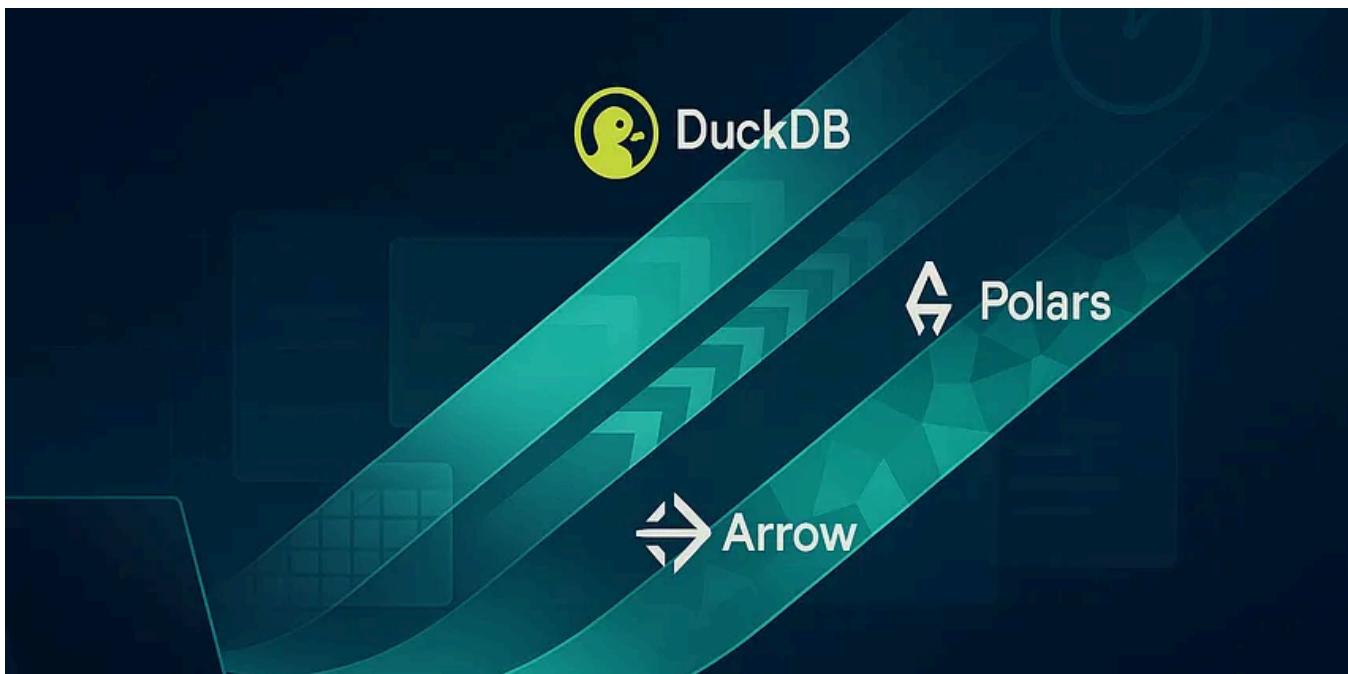


 Vijay Gadhav

Master SQL's QUALIFY Clause for Cleaner, Faster Window Queries

Note: If you're not a medium member, [CLICK HERE](#)

◆ May 20 ⌘ 12 💬 1





Thinking Loop

5 DuckDB–Arrow–Polars Workflows in Minutes

Turn day-long pipelines into small, local, reproducible runs without clusters or drama.

5d ago

14



...

See more recommendations