

[Open in app ↗](#)

Medium



Get unlimited access to the best of Medium for less than \$1/week. [Become a member](#) X

24 - PostgreSQL 17 Performance Tuning: Monitoring Table-Level Statistics with pg_stat_user_tables

4 min read · Sep 7, 2025

J [Jeyaram Ayyalusamy](#)  Following ▾

[Listen](#) [Share](#) [More](#)

J [Follow](#)

Jeyaram Ayyalusamy
158 followers
Oracle DBA | AWS Certified | 18+ yrs experience | DB Admin, RAC, GoldenGate, RDS, MySQL, PostgreSQL | Author of MySQL & RDS books | Cloud & Performance Tuning Expert

 [Book Author](#)



When tuning PostgreSQL, one of the most important steps is to **observe table-level statistics**. You cannot optimize what you cannot measure. PostgreSQL provides the

system view `pg_stat_user_tables`, which tracks how your tables are accessed, how many tuples (rows) are alive or dead, and when maintenance operations like VACUUM last occurred.

This system view is invaluable for performance tuning — when used properly, it can reveal how your queries and workloads truly interact with your data.

Step 1: Setup Demo with a Products Table

Let's first create a large table (`products`) with multiple columns and insert 10 million rows to simulate a real workload.

```
CREATE TABLE products (
    product_id    BIGSERIAL PRIMARY KEY,
    product_name  TEXT,
    category      TEXT,
    price         NUMERIC(10,2),
    stock_qty     INT
);
```



Jeyaram Ayyalusamy

158 followers

Oracle DBA | AWS Certified | 18+ yrs
DB Admin, RAC, GoldenGate, RDSS, MySQL, PostgreSQL | Author of MySQL & RDS books | Cloud & Performance Tuning

 Book Author

```
postgres=# CREATE TABLE products (
    product_id    BIGSERIAL PRIMARY KEY,
    product_name  TEXT,
    category      TEXT,
    price         NUMERIC(10,2),
    stock_qty     INT
);
CREATE TABLE
postgres=#
```

```
-- Insert 10 million rows
INSERT INTO products (product_name, category, price, stock_qty)
SELECT
    'Product_' || g,
    'Category_' || (g % 50),      -- 50 categories
    (random()*500)::NUMERIC(10,2),
    (random()*100)::INT
FROM generate_series(1, 10000000) g;
ANALYZE products;
```

```
postgres=# -- Insert 10 million rows
INSERT INTO products (product_name, category, price, st
SELECT
    'Product_' || g,
    'Category_' || (g % 50),      -- 50 categories
    (random()*500)::NUMERIC(10,2),
    (random()*100)::INT
FROM generate_series(1, 10000000) g;
ANALYZE products;
INSERT 0 10000000
ANALYZE
postgres=#
```



Jeyaram Ayyalusamy

158 followers

Oracle DBA | AWS Certified | 18+ yrs
DB Admin, RAC, GoldenGate, RDSS, MySQL, PostgreSQL | Author of MySQL & RDS books | Cloud & Performance Tuning



Book Author

Now we have a realistic table large enough to observe **sequential scans, index scans, and vacuum behavior**.

Step 2: Generate Some Load

To simulate activity, run queries that update and read from this table:

```
-- Update stock for a specific category
UPDATE products SET stock_qty = stock_qty + 1
```

```
WHERE category = 'Category_10';
```

```
postgres=# -- Update stock for a specific category
UPDATE products SET stock_qty = stock_qty + 1
WHERE category = 'Category_10';
UPDATE 200000
postgres=#

```

J

Jeyaram Ayyalusamy

158 followers

Oracle DBA | AWS Certified | 18+ yrs
DB Admin, RAC, GoldenGate, RDSS, MySQL, PostgreSQL | Author of MySQL & RDS books | Cloud & Performance Tuning

 Book Author

```
postgres=# SELECT * FROM products WHERE product_id = 5000000;
product_id | product_name | category | price | stock_qty
-----+-----+-----+-----+-----+
5000000 | Product_5000000 | Category_0 | 194.44 |      85
(1 row)
```

```
postgres=#

```

```
postgres=# -- Select by category (likely sequential scan if no index exists)
SELECT * FROM products WHERE category = 'Category_2';
```

product_id	product_name	category	price	stock_qty
252	Product_252	Category_2	236.69	73
302	Product_302	Category_2	9.44	50
352	Product_352	Category_2	310.24	24
402	Product_402	Category_2	175.06	52
452	Product_452	Category_2	74.66	91
502	Product_502	Category_2	73.18	98
2402	Product_2402	Category_2	57.82	87
2452	Product_2452	Category_2	136.81	25
2502	Product_2502	Category_2	490.25	12
2552	Product_2552	Category_2	456.85	6
2602	Product_2602	Category_2	215.48	55
2652	Product_2652	Category_2	278.45	53
2702	Product_2702	Category_2	358.29	64
2752	Product_2752	Category_2	178.38	24
2802	Product_2802	Category_2	77.39	23

Running these repeatedly (or using a tool like `pgbench` for load) will generate useful statistics in `pg_stat_user_tables`.

Step 3: Inspect `pg_stat_user_tables`

Now, query the system view to see how PostgreSQL is handling the table.

```
SELECT relname,
       n_live_tup,
       n_dead_tup,
       seq_scan,
       seq_tup_read,
       idx_scan,
       idx_tup_fetch,
       last_vacuum,
       last_autovacuum
  FROM pg_stat_user_tables
 WHERE relname = 'products';
```



Jeyaram Ayyalusamy

158 followers

Oracle DBA | AWS Certified | 18+ yrs experience
DB Admin, RAC, GoldenGate, Redshift, MySQL, PostgreSQL | Author of MySQL & RDS books | Cloud & Performance Tuning

Book Author

Step 4: Interpreting the Results

- `n_live_tup` → current estimate of live rows in the table.
- `n_dead_tup` → number of dead tuples waiting for cleanup (important for monitoring VACUUM effectiveness).
- `seq_scan` / `seq_tup_read` → how many sequential scans occurred and how many rows they processed.
- `idx_scan` / `idx_tup_fetch` → how many index scans occurred and how many rows they fetched.
- `last_vacuum` / `last_autovacuum` → when the table was last vacuumed.

For example, you may observe:

```
postgres=# SELECT relname,
    n_live_tup,
    n_dead_tup,
    seq_scan,
    seq_tup_read,
    idx_scan,
    idx_tup_fetch,
    last_vacuum,
    last_autovacuum
  FROM pg_stat_user_tables
 WHERE relname = 'products';
   relname  | n_live_tup | n_dead_tup | seq_scan | seq_tup_read | idx_scan | idx_tup_fetch | last_vacuum | last_autovacuum
-----+-----+-----+-----+-----+-----+-----+-----+-----+
 products | 9798007 | 606000 | 10 | 60398000 | 5 | 5 |
(1 row)

postgres=#

```

J

Jeyaram Ayyalusamy
 158 followers
 Oracle DBA | AWS Certified | 18+ yrs
 DB Admin, RAC, GoldenGate, RDSS, MySQL, PostgreSQL | Author of MySQL & RDS books | Cloud & Performance Tuning

Book Author

👉 This means:

- Around 50,000 dead tuples exist, which VACUUM will eventually reclaim.
- Sequential scans processed millions of rows — showing that some queries are likely **not using indexes**.

- Index scans fetched 600,000 rows across 6,000 scans — showing frequent lookups by indexed columns.

Why This Matters

- If **dead tuples keep growing** and VACUUM isn't cleaning them efficiently, you risk **table bloat** and performance degradation.
- If **sequential scans dominate**, queries may be missing useful indexes or the planner finds them cheaper due to data distribution.
- Both **number of scans** and **rows fetched** are critical metrics — they reveal not just how queries run, but their total impact on your database.

 By continuously monitoring `pg_stat_user_tables`, you can:

- When to VACUUM or ANALYZE more aggressively.
- Which queries need indexing help.
- Whether workloads are efficiently using PostgreSQL's statistics decisions.



Jeyaram Ayyalusamy

158 followers

Oracle DBA | AWS Certified | 18+ yrs
DB Admin, RAC, GoldenGate, RDSS, MySQL, PostgreSQL | Author of MySQL & RDS books | Cloud & Performance Tuning

 Book Author

📢 Stay Updated with Daily PostgreSQL & Cloud Tips!

If you've been finding my blog posts helpful and want to stay ahead with daily insights on PostgreSQL, Cloud Infrastructure, Performance Tuning, and DBA Best Practices — I invite you to subscribe to my Medium account.

 [Subscribe here](https://medium.com/@jramcloud1/subscribe) 

Your support means a lot — and you'll never miss a practical guide again!

🔗 Let's Connect!

If you enjoyed this post or would like to connect professionally, feel free to reach out to me on LinkedIn:

 [Jeyaram Ayyalusamy](https://medium.com/@jramcloud1/24-postgresql-17-performance-tuning-monitoring-table-level-statistics-with-pg-stat-user-tables-9b281933b03e)

I regularly share content on **PostgreSQL, database administration, cloud technologies, and data engineering**. Always happy to connect, collaborate, and discuss ideas!

[Postgresql](#)[Open Source](#)[Oracle](#)[MySQL](#)[AWS](#)[Following](#)

Written by Jeyaram Ayyalusamy

158 followers · 2 following

Oracle DBA | AWS Certified | 18+ yrs in DB Admin, RAC, GoldenGate, RDS, MySQL & RDS books | Cloud & Performance Expert

**Jeyaram Ayyalusamy**

158 followers

Oracle DBA | AWS Certified | 18+ yrs in DB Admin, RAC, GoldenGate, RDS, MySQL, PostgreSQL | Author of MySQL & RDS books | Cloud & Performance Expert

 Book Author

Responses (1)



Gvadakte

What are your thoughts?



Awsrn
Sep 14

...

Very insightful

 1 [Reply](#)

More from Jeyaram Ayyalusamy

The screenshot shows the AWS EC2 Instances page. The browser tab bar includes 'us-west-2' (active), 'Launch an instance | EC2 | us-west-2', 'Instances | EC2 | us-east-1', and '+'. The main content area is titled 'Instances Info' with a search bar and filters for 'Name', 'Instance ID', 'Instance state', 'Instance type', 'Status check', 'Alarm status', 'Availability Zone', 'Public IPv4 DNS', 'Public IPv4 IP', 'Elastic IP', and 'IPs'. A message states 'No instances' and 'You do not have any instances in this region'. A blue 'Launch instances' button is visible. Below this, a section titled 'Select an instance' is shown.

 Jeyaram Ayyalusamy 

Upgrading PostgreSQL from Version 16 to Version 17 on a Linux Server AWS EC2...

 A Complete Step-by-Step Guide to Installing PostgreSQL 16 on a Linux Server (AWS EC2) with Detailed Explanations)

Aug 4  40

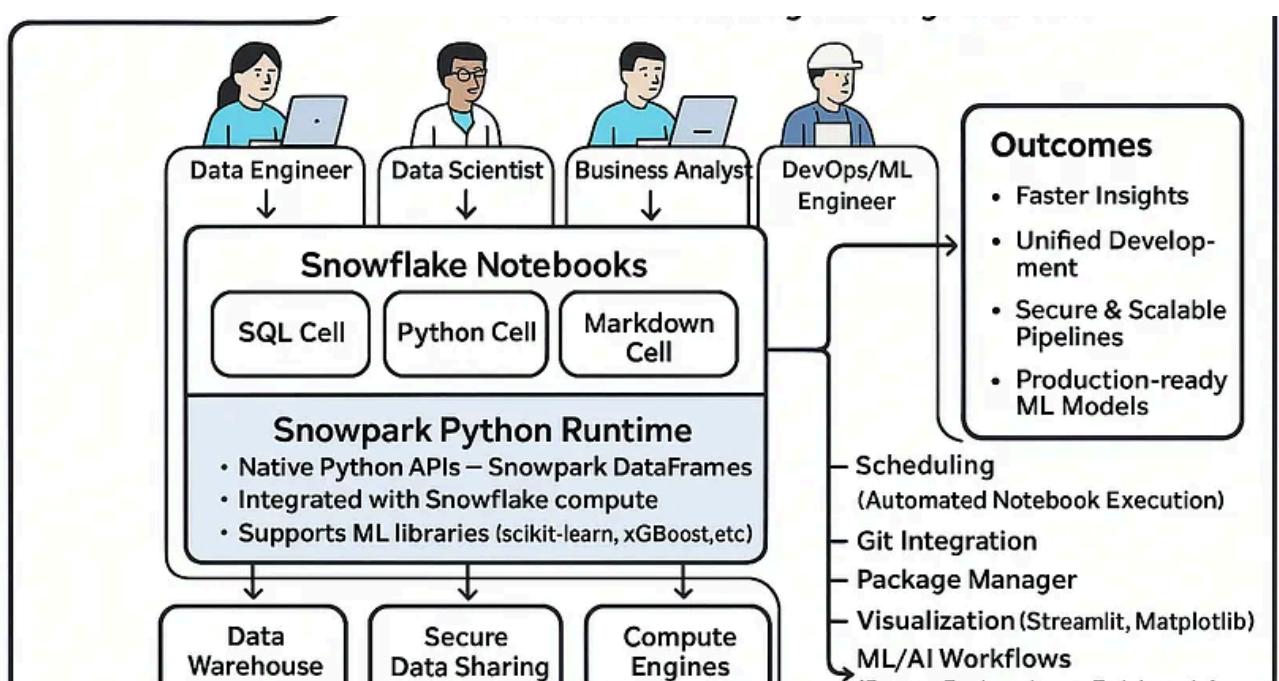


Jeyaram Ayyalusamy

158 followers

Oracle DBA | AWS Certified | 18+ yrs experience | DB Admin, RAC, GoldenGate, RDS, MySQL, PostgreSQL | Author of MySQL & RDS books | Cloud & Performance Tuning Expert

 Book Author

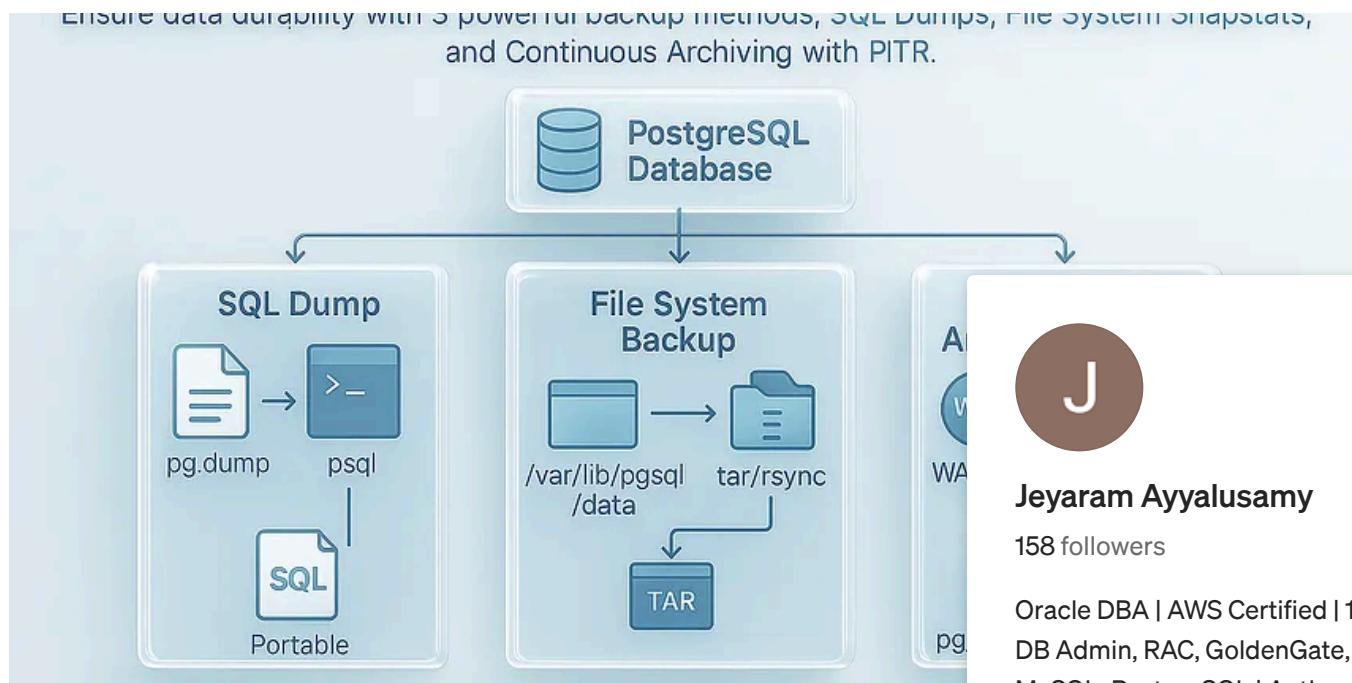


J Jeyaram Ayyalusamy 

16—Experience Snowflake with Notebooks and Snowpark Python: A Unified Data Engineering Platform

In the fast-moving world of data, organizations are no longer just collecting information—they're leveraging it to drive business...

Jul 13

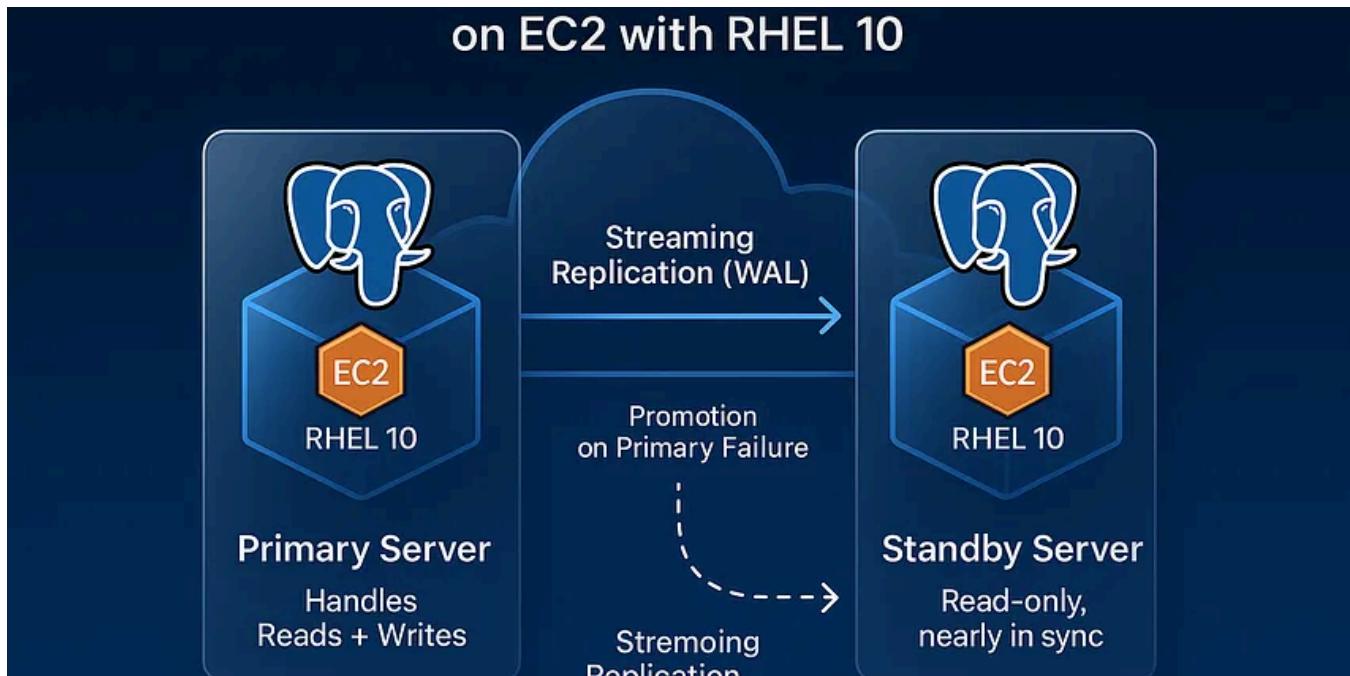


J Jeyaram Ayyalusamy 

💡 Essential Techniques Every DBA Should Know: Performance Strategies

In the world of databases, data is everything—and losing it can cost your business time, money, and trust. Whether you're managing a...

Aug 20  26



J Jeyaram Ayyalusamy

🚀 Streaming Replication in PostgreSQL 17 on EC2 with Deep Dive

🐘 Introduction: Why Streaming Replication Matters

Aug 20 50

[See all from Jeyaram Ayyalusamy](#)



Jeyaram Ayyalusamy

158 followers

Oracle DBA | AWS Certified | 18+ yrs
DB Admin, RAC, GoldenGate, RDSS, MySQL, PostgreSQL | Author of MySQL & RDS books | Cloud & Performance Tuning

Book Author

Recommended from Medium



Rizqi Mulki

PostgreSQL Index Bloat: The Silent Killer of Performance

How a 10TB database became 3x faster by fixing the invisible problem in production PostgreSQL deployments

★ Sep 15 11 1

Jeyaram Ayyalusamy

158 followers

Oracle DBA | AWS Certified | 18+ yrs experience
DB Admin, RAC, GoldenGate, RDS, MySQL, PostgreSQL | Author of MySQL & RDS books | Cloud & Performance Tuning

Book Author



Tomasz Gintowt

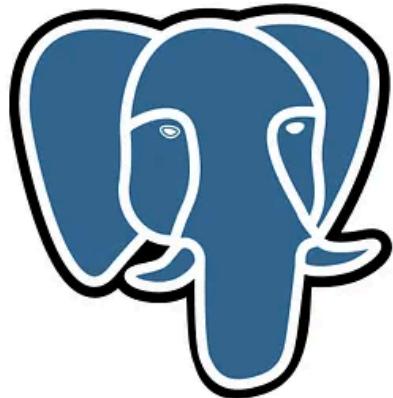
Security in PostgreSQL

PostgreSQL is a powerful open-source database. Security is very important when working with sensitive data like passwords, customer...

6d ago



5



Beyond Basic PostgreSQL Programmable Objects

In Stackademic by bektiaw

Beyond Basics: PostgreSQL Programmable Objects (Optimize, Control)

Tired of repeating the same SQL logic? Learn how PostgreSQL can c

Sep 1 68 1

J

Jeyaram Ayyalusamy

158 followers

Oracle DBA | AWS Certified | 18+ yrs
DB Admin, RAC, GoldenGate, RDSS, MySQL, PostgreSQL | Author of MySQL & RDS books | Cloud & Performance

Book Author



In CodeX by MayhemCode

The Secret Weapons That Separate Bash Beginners from Command Line Legends

Ever wondered how senior developers seem to have magical powers when it comes to bash? They write scripts that validate data, find patterns...

4d ago 7



...



Thinking Loop

5 DuckDB–Arrow–Polars Workflows in Minutes

Turn day-long pipelines into small, local, reproducible runs without c

5d ago 14



Jeyaram Ayyalusamy

158 followers

Oracle DBA | AWS Certified | 18+ yrs
DB Admin, RAC, GoldenGate, Redshift
MySQL, PostgreSQL | Author of MySQL & RDS books | Cloud & Performance

Book Author



R Rohan

JSONB in PostgreSQL: The Fast Lane to Flexible Data

“Why spin up yet another database just to store semi-structured data? discovered JSONB

4 Jul 18 12 1

See more recommendations



Jeyaram Ayyalusamy

158 followers

Oracle DBA | AWS Certified | 18+ yrs
DB Admin, RAC, GoldenGate, RDSS, MySQL, PostgreSQL | Author of MySQL & RDS books | Cloud & Performance Tuning

 Book Author