

[Open in app ↗](#)

Search



Get unlimited access to the best of Medium for less than \$1/week. [Become a member](#)



# PostgreSQL 17 Authentication: How pg\_hba.conf Controls Access Like a Firewall

8 min read · Jun 8, 2025



Jeyaram Ayyalusamy

Following



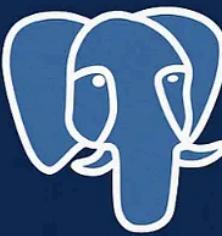
Listen



Share



More



## pg\_hba.conf



### A GUIDE TO POSTGRESQL AUTHENTICATION & SECURITY

- ✓ What pg\_hba.conf does
- ✓ How PostgreSQL authentication works
- ✓ Types of authentication: trust, md5 and scram-sha-256
- ✓ Best practices for securing PostgreSQL access

PostgreSQL is one of the most powerful open-source relational databases, but with great power comes great responsibility — especially when it comes to security. 

In PostgreSQL 17, **internal authentication** is managed primarily through one file:

👉 `pg_hba.conf` (PostgreSQL Host-Based Authentication)

In this guide, we'll explore:

- What `pg_hba.conf` does
- How PostgreSQL authentication works
- Types of authentication: trust, md5, and scram-sha-256
- Best practices for securing PostgreSQL access

Let's dive in 👇

## 📁 What is `pg_hba.conf` ?

You can think of `pg_hba.conf` as **PostgreSQL's built-in firewall and access control list (ACL)**.

Whenever any client tries to connect to your PostgreSQL database, PostgreSQL reads this file to answer four questions:

- 1 Who is trying to connect? → (User account)
- 2 Which database are they trying to access? → (Database name)
- 3 Where are they connecting from? → (Host or IP address)
- 4 How should they authenticate? → (Authentication method)

If the connection request matches a rule in `pg_hba.conf`, PostgreSQL proceeds to authenticate the user accordingly. If no matching rule exists, the connection is refused.

## 🔍 Where is `pg_hba.conf` Located in PostgreSQL?

When you're managing a PostgreSQL server, one of the most critical files you'll encounter is `pg_hba.conf` — the Host-Based Authentication configuration file that controls who can connect, from where, and how.

But where exactly is this file located? Let's break it down for you:

## Default Location of pg\_hba.conf

By default, PostgreSQL stores the `pg_hba.conf` file inside the server's **data directory**. This data directory is where PostgreSQL keeps all of its internal files, configurations, and database data files.

On most Linux-based installations (such as RHEL, CentOS, or Rocky Linux), if you're using PostgreSQL 17 installed from the official PostgreSQL repository, the typical path will be:

```
/var/lib/pgsql/17/data/pg_hba.conf
```

 Here:

- `/var/lib/pgsql/17/data/` is your data directory.
- `pg_hba.conf` is located right inside that folder.

## How to View the File

To view or edit the file, you typically need to switch to the `postgres` user (the dedicated PostgreSQL service account) and use any text viewer or editor:

```
sudo su - postgres
cat /var/lib/pgsql/17/data/pg_hba.conf
```

Or to edit:

```
sudo vi /var/lib/pgsql/17/data/pg_hba.conf
```

## Verify the Exact Path on Your Server

Since different Linux distributions and custom PostgreSQL installations might store data in different directories, you can safely retrieve the actual path PostgreSQL is using by running:

```
SHOW hba_file;
```

Example output:

```
/var/lib/pgsql/17/data/pg_hba.conf
```

- This ensures you're always editing the correct file.

## 🔍 Checking Active Authentication Methods

If you want to quickly see which authentication methods are currently configured, you can filter the file using `grep`:

```
cat /var/lib/pgsql/17/data/pg_hba.conf | grep -i method
```

Or more simply:

```
grep -i method /var/lib/pgsql/17/data/pg_hba.conf
```

This will return lines that mention the word “method”, helping you focus directly on the authentication configurations (like `trust`, `md5`, `scram-sha-256`, etc.).

## 🔴 Don't Forget: Reload Configuration After Changes

After editing `pg_hba.conf`, PostgreSQL won't automatically pick up your changes until you reload the configuration:

```
SELECT pg_reload_conf();
```

Or from shell:

```
sudo systemctl reload postgresql-17
```

 This safely applies your changes without requiring a full server restart.

## Quick Summary

Action	Command
Check default file path	/var/lib/pgsql/17/data/pg_hba.conf
Switch to postgres user	sudo su - postgres
View file contents	cat /var/lib/pgsql/17/data/pg_hba.conf
Locate exact file	SHOW hba_file;
Filter for methods	grep -i method /var/lib/pgsql/17/data/pg_hba.conf
Apply changes	pg_reload_conf() or `systemctl reload postgresql-17`

 *Pro Tip:* Always double-check your file path using SHOW hba\_file; — especially on custom installations, Docker containers, or managed cloud PostgreSQL services.

## The Structure of pg\_hba.conf in PostgreSQL: How Access Rules Work

In PostgreSQL, your first line of defense for controlling who can access your database is a single file: pg\_hba.conf . But to use it effectively, you need to fully understand its structure and how each column defines the rules for database access.

Let's break it down in simple terms:

## The 5 Columns of pg\_hba.conf

Every rule inside `pg_hba.conf` follows a consistent five-column structure:

```
# TYPE  DATABASE  USER  ADDRESS  METHOD
```

Each column plays a crucial role:

Column	Description
<b>TYPE</b>	The type of connection: <code>local</code> (Unix socket) or <code>host</code> (TCP/IP)
<b>DATABASE</b>	The database name this rule applies to, or <code>all</code> for any database
<b>USER</b>	The PostgreSQL user(s) this rule applies to, or <code>all</code> for any user
<b>ADDRESS</b>	The client's IP address (for <code>host</code> connections)
<b>METHOD</b>	The authentication method used (e.g. <code>trust</code> , <code>md5</code> , <code>scram-sha-256</code> , etc.)

## 💡 Let's Look at Each Column in Detail

### 1 TYPE — Connection Type

- `local` : Unix-domain socket connections (typically for connections from the same server).
- `host` : TCP/IP connections (for remote or network-based connections).
- `hostssl` : TCP/IP connections encrypted with SSL.
- `hostnossal` : TCP/IP connections without SSL encryption.

### 👉 Example:

```
host      all      all      127.0.0.1/32      md5
```

This rule applies to any client connecting via TCP/IP (host).

## 2 DATABASE — Which Databases the Rule Applies To

- You can specify:
- A specific database name.
- The keyword `all` to apply the rule to every database.
- `pguser` to match databases named the same as the connecting user.
- `replication` to control replication connections.

👉 Example:

```
host      dvdrental      all      127.0.0.1/32      md5
```

This rule allows connections only to the `dvdrental` database.

## 3 USER — Who Can Connect

- You can list:
- A single PostgreSQL username.
- Multiple usernames separated by commas.
- The keyword `all` for any user.

👉 Example:

```
host      all      appuser,readonlyuser      127.0.0.1/32      md5
```

This allows only `appuser` and `readonlyuser` to connect.

## 4 ADDRESS — The Client's IP Address or Subnet

- For host rules, this defines:
  - The client IP address ( 127.0.0.1 for localhost).
  - The subnet ( 192.168.11.0/24 for entire network ranges).
  - 0.0.0.0/0 for all IPv4 addresses.
  - ::1/128 for IPv6 localhost.

👉 Example:

```
host      all      all      192.168.11.0/24      md5
```

This allows connections from any client in the 192.168.11.x subnet.

## 5 METHOD — Authentication Method

- trust — No password required (not secure for production).
- reject — Always deny access.
- md5 — Use MD5 password-based authentication.
- scram-sha-256 — More secure, recommended for PostgreSQL 10+.
- peer — Use system user identity (for local connections).
- password — Send password in clear text (rarely used anymore).

👉 Example:

```
host      all      all      127.0.0.1/32      scram-sha-256
```

Requires SCRAM-based authentication.

## ✍ Full Example Rule

Let's look at a full rule:

```
host      all      all      127.0.0.1/32      md5
```

📝 This means:

- Allow connections over TCP/IP ( host )
- For all databases
- From all users
- Only from localhost (127.0.0.1)
- Using MD5 password authentication

## 💡 Important Notes

- Rules are evaluated **top-down** — PostgreSQL applies the first rule that matches.
- Always place more restrictive rules first.
- After modifying `pg_hba.conf`, reload the configuration:

```
SELECT pg_reload_conf();
```

Or via shell:

```
sudo systemctl reload postgresql-17
```

## Summary Table

Column	Example	Meaning
TYPE	host	TCP/IP connections
DATABASE	all	Applies to all databases
USER	all	Applies to all users
ADDRESS	127.0.0.1/32	Localhost only
METHOD	md5	Requires MD5 password

 **Pro tip:** pg\_hba.conf is one of the most powerful access control tools in PostgreSQL. Mastering its structure is key to securing your database.

 **PostgreSQL Authentication Methods Explained (for Beginners & DBAs)**

PostgreSQL gives you full control over who can access your database, from where, and with which credentials. This flexibility is managed using the `pg_hba.conf` file, where you configure your **authentication methods**.

Let's explore the most common authentication methods in PostgreSQL — when to use them, how to configure them, and how to test connections.



## 1 Trust Authentication — “No Password Mode”

**Summary:**

- No password required
- Extremely risky for production — use only in local development or isolated environments

**How it works:**

With `trust` authentication, PostgreSQL simply trusts any client connecting that matches the rule. No password is ever requested.

**Example Configuration in `pg_hba.conf`:**

```
host      all      all      127.0.0.1/32      trust
```

- This rule allows *any* user to connect to *any* database from `127.0.0.1` (localhost), without providing a password.

**Testing the Connection:**

```
psql postgres -h 127.0.0.1 -U postgres
```

You will not be prompted for a password.

**What happens if remote clients try?**

```
psql postgres -h <remote_IP> -U postgres
```

👉 Result:

```
FATAL: pg_hba.conf rejects connection for host "<remote_IP>", user "postgres",
```



### 🔒 Important Note:

Never use `trust` in production environments unless you're 100% certain the machine is isolated and fully protected.

## 🔑 2 MD5 Authentication — “Password with MD5 Hashing”

### Summary:

- Password required
- Uses MD5 hashing to store & verify passwords
- Safer than `trust`, but older and less secure than modern alternatives

### How it works:

PostgreSQL stores user passwords hashed with MD5, and verifies them on login.

### Example Configuration in `pg_hba.conf`:

```
host      all      all      127.0.0.1/32      md5
```

- Requires clients to supply a password.

### Testing the Connection:

```
psql postgres -h <remote_IP> -U postgres
```



Result:  
You'll be prompted for a password:

```
Password for user postgres:
```

## ⚠ MD5 Considerations:

- MD5 is widely supported but considered outdated.
- Modern security best practices prefer stronger hashing methods like SCRAM.

## 🔒 3 SCRAM-SHA-256 Authentication — “Modern, Recommended Approach”

Summary:

- Strong, salted password hashing
- Supported since PostgreSQL 10 (default in PostgreSQL 13+)
- Recommended for production use

## Why SCRAM?

SCRAM-SHA-256 (Salted Challenge Response Authentication Mechanism) uses modern hashing algorithms, adding stronger protection against password leaks and attacks.

## 🔧 How to Set Up SCRAM Authentication

### Step 1 : Enable SCRAM Password Encryption

Inside your PostgreSQL session:

```
SET password_encryption = 'scram-sha-256';
```

- ✓ This tells PostgreSQL to store future passwords using SCRAM-SHA-256.

## Step 2 : Set (or reset) the User Password

For example, to update the `postgres` user:

```
ALTER USER postgres WITH PASSWORD 'your_password';
```

- 👉 The new password is now hashed with SCRAM-SHA-256.

## Step 3 : Update `pg_hba.conf` to Require SCRAM

```
host      all      all      <your_IP>/32      scram-sha-256
```

- Replace `<your_IP>` with the client's IP address or subnet.

## Step 4 : Reload Configuration

```
SELECT pg_reload_conf();
```

or from shell:

```
sudo systemctl reload postgresql-17
```

## Testing SCRAM Authentication:

```
psql postgres -h <remote_IP> -U postgres
```

👉 You'll be prompted for a password just like MD5 — but now using strong SCRAM-based verification.



## Quick Comparison Table

Method	Password?	Security Level	Use Case
trust	✗	🚫 Low	Only dev/test
md5	✓	⚠️ Medium	Legacy systems
scram-sha-256	✓	✓ High	Recommended production

🔒 **Best Practice:** Always use `scram-sha-256` for production systems on PostgreSQL 17 for better password security.

## ▶ Best Practices for PostgreSQL Authentication

- 🔒 Always prefer `scram-sha-256` for production.
- 🚫 Avoid `trust` outside of isolated dev machines.
- 🔎 Review `pg_hba.conf` regularly as part of security audits.
- 🔍 Reload configuration after changes (no downtime required).
- 📜 Restrict IP addresses carefully in the `ADDRESS` field.

## 🔒 Why PostgreSQL Authentication Is So Powerful

Unlike many databases, PostgreSQL gives you full flexibility to:

- 🔒 Control every incoming connection
- 🔒 Define granular user/database rules
- 🔒 Combine with SSL/TLS encryption
- 🔒 Prevent unauthorized access at the connection level

Properly configured `pg_hba.conf` is your **first line of defense** before authentication even happens.

## 🏁 Conclusion

PostgreSQL's internal authentication via `pg_hba.conf` is simple, powerful, and incredibly flexible — but only when used correctly.

- ✅ Use `scram-sha-256` for production-grade password security.
- ✅ Keep `pg_hba.conf` clean, specific, and locked down.
- ✅ Reload configuration after edits, no restart needed.

-  Remember: Database security starts before the connection even begins!

 PostgreSQL 17 continues to deliver serious enterprise-grade security – but only if you configure it properly. Master `pg_hba.conf` and you master your database perimeter.

 If you found this guide helpful, follow me(medium) for more practical PostgreSQL tutorials, database architecture guides, and hands-on DBA content.

## Let's Connect!

If you enjoyed this post or would like to connect professionally, feel free to reach out to me on LinkedIn:

 [Jeyaram Ayyalusamy](#)

I regularly share content on PostgreSQL, database administration, cloud technologies, and data engineering. Always happy to connect, collaborate, and discuss ideas!

Postgresql

Sql

Open Source

Security

Database



Following 

## Written by Jeyaram Ayyalusamy

76 followers · 2 following

Oracle DBA | AWS Certified | 18+ yrs in DB Admin, RAC, GoldenGate, RDS, MySQL, PostgreSQL | Author of MySQL & RDS books | Cloud & Performance Expert



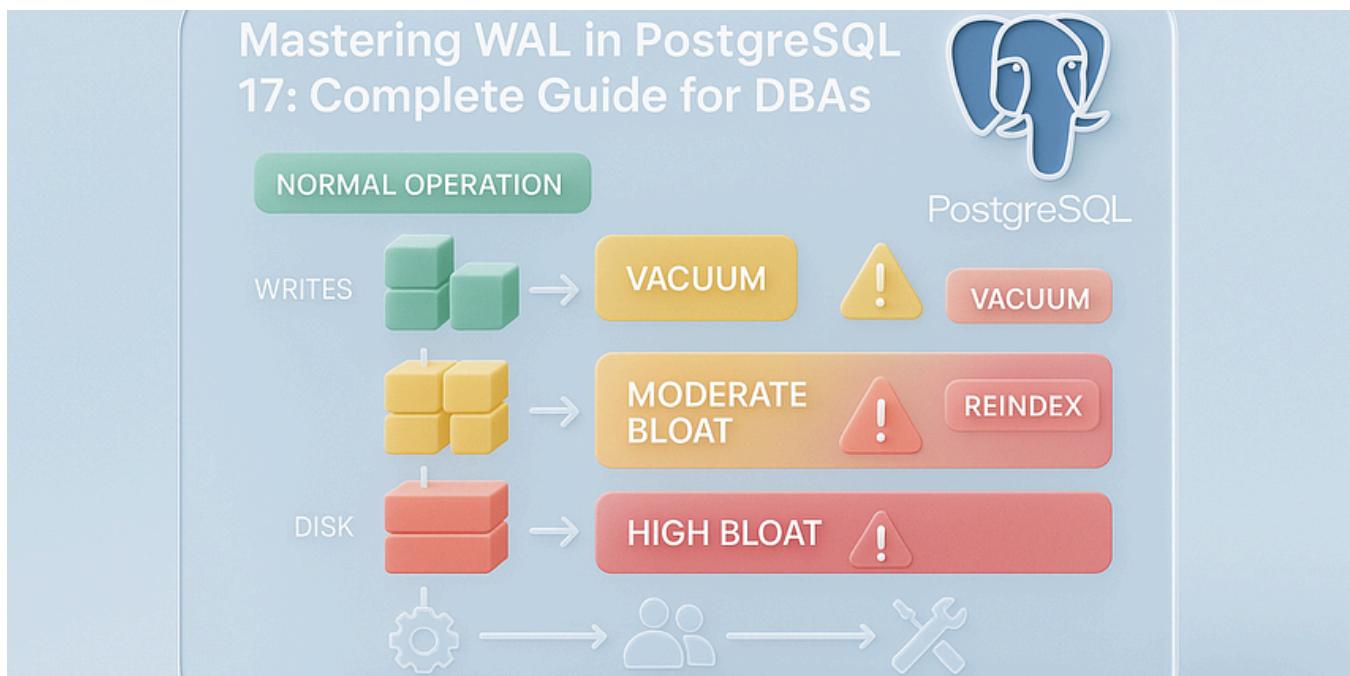
## No responses yet



Gvadakte

What are your thoughts?

## More from Jeyaram Ayyalusamy



J Jeyaram Ayyalusamy

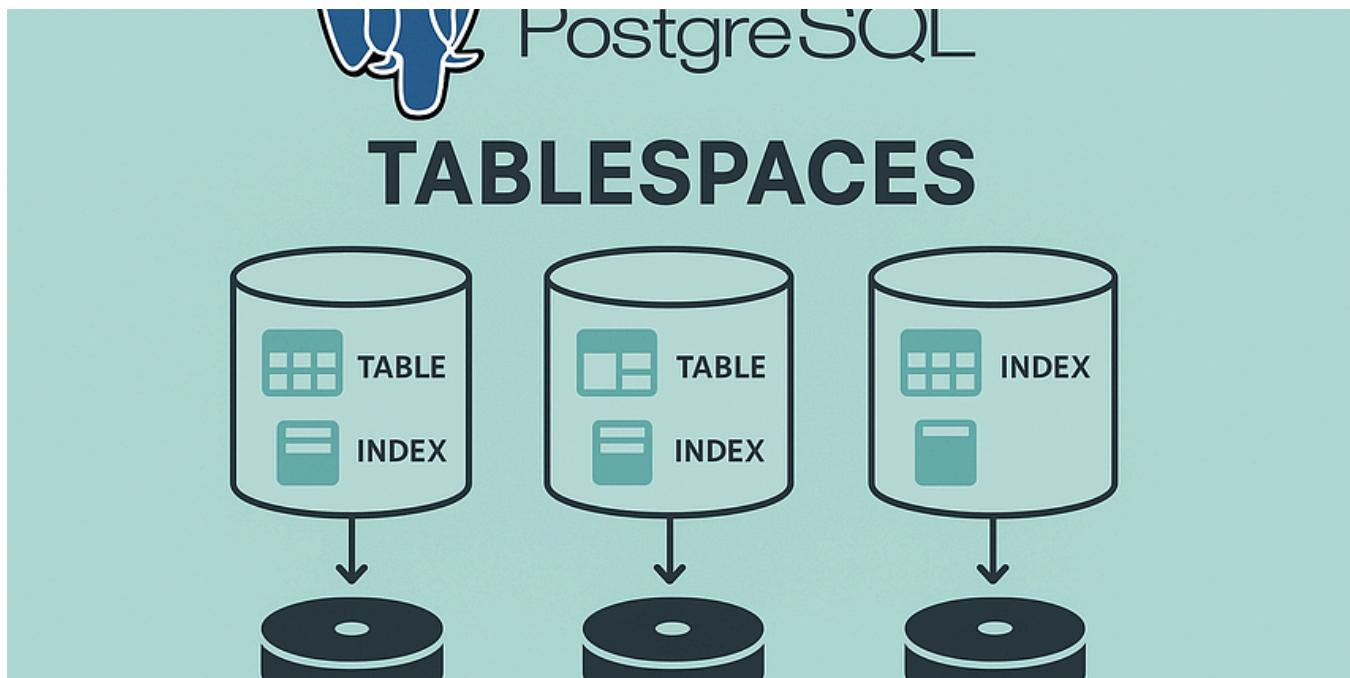
## Understanding and Managing Bloating in PostgreSQL: A Complete Guide for DBAs

PostgreSQL is a powerful, reliable, and feature-rich open-source relational database system. It's praised for its extensibility, ACID...

Jun 25 52



...



J Jeyaram Ayyalusamy

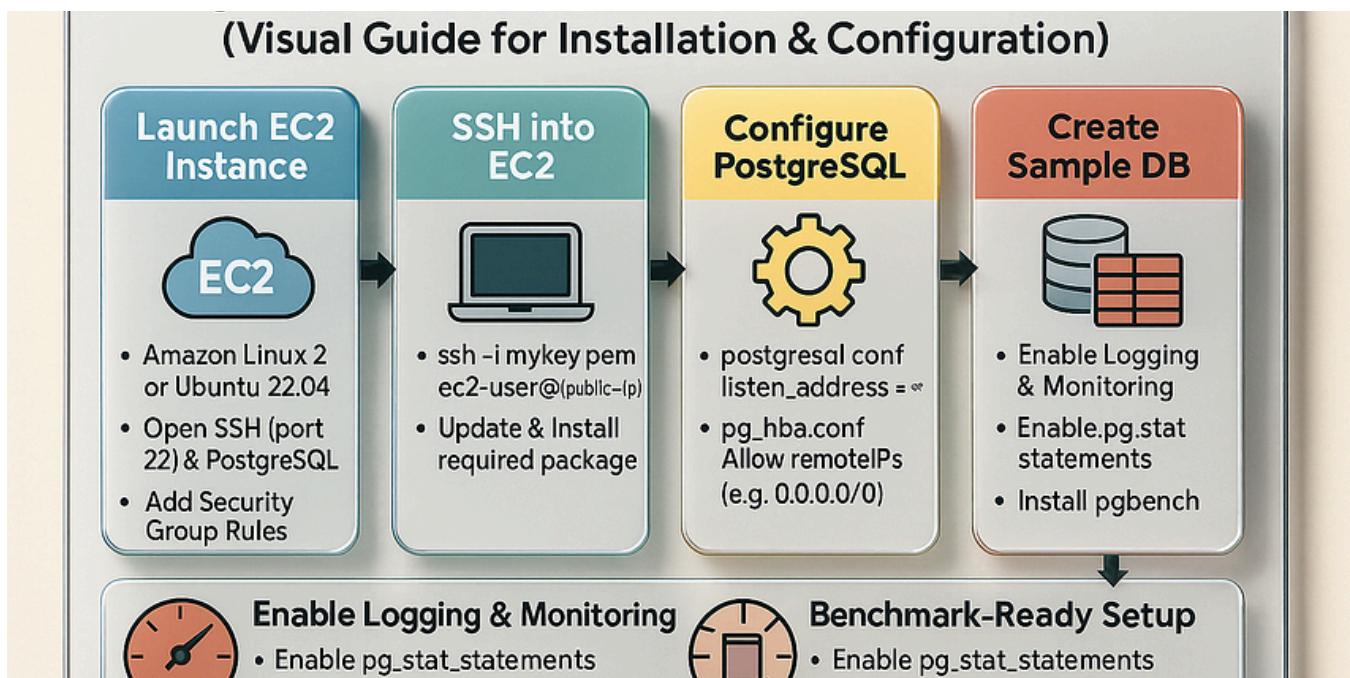
## PostgreSQL Tablespaces Explained: Complete Guide for PostgreSQL 17 DBAs

PostgreSQL offers a powerful feature called tablespaces, allowing database administrators to take control of how and where data is...

Jun 12 8



...



J Jeyaram Ayyalusamy

## PostgreSQL 17 on AWS EC2—Full Installation & Configuration Walkthrough

Setting up PostgreSQL 17 on AWS EC2 might seem complex at first—but once you break down each component, it becomes an efficient and...

1d ago 50



...



J Jeyaram Ayyalusamy

## How to Install PostgreSQL 17 on Red Hat, Rocky, AlmaLinux, and Oracle Linux (Step-by-Step Guide)

PostgreSQL 17 is the latest release of one of the world's most advanced open-source relational databases. If you're using a Red Hat-based...

Jun 3



...

See all from Jeyaram Ayyalusamy

## Recommended from Medium

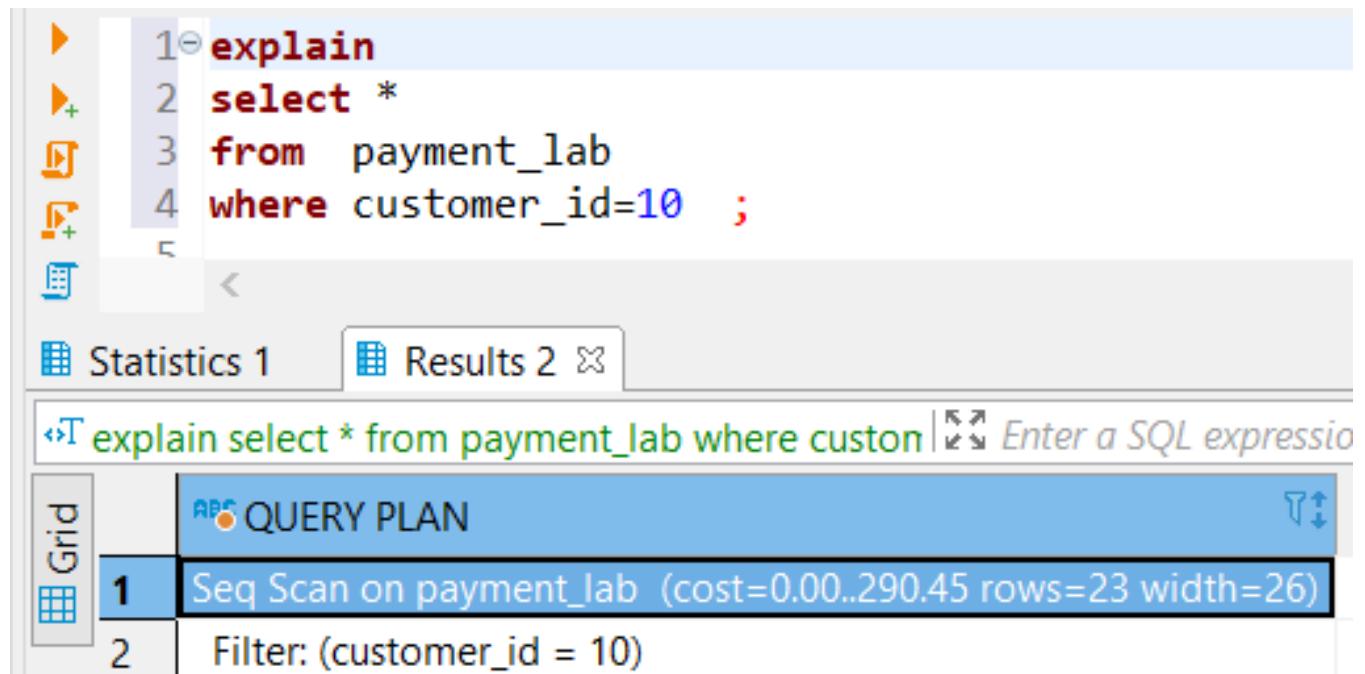
 Rizqi Mulki

## Full-Text Search in PostgreSQL: Better Than You Think

Why developers are abandoning Elasticsearch for PostgreSQL's built-in search—and how it handles 10M records without breaking a sweat

5d ago 2

[+]



The screenshot shows a PostgreSQL query editor interface. At the top, there is a code editor with the following SQL query:

```
1 explain
2 select *
3 from payment_lab
4 where customer_id=10 ;
```

Below the code editor are two tabs: "Statistics 1" and "Results 2". The "Results 2" tab is currently selected, showing the "QUERY PLAN" section. The plan details a "Seq Scan on payment\_lab (cost=0.00..290.45 rows=23 width=26)" followed by a "Filter: (customer\_id = 10)".

At the bottom of the interface, there is a search bar containing the text "explain select \* from payment\_lab where custom" and a placeholder "Enter a SQL expression".

 Muhammet Kurtoglu

## Postgresql Query Performance Analysis

SQL tuning is critically important for ensuring database performance, reliability, and scalability. In most cases, performance issues in a...

6d ago



10



Azlan Jamal

## Stop Using SERIAL in PostgreSQL: Here's What You Should Do Instead

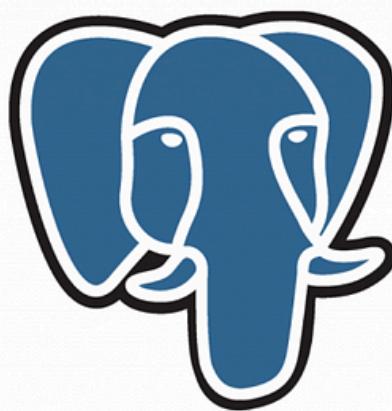
In PostgreSQL, there are several ways to create a PRIMARY KEY for an id column, and they usually involve different ways of generating the...



Jul 12



33



# PostgreSQL

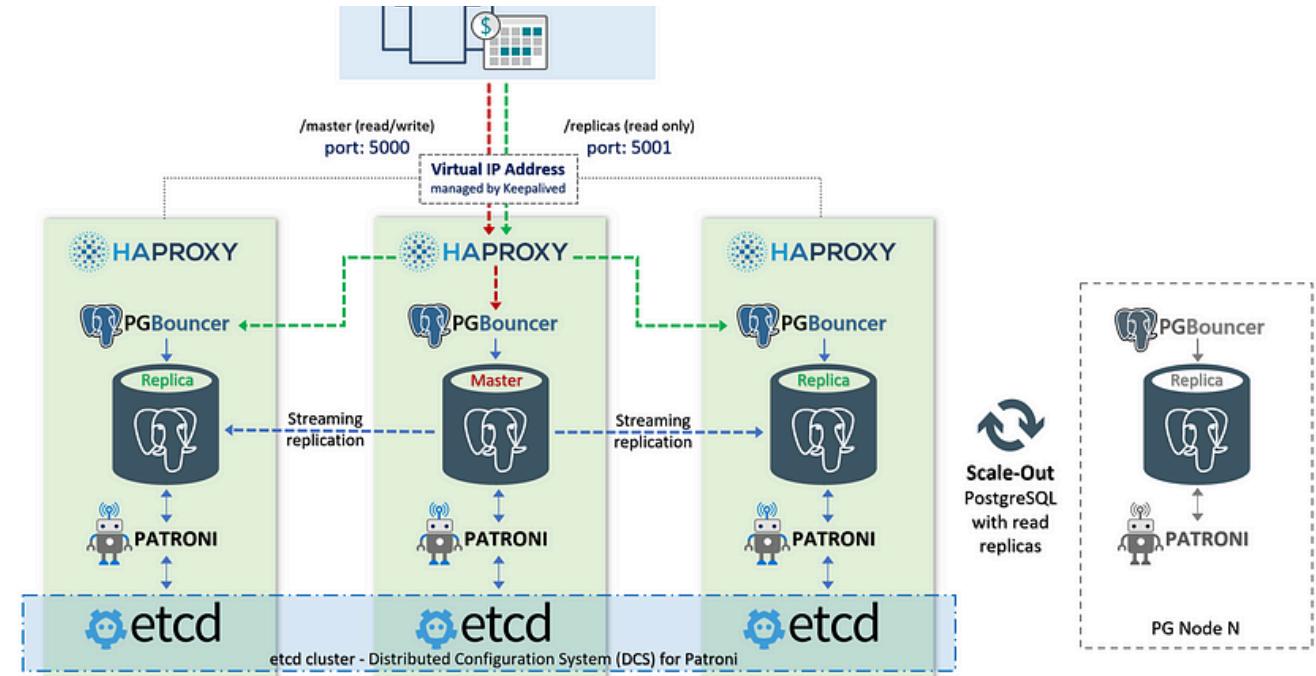


Harishsingh

# PostgreSQL 18 in Microservices: You Don't Need a Separate DB for Everything

# Introduction: The Myth of Database-Per-Service

Jul 13 11 1



 Kanat Akylson

## **How to Deploy a High-Availability PostgreSQL Cluster in 5 Minutes Using Ansible and Patroni**

This tutorial shows how to spin up a production-grade HA PostgreSQL cluster in just 5 minutes using m2y ready-made GitHub repository with...

Jun 9 65 1

 DevTrendsDaily

## How We Tuned PgBouncer for 10,000 RPS with Just 50 Connections

(And Why It Worked Better Than pgxpool Ever Did)

 Jun 18  11

...

[See more recommendations](#)