

[Open in app ↗](#)**Medium**

Search



Get unlimited access to the best of Medium for less than \$1/week. [Become a member](#)



Understanding and Managing Bloating in PostgreSQL: A Complete Guide for DBAs

16 min read · Jun 24, 2025



Jeyaram Ayyalusamy



Following



Listen



Share



More



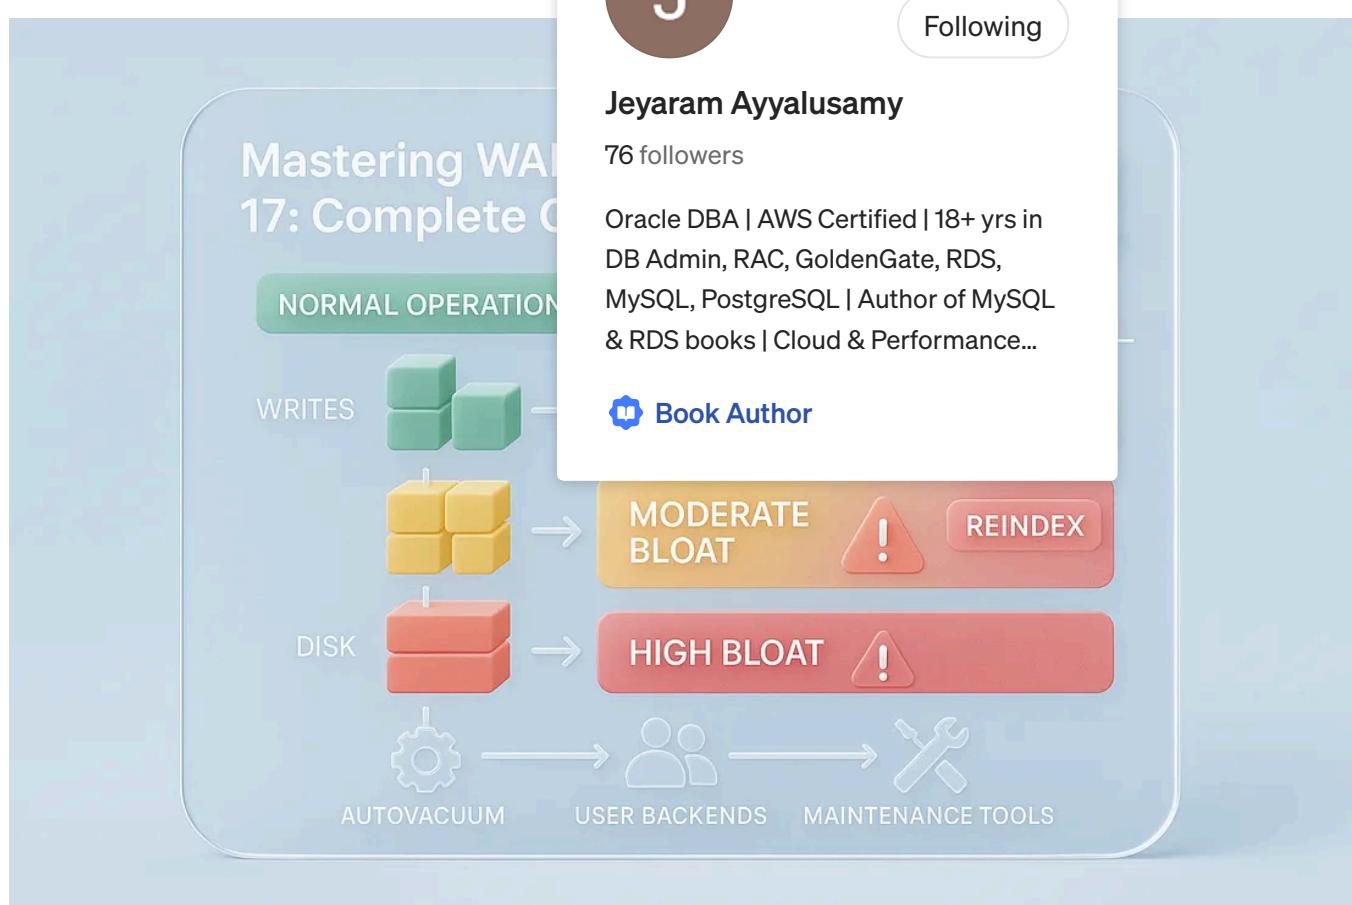
Following

Jeyaram Ayyalusamy

76 followers

Oracle DBA | AWS Certified | 18+ yrs in DB Admin, RAC, GoldenGate, RDS, MySQL, PostgreSQL | Author of MySQL & RDS books | Cloud & Performance...

Book Author



PostgreSQL is a powerful, reliable, and feature-rich open-source relational database system. It's praised for its extensibility, ACID compliance, and performance under large-scale workloads. However, like any complex database engine, PostgreSQL also

has its internal housekeeping challenges — and one of the most common among them is **bloating**.

What is Bloating?

Bloating in PostgreSQL refers to the **excess and unnecessary disk space consumption** in tables and indexes due to PostgreSQL's internal architecture and data modification processes. It's not a bug or malfunction — it's a byproduct of how PostgreSQL handles updates and deletes using its **MVCC (Multi-Version Concurrency Control)** model.

When you delete or update rows in PostgreSQL:

- The original rows are not immediately removed from disk.
- Instead, they're marked as dead and remain in the table or index pages until reclaimed.

Over time, these dead tuples accu...

- Increased table and index size
- Decreased query performance
- Less effective use of cache and...

This is known as **table bloat** or **index bloat**.



Jeyaram Ayyalusamy

76 followers

Oracle DBA | AWS Certified | 18+ yrs in
DB Admin, RAC, GoldenGate, RDS,
MySQL, PostgreSQL | Author of MySQL
& RDS books | Cloud & Performance...

 Book Author

Why Does Bloating Occur?

PostgreSQL uses MVCC to ensure consistency in concurrent transactions. While this design provides excellent performance and isolation, it also means PostgreSQL keeps multiple versions of rows:

- A **new version** is written during every update.
- The **old version** is retained until it is cleaned up by a **VACUUM** process.
- Indexes still point to old/dead tuples until they are rebuilt.

If regular VACUUM operations (manual or autovacuum) are not tuned properly or cannot keep up with write activity, bloat can grow unchecked.

Common causes include:

- Heavy update/delete workloads.
- Long-running transactions preventing VACUUM from cleaning.
- Poor autovacuum settings.
- Lack of regular maintenance routines like reindexing.

What This Article Covers

In this article, we'll walk you through effectively:

1. **What is Bloat?** — A technical deep dive into indexes.
2. **Why Bloat Matters** — Understanding the implications.
3. **How to Detect Bloat** — Using PostgreSQL's pgstattuple .
4. **How to Fix and Prevent It** — Including VACUUM , REINDEX , and best practices.

By the end of this article, you'll be equipped with:

- The knowledge to identify bloat before it becomes a problem.
- Practical tools and commands to manage it.
- Strategies to maintain long-term PostgreSQL performance and stability.

Bloat may be subtle, but its impact on performance and storage is real. Fortunately, with the right awareness and maintenance strategies, it's entirely manageable.



Jeyaram Ayyalusamy

76 followers

Oracle DBA | AWS Certified | 18+ yrs in DB Admin, RAC, GoldenGate, RDS, MySQL, PostgreSQL | Author of MySQL & RDS books | Cloud & Performance...

Book Author

ile bloating

oles and

implications.

sions like

Let's dive deeper into the causes and solutions for PostgreSQL bloat.

► Primary Causes of Bloating in PostgreSQL

While PostgreSQL offers exceptional performance and data integrity features, its internal architecture can lead to **bloating** — an inefficient use of disk space in both tables and indexes. Understanding what causes this bloat is essential for maintaining a healthy and high-performing database.

Let's dive into the **primary culprits behind PostgreSQL bloat**:

1 Dead Tuples: The Silent Accomplice

PostgreSQL uses **Multi-Version Concurrency Control** (MVCC) to allow concurrent access to data without locking. While this maintains data consistency, it also introduces sources of bloat.

How it works:

- When a row is **updated**, PostgreSQL creates a **new version** of the row instead of overwriting the old one. Instead, it
- When a row is **deleted**, it is not immediately removed from disk; it is simply marked as **invisible** to future transactions.
- These dead rows are only cleaned up later by **VACUUM** or **autovacuum**.

The problem:

If **VACUUM** cannot run frequently enough (e.g., due to long-running transactions or misconfigured autovacuum settings), dead tuples accumulate in tables and indexes. These leftover versions inflate table size, reduce index effectiveness, and degrade performance.



Jeyaram Ayyalusamy

76 followers

Oracle DBA | AWS Certified | 18+ yrs in
DB Admin, RAC, GoldenGate, RDS,
MySQL, PostgreSQL | Author of MySQL
& RDS books | Cloud & Performance...

Book Author

e concurrent
urrency and
ources of

. Instead, it

2 Free Space Map Fragmentation

PostgreSQL uses a data structure called the **Free Space Map (FSM)** to keep track of available space in table pages. When rows are deleted or updated, the database relies on the FSM to identify where new rows can be inserted without allocating new pages.

The issue:

- Over time, as rows are frequently added and removed, the FSM becomes **fragmented**.
- This fragmentation leads to **suboptimal space utilization**, where many pages may have small amounts of unused space that are too fragmented to be reused efficiently.
- The result is a bloated table file, even if the logical number of rows remains the same.

Why it matters:

Fragmentation in FSM doesn't just affect performance because more blocks need to be read or written for each operation.



Jeyaram Ayyalusamy

76 followers

Oracle DBA | AWS Certified | 18+ yrs in
DB Admin, RAC, GoldenGate, RDS,
MySQL, PostgreSQL | Author of MySQL
& RDS books | Cloud & Performance...



Book Author

because more

critical for
hidden

3 Unindexed Foreign Keys

Foreign keys ensure referential integrity, but they can also lead to performance and space problems.

What happens:

- When a referenced row is deleted or updated in the parent table, PostgreSQL must **check the child table** to enforce the foreign key constraint.
- If the foreign key column in the child table is **not indexed**, this check results in **full table scans**, which are expensive and create **temporary bloat**, especially under high DML workloads.
- It may also slow down `VACUUM` or `DELETE` operations, indirectly contributing to the accumulation of dead tuples.

Best practice:

Always create an index on any column used in a foreign key constraint. This not only improves performance but also reduces the risk of bloat through unnecessary

full scans.

4 High Update Frequency: The Transactional Pressure Cooker

In highly transactional systems — like e-commerce platforms, financial services apps, or event logging systems — certain tables receive a **continuous stream of updates**. These can be subtle changes like modifying a status flag or incrementing a counter.

The outcome:

- Every `UPDATE` operation creates a **new row version** and leaves the old one behind.
- As the frequency increases, so does the number of versions.
- Even if autovacuum is enabled, it may not keep up, especially if transactions are long-running.

Real-world example:

A table storing `user_sessions` or similar data might receive several updates per minute. Without regular `VACUUM` or `ANALYZE` operations, it can quickly become bloated over time, a matter of hours or days.



Jeyaram Ayyalusamy

76 followers

Oracle DBA | AWS Certified | 18+ yrs in DB Admin, RAC, GoldenGate, RDS, MySQL, PostgreSQL | Author of MySQL & RDS books | Cloud & Performance...

 Book Author

les.

pace,

nds of times
table bloat in

Summary

Understanding the **root causes of bloat** helps you stay ahead of performance issues and wasted disk space. Here's a quick recap:

Cause	Description
Dead Tuples	Old row versions left behind by <code>UPDATE/DELETE</code> operations
Free Space Fragmentation	Inefficient reuse of space within pages due to FSM fragmentation
Unindexed Foreign Keys	Lack of indexes on FK columns leads to table scans and vacuum delays
High Update Frequency	Frequent row updates rapidly generate bloat if not controlled by autovacuum

By proactively addressing these issues, you'll ensure that your PostgreSQL environment remains **lean, fast, and scalable**.

🛠 Strategies to Control and Prevent Bloating in PostgreSQL

While bloating is a natural byproduct of PostgreSQL's MVCC (Multi-Version Concurrency Control) model, it doesn't have to become a performance bottleneck. PostgreSQL provides a robust set of tools and practices to help you **control, reduce, and prevent** table and index bloat.

Let's explore the most effective strategies to manage bloat proactively:

✓ VACUUM: Clean Up Dead T

VACUUM is PostgreSQL's built-in tool to clean up dead rows and free space. It scans through the database to identify rows that have been updated or deleted but still occupy space in the table's data pages. It then marks these rows as "dead" and removes them from the table's data map and the free space map (FSM).

Key Benefits:

- Reclaims storage from dead rows.
- Prevents transaction ID wraparound issues.
- Improves performance by reducing the amount of “junk” PostgreSQL has to scan.

Example:

```
VACUUM user_activity_log_1;
```

For a more aggressive cleanup (which also updates statistics), use:



Jeyaram Ayyalusamy

76 followers

Oracle DBA | AWS Certified | 18+ yrs in DB Admin, RAC, GoldenGate, RDS, MySQL, PostgreSQL | Author of MySQL & RDS books | Cloud & Performance...

Book Author

old row
es visibility
e inserts.
e size).

```
VACUUM FULL user_activity_log_1;
```

⚠ Note: `VACUUM FULL` requires an exclusive lock and rewrites the table, so use it during maintenance windows.

```
postgres=# VACUUM user_activity_log_1;
VACUUM
postgres=#
postgres=# VACUUM FULL user_activity_log_1;
VACUUM
postgres=#
```



Jeyaram Ayyalusamy

76 followers

ns VACUUM

Oracle DBA | AWS Certified | 18+ yrs in
DB Admin, RAC, GoldenGate, RDS,
MySQL, PostgreSQL | Author of MySQL
& RDS books | Cloud & Performance...

Book Author

✓ AUTOVACUUM: Set It and Forget It

PostgreSQL includes an `autovacuum` (and `ANALYZE`) based on activity thresholds.

Why It Matters:

- Runs in the background without blocking other queries.
- Ensures dead tuples are cleaned up.
- Keeps statistics fresh for the query planner.

However, **autovacuum settings must be properly tuned**. On high-write tables, the default thresholds may be too conservative.

Recommended Actions:

- Ensure `autovacuum` is enabled (`track_counts = on`, `autovacuum = on`).
- Adjust settings like:
 - `autovacuum_vacuum_threshold`
 - `autovacuum_vacuum_scale_factor`
 - `autovacuum_naptime`

- Monitor logs and `pg_stat_user_tables` to ensure autovacuum is triggering as expected.

REINDEX: Fight Index Bloat Directly

While `VACUUM` addresses table bloat, it **does not reclaim space** in indexes. Over time, especially with frequent inserts and deletes, index structures become **bloated and inefficient**.

Solution: Use `REINDEX`

`REINDEX` rebuilds the index from scratch, eliminating fragmentation and unused space.

Example:

```
REINDEX INDEX idx_user_activit
```

output:

```
postgres=# REINDEX INDEX idx_u
REINDEX
postgres=#
```



Jeyaram Ayyalusamy

76 followers

Oracle DBA | AWS Certified | 18+ yrs in
DB Admin, RAC, GoldenGate, RDS,
MySQL, PostgreSQL | Author of MySQL
& RDS books | Cloud & Performance...

 Book Author

You can also reindex an entire table or database:

```
REINDEX TABLE user_activity_log_1;
REINDEX DATABASE my_database;
```

output:

```
postgres=# REINDEX TABLE user_activity_log_1;
REINDEX
postgres=# REINDEX DATABASE postgres;
REINDEX
postgres=#
```

For production environments, consider:

```
REINDEX INDEX CONCURRENTLY idx_user_activity_log_1_id;
```

output:

```
postgres=# REINDEX INDEX CONCU
REINDEX
postgres=#
```

This allows reindexing without blo



Jeyaram Ayyalusamy

76 followers

Oracle DBA | AWS Certified | 18+ yrs in
DB Admin, RAC, GoldenGate, RDS,
MySQL, PostgreSQL | Author of MySQL
& RDS books | Cloud & Performance...

Book Author

Proper Indexing: Avoid Hidden Bloat Triggers

A common source of performance inefficiency (and indirect bloat) is **missing indexes**, particularly on foreign key columns.

Why It Matters:

- PostgreSQL must scan the referencing (child) table when the referenced (parent) row is modified.
- Without an index, this results in full table scans, increasing disk activity and maintenance overhead.
- It can slow down `VACUUM` and increase the chance of dead tuple accumulation.

Best Practice:

Ensure that every foreign key column is backed by an index.

Example:

```
CREATE INDEX idx_orders_customer_id ON orders(customer_id);
```

output:

```
postgres=# CREATE INDEX idx_orders_customer_id ON orders(customer_id);
CREATE INDEX
postgres=#
```

Monitoring and Scheduled

The best way to manage bloat is to
becomes critical.

Steps for Proactive Monitoring:

Use PostgreSQL system views:

- pg_stat_user_tables
- pg_stat_user_indexes
- pg_class

Install and use extensions like:

- pgstattuple
- pg_bloat_check
- Set up disk space alerts and monitor index-to-table size ratios.

Maintenance Checklist:

- Regularly run VACUUM and analyze autovacuum effectiveness.
- Schedule REINDEX operations during low-traffic periods.



Jeyaram Ayyalusamy

76 followers

Oracle DBA | AWS Certified | 18+ yrs in
DB Admin, RAC, GoldenGate, RDS,
MySQL, PostgreSQL | Author of MySQL
& RDS books | Cloud & Performance...



- Create a health dashboard or automated report for bloat statistics.
- Combine all these steps into a monthly or quarterly DBA routine.

💡 Final Thoughts

PostgreSQL gives you the tools to fight bloat — but **you must use them proactively**. By combining autovacuum tuning, proper indexing, and regular reindexing with ongoing monitoring, you can ensure your database remains:

- Efficient
- Lean on disk
- Fast in query performance
- Stable for long-term use

Bloat is inevitable — but its impact



Jeyaram Ayyalusamy

76 followers

Oracle DBA | AWS Certified | 18+ yrs in
DB Admin, RAC, GoldenGate, RDS,
MySQL, PostgreSQL | Author of MySQL
& RDS books | Cloud & Performance...

 Book Author

PostgreSQL

detect and
ugh, we'll
apply cleanup

💡 Practical Demo: Simulating Bloat

To truly understand how **bloating** happens and how to resolve it, nothing beats a hands-on practical demo. We'll simulate bloat step-by-step, visualizing the results along the way.

For this demo, we'll use **customized table and column names** to make it more meaningful and easier to follow.

🔧 Step 1 — Simulate Bloat by Creating a Large Table

We'll begin by creating a large table named `employee_activity_log` to simulate real-world transactional data. This table will have 10 million rows with a single integer column named `activity_code`.

Commands:

```
sudo su - postgres
psql -c "CREATE TABLE employee_activity_log AS SELECT generate_series AS activity_id, employee_code, activity_type, activity_desc, activity_start_time, activity_end_time, duration, location FROM generate_series(1, 1000000) CROSS JOIN generate_series('2023-01-01'::date, '2023-01-01'::date, '1 minute')"
psql -c "CREATE INDEX idx_employee_activity_code ON employee_activity_log (activity_id, employee_code)"
```

output:

```
[ec2-user@ip-172-31-20-155 ~]$ sudo su - postgres
Last login: Tue Jun 24 20:19:57 UTC 2025 on pts/2
[postgres@ip-172-31-20-155 ~]$
[postgres@ip-172-31-20-155 ~]$
[postgres@ip-172-31-20-155 ~]$/psql -c "CREATE TABLE employee_activity_log AS $"
SELECT 10000000
[postgres@ip-172-31-20-155 ~]$
[postgres@ip-172-31-20-155 ~]$
[postgres@ip-172-31-20-155 ~]$
CREATE INDEX
[postgres@ip-172-31-20-155 ~]$
```

Once the data is loaded and indexed

```
psql -c "\dt+"
psql -c "\di+"
```



Jeyaram Ayyalusamy

76 followers

Oracle DBA | AWS Certified | 18+ yrs in
DB Admin, RAC, GoldenGate, RDS,
MySQL, PostgreSQL | Author of MySQL
& RDS books | Cloud & Performance...



output:

```
[postgres@ip-172-31-20-155 ~]$ psql -c "\dt+"
                                         List of relations
 Schema |           Name            | Type  | Owner   | Persistence | Access method
-----+---------------------+-----+-----+-----+-----+
 public | course              | table | postgres | permanent | heap
 public | employee_activity_log | table | postgres | permanent | heap
 public | orders               | table | postgres | permanent | heap
 public | user_activity_log_1   | table | postgres | permanent | heap
 public | user_activity_log_2   | table | postgres | permanent | heap
 public | user_activity_log_3   | table | postgres | permanent | heap
(6 rows)
```

```
[postgres@ip-172-31-20-155 ~]$
[postgres@ip-172-31-20-155 ~]$ psql -c "\di+"
```

List of relations					
Schema	Name	Type	Owner	Table	
public	idx_employee_activity_code	index	postgres	employee_activity_log	
public	idx_orders_customer_id	index	postgres	orders	
public	idx_user_activity_log_1_id	index	postgres	user_activity_log_1	
public	idx_user_activity_log_2_id	index	postgres	user_activity_log_2	
public	idx_user_activity_log_3_id	index	postgres	user_activity_log_3	
public	orders_pkey	index	postgres	orders	

(6 rows)

```
[postgres@ip-172-31-20-155 ~]$
```

These commands will display the current state of your database — your baseline before bloating.



Jeyaram Ayyalusamy

76 followers

Oracle DBA | AWS Certified | 18+ yrs in
DB Admin, RAC, GoldenGate, RDS,
MySQL, PostgreSQL | Author of MySQL
& RDS books | Cloud & Performance...

Book Author

ables, each
workload

Step 2 — Generate More Tables

We'll now simulate a busy environment filled with 1 million rows. These tables will be snapshots.

Commands:

```
createdb metrics_lab
for i in {1..9}; do
    psql -d metrics_lab -c "CREATE TABLE project_task_records_${i} AS SELECT ge
done
```

Each table (`project_task_records_1` to `project_task_records_9`) has a column named `task_number` populated with integers.

```
[postgres@ip-172-31-20-155 ~]$ createdb metrics_lab
for i in {1..9}; do
    psql -d metrics_lab -c "CREATE TABLE project_task_records_{i} AS SELECT ge
done
SELECT 1000000
[postgres@ip-172-31-20-155 ~]$
```



Jeyaram Ayyalusamy

76 followers

Oracle DBA | AWS Certified | 18+ yrs in
DB Admin, RAC, GoldenGate, RDS,
MySQL, PostgreSQL | Author of MySQL
& RDS books | Cloud & Performance...



Step 3 — Create Dead Tuples

To simulate dead tuples (which are rows of data from each of the smaller tables).

Commands:

```
for i in {1..9}; do
    psql -d metrics_lab -c "DELE
done
```

... delete most

} WHERE task_nu

Each table now has **only one row left**, but PostgreSQL hasn't reclaimed the space from deleted rows yet — that's where bloat begins.

```
[postgres@ip-172-31-20-155 ~]$ for i in {1..9}; do
    psql -d metrics_lab -c "DELETE FROM project_task_records_{i} WHERE task_nu
done
DELETE 999999
DELETE 999999
DELETE 999999
DELETE 999999
```

```
DELETE 999999
DELETE 999999
DELETE 999999
DELETE 999999
DELETE 999999
[postgres@ip-172-31-20-155 ~]$
```



Step 4 — Check Bloat Using pgstattuple

PostgreSQL's `pgstattuple` extension is a powerful tool used to inspect **table and index bloat**, **dead tuples**, and overall storage efficiency. However, this extension isn't enabled by default – it must be installed on the system first, especially on Red Hat Enterprise Linux (RHEL) or its derivatives.

Here's a step-by-step breakdown of how to install and use `pgstattuple` in a real-world example on a PostgreSQL 17 instance.



Step 1: Install the postgresql17-contrib package

The `pgstattuple` extension is part of the `postgresql17-contrib` package provided by the PostgreSQL Global Development Group (PGDG) repository.

Command:

```
sudo yum install postgresql17-contrib
```

Explanation:

- This command begins by resolving dependencies and retrieving two main packages:
- `postgresql17-contrib` (732 KB): Contains optional PostgreSQL extensions like `pgstattuple`, `tablefunc`, and more.



Jeyaram Ayyalusamy

76 followers

Oracle DBA | AWS Certified | 18+ yrs in DB Admin, RAC, GoldenGate, RDS, MySQL, PostgreSQL | Author of MySQL & RDS books | Cloud & Performance...



Book Author

cess, as seen in

luded in the elopment

- `libxslt` (190 KB): A required dependency for certain contrib modules.

Even if the system shows a message like:

```
This system is not registered with an entitlement server...
```

...it doesn't prevent the PGDG repository from working. The install still proceeds normally, downloading and verifying packages.

Outcome:

The packages were downloaded successfully and the transaction completed with:

```
Updating Subscription Management
Unable to read consumer identity
```

```
This system is not registered
```

```
Last metadata expiration check
Dependencies resolved.
```

```
=====
Package
```

```
Installing:
postgresql17-contrib
```

```
Installing dependencies:
```

```
libxslt
```

Transaction Summary

```
=====
Install 2 Packages
```

```
Total download size: 922 k
```

```
Installed size: 3.2 M
```

```
Is this ok [y/N]: y
```

```
Downloading Packages:
```

```
(1/2): libxslt-1.1.39-7.el10_0.x86_64.rpm
```

```
(2/2): postgresql17-contrib-17.5-3PGDG.rhel10.x86_64.rpm
```

```
=====
Total
```

```
Running transaction check
```

```
Transaction check succeeded.
```

```
Running transaction test
```

```
Transaction test succeeded.
```



Jeyaram Ayyalusamy

76 followers

Oracle DBA | AWS Certified | 18+ yrs in
DB Admin, RAC, GoldenGate, RDS,
MySQL, PostgreSQL | Author of MySQL
& RDS books | Cloud & Performance...

Book Author

use "rhc" or

17 2025.

x86_64

Running transaction

Preparing :

Installing : libxslt-1.1.39-7.el10_0.x86_64

Installing : postgresql17-contrib-17.5-3PGDG.rhel10.x86_64

Running scriptlet: postgresql17-contrib-17.5-3PGDG.rhel10.x86_64

Installed products updated.

Installed:

libxslt-1.1.39-7.el10_0.x86_64

Complete!

[postgres@ip-172-31-20-155 ~]\$

This confirms the contrib package and its dependencies were installed correctly.

Step II: Restart the PostgreSQL Server

To make sure the extension module is properly loaded after installation, it's best to restart the PostgreSQL server.

Command:

```
sudo systemctl restart postgresql
```



Jeyaram Ayyalusamy

76 followers

Oracle DBA | AWS Certified | 18+ yrs in
DB Admin, RAC, GoldenGate, RDS,
MySQL, PostgreSQL | Author of MySQL
& RDS books | Cloud & Performance...

Book Author

This ensures the server loads the new extension libraries and is ready to activate them within a database.

Step III: Enable the pgstattuple Extension in Your Database

With the contrib package installed and PostgreSQL restarted, the extension is now available to be created inside any database.

Command:

```
psql -d postgres -c "CREATE EXTENSION pgstattuple;"
```

This connects to the `postgres` database and activates the `pgstattuple` extension.

Result:

```
[postgres@ip-172-31-20-155 ~]$  
[postgres@ip-172-31-20-155 ~]$ psql -d metrics_lab -c "CREATE EXTENSION pgstattuple;"  
CREATE EXTENSION  
[postgres@ip-172-31-20-155 ~]$
```

This confirms that the extension was successfully installed and is now ready to use.



Jeyaram Ayyalusamy

76 followers

Oracle DBA | AWS Certified | 18+ yrs in
DB Admin, RAC, GoldenGate, RDS,
MySQL, PostgreSQL | Author of MySQL
& RDS books | Cloud & Performance...

Book Author

You can now begin analyzing table sizes.

```
SELECT * FROM pgstattuple('your_table_name');
```

Run a sample bloat analysis on the main table:

```
SELECT  
    pg_size.pretty(pg_relation_size('employee_activity_log')) AS table_size,  
    pg_size.pretty(pg_relation_size('idx_employee_activity_code')) AS index_size,  
    (pgstattuple('idx_employee_activity_code')).dead_tuple_percent;
```

This query provides:

- Total table size
- Index size

- Percentage of dead tuples in the index (a direct sign of index bloat)

```
postgres=# 
postgres=# SELECT
    pg_size.pretty(pg_relation_size('employee_activity_log')) AS table_size,
    pg_size.pretty(pg_relation_size('idx_employee_activity_code')) AS index_size,
    (pgstattuple('idx_employee_activity_code')).dead_tuple_percent;
table_size | index_size | dead_tuple_percent
-----+-----+-----+
 346 MB   | 214 MB   |          0
(1 row)

postgres=#

```



Jeyaram Ayyalusamy

76 followers

Oracle DBA | AWS Certified | 18+ yrs in
DB Admin, RAC, GoldenGate, RDS,
MySQL, PostgreSQL | Author of MySQL
& RDS books | Cloud & Performance...

 Book Author

- `VACUUM` clears the dead rows, making space reusable.
- `ANALYZE` updates the planner's statistics to help PostgreSQL make smarter execution decisions.

```
[postgres@ip-172-31-20-155 ~]$ psql -c "VACUUM employee_activity_log;" 
VACUUM
[postgres@ip-172-31-20-155 ~]$
[postgres@ip-172-31-20-155 ~]$ psql -c "ANALYZE employee_activity_log;" 
ANALYZE
[postgres@ip-172-31-20-155 ~]$
```

 Note: For full cleanup and disk space reclamation, `VACUUM FULL` may be required—but it requires an exclusive lock.

Step 6 — Advanced Bloat Analysis

You can run a broader scan across all user tables to identify bloated tables and indexes by comparing total size to index size.

Query:

```
SELECT
    relname AS table_name,
    pg_total_relation_size(relid) AS total_size,
    pg_indexes_size(relid) AS index_size,
    pg_total_relation_size(relid) - pg_indexes_size(relid) AS table_only_size
FROM pg_catalog.pg_statio_user_tables
ORDER BY total_size DESC;
```



Jeyaram Ayyalusamy

76 followers

Oracle DBA | AWS Certified | 18+ yrs in
DB Admin, RAC, GoldenGate, RDS,
MySQL, PostgreSQL | Author of MySQL
& RDS books | Cloud & Performance...

 Book Author

This will list all tables sorted by size.

- Total size (table + indexes)
- Size of indexes alone
- Size of the table without indexes

```
postgres=# SELECT
  relname AS table_name,
  pg_total_relation_size(relid) AS total_size,
  pg_indexes_size(relid) AS index_size,
  pg_total_relation_size(relid) - pg_indexes_size(relid) AS table_only_size
FROM pg_catalog.pg_statio_user_tables
ORDER BY total_size DESC;
table_name      | total_size | index_size | table_only_size
-----+-----+-----+-----
employee_activity_log | 587243520 | 224641024 | 362602496
orders           | 74088448  | 29761536  | 44326912
user_activity_log_2 | 24576     | 8192       | 16384
user_activity_log_3 | 24576     | 8192       | 16384
```

course		16384		0		16384
user_activity_log_1		8192		8192		0
(6 rows)						

postgres=#

By analyzing this data, you can **prioritize which tables need reindexing or vacuuming** based on size and bloat characteristics.

💡 Summary

In this practical demo, we:

- Simulated bloating with large character columns
- Identified bloat using PostgreSQL's ANALYZE command
- Cleaned up unused space using VACUUM
- Conducted a full-space analysis with pgstattuple

This kind of proactive monitoring ensures that your PostgreSQL performance stays high and efficient.



Jeyaram Ayyalusamy

76 followers

Oracle DBA | AWS Certified | 18+ yrs in DB Admin, RAC, GoldenGate, RDS, MySQL, PostgreSQL | Author of MySQL & RDS books | Cloud & Performance...

eping

Book Author

📊 Why You Should Actively Manage Bloat in PostgreSQL

Bloat is one of the most common, yet often overlooked, performance and storage issues in PostgreSQL databases. While it naturally occurs due to PostgreSQL's MVCC (Multi-Version Concurrency Control) mechanism, failing to control it can quietly degrade your database over time.

Here's why active bloat management is essential:

✅ Faster Query Performance

As indexes and tables grow due to bloat, PostgreSQL needs to scan more pages to find relevant rows. This leads to:

- Increased I/O latency
- Slower index scans and sequential reads
- Less efficient use of shared buffers and cache

By controlling bloat through regular vacuuming and reindexing, you ensure that the database engine processes only the data it needs — leading to faster query execution and a smoother user experience.

Lower Storage Consumption

Dead tuples and fragmented indexes consume real disk space. Over time, they can grow to multiple times their logical size.

Benefits of bloat control on storage:

- Freed up wasted disk space
- Avoids unnecessary scaling of hardware
- Reduces the cost of cloud storage



Jeyaram Ayyalusamy

76 followers

Oracle DBA | AWS Certified | 18+ yrs in DB Admin, RAC, GoldenGate, RDS, MySQL, PostgreSQL | Author of MySQL & RDS books | Cloud & Performance...

 Book Author

This is especially critical in environments with large datasets or multi-terabyte databases where storage costs can be significant.

More Efficient Backups

Backup tools like `pg_basebackup` or logical dump utilities copy the physical size of the database, not just the active data.

When bloat inflates your tables and indexes:

- Backups take longer to complete

- Backup file sizes increase
- More bandwidth is used for remote or cloud backups

By reducing bloat, you make your backup processes faster, lighter, and less error-prone.

Better Replication Stability

In PostgreSQL streaming replication, the WAL (Write-Ahead Log) is used to synchronize changes to standby nodes. Bloat affects replication in the following ways:

- Large bloat leads to excessive WAL traffic
- Autovacuum or manual vacuum operations may be delayed
- Increased WAL traffic can lag synchronization

Regular bloat control minimizes unexpected replication issues and ensures predictable replication behavior across all nodes.



Jeyaram Ayyalusamy
76 followers
Oracle DBA | AWS Certified | 18+ yrs in DB Admin, RAC, GoldenGate, RDS, MySQL, PostgreSQL | Author of MySQL & RDS books | Cloud & Performance...

 Book Author

leanups)

faster, more

Reduced Downtime During vacuuming

If bloat is ignored for too long, a full VACUUM or even VACUUM FULL becomes unavoidable. These operations can:

- Lock tables for extended periods
- Delay query processing
- Increase the risk of blocking and deadlocks

However, proactive vacuuming and bloat cleanup help avoid emergency scenarios where downtime is required for table rewrites or aggressive maintenance.

🏁 Conclusion

Bloating in PostgreSQL is not a bug — it's a byproduct of how the database manages concurrency and updates. But when ignored, it grows silently and steadily until it impacts performance, storage, replication, and maintainability.

The good news? Bloat is preventable and manageable.

By implementing a regular routine of:

- 🖌 VACUUM and autovacuum tuning
- 🏚 REINDEX for index maintenance
- 📈 Monitoring with tools like `pgstattuple`, `pg_stat_user_tables`, or custom dashboards

...you ensure that your PostgreSQL

- ✅ High-performing
- ✅ Scalable
- ✅ Resource-efficient
- ✅ Reliable in production

⌚ Bloat control is not optional — every professional PostgreSQL DBA or engineer.



Jeyaram Ayyalusamy

76 followers

Oracle DBA | AWS Certified | 18+ yrs in DB Admin, RAC, GoldenGate, RDS, MySQL, PostgreSQL | Author of MySQL & RDS books | Cloud & Performance...

📘 Book Author

Make it a habit, not a reaction.

🔔 Stay Updated with Daily PostgreSQL & Cloud Tips!

If you've been finding my blog posts helpful and want to stay ahead with daily insights on PostgreSQL, Cloud Infrastructure, Performance Tuning, and DBA Best Practices — I invite you to subscribe to my Medium account.

🔔 Subscribe here 👉 <https://medium.com/@jramcloud1/subscribe>

Your support means a lot — and you'll never miss a practical guide again!

🔗 Let's Connect!

If you enjoyed this post or would like to connect professionally, feel free to reach out to me on LinkedIn:

👉 [Jeyaram Ayyalusamy](#)

I regularly share content on **PostgreSQL, database administration, cloud technologies, and data engineering**. Always happy to connect, collaborate, and discuss ideas!

Postgresql

Oracle

Sql

AWS

Open Source

J

Written by Jeyaram Ayyalusamy

76 followers · 2 following

Oracle DBA | AWS Certified | 18+ yrs in DB Admin, MySQL & RDS books | Cloud & Performance...



Jeyaram Ayyalusamy

76 followers

Oracle DBA | AWS Certified | 18+ yrs in DB Admin, RAC, GoldenGate, RDS, MySQL, PostgreSQL | Author of MySQL & RDS books | Cloud & Performance...

Following

Book Author

No responses yet



Gvadakte

What are your thoughts?

More from Jeyaram Ayyalusamy

J Jeyaram Ayyalusamy

PostgreSQL Tablespaces Explained for DBAs

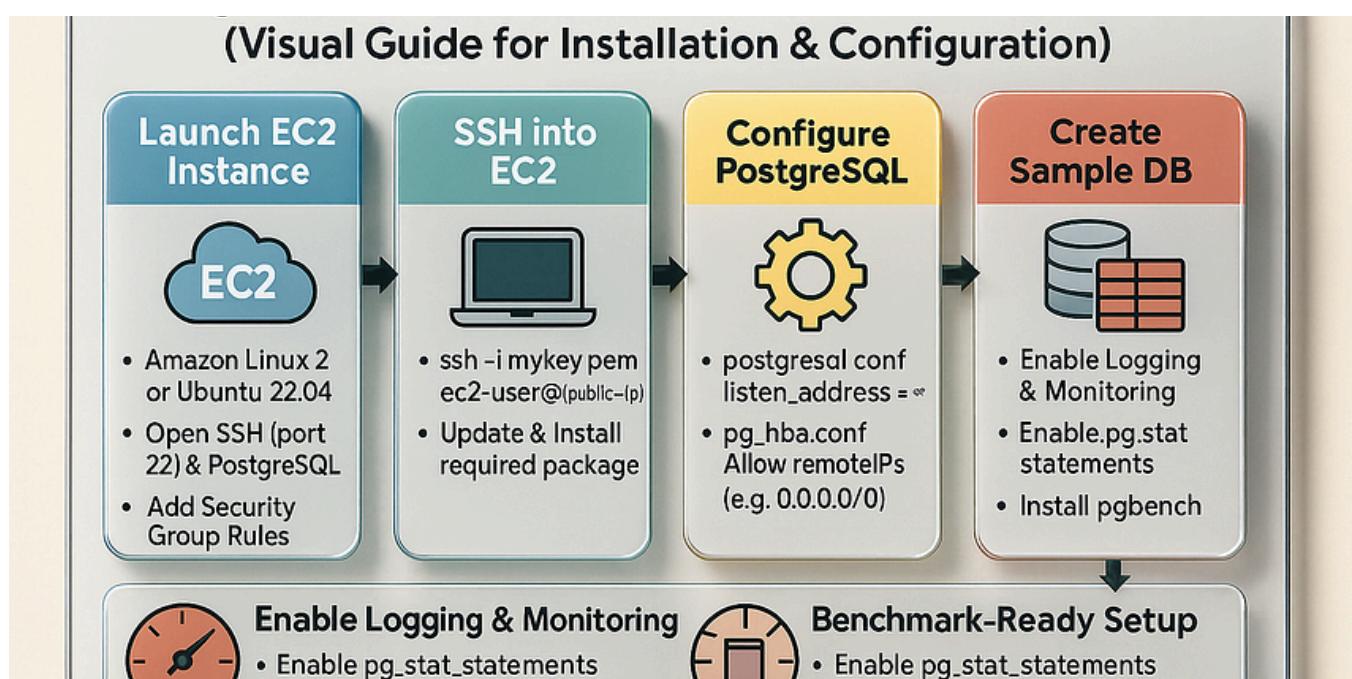
PostgreSQL offers a powerful feature called tablespaces that allows DBAs to take control of how and where data is stored.

Jun 12 8

J Jeyaram Ayyalusamy
76 followers
Oracle DBA | AWS Certified | 18+ yrs in DB Admin, RAC, GoldenGate, RDS, MySQL, PostgreSQL | Author of MySQL & RDS books | Cloud & Performance...

Book Author

tgreSQL 17
ministrators to



J Jeyaram Ayyalusamy 

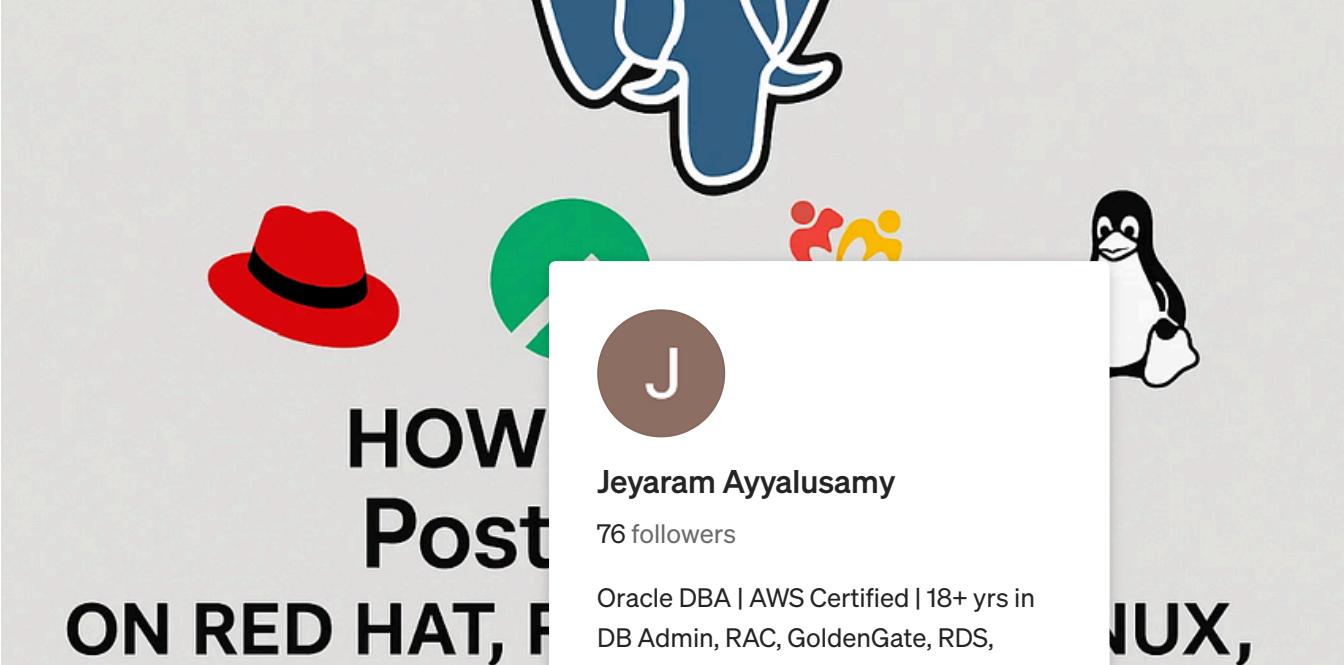
PostgreSQL 17 on AWS EC2—Full Installation & Configuration Walkthrough

Setting up PostgreSQL 17 on AWS EC2 might seem complex at first—but once you break down each component, it becomes an efficient and...

1d ago  50



...



HOW PostgreSQL 17 on Red Hat, RHEL, and Oracle

J
Jeyaram Ayyalusamy
76 followers
Oracle DBA | AWS Certified | 18+ yrs in DB Admin, RAC, GoldenGate, RDS, MySQL, PostgreSQL | Author of MySQL & RDS books | Cloud & Performance...

 Jeyaram Ayyalusamy 

J Jeyaram Ayyalusamy 

How to Install PostgreSQL 17 on Linux (Step-by-Step Guide)

PostgreSQL 17 is the latest release of one of the world's most advanced open-source relational databases. If you're using a Red Hat-based...

Jun 3



...

```
sudo systemctl restart etcd
sudo systemctl status etcd
etcd - highly-available key value store
ded (/lib/systemd/system/etcd.service; enabled; vendor preset: enabled)
  ● etcd (running) since Thu 2025-05-29 00:39:20 UTC; 5s ago
    ps://etcd.10/docs
      ● etcd
        ● 3 (etcd)
          limit: 4558
        5M
        ms
      ● stem.slice/etcd.service
        463 /usr/bin/etcd

etcdnode etcd[2463]: 8e9e05c52164694d became candidate at term 5
etcdnode etcd[2463]: 8e9e05c52164694d received MsgVoteResp from 8e9e05c52164694d at term 5
etcdnode etcd[2463]: 8e9e05c52164694d became leader at term 5
etcdnode etcd[2463]: raft.node: 8e9e05c52164694d elected leader 8e9e05c52164694d at term 5
etcdnode etcd[2463]: published {Name:etcdnode ClientURLs:[http://192.168.32.140:2379]} to cluster cdf8
etcdnode etcd[2463]: ready to serve client requests
etcdnode etcd[2463]: ready to serve client requests
etcdnode etcd[2463]: serving insecure client requests on 192.168.32.140:2379, this is strongly discour
etcdnode etcd[2463]: serving insecure client requests on 127.0.0.1:2379, this is strongly discouraged!
etcdnode systemd[1]: Started etcd - highly-available key value store.
```

J Jeyaram Ayyalusamy 

Building a High Availability Po Step Guide

This will follow the standard instruction
your nodes: `node1`, `node2`, `etcdn

May 29  1

See all



Jeyaram Ayyalusamy

76 followers

Oracle DBA | AWS Certified | 18+ yrs in
DB Admin, RAC, GoldenGate, RDS,
MySQL, PostgreSQL | Author of MySQL
& RDS books | Cloud & Performance...

 Book Author

Step-by-

ation based on



...

Recommended from Medium

The screenshot shows a dark-themed dashboard titled "Postgres Performance Tuning". It features a large blue hexagonal logo in the center. On the left, there are several cards: "Postgres Features", "Postgres Statistics", "Postgres Metrics", "Postgres Configuration", "Postgres Tracing", and "Postgres Monitoring". The "Postgres Configuration" card has sections for "Postgres Configuration", "Postgres Metrics", and "Postgres Monitoring". The "Postgres Metrics" section includes a "CPU Usage" chart and a "Memory Usage" chart. At the bottom, there's a "Postgres Configuration" card with a "Configure" button.

Rizqi Mulki

Postgres Performance Tuning

The advanced techniques that separate...

6d ago 55

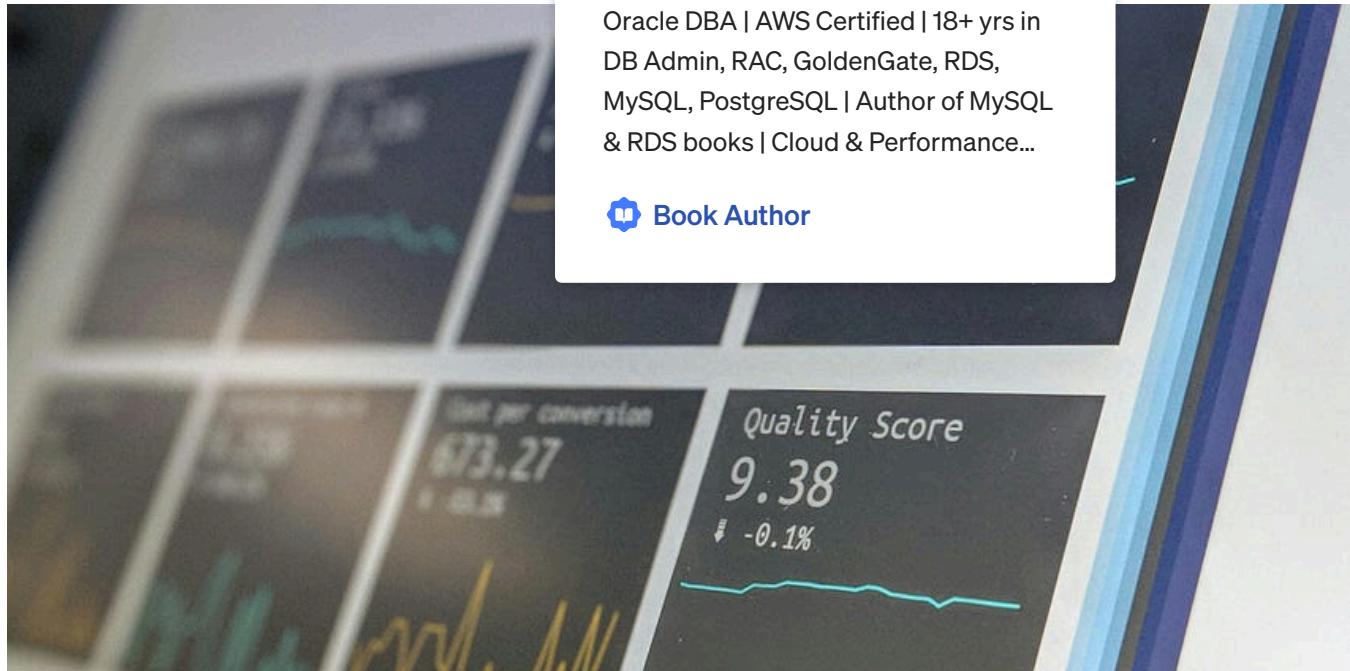


Jeyaram Ayyalusamy

76 followers

Oracle DBA | AWS Certified | 18+ yrs in
DB Admin, RAC, GoldenGate, RDS,
MySQL, PostgreSQL | Author of MySQL
& RDS books | Cloud & Performance...

Book Author

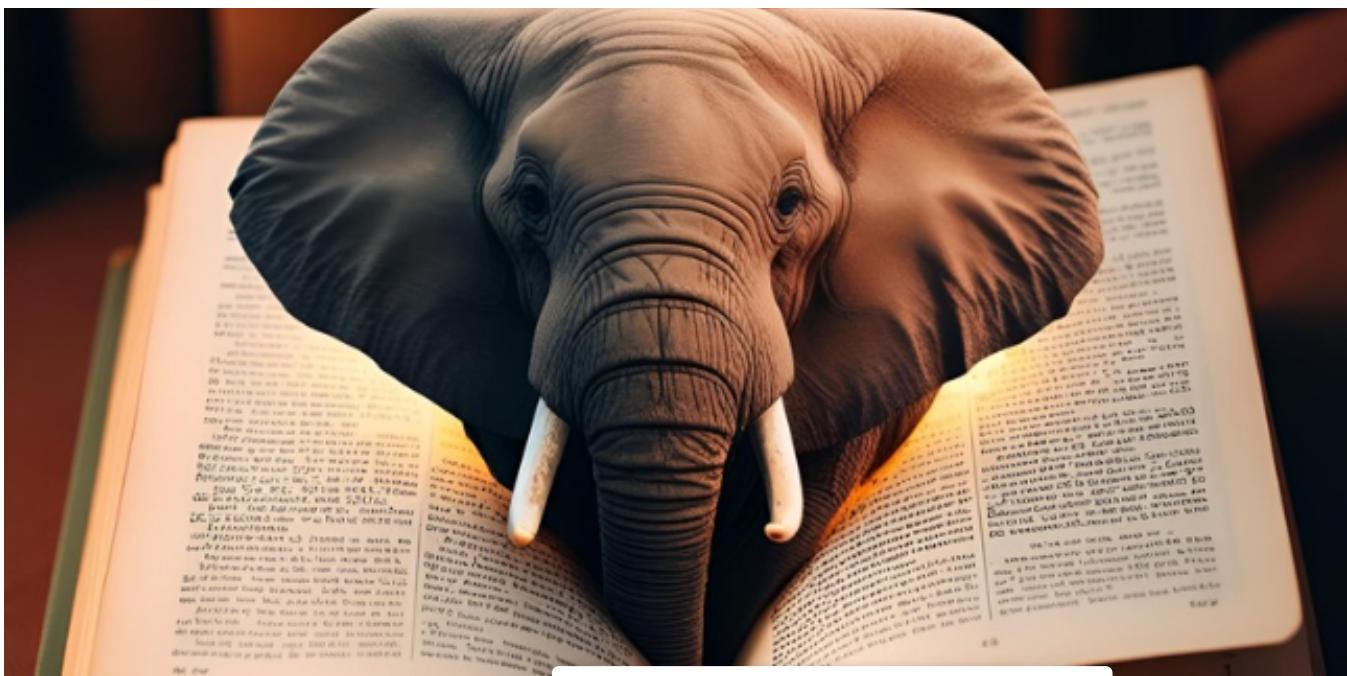


Azlan Jamal

Stop Using SERIAL in PostgreSQL: Here's What You Should Do Instead

In PostgreSQL, there are several ways to create a PRIMARY KEY for an id column, and they usually involve different ways of generating the...

Jul 12 33



Understanding PostgreSQL Part 1

PostgreSQL stores data on disk in fixed-size pages. In this post, we'll explore how PostgreSQL organizes data.

May 14 58 1



Jeyaram Ayyalusamy

76 followers

Oracle DBA | AWS Certified | 18+ yrs in DB Admin, RAC, GoldenGate, RDS, MySQL, PostgreSQL | Author of MySQL & RDS books | Cloud & Performance...

Book Author

The screenshot shows a PostgreSQL query editor interface. At the top, there is a code editor with the following SQL query:

```

1 explain
2 select *
3 from payment_lab
4 where customer_id=10 ;

```

Below the code editor, there are two tabs: "Statistics 1" and "Results 2". The "Results 2" tab is currently selected, displaying the "QUERY PLAN" for the query. The plan shows the following steps:

- Seq Scan on payment_lab (cost=0.00..290.45 rows=23 width=26)
- Filter: (customer_id = 10)

M Muhammet Kurtoglu

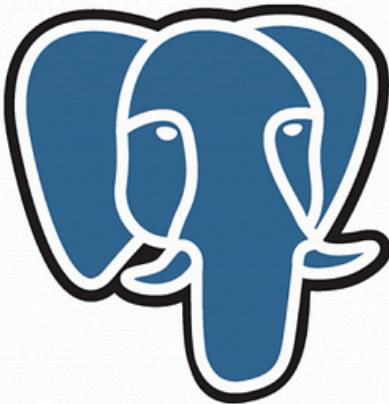
Postgresql Query Performance Analysis

SQL tuning is critically important for ensuring database performance, reliability, and scalability. In most cases, performance issues in a...

6d ago 👏 10



...



PostgreSQL

 Harishsingh

PostgreSQL 18 in Microservices Everything

Introduction: The Myth of Database-Per-

Jul 13 👏 11 1



Jeyaram Ayyalusamy

76 followers

Oracle DBA | AWS Certified | 18+ yrs in
DB Admin, RAC, GoldenGate, RDS,
MySQL, PostgreSQL | Author of MySQL
& RDS books | Cloud & Performance...

 Book Author

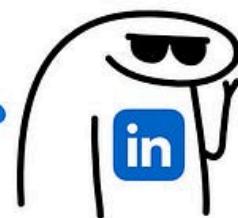
DB for



...



LinkedIn is moving from Kafka to this



The company that created Kafka is replacing
it with a new solution



In Data Engineer Things by Vu Trinh

The company that created Kafka is replacing it with a new solution

How did LinkedIn build Northguard, the new scalable log storage

★ Jul 17

330

6



...

[See more recommendations](#)**Jeyaram Ayyalusamy**

76 followers

Oracle DBA | AWS Certified | 18+ yrs in
DB Admin, RAC, GoldenGate, RDS,
MySQL, PostgreSQL | Author of MySQL
& RDS books | Cloud & Performance...

 Book Author