

# PostgreSQL Monitoring: A Complete Guide to Database Observability

Effective monitoring is crucial for maintaining healthy PostgreSQL databases. This comprehensive guide covers everything you need to know about monitoring PostgreSQL, from basic metrics to advanced monitoring strategies.

## Table of Contents

- [Essential PostgreSQL Metrics](#)
- [Monitoring Tools](#)
- [Query Performance Monitoring](#)
- [System Resource Monitoring](#)
- [Alerting Strategies](#)
- [Best Practices](#)

## Essential PostgreSQL Metrics {#essential-metrics}

### 1. Connection Metrics

```
-- Monitor active connections
SELECT count(*), state FROM pg_stat_activity GROUP BY state;

-- Check connection utilization

SELECT current_setting('max_connections')::int as max_conn, count(*) as current_conn,
(count(*)::float/current_setting('max_connections')::int * 100)::int as conn_percent FROM
pg_stat_activity;
```

### 2. Database Size and Growth

```
-- Monitor database sizes

SELECT  datname, pg_size_pretty(pg_database_size(datname)) as size,
pg_size_pretty(pg_database_size(datname) - coalesce(previous_size, 0)) as growthFROM
pg_database;
```

*-- Table growth monitoring*

```
SELECT  schemaname, relname, n_live_tup as row_count,
pg_size_pretty(pg_total_relation_size(schemaname || '.' || relname)) as total_size
FROM pg_stat_user_tables
ORDER BY pg_total_relation_size(schemaname || '.' || relname) DESC;
```

### 3. Transaction Metrics

*-- Transaction statistics*

```
SELECT  datname, xact_commit, xact_rollback, (xact_rollback::float / (xact_commit +
xact_rollback) * 100)::numeric(10,2)  as rollback_ratio
FROM pg_stat_database;
```

## Monitoring Tools {#monitoring-tools}

### 1. Prometheus with PostgreSQL Exporter

Example `prometheus.yml` configuration:

```
scrape_configs:
- job_name: 'postgresql'
  static_configs:
    - targets: ['localhost:9187']
  metrics_path: '/metrics'
```

**Example queries:**

*# Query rate of transactions*

```
rate(pg_stat_database_xact_commit{datname="mydb"}[5m])
```

*# Database size*

```
pg_database_size_bytes{datname="mydb"}
```

*# Connection utilization*

```
pg_stat_activity_count / pg_settings_max_connections * 100
# Query rate of transactions
rate(pg_stat_database_xact_commit{datname="mydb"}[5m])
# Database size
pg_database_size_bytes{datname="mydb"}
# Connection utilization
pg_stat_activity_count / pg_settings_max_connections * 100
```

**PROMQL**

## 2. Grafana Dashboard Setup

Example dashboard JSON:

```
{
  "panels": [
    {
      "title": "Active Connections",
      "targets": [
        {
          "expr": "pg_stat_activity_count{datname='mydb',state='active'}"
        }
      ]
    },
    {
      "title": "Transaction Rate",
      "targets": [
        {
          "expr": "rate(pg_stat_database_xact_commit{datname='mydb'}[5m])"
        }
      ]
    }
  ]
}
```

## 3. pgMonitor Configuration

Example pgMonitor setup:

```
# Install pgMonitor
git clone https://github.com/CrunchyData/pgmonitor.git
cd pgmonitor/exporter
./install-pg-exporter.sh

# Configure metrics collection
cat << EOF > /etc/pgmonitor/exporter/postgres/queries.yml
pg_replication:
  query: "SELECT * FROM pg_stat_replication"
metrics:
  - write_lag_bytes
  - replay_lag_bytes

EOF
```

<b>Query Performance Monitoring {#query-performance}</b>
--

## 1. Slow Query Analysis

```
-- Create extension if not exists
CREATE EXTENSION pg_stat_statements;

-- Find slow queries
SELECT
    substring(query, 1, 50) as short_query,
    round(total_time::numeric, 2) as total_time,
    calls,
    round(mean_time::numeric, 2) as mean_time,
    round((100 * total_time / sum(total_time::numeric) over ()):numeric, 2) as percentage
FROM pg_stat_statements
ORDER BY total_time DESC
LIMIT 10;
```

## 2. Index Usage Monitoring

```
-- Monitor index usage
SELECT  schemaname, tablename, indexname, idx_scan as number_of_scans, idx_tup_read as
tuples_read, idx_tup_fetch as tuples_fetched
FROM pg_stat_user_indexes
ORDER BY idx_scan DESC;

SQL
```

# System Resource Monitoring {#system-resources}

## 1. CPU Usage Monitoring

- Check CPU-intensive queries

```
SELECT  pid, datname, username, query_start, state, query
FROM pg_stat_activity
WHERE state = 'active'
ORDER BY query_start;
```

SQL

-

## 2. Memory Usage

-- Monitor shared buffer usage

```
SELECT  buffers_clean, buffers_backend, buffers_backend_fsync
FROM pg_stat_bgwriter;

-- Check memory context usage
SELECT  name, setting, unit, context
FROM pg_settings
WHERE name LIKE '%memory%' OR name LIKE '%buffer%';
```

SQL

# Alerting Strategies {#alerting}

## 1. Connection Alerts

```
-- Create alert function for connection threshold

CREATE OR REPLACE FUNCTION check_connection_threshold()

RETURNS trigger AS $$

BEGIN

    IF (SELECT count(*) FROM pg_stat_activity) >

        (current_setting('max_connections')::int * 0.8) THEN

        PERFORM pg_notify(

            'connection_alert',

            'Connection threshold exceeded: ' || count(*) || ' connections'

        );

    END IF;

    RETURN NULL;

END;

$$ LANGUAGE plpgsql;
```

## 2. Replication Lag Alerts

```
-- Monitor replication lag

SELECT

    client_addr,

    state,

    sent_lsn,

    write_lsn,

    flush_lsn,

    replay_lsn,

    pg_wal_lsn_diff(sent_lsn, replay_lsn) as lag_bytes

FROM pg_stat_replication
```

# Best Practices {#best-practices}

## 1. Regular Maintenance

```
-- Monitor table bloat
WITH constants AS (
    SELECT current_setting('block_size')::numeric AS bs
)
SELECT
    schemaname,
    tablename,
    pg_size_pretty(bloat_size) as bloat_size,
    round(bloat_ratio::numeric, 2) as bloat_ratio
FROM (
    SELECT
        schemaname,
        tablename,
        bs*n_live_tup as actual_size,
        bs*reltuples as expected_size,
        bs*(reltuples-n_live_tup) as bloat_size,
        100.0 * (reltuples-n_live_tup)/reltuples as bloat_ratio
    FROM pg_stat_user_tables t
    JOIN pg_class c ON t.relname = c.relname
    CROSS JOIN constants
    WHERE reltuples > 0
) s
WHERE bloat_ratio > 10
ORDER BY bloat_size DESC;
```

## 2. Vacuum Monitoring

```
-- Monitor autovacuum activity

SELECT  schemaname, relname, last_vacuum, last_autovacuum, vacuum_count,
        autovacuum_countFROM pg_stat_user_tables;

SQL
```

### 3. Logging Configuration

```
# postgresql.conf logging settings
log_min_duration_statement = 1000 # Log queries taking more than 1 second
log_checkpoints = on
log_connections = on
log_disconnections = on
log_lock_waits = on
log_temp_files = 0
log_autovacuum_min_duration = 0
INI
```

## Monitoring Automation Script

Here's a sample Python script for automated monitoring:

```
import psycopg2

import time

from datetime import datetime

def monitor_database():

    conn = psycopg2.connect("dbname=postgres user=postgres")

    cur = conn.cursor()

    # Collect metrics

    metrics = {

        'timestamp': datetime.now(),

        'connections': {},

        'performance': {},

        'storage': {}

    }

    # Connection metrics
```

```

cur.execute("""

    SELECT count(*), state

    FROM pg_stat_activity

    GROUP BY state

""")

metrics['connections']['by_state'] = dict(cur.fetchall())


# Performance metrics

cur.execute("""

    SELECT datname, xact_commit, xact_rollback

    FROM pg_stat_database

    WHERE datname = current_database()

""")

metrics['performance']['transactions'] = cur.fetchone()


# Storage metrics

cur.execute("""

    SELECT pg_size_pretty(pg_database_size(current_database()))

""")

metrics['storage']['database_size'] = cur.fetchone()[0]


return metrics


def main():

    while True:

        try:

            metrics = monitor_database()

```



```
# Process metrics (e.g., send to monitoring system)

print(f'Metrics collected at {metrics['timestamp']}')

time.sleep(60) # Collect metrics every minute

except Exception as e:

    print(f'Error collecting metrics: {e}')

    time.sleep(5)

if __name__ == "__main__":

    main()
```

## Conclusion

Effective PostgreSQL monitoring is essential for maintaining database health and performance. By implementing the monitoring strategies and tools discussed in this guide, you can:

- Proactively identify and resolve issues
- Optimize database performance
- Plan capacity effectively
- Ensure high availability
- Maintain data integrity

Remember to regularly review and adjust your monitoring setup as your database requirements evolve. The key is to find the right balance between comprehensive monitoring and system overhead.