

Setup a Simple HA PostgreSQL Cluster Using Patroni, Etcd, HAProxy, and Keepalived

by Biondi Septian S

Objective

To simulate the building of a simple highly available PostgreSQL Patroni Cluster from scratch in a development environment.

Assumptions

- I don't care about the firewall in this demo environment, therefore I'm turning off the ufw service in this demo.
- I'm using default values of PostgreSQL Memory Parameters because the VMs that I'm using are small in terms of specification.
- I'm doing this demo in my VirtualBox, but I won't specify step by step related to the virtualbox management. I'm focusing only on the setup of the PostgreSQL Cluster.

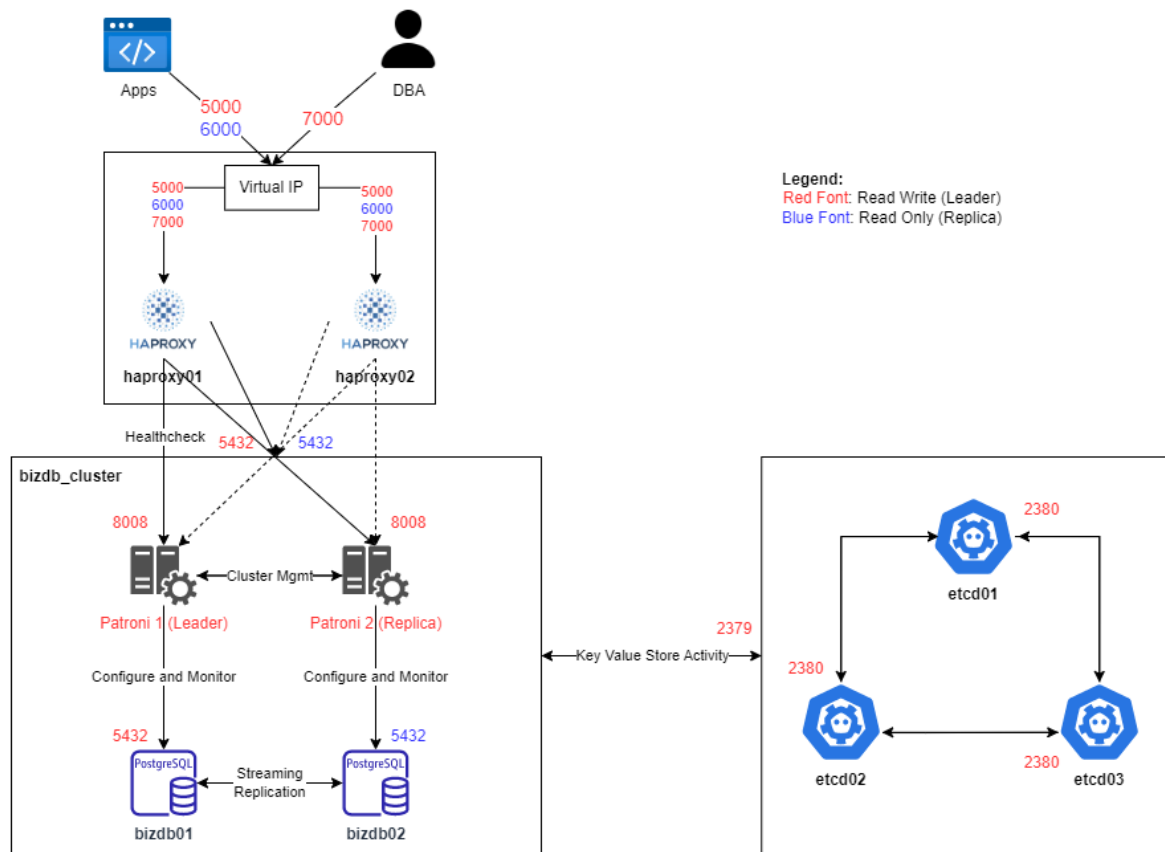
Prerequisites

- 2 VMs for PostgreSQL and Patroni:
bizdb01, bizdb02
- 2 VMs for HAProxy:
haproxy01, haproxy02
- 3 Etcd nodes will be installed on shared VMs:
bizdb01, bizdb02, haproxy01
- OS: Ubuntu 20.04
- PostgreSQL version: 14
- PostgreSQL cluster name: bizdb_cluster
- RAM:
bizdb01: 4 GB
bizdb02: 4 GB
haproxy01: 4 GB

haproxy02: 2 GB

- CPU:
bizdb01: 2 CPU
bizdb02: 2 CPU
haproxy01: 2 CPU
haproxy02: 2 CPU
- Storage:
bizdb01: 20 GB
bizdb02: 20 GB
haproxy01: 20 GB
haproxy02: 15 GB
- IP Address:
bizdb01 (PG Node 1): 192.168.150.100
bizdb02 (PG Node 2): 192.168.150.101
haproxy01 (HAProxy Node 1): 192.168.150.110
haproxy02 (HAProxy Node 2): 192.168.150.111
Virtual IP Address Reserved for HAProxy: 192.168.150.200

Architecture



Steps

1. Install Etcd on 3 Nodes (bizdb01, bizdb02, haproxy01)

Do Below Steps on 3 Nodes (bizdb01, bizdb02, haproxy01), Change the Hostname and IP Address according to the Hosts:

Install Required Tools

```
root@bizdb01:~# sudo apt-get install wget curl -y
```

Download etcd Binary From Github Using wget

```
root@bizdb01:~# wget
```

<https://github.com/etcd-io/etcd/releases/download/v3.5.7/etcd-v3.5.7-linux-amd64.tar.gz>

....

```
....  
etcd-v3.5.7-linux-amd64.tar.gz  
100%[=====>]  
17.60M 4.46MB/s in 5.0s
```

2023-03-15 09:23:11 (3.55 MB/s) - 'etcd-v3.5.7-linux-amd64.tar.gz' saved [18458320/18458320]

Unzip the etcd Binary Files and Rename Directory to etcd

```
root@bizdb01:~# tar xvf etcd-v3.5.7-linux-amd64.tar.gz  
root@bizdb01:~# mv etcd-v3.5.7-linux-amd64 etcd
```

Move All Binary Files Into /usr/local/bin Directory

```
root@bizdb01:~# cd etcd/  
root@bizdb01:~/etcd# sudo mv etcd etcdctl etcdutl /usr/local/bin/
```

Check etcd Binaries Version

```
root@bizdb01:~/etcd# etcd --version  
etcd Version: 3.5.7  
Git SHA: 215b53cf3  
Go Version: go1.17.13  
Go OS/Arch: linux/amd64
```

```
root@bizdb01:~/etcd# etcdctl version  
etcdctl version: 3.5.7  
API version: 3.5
```

```
root@bizdb01:~/etcd# etcdutl version  
etcdutl version: 3.5.7  
API version: 3.5
```

Setup Etcd Service

Edit /etc/hosts to Give Alias For Each Etcd Cluster Members

```
root@bizdb01:~# nano /etc/hosts  
192.168.150.100 etcd01  
192.168.150.101 etcd02  
192.168.150.110 etcd03
```

Create Group and User for Etcd

```
root@bizdb01:~# sudo groupadd --system etcd
```

```
root@bizdb01:~# sudo useradd -s /sbin/nologin --system -g etcd etcd
```

Create Data Directory and Config Directory

```
root@bizdb01:~# sudo mkdir -p /var/lib/etcd/
```

```
root@bizdb01:~# sudo mkdir /etc/etcd/
```

Change Ownership and Permission Of /var/lib/etcd Directory To Etcd User

```
root@bizdb01:~# sudo chown -R etcd:etcd /var/lib/etcd/
```

```
root@bizdb01:~# sudo chmod -R 700 /var/lib/etcd/
```

Identify Primary Network Interface To Run Etcd

```
root@bizdb01:~# ip ad
```

```
1: lo: <LOOPBACK,UP,LOWER_UP> mtu 65536 qdisc noqueue state UNKNOWN group default qlen 1000
```

```
    link/loopback 00:00:00:00:00:00 brd 00:00:00:00:00:00
```

```
    inet 127.0.0.1/8 scope host lo
```

```
        valid_lft forever preferred_lft forever
```

```
    inet6 ::1/128 scope host
```

```
        valid_lft forever preferred_lft forever
```

```
2: enp0s3: <BROADCAST,MULTICAST,UP,LOWER_UP> mtu 1500 qdisc fq_codel state UP group default qlen 1000
```

```
    link/ether 08:00:27:dd:a2:da brd ff:ff:ff:ff:ff:ff
```

```
    inet 10.0.2.15/24 brd 10.0.2.255 scope global dynamic enp0s3
```

```
        valid_lft 78324sec preferred_lft 78324sec
```

```
    inet6 fe80::a00:27ff:fedd:a2da/64 scope link
```

```
        valid_lft forever preferred_lft forever
```

```
3: enp0s8: <BROADCAST,MULTICAST,UP,LOWER_UP> mtu 1500 qdisc fq_codel state UP group default qlen 1000
```

```
    link/ether 08:00:27:39:88:62 brd ff:ff:ff:ff:ff:ff
```

```
    inet 192.168.150.100/24 brd 192.168.150.255 scope global enp0s8
```

```
        valid_lft forever preferred_lft forever
```

```
    inet6 fe80::a00:27ff:fe39:8862/64 scope link
```

```
        valid_lft forever preferred_lft forever
```

Load and Check Some Environment Variables For Used in Service Daemon

```
root@bizdb01:~# NET_NAME="enp0s8"
```

```
root@bizdb01:~# ETCD_HOST_IP=$(ip addr show $NET_NAME | grep "inet\b" | awk '{print $2}' | cut -d/ -f1)
```

```
root@bizdb01:~# ETCD_NAME=$(cat /etc/hosts | grep "$(ip addr show $NET_NAME | grep "inet\b" | awk '{print $2}' | cut -d/ -f1)" | awk '{print $2}')
```

```
root@bizdb01:~# echo $NET_NAME
```

```
enp0s8
```

```
root@bizdb01:~# echo $ETCD_HOST_IP
```

192.168.150.100

```
root@bizdb01:~# echo $ETCD_NAME
etcd01
```

#NET_NAME is used to store the name of the primary network interface of etcd server
#ETCD_HOST_IP is used to store the IP of the host of the Etcd member in the clusters.
#ETCD_NAME is used to store the unique name of the Etcd member by using hostname.

Create Etcd Unit File in Systemd To Manage Etcd Service Using Following Configuration

```
root@bizdb01:~# cat <<EOF | sudo tee /etc/systemd/system/etcd.service
```

[Unit]

Description=etcd - highly-available key value store

Documentation=https://etcd.io/docs

Documentation=man:etcd

[Service]

Type=notify

User=etcd

**ExecStart=/usr/local/bin/etcd **

**--name \${ETCD_NAME} **

**--data-dir=/var/lib/etcd **

**--initial-advertise-peer-urls http://\${ETCD_HOST_IP}:2380 **

**--listen-peer-urls http://\${ETCD_HOST_IP}:2380 **

**--listen-client-urls http://\${ETCD_HOST_IP}:2379,http://127.0.0.1:2379 **

**--advertise-client-urls http://\${ETCD_HOST_IP}:2379 **

**--initial-cluster-token etcd-cluster **

--initial-cluster etcd01=http://etcd01:2380,etcd02=http://etcd02:2380,etcd03=http://etcd03:2380

**--initial-cluster-state new **

**--enable-v2=true **

[Install]

WantedBy=multi-user.target

EOF

Reload Daemon and Enable The Service To Start On Boot

```
root@bizdb01:~# sudo systemctl daemon-reload
```

```
root@bizdb01:~# sudo systemctl enable etcd.service
```

Created symlink /etc/systemd/system/multi-user.target.wants/etcd.service →
/etc/systemd/system/etcd.service.

Start The Etcd Service on 3 Nodes

```
root@bizdb01:~# sudo systemctl start etcd
```

Check Status of The Etcd Services

```
root@bizdb01:~# sudo systemctl status etcd
```

- etcd.service - etcd - highly-available key value store

Loaded: loaded (/etc/systemd/system/etcd.service; enabled; vendor preset: enabled)

Active: active (running) since Fri 2023-03-17 11:21:07 UTC; 15min ago

Docs: <https://etcd.io/docs>

man:etcd

Main PID: 3711 (etcd)

Tasks: 8 (limit: 4612)

Memory: 16.1M

CGroup: /system.slice/etcd.service

└─3711 /usr/local/bin/etcd --name etcd01 --data-dir=/var/lib/etcd --initial-advertise-peer-urls

http://192.168.150.100:2380

```
root@bizdb01:~# etcdctl endpoint status --write-out=table
```

```
--endpoints=http://etcd01:2379,http://etcd02:2379,http://etcd03:2379
```

ENDPOINT	ID	VERSION	DB SIZE	IS LEADER	IS LEARNER	RAFT TERM	RAFT INDEX	RAFT APPLIED INDEX	ERRORS
http://etcd01:2379	b8b747c74aaea686	3.5.7	20 kB	true	false	7	26		
http://etcd02:2379	b3504381e8ba3cb	3.5.7	20 kB	false	false	7	26		
http://etcd03:2379	f572fdcf5cb68406	3.5.7	20 kB	false	false	7	26		

2. Install PostgreSQL 14 on 2 Nodes (bizdb01, bizdb02)

Install GNUPG

```
root@bizdb01:~# sudo apt install gnupg2
```

Add PostgreSQL 14 Repository for Ubuntu 20.04 Release

```
root@bizdb01:~# sudo sh -c 'echo "deb http://apt.postgresql.org/pub/repos/apt $(lsb_release -cs)-pgdg main" > /etc/apt/sources.list.d/pgdg.list'
```

Import GPG Signing Key For The Repository

```
root@bizdb01:~# wget --quiet -O - https://www.postgresql.org/media/keys/ACCC4CF8.asc | sudo apt-key add -  
OK
```

Update APT Package List

```
root@bizdb01:~# sudo apt update  
Hit:1 http://id.archive.ubuntu.com/ubuntu focal InRelease  
Get:2 http://id.archive.ubuntu.com/ubuntu focal-updates InRelease [114 kB]  
Get:3 http://id.archive.ubuntu.com/ubuntu focal-backports InRelease [108 kB]  
Get:4 http://id.archive.ubuntu.com/ubuntu focal-security InRelease [114 kB]  
Get:5 http://id.archive.ubuntu.com/ubuntu focal-updates/main amd64 Packages [2,425 kB]  
Get:6 http://id.archive.ubuntu.com/ubuntu focal-updates/universe amd64 Packages [1,038 kB]  
Get:7 http://apt.postgresql.org/pub/repos/apt focal-pgdg InRelease [91.6 kB]  
Get:8 http://apt.postgresql.org/pub/repos/apt focal-pgdg/main amd64 Packages [258 kB]  
Fetched 4,149 kB in 3s (1,579 kB/s)  
Reading package lists... Done  
Building dependency tree  
Reading state information... Done  
28 packages can be upgraded. Run 'apt list --upgradable' to see them.
```

Install PostgreSQL 14 Package

```
root@bizdb01:~# sudo apt -y install postgresql-14  
....  
....  
Creating config file /etc/postgresql-common/createcluster.conf with new version  
Building PostgreSQL dictionaries from installed myspell/hunspell packages...  
Removing obsolete dictionary files:  
'/etc/apt/trusted.gpg.d/apt.postgresql.org.gpg' ->  
'/usr/share/postgresql-common/pgdg/apt.postgresql.org.gpg'  
Created symlink /etc/systemd/system/multi-user.target.wants/postgresql.service →  
/lib/systemd/system/postgresql.service.  
Setting up postgresql-14 (14.7-1.pgdg20.04+1) ...  
Creating new PostgreSQL cluster 14/main ...  
/usr/lib/postgresql/14/bin/initdb -D /var/lib/postgresql/14/main --auth-local peer --auth-host  
scram-sha-256 --no-instructions  
The files belonging to this database system will be owned by user "postgres".  
This user must also own the server process.
```


The database cluster will be initialized with locale "en_US.UTF-8".
The default database encoding has accordingly been set to "UTF8".
The default text search configuration will be set to "english".

Data page checksums are disabled.

```
fixing permissions on existing directory /var/lib/postgresql/14/main ... ok
creating subdirectories ... ok
selecting dynamic shared memory implementation ... posix
selecting default max_connections ... 100
selecting default shared_buffers ... 128MB
selecting default time zone ... Asia/Jakarta
creating configuration files ... ok
running bootstrap script ... ok
performing post-bootstrap initialization ... ok
syncing data to disk ... ok
update-alternatives: using /usr/share/postgresql/14/man/man1/postmaster.1.gz to provide
/usr/share/man/man1/postmaster.1.gz (postmaster.1.gz) in auto mode
Processing triggers for systemd (245.4-4ubuntu3.20) ...
Processing triggers for man-db (2.9.1-1) ...
Processing triggers for libc-bin (2.31-0ubuntu9.9) ...
```

Drop Existing main Cluster From PostgreSQL Default Installation

We don't need this built-in cluster from installation, because later Patroni will take care of new database cluster initialization

```
root@bizdb01:~# pg_dropcluster --stop 14 main
```

Edit PostgreSQL Default Config File, Change listen_addresses

```
root@bizdb01:~# cd /etc/postgresql/14/main
root@bizdb01:/etc/postgresql/14/main# nano postgresql.conf
# - Connection Settings -
```

```
listen_addresses = '*'
```

```
root@bizdb01:/etc/postgresql/14/main# sudo systemctl stop postgresql
root@bizdb01:/etc/postgresql/14/main# sudo systemctl start postgresql
```

3. Setup Patroni on 2 Postgresql Nodes (bizdb01, bizdb02)

Install Patroni

```
root@bizdb01:~# curl https://bootstrap.pypa.io/pip/3.6/get-pip.py -o /tmp/get-pip.py -k
```


Note: /CN filled with Shared Virtual IP address of the HAProxy we created before.

Copy Cert and Key To Directory Created Before (/usr/patroni/conf) On Node 1

```
root@bizdb01:~# cp server.crt /usr/patroni/conf/server.crt
root@bizdb01:~# cp server.key /usr/patroni/conf/server.key
```

Change Cert and Key Files Mode and Owner On Node 1

```
root@bizdb01:~# cd /usr/patroni/conf
root@bizdb01:/usr/patroni/conf# chmod 400 server.crt
root@bizdb01:/usr/patroni/conf# chmod 400 server.key
root@bizdb01:/usr/patroni/conf# chown postgres:postgres server.key
root@bizdb01:/usr/patroni/conf# chown postgres:postgres server.crt
```

Setup Passwordless SSH Connection Between bizdb01 And bizdb02

Edit Etc Hosts File on Both Nodes To Include bizdb01 and bizdb02:

```
root@bizdb01:~# nano /etc/hosts
```

```
....
....
192.168.150.100 etcd01 bizdb01
192.168.150.101 etcd02 bizdb02
```

```
root@bizdb02:~# nano /etc/hosts
```

```
....
....
192.168.150.100 etcd01 bizdb01
192.168.150.101 etcd02 bizdb02
```

Change Postgres OS User Password on Both Nodes:

```
root@bizdb01:~# passwd postgres
New password:
Retype new password:
passwd: password updated successfully
```

```
root@bizdb02:~# passwd postgres
New password:
Retype new password:
passwd: password updated successfully
```

Generate Public/Private Key Pair Using Postgres User on Both Nodes:

```
root@bizdb01:~# sudo -i -u postgres
postgres@bizdb01:~$ ssh-keygen -t rsa -N ""
Generating public/private rsa key pair.
Enter file in which to save the key (/var/lib/postgresql/.ssh/id_rsa):
Created directory '/var/lib/postgresql/.ssh'.
Your identification has been saved in /var/lib/postgresql/.ssh/id_rsa
Your public key has been saved in /var/lib/postgresql/.ssh/id_rsa.pub
The key fingerprint is:
SHA256:dpOBn0ZdxsM711Vbr5NJg3Jg3QTKHF7kp86nyEMkKmU postgres@bizdb01
The key's randomart image is:
+---[RSA 3072]-----+
|      oo+B* +|
|      o *.B= |
|      . + B ++=|
|      Eo.*..oO.|
|      oS.Bo *o |
|      ...o ..o .|
|      . . o .|
|      ... o |
|      o.. |
+----[SHA256]-----+
```

```
root@bizdb02:~# sudo -i -u postgres
postgres@bizdb02:~$ ssh-keygen -t rsa -N ""
Generating public/private rsa key pair.
Enter file in which to save the key (/var/lib/postgresql/.ssh/id_rsa):
Created directory '/var/lib/postgresql/.ssh'.
Your identification has been saved in /var/lib/postgresql/.ssh/id_rsa
Your public key has been saved in /var/lib/postgresql/.ssh/id_rsa.pub
The key fingerprint is:
SHA256:VVtzX5tOFFPkCm+l6g4iXgQTYGAON6k35HpRqcpDz4ok postgres@bizdb02
The key's randomart image is:
+---[RSA 3072]-----+
|**      . oo=|
|=o      ..o.+=|
|... . . . .+ +o|
|.+.o o .. o + .|
|..*+. S. . o |
|+.*.o . |
|o *..+.. . |
|Eo. +.. + |
| . o |
```

+----[SHA256]-----+

Copy the Public Key From bizdb01 to bizdb02:

```
postgres@bizdb01:~$ ssh-copy-id postgres@bizdb02
```

```
/usr/bin/ssh-copy-id: INFO: Source of key(s) to be installed: "/var/lib/postgresql/.ssh/id_rsa.pub"
```

```
/usr/bin/ssh-copy-id: INFO: attempting to log in with the new key(s), to filter out any that are already installed
```

```
/usr/bin/ssh-copy-id: INFO: 1 key(s) remain to be installed -- if you are prompted now it is to install the new keys
```

```
postgres@bizdb02's password:
```

Number of key(s) added: 1

Now try logging into the machine, with: "ssh 'postgres@bizdb02'"
and check to make sure that only the key(s) you wanted were added.

Copy the Public Key From bizdb02 to bizdb01:

```
postgres@bizdb02:~$ ssh-copy-id postgres@bizdb01
```

```
/usr/bin/ssh-copy-id: INFO: Source of key(s) to be installed: "/var/lib/postgresql/.ssh/id_rsa.pub"
```

```
/usr/bin/ssh-copy-id: INFO: attempting to log in with the new key(s), to filter out any that are already installed
```

```
/usr/bin/ssh-copy-id: INFO: 1 key(s) remain to be installed -- if you are prompted now it is to install the new keys
```

```
postgres@bizdb01's password:
```

Number of key(s) added: 1

Now try logging into the machine, with: "ssh 'postgres@bizdb01'"
and check to make sure that only the key(s) you wanted were added.

Test Connect From Both Nodes to Both Nodes:

```
postgres@bizdb01:~$ ssh postgres@bizdb02
```

```
postgres@bizdb02:~$ ssh postgres@bizdb01
```

Copy Server.crt and Server.key To Node 2

Create Temporary Directory On Node 2 To Store Cert And Key:

```
root@bizdb02:/usr/patroni# sudo -i -u postgres
```

```
postgres@bizdb02:~$ pwd
```

```
/var/lib/postgresql
```

```
postgres@bizdb02:~$ mkdir tempcrt
```

SCP Server.crt and Server.key From Node 1 To Temporary Directory On Node 2:

```
root@bizdb01:~# sudo -i -u postgres
postgres@bizdb01:~$ scp -p /usr/patroni/conf/server.crt
postgres@bizdb02:/var/lib/postgresql/tmpcrt/server.crt
server.crt                                100% 1164  439.7KB/s
00:00
postgres@bizdb01:~$ scp -p /usr/patroni/conf/server.key
postgres@bizdb02:/var/lib/postgresql/tmpcrt/server.key
server.key                                100% 1675  233.1KB/s
00:00
```

On Node 2 As Root User, Copy Server.crt and Server.key To /usr/patroni/conf Directory:

```
root@bizdb02:~# cd /var/lib/postgresql/tmpcrt/
root@bizdb02:/var/lib/postgresql/tmpcrt# ls -ltr
total 8
-r----- 1 postgres postgres 1164 Mar 19 11:35 server.crt
-r----- 1 postgres postgres 1675 Mar 19 11:35 server.key
root@bizdb02:/var/lib/postgresql/tmpcrt# cp -p server.crt /usr/patroni/conf/
root@bizdb02:/var/lib/postgresql/tmpcrt# cp -p server.key /usr/patroni/conf/
```

Create Config File for Patroni (YAML)

Create Configuration File By Touch

```
root@bizdb01:~# touch /usr/patroni/conf/postgresql.yml
```

Prepare The YAML Config On Node 1 (bizdb01), The Green Highlight is Unique Value on Each Node:

```
scope: postgres
namespace: /bizdb_cluster/
name: bizdb01
restapi:
  listen: 192.168.150.100:8008
  connect_address: 192.168.150.100:8008
etcd:
  hosts: 192.168.150.100:2379, 192.168.150.101:2379, 192.168.150.110:2379
bootstrap:
  dcs:
    ttl: 30
    loop_wait: 10
    retry_timeout: 10
```

maximum_lag_on_failover: 1048576
 maximum_lag_on_syncnode: 15000000
 synchronous_mode: false
 postgresql:
 use_pg_rewind: true
 use_slots: true
 parameters:
 shared_buffers: 128MB
 work_mem: 4MB
 maintenance_work_mem: 64MB
 max_worker_processes: 8
 wal_buffers: 4MB
 max_wal_size: 1GB
 min_wal_size: 80MB
 effective_cache_size: 4GB
 fsync: on
 checkpoint_completion_target: 0.9
 log_rotation_size: 10MB
 listen_addresses: "*"
 max_connections: 100
 temp_buffers: 4MB
 ssl: true
 ssl_cert_file: /usr/patroni/conf/server.crt
 ssl_key_file: /usr/patroni/conf/server.key
 initdb:
 - encoding: UTF8
 - data-checksums
 pg_hba:
 - host replication replicator 127.0.0.1/32 md5
 - host replication replicator 192.168.150.100/32 md5
 - host replication replicator 192.168.150.101/32 md5
 - host all all 0.0.0.0/0 md5
 users:
 admin:
 password: admin
 options:
 - createrole
 - createdb

 postgresql:
 listen: 192.168.150.100:5432
 connect_address: 192.168.150.100:5432
 data_dir: /var/lib/postgresql/14/bizdb_cluster
 bin_dir: /usr/lib/postgresql/14/bin

pgpass: /tmp/pgpass
authentication:
 replication:
 username: replicator
 password: replicator
 superuser:
 username: postgres
 password: postgres
 rewind:
 username: pgrewind
 password: pgrewind
tags:
 nofailover: false
 noloadbalance: false
 clonefrom: false
 nosync: true

Prepare The YAML File On Node 2 (bizdb02), The Green Highlight is Unique Value on Each Node:

scope: postgres
namespace: /bizdb_cluster/
name: bizdb02
restapi:
 listen: 192.168.150.101:8008
 connect_address: 192.168.150.101:8008
etcd:
 hosts: 192.168.150.100:2379, 192.168.150.101:2379, 192.168.150.110:2379
bootstrap:
 dcs:
 ttl: 30
 loop_wait: 10
 retry_timeout: 10
 maximum_lag_on_failover: 1048576
 maximum_lag_on_syncnode: 15000000
 synchronous_mode: false
 postgresql:
 use_pg_rewind: true
 use_slots: true
 parameters:
 shared_buffers: 128MB
 work_mem: 4MB
 maintenance_work_mem: 64MB
 max_worker_processes: 8
 wal_buffers: 4MB

```

max_wal_size: 1GB
min_wal_size: 80MB
effective_cache_size: 4GB
fsync: on
checkpoint_completion_target: 0.9
log_rotation_size: 10MB
listen_addresses: "*"
max_connections: 100
temp_buffers: 4MB
ssl: true
ssl_cert_file: /usr/patroni/conf/server.crt
ssl_key_file: /usr/patroni/conf/server.key
initdb:
  - encoding: UTF8
  - data-checksums
pg_hba:
  - host replication replicator 127.0.0.1/32 md5
  - host replication replicator 192.168.150.100/32 md5
  - host replication replicator 192.168.150.101/32 md5
  - host all all 0.0.0.0/0 md5
users:
  admin:
    password: admin
    options:
      - createrole
      - createdb
postgresql:
  listen: 192.168.150.101:5432
  connect_address: 192.168.150.101:5432
  data_dir: /var/lib/postgresql/14/bizdb_cluster
  bin_dir: /usr/lib/postgresql/14/bin
  pgpass: /tmp/pgpass
  authentication:
    replication:
      username: replicator
      password: replicator
    superuser:
      username: postgres
      password: postgres
  rewind:
    username: pgrewind
    password: pgrewind
tags:

```

nofailover: false
noloadbalance: false
clonefrom: false
nosync: true

Edit Configuration File On Both Node 1 and Node 2 (bizdb01, bizdb02) Using Nano and Fill Each With YAML Config Node 1 and Node 2 Above:

root@bizdb01:~# **nano /usr/patroni/conf/postgresql.yml**
root@bizdb02:~# **nano /usr/patroni/conf/postgresql.yml**

Create Patroni Service

root@bizdb01:~# **nano /usr/lib/systemd/system/patroni.service**

[Unit]

Description=patroni
Documentation=https://patroni.readthedocs.io/en/latest/index.html
After=syslog.target network.target etcd.target
Wants=network-online.target

[Service]

Type=simple
User=postgres
Group=postgres
PermissionsStartOnly=true
ExecStart=/usr/local/bin/patroni /usr/patroni/conf/postgresql.yml
ExecReload=/bin/kill -HUP \$MAINPID
LimitNOFILE=65536
KillMode=process
KillSignal=SIGINT
Restart=on-abnormal
RestartSec=30s
TimeoutSec=0

[Install]

WantedBy=multi-user.target

Start Patroni on Both Nodes (bizdb01, bizdb02)

root@bizdb01:~# **sudo systemctl daemon-reload**
root@bizdb01:~# **sudo systemctl start patroni**
root@bizdb01:~# **sudo systemctl status patroni**

- patroni.service - patroni
Loaded: loaded (/lib/systemd/system/patroni.service; disabled; vendor preset: enabled)

Active: active (running) since Sun 2023-03-19 14:54:12 WIB; 2min 59s ago
 Docs: <https://patroni.readthedocs.io/en/latest/index.html>
 Main PID: 4301 (patroni)
 Tasks: 14 (limit: 4609)
 Memory: 136.4M
 CGroup: /system.slice/patroni.service

```

├─4301 /usr/bin/python3 /usr/local/bin/patroni /usr/patroni/conf/postgresql.yml
├─4342 /usr/lib/postgresql/14/bin/postgres -D /var/lib/postgresql/14/bizdb_cluster
--config-file=/var/lib/postgresql/14/bizdb_cluster>
├─4346 postgres: postgres: checkpointer
├─4347 postgres: postgres: background writer
├─4348 postgres: postgres: walwriter
├─4349 postgres: postgres: autovacuum launcher
├─4350 postgres: postgres: stats collector
├─4351 postgres: postgres: logical replication launcher
├─4354 postgres: postgres: postgres postgres 192.168.150.100(49336) idle
└─4365 postgres: postgres: walsender replicator 192.168.150.101(48170) streaming
0/3000060

Mar 19 14:55:34 bizdb01 patroni[4301]: 2023-03-19 14:55:34,383 INFO: no action. I am (bizdb01), the
leader with the lock

```

```

root@bizdb02:~# sudo systemctl daemon-reload
root@bizdb02:~# sudo systemctl start patroni
root@bizdb02:~# sudo systemctl enable patroni
root@bizdb02:~# sudo systemctl status patroni
● patroni.service - patroni
   Loaded: loaded (/lib/systemd/system/patroni.service; disabled; vendor preset: enabled)
   Active: active (running) since Sun 2023-03-19 14:54:30 WIB; 3min 26s ago
     Docs: https://patroni.readthedocs.io/en/latest/index.html
   Main PID: 3190 (patroni)
     Tasks: 12 (limit: 4609)
    Memory: 108.6M
    CGroup: /system.slice/patroni.service
           └─3190 /usr/bin/python3 /usr/local/bin/patroni /usr/patroni/conf/postgresql.yml
           └─3213 /usr/lib/postgresql/14/bin/postgres -D /var/lib/postgresql/14/bizdb_cluster
--config-file=/var/lib/postgresql/14/bizdb_cluster>
           └─3215 postgres: postgres: startup recovering 00000001000000000000000000000003
           └─3219 postgres: postgres: checkpointer
           └─3220 postgres: postgres: background writer
           └─3221 postgres: postgres: stats collector
           └─3227 postgres: postgres: postgres postgres 192.168.150.101(54304) idle
           └─3232 postgres: postgres: walreceiver streaming 0/3000060

```

Mar 19 14:56:24 bizdb02 patroni[3190]: 2023-03-19 14:56:24,391 INFO: no action. I am (bizdb02), a secondary, and following a leader (bizdb01)

4. Install HAProxy on 2 Nodes (haproxy01, haproxy02)

Enable PPA

```
root@haproxy01:~# sudo apt-get install --no-install-recommends software-properties-common
```

```
root@haproxy01:~# sudo add-apt-repository ppa:vbernat/haproxy-2.6
```

HAProxy is a free, very fast and reliable solution offering high availability, load balancing, and proxying for TCP and HTTP-based applications. It is particularly suited for web sites crawling under very high loads while needing persistence or Layer7 processing. Supporting tens of thousands of connections is clearly realistic with todays hardware. Its mode of operation makes its integration into existing architectures very easy and riskless, while still offering the possibility not to expose fragile web servers to the Net.

This PPA contains packages for HAProxy 2.6.

More info: <https://launchpad.net/~vbernat/+archive/ubuntu/haproxy-2.6>

Press [ENTER] to continue or Ctrl-c to cancel adding it.

Install HAProxy

```
root@haproxy01:~# sudo apt-get install haproxy=2.6.*
```

Reading package lists... Done

Building dependency tree

Reading state information... Done

Selected version '2.6.11-1ppa1~focal' (HAProxy 2.6:20.04/focal [amd64]) for 'haproxy'

The following additional packages will be installed:

liblua5.3-0

Suggested packages:

vim-haproxy haproxy-doc

The following NEW packages will be installed:

haproxy liblua5.3-0

0 upgraded, 2 newly installed, 0 to remove and 26 not upgraded.

Need to get 1,806 kB of archives.

After this operation, 4,480 kB of additional disk space will be used.

Do you want to continue? [Y/n] **y**

Check HAProxy Version

```
root@haproxy01:~# sudo haproxy -v
```

HAProxy version 2.6.11-1ppa1~focal 2023/03/18 - <https://haproxy.org/>

Status: long-term supported branch - will stop receiving fixes around Q2 2027.

Known bugs: <http://www.haproxy.org/bugs/bugs-2.6.11.html>

Running on: Linux 5.4.0-144-generic #161-Ubuntu SMP Fri Feb 3 14:49:04 UTC 2023 x86_64

Backup Existing Config File of HAProxy

```
root@haproxy01:~# cp -p /etc/haproxy/haproxy.cfg /etc/haproxy/haproxy.cfg.bak
```

Replace The Content of HAProxy Config File With Below

```
oot@haproxy01:~# cat /dev/null > /etc/haproxy/haproxy.cfg
```

```
root@haproxy01:~# nano /etc/haproxy/haproxy.cfg
```

global

```
log 127.0.0.1 local2
chroot /var/lib/haproxy
pidfile /var/run/haproxy.pid
maxconn 6000
user haproxy
group haproxy
daemon
stats socket /var/lib/haproxy/stats
```

defaults

```
mode tcp
log global
retries 3
timeout queue 1m
timeout connect 10s
timeout client 31m
timeout server 31m
timeout check 10s
maxconn 3000
```

listen stats

```
mode http
bind *:7000
stats enable
stats uri /
```

listen postgres

```
bind *:5000
option httpchk
http-check expect status 200
default-server inter 3s fall 3 rise 2 on-marked-down shutdown-sessions
server bizdb01 192.168.150.100:5432 maxconn 100 check port 8008
server bizdb02 192.168.150.101:5432 maxconn 100 check port 8008
```

listen postgres-readonly

```
bind *:6000
option httpchk GET /replica
```

```
http-check expect status 200
default-server inter 3s fall 3 rise 2 on-marked-down shutdown-sessions
server bizdb01 192.168.150.100:5432 maxconn 100 check port 8008
server bizdb02 192.168.150.101:5432 maxconn 100 check port 8008
```

Edit HAProxy Service Script and Add LimitNOFILE Param

```
root@haproxy01:~# nano /usr/lib/systemd/system/haproxy.service
```

```
....
....
[Service]
LimitNOFILE=6000
```

Reload Daemon and Restart HAProxy Service

```
root@haproxy01:~# sudo systemctl daemon-reload
root@haproxy01:~# sudo systemctl stop haproxy
root@haproxy01:~# sudo systemctl start haproxy
```

5. Test Connect To PostgreSQL Cluster Through HAProxy

Using Psql (Connect Without SSL)

```
root@bizdb01:~# sudo -i -u postgres
postgres@bizdb01:~$ psql -h 192.168.150.110 -p 5000 -d postgres -U postgres
Password for user postgres:
psql (14.7 (Ubuntu 14.7-1.pgdg20.04+1))
SSL connection (protocol: TLSv1.3, cipher: TLS_AES_256_GCM_SHA384, bits: 256, compression: off)
Type "help" for help.
```

```
postgres=# \q
postgres@bizdb01:~$ psql -h 192.168.150.111 -p 5000 -d postgres -U postgres
Password for user postgres:
psql (14.7 (Ubuntu 14.7-1.pgdg20.04+1))
SSL connection (protocol: TLSv1.3, cipher: TLS_AES_256_GCM_SHA384, bits: 256, compression: off)
Type "help" for help.
```

```
postgres=# \q
```

Using Psql (Connect With SSL)

```
postgres@bizdb01:~$ psql "host=192.168.150.110 port=5000 dbname=postgres user=postgres
password=postgres sslmode=verify-ca sslrootcert=/usr/patroni/conf/server.crt"
psql (14.7 (Ubuntu 14.7-1.pgdg20.04+1))
SSL connection (protocol: TLSv1.3, cipher: TLS_AES_256_GCM_SHA384, bits: 256, compression: off)
```

Type "help" for help.

```
postgres=# \q
```

```
postgres@bizdb01:~$ psql "host=192.168.150.111 port=5000 dbname=postgres user=postgres password=postgres sslmode=verify-ca sslrootcert=/usr/patroni/conf/server.crt"
```

```
psql (14.7 (Ubuntu 14.7-1.pgdg20.04+1))
```

```
SSL connection (protocol: TLSv1.3, cipher: TLS_AES_256_GCM_SHA384, bits: 256, compression: off)
```

```
Type "help" for help.
```

```
postgres=# \q
```

Using Psql, Check Read-Write or Read-Only Connection on Different Ports

```
postgres@bizdb01:~$ psql -h 192.168.150.110 -U postgres -p 5000 -c "SELECT pg_is_in_recovery();"
```

```
Password for user postgres:
```

```
pg_is_in_recovery
```

```
-----  
f
```

```
(1 row)
```

```
postgres@bizdb01:~$ psql -h 192.168.150.110 -U postgres -p 6000 -c "SELECT pg_is_in_recovery();"
```

```
Password for user postgres:
```

```
pg_is_in_recovery
```

```
-----  
t
```

```
(1 row)
```

Check PostgreSQL Cluster Through HAProxy Stats Dashboard

<http://192.168.150.110:7000>

Statistics Report for HAProxy

HAProxy version 2.6.11-1ppa1~focal, released 2023/03/18

Statistics Report for pid 3305

> General process information

pid = 3305 (process #1, nbproc = 1, nbthread = 2)
 uptime = 0s 0h46m31s
 system limits: memmax = unlimited, ulimit-n = 12050
 maxsock = 12050, maxconn = 6000, maxpipes = 0
 current conns = 2, current pipes = 0/0, conn rate = 1/sec, bit rate = 2.175 kbps
 Running tasks: 0/21, idle = 100 %

active UP
 active UP, going down
 active DOWN, going up
 active or backup DOWN
 active or backup DOWN for maintenance (MAINT)
 active or backup SOFT STOPPED for maintenance
 Note: "NOLB"/"DRAIN" = UP with load-balancing disabled

backup UP
 backup UP, going down
 backup DOWN, going up
 not checked

Display option:
 • Scope:
 • Hide DOWN servers
 • Refresh now
 • CSV export
 • JSON export (schema)

External resources:
 • Primary site
 • Updates (v2.6)
 • Online manual

state		Queue		Session rate		Sessions		Sessions		Bytes		Denied		Errors		Warnings		Server														
		Cur	Max	Limit	Cur	Max	Limit	Cur	Max	Limit	Total	LbTot	Last	In	Out	Req	Resp	Req	Conn	Resp	Retr	Redis	Status	LastChk	Wght	Act	Bck	Chk	Dwn	Dwntme	Thrtle	
Frontend		0	0		1	2	-	2	3	3 000	3 000	5		13 616	27 686	0	0	0	0	0	0	0	0	OPEN								
Backend		0	0		0	0		0	0	300	300	0	0s	13 616	27 686	0	0	0	0	0	0	0	0	46m31s UP		0/0	0	0	0	0		

postgres		Queue		Session rate		Sessions		Sessions		Bytes		Denied		Errors		Warnings		Server														
		Cur	Max	Limit	Cur	Max	Limit	Cur	Max	Limit	Total	LbTot	Last	In	Out	Req	Resp	Req	Conn	Resp	Retr	Redis	Status	LastChk	Wght	Act	Bck	Chk	Dwn	Dwntme	Thrtle	
Frontend		0	0	-	0	2	-	0	2	3 000	3 000	14		13 616	27 686	0	0	0	0	0	0	0	0	OPEN								
bizdb01		0	0	-	0	2	-	0	2	100	100	14	14	6m59s	13 616	27 686	0	0	0	0	0	0	0	46m31s UP	L7OK/200 in 3ms	1/1	Y	-	0	0	0s	-
bizdb02		0	0	-	0	0	-	0	0	100	100	0	0	?	0	0	0	0	0	0	0	0	46m30s DOWN	L7STS/503 in 7ms	1/1	Y	-	1	1	46m30s	-	
Backend		0	0		0	2		0	2	300	300	14	14	6m59s	13 616	27 686	0	0	0	0	0	0	46m31s UP		1/1	1	0	0	0	0s		

postgres-readonly		Queue		Session rate		Sessions		Sessions		Bytes		Denied		Errors		Warnings		Server														
		Cur	Max	Limit	Cur	Max	Limit	Cur	Max	Limit	Total	LbTot	Last	In	Out	Req	Resp	Req	Conn	Resp	Retr	Redis	Status	LastChk	Wght	Act	Bck	Chk	Dwn	Dwntme	Thrtle	
Frontend		0	1	-	0	1	-	0	1	3 000	3 000	5		5 938	11 980	0	0	0	0	0	0	0	0	OPEN								
bizdb01		0	0	-	0	0	-	0	0	100	100	0	0	?	0	0	0	0	0	0	0	0	46m29s DOWN	L7STS/503 in 3ms	1/1	Y	-	1	1	46m29s	-	
bizdb02		0	0	-	0	1	-	0	1	100	100	5	5	6m46s	5 938	11 980	0	0	0	0	0	0	46m31s UP	L7OK/200 in 3ms	1/1	Y	-	0	0	0s	-	
Backend		0	0		0	1		0	1	300	300	5	5	6m46s	5 938	11 980	0	0	0	0	0	0	46m31s UP		1/1	1	0	0	0	0s		

6. Patronictl Commands

List Nodes:

```
root@bizdb01:~# patronictl -c /usr/patroni/conf/postgresql.yml list
+ Cluster: postgres -----+-----+-----+-----+-----+-----+
| Member | Host           | Role   | State | TL | Lag in MB | Tags          |
+-----+-----+-----+-----+-----+-----+
| bizdb01 | 192.168.150.100 | Leader | running | 4 |          | nosync: true |
| bizdb02 | 192.168.150.101 | Replica | running | 4 |          | nosync: true |
+-----+-----+-----+-----+-----+-----+

```

Failover:

```
root@bizdb01:~# patronictl -c /usr/patroni/conf/postgresql.yml failover
```

Current cluster topology

```
+ Cluster: postgres -----+-----+-----+-----+-----+-----+
| Member | Host           | Role   | State | TL | Lag in MB | Tags          |
+-----+-----+-----+-----+-----+-----+
| bizdb01 | 192.168.150.100 | Leader | running | 4 |          | nosync: true |
| bizdb02 | 192.168.150.101 | Replica | running | 4 |          | nosync: true |
+-----+-----+-----+-----+-----+-----+

```

Candidate ['bizdb02'] []: **bizdb02**

Are you sure you want to failover cluster postgres, demoting current leader bizdb01? [y/N]: **y**

2023-03-20 17:21:38.43431 Successfully failed over to "bizdb02"

```
+ Cluster: postgres -----+-----+-----+-----+-----+-----+
| Member | Host           | Role   | State | TL | Lag in MB | Tags          |
+-----+-----+-----+-----+-----+-----+

```

```

+-----+-----+-----+-----+---+-----+-----+
| bizdb01 | 192.168.150.100 | Replica | stopped |   | unknown | nosync: true |
| bizdb02 | 192.168.150.101 | Leader  | running | 4 |         | nosync: true |
+-----+-----+-----+-----+---+-----+-----+
root@bizdb01:~# patronictl -c /usr/patroni/conf/postgresql.yml list
+ Cluster: postgres -----+-----+-----+-----+-----+
| Member | Host           | Role   | State   | TL | Lag in MB | Tags          |
+-----+-----+-----+-----+---+-----+-----+
| bizdb01 | 192.168.150.100 | Replica | running | 5 | 0         | nosync: true |
| bizdb02 | 192.168.150.101 | Leader  | running | 5 |          | nosync: true |
+-----+-----+-----+-----+---+-----+-----+

```

7. Install Keepalived on HAProxy Nodes

Determine The Default Network Interface on Both Nodes For Virtual IP Address

```

root@haproxy01:~# ifconfig -a
enp0s3: flags=4163<UP,BROADCAST,RUNNING,MULTICAST> mtu 1500
    inet 10.0.2.15 netmask 255.255.255.0 broadcast 10.0.2.255
    inet6 fe80::a00:27ff:fe17:643a prefixlen 64 scopeid 0x20<link>
    ether 08:00:27:17:64:3a txqueuelen 1000 (Ethernet)
    RX packets 7018 bytes 10120224 (10.1 MB)
    RX errors 0 dropped 0 overruns 0 frame 0
    TX packets 1011 bytes 80841 (80.8 KB)
    TX errors 0 dropped 0 overruns 0 carrier 0 collisions 0

enp0s8: flags=4163<UP,BROADCAST,RUNNING,MULTICAST> mtu 1500
    inet 192.168.150.110 netmask 255.255.255.0 broadcast 192.168.150.255
    inet6 fe80::a00:27ff:fe75:5361 prefixlen 64 scopeid 0x20<link>
    ether 08:00:27:75:53:61 txqueuelen 1000 (Ethernet)
    RX packets 1017273 bytes 103387316 (103.3 MB)
    RX errors 0 dropped 0 overruns 0 frame 0
    TX packets 1005278 bytes 103836782 (103.8 MB)
    TX errors 0 dropped 0 overruns 0 carrier 0 collisions 0

lo: flags=73<UP,LOOPBACK,RUNNING> mtu 65536
    inet 127.0.0.1 netmask 255.0.0.0
    inet6 ::1 prefixlen 128 scopeid 0x10<host>
    loop txqueuelen 1000 (Local Loopback)
    RX packets 9364 bytes 506298 (506.2 KB)
    RX errors 0 dropped 0 overruns 0 frame 0
    TX packets 9364 bytes 506298 (506.2 KB)
    TX errors 0 dropped 0 overruns 0 carrier 0 collisions 0

```

In my case, I will use **enp0s8** because it is the interface where HAProxy IP resides.

Install Keepalived From APT Repository on Both Nodes

```
root@haproxy01:~# sudo apt-get install keepalived -y
```

```
root@haproxy02:~# sudo apt-get install keepalived -y
```

Using Nano, Create keepalived.conf on Node 1

```
root@haproxy01:~# nano /etc/keepalived/keepalived.conf
```

```
global_defs {
    router_id gd_ha01
    default_interface enp0s8
}
vrrp_script chk_haproxy {
    script "killall -0 haproxy" # check the haproxy process
    interval 2                # every 2 seconds
    fall 2
    rise 2
}
vrrp_instance VI_1 {
    interface enp0s8          # interface to monitor
    nopreempt
    state BACKUP               # use nopreempt, so BACKUP on both of haproxy01 and haproxy02
    virtual_router_id 91      # use a unique id shared between haproxy01 and haproxy02
    priority 101              # 101 on haproxy01, 100 on haproxy02
    advert_int 1
    unicast_src_ip 192.168.150.110
    unicast_peer {
        192.168.150.111
    }
    authentication {
        auth_type PASS
        auth_pass haproxy1234 # specify the same password for haproxy01 and haproxy02
    }
    virtual_ipaddress {
        192.168.150.200 # specify a virtual ip address agreed before when generating SSL for
    }
}
track_script {
    chk_haproxy
}
```

Using Nano, Create keepalived.conf on Node 2

root@haproxy02:~# nano /etc/keepalived/keepalived.conf

```
global_defs {
    router_id gd_ha02
    default_interface enp0s8
}
vrrp_script chk_haproxy {
    script "killall -0 haproxy" # check the haproxy process
    interval 2                # every 2 seconds
    fall 2
    rise 2
}
vrrp_instance VI_1 {
    interface enp0s8           # interface to monitor
    nopreempt
    state BACKUP               # use nopreempt, so BACKUP on both of haproxy01 and haproxy02
    virtual_router_id 91       # use a unique id shared between haproxy01 and haproxy02
    priority 100               # 101 on haproxy01, 100 on haproxy02
    advert_int 1
    authentication {
        auth_type PASS
        auth_pass haproxy1234 # specify the same password for haproxy01 and haproxy02
    }
    unicast_src_ip 192.168.150.111
    unicast_peer {
        192.168.150.110
    }
    virtual_ipaddress {
        192.168.150.200 # specify a virtual ip address agreed before when generating SSL for
    }
    track_script {
        chk_haproxy
    }
}
```

Configure IP Forwarding and Non-local Binding on Both HAProxy Nodes

root@haproxy01:~# nano /etc/sysctl.conf

....

....

net.ipv4.ip_forward = 1

net.ipv4.ip_nonlocal_bind = 1

```
root@haproxy01:~# sysctl -p
net.ipv4.ip_forward = 1
net.ipv4.ip_nonlocal_bind = 1
```

```
root@haproxy02:~# nano /etc/sysctl.conf
```

```
....
```

```
....
```

```
net.ipv4.ip_forward = 1
net.ipv4.ip_nonlocal_bind = 1
```

```
root@haproxy02:~# sysctl -p
net.ipv4.ip_forward = 1
net.ipv4.ip_nonlocal_bind = 1
```

Start Keepalived Services on Both HAProxy Nodes

```
root@haproxy01:~# sudo systemctl enable keepalived
```

Synchronizing state of keepalived.service with SysV service script with
/lib/systemd/systemd-sysv-install.

Executing: /lib/systemd/systemd-sysv-install enable keepalived

```
root@haproxy01:~# sudo systemctl start keepalived
```

```
root@haproxy01:~# sudo systemctl status keepalived
```

- keepalived.service - Keepalived Daemon (LVS and VRRP)
 - Loaded: loaded (/lib/systemd/system/keepalived.service; enabled; vendor preset: enabled)
 - Active: **active (running)** since Mon 2023-03-20 18:27:26 WIB; 13s ago
 - Main PID: 4970 (keepalived)
 - Tasks: 2 (limit: 4609)
 - Memory: 3.2M
 - CGroup: /system.slice/keepalived.service
 - └─4970 /usr/sbin/keepalived --dont-fork
 - └─4971 /usr/sbin/keepalived --dont-fork

Check The Virtual IP Address On Both Nodes

```
root@haproxy01:~# ip addr show dev enp0s8
```

```
3: enp0s8: <BROADCAST,MULTICAST,UP,LOWER_UP> mtu 1500 qdisc fq_codel state UP group
default qlen 1000
```

```
link/ether 08:00:27:75:53:61 brd ff:ff:ff:ff:ff:ff
```

```
inet 192.168.150.110/24 brd 192.168.150.255 scope global enp0s8
```

```
valid_lft forever preferred_lft forever
```

```
inet 192.168.150.200/32 scope global enp0s8
```

```
valid_lft forever preferred_lft forever
```

```
inet6 fe80::a00:27ff:fe75:5361/64 scope link
    valid_lft forever preferred_lft forever
```

```
root@haproxy02:~# ip addr show dev enp0s8
```

```
3: enp0s8: <BROADCAST,MULTICAST,UP,LOWER_UP> mtu 1500 qdisc fq_codel state UP group
default qlen 1000
    link/ether 08:00:27:9e:4a:3a brd ff:ff:ff:ff:ff:ff
    inet 192.168.150.111/24 brd 192.168.150.255 scope global enp0s8
        valid_lft forever preferred_lft forever
    inet6 fe80::a00:27ff:fe9e:4a3a/64 scope link
        valid_lft forever preferred_lft forever
```

The IP Address should only appear on one of the two nodes of the HAProxy hosts.

Change Bind Address From * to The Virtual IP Address in Both HAProxy Config

```
root@haproxy01:~# nano /etc/haproxy/haproxy.cfg
```

```
listen stats
....
bind 192.168.150.200:7000
....
listen postgres
....
bind 192.168.150.200:5000
....
listen postgres-readonly
....
bind 192.168.150.200:6000
....
```

```

defaults
  mode tcp
  log global
  retries 3
  timeout queue 1m
  timeout connect 10s
  timeout client 31m
  timeout server 31m
  timeout check 10s
  maxconn 3000
listen stats
  mode http
  bind 192.168.150.200:7000
  stats enable
  stats uri /
listen postgres
  bind 192.168.150.200:5000
  option httpchk
  http-check expect status 200
  default-server inter 3s fall 3 rise 2 on-marked-down shutdown-sessions
  server bizdb01 192.168.150.100:5432 maxconn 100 check port 8008
  server bizdb02 192.168.150.101:5432 maxconn 100 check port 8008
listen postgres-readonly
  bind 192.168.150.200:6000
  option httpchk GET /replica
  http-check expect status 200
  default-server inter 3s fall 3 rise 2 on-marked-down shutdown-sessions
  server bizdb01 192.168.150.100:5432 maxconn 100 check port 8008
  server bizdb02 192.168.150.101:5432 maxconn 100 check port 8008

```

root@haproxy01:~# **sudo systemctl stop haproxy**

root@haproxy01:~# **sudo systemctl start haproxy**

Test Connect to PostgreSQL Cluster Using Virtual IP Address

Without SSL:

root@bizdb01:~# **sudo -i -u postgres**

postgres@bizdb01:~\$ **psql -h 192.168.150.200 -p 5000 -d postgres -U postgres**

Password for user postgres:

psql (14.7 (Ubuntu 14.7-1.pgdg20.04+1))

SSL connection (protocol: TLSv1.3, cipher: TLS_AES_256_GCM_SHA384, bits: 256, compression: off)

Type "help" for help.

postgres=# \q

With SSL:

postgres@bizdb01:~\$ **psql "host=192.168.150.200 port=5000 dbname=postgres user=postgres password=postgres sslmode=verify-full sslrootcert=/usr/patroni/conf/server.crt"**

psql (14.7 (Ubuntu 14.7-1.pgdg20.04+1))

SSL connection (protocol: TLSv1.3, cipher: TLS_AES_256_GCM_SHA384, bits: 256, compression: off)
Type "help" for help.

postgres=# \q

Check PostgreSQL Cluster Through HAProxy Stats Dashboard Using Virtual IP Address

http://192.168.150.200:7000

8. Test Keepalived Service For HAProxy

Check Keepalived Service Logs on Both Nodes

root@haproxy01:~# journalctl -u keepalived -f

-- Logs begin at Fri 2023-03-17 19:21:01 WIB. --

Mar 20 18:27:26 haproxy01 Keepalived_vrrp[4971]: (VI_1) Entering BACKUP STATE

Mar 20 18:27:29 haproxy01 Keepalived_vrrp[4971]: (VI_1) Entering MASTER STATE

Mar 20 19:02:37 haproxy01 Keepalived_vrrp[4971]: Script `chk_haproxy` now returning 1

Mar 20 19:02:39 haproxy01 Keepalived_vrrp[4971]: VRRP_Script(chk_haproxy) failed (exited with status 1)

Mar 20 19:02:39 haproxy01 Keepalived_vrrp[4971]: (VI_1) Entering FAULT STATE

Mar 20 19:02:39 haproxy01 Keepalived_vrrp[4971]: (VI_1) sent 0 priority

Mar 20 19:02:47 haproxy01 Keepalived_vrrp[4971]: Script `chk_haproxy` now returning 0

Mar 20 19:02:49 haproxy01 Keepalived_vrrp[4971]: VRRP_Script(chk_haproxy) succeeded

Mar 20 19:02:49 haproxy01 Keepalived_vrrp[4971]: (VI_1) Entering BACKUP STATE

Mar 20 19:02:53 haproxy01 Keepalived_vrrp[4971]: (VI_1) Entering MASTER STATE

root@haproxy02:~# journalctl -u keepalived -f

-- Logs begin at Fri 2023-03-17 19:21:01 WIB. --

Mar 20 18:27:33 haproxy02 Keepalived_vrrp[4083]: SECURITY VIOLATION - scripts are being executed but script_security not enabled.

Mar 20 18:27:33 haproxy02 Keepalived_vrrp[4083]: Registering gratuitous ARP shared channel

Mar 20 18:27:33 haproxy02 Keepalived_vrrp[4083]: VRRP_Script(chk_haproxy) succeeded

Mar 20 18:27:33 haproxy02 Keepalived_vrrp[4083]: (VI_1) Entering BACKUP STATE

Mar 20 19:02:32 haproxy02 Keepalived_vrrp[4083]: Script `chk_haproxy` now returning 1

Mar 20 19:02:34 haproxy02 Keepalived_vrrp[4083]: VRRP_Script(chk_haproxy) failed (exited with status 1)

Mar 20 19:02:34 haproxy02 Keepalived_vrrp[4083]: (VI_1) Entering FAULT STATE

Mar 20 19:02:52 haproxy02 Keepalived_vrrp[4083]: Script `chk_haproxy` now returning 0

Mar 20 19:02:54 haproxy02 Keepalived_vrrp[4083]: VRRP_Script(chk_haproxy) succeeded

Mar 20 19:02:54 haproxy02 Keepalived_vrrp[4083]: (VI_1) Entering BACKUP STATE

Simulate HAProxy Service Problem on Node 1

root@haproxy01:~# **sudo systemctl stop haproxy**

Keepalived Log on Node 1 (Now Node 1 becomes FAULT):

```
root@haproxy01:~# journalctl -u keepalived -f
-- Logs begin at Fri 2023-03-17 19:21:01 WIB. --
Mar 20 18:27:26 haproxy01 Keepalived_vrrp[4971]: (VI_1) Entering BACKUP STATE
Mar 20 18:27:29 haproxy01 Keepalived_vrrp[4971]: (VI_1) Entering MASTER STATE
Mar 20 19:02:37 haproxy01 Keepalived_vrrp[4971]: Script `chk_haproxy` now returning 1
Mar 20 19:02:39 haproxy01 Keepalived_vrrp[4971]: VRRP_Script(chk_haproxy) failed (exited with status 1)
Mar 20 19:02:39 haproxy01 Keepalived_vrrp[4971]: (VI_1) Entering FAULT STATE
Mar 20 19:02:39 haproxy01 Keepalived_vrrp[4971]: (VI_1) sent 0 priority
Mar 20 19:02:47 haproxy01 Keepalived_vrrp[4971]: Script `chk_haproxy` now returning 0
Mar 20 19:02:49 haproxy01 Keepalived_vrrp[4971]: VRRP_Script(chk_haproxy) succeeded
Mar 20 19:02:49 haproxy01 Keepalived_vrrp[4971]: (VI_1) Entering BACKUP STATE
Mar 20 19:02:53 haproxy01 Keepalived_vrrp[4971]: (VI_1) Entering MASTER STATE
Mar 20 19:21:43 haproxy01 Keepalived_vrrp[4971]: Script `chk_haproxy` now returning 1
Mar 20 19:21:45 haproxy01 Keepalived_vrrp[4971]: VRRP_Script(chk_haproxy) failed (exited with status 1)
Mar 20 19:21:45 haproxy01 Keepalived_vrrp[4971]: (VI_1) Entering FAULT STATE
Mar 20 19:21:45 haproxy01 Keepalived_vrrp[4971]: (VI_1) sent 0 priority
```

Keepalived Log on Node 2 (Now Node 2 becomes MASTER):

```
root@haproxy02:~# journalctl -u keepalived -f
-- Logs begin at Fri 2023-03-17 19:21:01 WIB. --
Mar 20 18:27:33 haproxy02 Keepalived_vrrp[4083]: SECURITY VIOLATION - scripts are being executed but script_security not enabled.
Mar 20 18:27:33 haproxy02 Keepalived_vrrp[4083]: Registering gratuitous ARP shared channel
Mar 20 18:27:33 haproxy02 Keepalived_vrrp[4083]: VRRP_Script(chk_haproxy) succeeded
Mar 20 18:27:33 haproxy02 Keepalived_vrrp[4083]: (VI_1) Entering BACKUP STATE
Mar 20 19:02:32 haproxy02 Keepalived_vrrp[4083]: Script `chk_haproxy` now returning 1
Mar 20 19:02:34 haproxy02 Keepalived_vrrp[4083]: VRRP_Script(chk_haproxy) failed (exited with status 1)
Mar 20 19:02:34 haproxy02 Keepalived_vrrp[4083]: (VI_1) Entering FAULT STATE
Mar 20 19:02:52 haproxy02 Keepalived_vrrp[4083]: Script `chk_haproxy` now returning 0
Mar 20 19:02:54 haproxy02 Keepalived_vrrp[4083]: VRRP_Script(chk_haproxy) succeeded
Mar 20 19:02:54 haproxy02 Keepalived_vrrp[4083]: (VI_1) Entering BACKUP STATE
Mar 20 19:21:45 haproxy02 Keepalived_vrrp[4083]: (VI_1) Backup received priority 0 advertisement
Mar 20 19:21:45 haproxy02 Keepalived_vrrp[4083]: (VI_1) Backup received priority 0 advertisement
Mar 20 19:21:46 haproxy02 Keepalived_vrrp[4083]: (VI_1) Entering MASTER STATE
```

Check Virtual IP Address on HAProxy Node 2, The Virtual IP now is switch to there:

root@haproxy02:~# **ip addr show dev enp0s8**

3: enp0s8: <BROADCAST,MULTICAST,UP,LOWER_UP> mtu 1500 qdisc fq_codel state UP group default qlen 1000

```
link/ether 08:00:27:9e:4a:3a brd ff:ff:ff:ff:ff:ff
inet 192.168.150.111/24 brd 192.168.150.255 scope global enp0s8
    valid_lft forever preferred_lft forever
inet 192.168.150.200/32 scope global enp0s8
    valid_lft forever preferred_lft forever
inet6 fe80::a00:27ff:fe9e:4a3a/64 scope link
    valid_lft forever preferred_lft forever
```

Simulate HAProxy Service on Node 1 is Backed Up Again

```
root@haproxy01:~# sudo systemctl start haproxy
```

Keepalived Log on Node 1 (Now Node 1 becomes BACKUP):

```
root@haproxy01:~# journalctl -u keepalived -f
-- Logs begin at Fri 2023-03-17 19:21:01 WIB. --
Mar 20 18:27:26 haproxy01 Keepalived_vrrp[4971]: (VI_1) Entering BACKUP STATE
Mar 20 18:27:29 haproxy01 Keepalived_vrrp[4971]: (VI_1) Entering MASTER STATE
Mar 20 19:02:37 haproxy01 Keepalived_vrrp[4971]: Script `chk_haproxy` now returning 1
Mar 20 19:02:39 haproxy01 Keepalived_vrrp[4971]: VRRP_Script(chk_haproxy) failed (exited with status 1)
Mar 20 19:02:39 haproxy01 Keepalived_vrrp[4971]: (VI_1) Entering FAULT STATE
Mar 20 19:02:39 haproxy01 Keepalived_vrrp[4971]: (VI_1) sent 0 priority
Mar 20 19:02:47 haproxy01 Keepalived_vrrp[4971]: Script `chk_haproxy` now returning 0
Mar 20 19:02:49 haproxy01 Keepalived_vrrp[4971]: VRRP_Script(chk_haproxy) succeeded
Mar 20 19:02:49 haproxy01 Keepalived_vrrp[4971]: (VI_1) Entering BACKUP STATE
Mar 20 19:02:53 haproxy01 Keepalived_vrrp[4971]: (VI_1) Entering MASTER STATE
Mar 20 19:21:43 haproxy01 Keepalived_vrrp[4971]: Script `chk_haproxy` now returning 1
Mar 20 19:21:45 haproxy01 Keepalived_vrrp[4971]: VRRP_Script(chk_haproxy) failed (exited with status 1)
Mar 20 19:21:45 haproxy01 Keepalived_vrrp[4971]: (VI_1) Entering FAULT STATE
Mar 20 19:21:45 haproxy01 Keepalived_vrrp[4971]: (VI_1) sent 0 priority
Mar 20 19:31:18 haproxy01 Keepalived_vrrp[4971]: Script `chk_haproxy` now returning 0
Mar 20 19:31:20 haproxy01 Keepalived_vrrp[4971]: VRRP_Script(chk_haproxy) succeeded
Mar 20 19:31:20 haproxy01 Keepalived_vrrp[4971]: (VI_1) Entering BACKUP STATE
```