

50 Essential PostgreSQL Queries by Category: A Complete Guide

Table of Contents

1. Database Administration
2. Performance Monitoring
3. Security Management
4. Data Manipulation
5. Advanced Features

Database Administration

Database and Table Information

-
<ul style="list-style-type: none">● - List all databases <pre>SELECT datname, datdba, encoding, datcollate, datctype FROM pg_database;</pre>
<ul style="list-style-type: none">● -- List all schemas <pre>SELECT schema_name, schema_owner FROM information_schema.schemata;</pre>
<ul style="list-style-type: none">● -- List tables with sizes <pre>SELECT schemaname as schema_name, relname as table_name, pg_size_pretty(pg_total_relation_size(relid)) as total_size, pg_size_pretty(pg_table_size(relid)) as table_size, pg_size_pretty(pg_indexes_size(relid)) as index_size FROM pg_catalog.pg_statio_user_tables ORDER BY pg_total_relation_size(relid) DESC;</pre>
<ul style="list-style-type: none">● -- Show table column information <pre>SELECT column_name, data_type, character_maximum_length, column_default, is_nullable FROM information_schema.columns WHERE table_schema = 'public' AND table_name = 'your_table';</pre>

-

- - List all indexes

```
SELECT
    schemaname,
    tablename,
    indexname,
    indexdef
FROM pg_indexes
WHERE schemaname = 'public'
ORDER BY tablename, indexname;
```

sql

● Backup and Maintenance

```
-- Analyze table statistics
ANALYZE VERBOSE your_table;

-- Vacuum table
VACUUM (VERBOSE, ANALYZE) your_table;

-- Reindex table
REINDEX TABLE your_table;

-- Show last vacuum and analyze time
SELECT
    relname as table_name,
    last_vacuum,
    last_autovacuum,
    last_analyze,
    last_autoanalyze
FROM pg_stat_user_tables;

-- Show dead tuples and autovacuum status
SELECT
    schemaname,
    relname,
    n_dead_tup,
    n_live_tup,
    n_mod_since_analyze
FROM pg_stat_user_tables;
```

sql

Performance Monitoring

● Query Performance

```

-- Show slow queries
SELECT
    pid,
    age(clock_timestamp(), query_start) as duration,
    query
FROM pg_stat_activity
WHERE state != 'idle'
ORDER BY duration DESC;

-- Identify queries with high total time
SELECT
    query,
    calls,
    total_time,
    rows,
    mean_time
FROM pg_stat_statements
ORDER BY total_time DESC
LIMIT 10;

-- Find queries with most block I/O
SELECT
    query,
    shared_blks_hit,
    shared_blks_read,
    shared_blks_written
FROM pg_stat_statements
ORDER BY shared_blks_read + shared_blks_written DESC
LIMIT 10;

-- Show table cache hit ratios
SELECT
    relname as table_name,
    heap_blks_read as blocks_read,
    heap_blks_hit as blocks_hit,
    CASE WHEN heap_blks_hit + heap_blks_read = 0
        THEN 0
        ELSE round(heap_blks_hit::numeric / (heap_blks_hit + heap_blks_read), 4)
    END as cache_hit_ratio
FROM pg_statio_user_tables
ORDER BY cache_hit_ratio DESC;

-- Index usage statistics
SELECT
    schemaname,
    relname,
    indexrelname,
    idx_scan,
    idx_tup_read,
    idx_tup_fetch
FROM pg_stat_user_indexes
ORDER BY idx_scan DESC;

```

Security Management

User and Role Management

```
-- List all roles
SELECT
    rolname,
    rolsuper,
    rolcreatorole,
    rolcreatedb,
    rolcanlogin
FROM pg_roles;

-- Show role permissions
SELECT
    grantee, table_schema, table_name, privilege_type
FROM information_schema.role_table_grants
WHERE grantee = 'your_role';

-- List active sessions by user
SELECT
    username,
    count(*) as session_count
FROM pg_stat_activity
GROUP BY username;

-- Show connection limits per role
SELECT
    rolname,
    rolconlimit
FROM pg_roles
WHERE rolconlimit <> -1;

-- Audit role memberships
SELECT
    pg_get_userbyid(member) as member,
    pg_get_userbyid(roleid) as role
FROM pg_auth_members;
```

Data Manipulation

Advanced Queries

```
-- Common Table Expression (CTE)
WITH ranked_orders AS (
    SELECT
        customer_id,
        order_date,
```

```

        total_amount,
        RANK() OVER (PARTITION BY customer_id ORDER BY total_amount DESC) as rank
    FROM orders
)
SELECT * FROM ranked_orders WHERE rank = 1;

-- Window Functions
SELECT
    department,
    employee_name,
    salary,
    AVG(salary) OVER (PARTITION BY department) as dept_avg,
    salary - AVG(salary) OVER (PARTITION BY department) as diff_from_avg
FROM employees;

-- JSON Operations
SELECT
    id,
    data->>'name' as name,
    (data->>'age')::int as age,
    jsonb_array_elements(data->'hobbies') as hobby
FROM users
WHERE (data->>'age')::int > 25;

-- Full Text Search
SELECT
    title,
    ts_rank_cd(to_tsvector('english', content), query) as rank
FROM articles, plainto_tsquery('english', 'your search terms') query
WHERE to_tsvector('english', content) @@ query
ORDER BY rank DESC;

-- Recursive Queries
WITH RECURSIVE subordinates AS (
    -- Base case
    SELECT employee_id, manager_id, name, 1 as level
    FROM employees
    WHERE employee_id = 1

    UNION ALL

    -- Recursive case
    SELECT e.employee_id, e.manager_id, e.name, s.level + 1
    FROM employees e
    INNER JOIN subordinates s ON s.employee_id = e.manager_id
)
SELECT * FROM subordinates;

```

sql

Advanced Features

Partitioning and Inheritance

```

-- Create partitioned table
CREATE TABLE measurements (

```

```

    city_id    int not null,
    logdate    date not null,
    peaktemp   int,
    unitsales  int
) PARTITION BY RANGE (logdate);

-- Create partition
CREATE TABLE measurements_y2025m01 PARTITION OF measurements
    FOR VALUES FROM ('2025-01-01') TO ('2025-02-01');

-- Show partitions
SELECT
    parent.relname as table_name,
    child.relname as partition_name,
    pg_get_expr(child.relpartbound, child.oid) as partition_expression
FROM pg_inherits
JOIN pg_class parent ON pg_inherits.inhparent = parent.oid
JOIN pg_class child ON pg_inherits.inhrelid = child.oid
WHERE parent.relname = 'measurements';

```

Materialized Views

```

-- Create materialized view

CREATE MATERIALIZED VIEW sales_summary AS
SELECT
    date_trunc('month', order_date) as month,
    product_category,
    SUM(amount) as total_sales,
    COUNT(*) as order_count
FROM orders
GROUP BY 1, 2
WITH DATA;

-- Refresh materialized view

REFRESH MATERIALIZED VIEW CONCURRENTLY sales_summary;

-- Show materialized view status

SELECT
    schemaname,
    matviewname,
    matviewowner,
    ispopulated
FROM pg_matviews;

```

Extensions and Custom Functions

```

-- List available extensions
SELECT * FROM pg_available_extensions;

-- Create custom function
CREATE OR REPLACE FUNCTION calculate_age(birthdate date)

```

```
RETURNS integer AS $$
BEGIN
    RETURN extract(year from age(current_date, birthdate));
END;
$$ LANGUAGE plpgsql;

-- Show function definitions
SELECT
    p.proname as function_name,
    pg_get_functiondef(p.oid) as definition
FROM pg_proc p
JOIN pg_namespace n ON p.pronamespace = n.oid
WHERE n.nspname = 'public';
```

Best Practices

Always use appropriate indexes

Create indexes for frequently queried columns

1. Monitor index usage
2. Remove unused indexes
3. Regular maintenance

● Schedule VACUUM and ANALYZE

1. Monitor table bloat
2. Keep statistics up to date
3. Query optimization

● Use EXPLAIN ANALYZE

1. Avoid SELECT *
2. Use appropriate JOIN types
3. Security

● Regular security audits

1. Principle of least privilege
2. Strong password policies