

[Back to Blog](#)

50 Essential PostgreSQL Queries by Category: A Complete Guide

TheDBAdmin Team • January 30, 2025

[postgresql](#)[sql](#)[database-administration](#)[performance](#)[tutorials](#)

PostgreSQL is a powerful open-source database system with extensive features. This guide presents 50 essential SQL queries organized by category to help you master PostgreSQL administration and development.

Table of Contents

[Database Administration](#)[Performance Monitoring](#)[Security Management](#)[Data Manipulation](#)[Advanced Features](#)

Database Administration

Database and Table Information

```
-- List all databases
SELECT datname, datdba, encoding, datcollate, datctype
FROM pg_database;
```

```
-- List all schemas
```

```
SELECT schema_name, schema_owner
FROM information_schema.schemata;

-- List tables with sizes
SELECT
    schemaname AS schema_name,
    relname AS table_name,
    pg_size.pretty(pg_total_relation_size(relid)) AS total_size,
    pg_size.pretty(pg_table_size(relid)) AS table_size,
    pg_size.pretty(pg_indexes_size(relid)) AS index_size
FROM pg_catalog.pg_stat_user_tables
ORDER BY pg_total_relation_size(relid) DESC;

-- Show table column information
SELECT
    column_name,
    data_type,
    character_maximum_length,
    column_default,
    is_nullable
FROM information_schema.columns
WHERE table_schema = 'public'
AND table_name = 'your_table';

-- List all indexes
SELECT
    schemaname,
    tablename,
    indexname,
    indexdef
FROM pg_indexes
WHERE schemaname = 'public'
ORDER BY tablename, indexname;
```

Backup and Maintenance

```
-- Analyze table statistics
ANALYZE VERBOSE your_table;

-- Vacuum table
VACUUM (VERBOSE, ANALYZE) your_table;
```

```
-- Reindex table
REINDEX TABLE your_table;

-- Show last vacuum and analyze time
SELECT
    relname AS table_name,
    last_vacuum,
    last_autovacuum,
    last_analyze,
    last_autoanalyze
FROM pg_stat_user_tables;

-- Show dead tuples and autovacuum status
SELECT
    schemaname,
    relname,
    n_dead_tup,
    n_live_tup,
    n_mod_since_analyze
FROM pg_stat_user_tables;
```

Performance Monitoring

Query Performance

```
-- Show slow queries
SELECT
    pid,
    age(clock_timestamp(), query_start) AS duration,
    query
FROM pg_stat_activity
WHERE state != 'idle'
ORDER BY duration DESC;

-- Identify queries with high total time
SELECT
    query,
    calls,
    total_time,
    rows,
    mean_time
```

```
FROM pg_stat_statements
ORDER BY total_time DESC
LIMIT 10;

-- Find queries with most block I/O
SELECT
    query,
    shared_blk_hit,
    shared_blk_read,
    shared_blk_written
FROM pg_stat_statements
ORDER BY shared_blk_read + shared_blk_written DESC
LIMIT 10;

-- Show table cache hit ratios
SELECT
    relname AS table_name,
    heap_blk_read AS blocks_read,
    heap_blk_hit AS blocks_hit,
    CASE WHEN heap_blk_hit + heap_blk_read = 0
        THEN 0
        ELSE round(heap_blk_hit::numeric / (heap_blk_hit + heap_blk_read) * 100)
    END AS cache_hit_ratio
FROM pg_stat_user_tables
ORDER BY cache_hit_ratio DESC;

-- Index usage statistics
SELECT
    schemaname,
    relname,
    indexrelname,
    idx_scan,
    idx_tup_read,
    idx_tup_fetch
FROM pg_stat_user_indexes
ORDER BY idx_scan DESC;
```

Security Management

User and Role Management

```
-- List all roles
SELECT
    rolname,
    rolsuper,
    rolcreaterole,
    rolcreatedb,
    rolcanlogin
FROM pg_roles;

-- Show role permissions
SELECT
    grantee, table_schema, table_name, privilege_type
FROM information_schema.role_table_grants
WHERE grantee = 'your_role';

-- List active sessions by user
SELECT
    username,
    count(*) as session_count
FROM pg_stat_activity
GROUP BY username;

-- Show connection limits per role
SELECT
    rolname,
    rolconnlimit
FROM pg_roles
WHERE rolconnlimit <> -1;

-- Audit role memberships
SELECT
    pg_get_userbyid(member) as member,
    pg_get_userbyid(roleid) as role
FROM pg_auth_members;
```

Data Manipulation

Advanced Queries

```
-- Common Table Expression (CTE)
WITH ranked_orders AS (
    SELECT
        customer_id,
        order_date,
        total_amount,
        RANK() OVER (PARTITION BY customer_id ORDER BY total_amount DESC)
    FROM orders
)
SELECT * FROM ranked_orders WHERE rank = 1;

-- Window Functions
SELECT
    department,
    employee_name,
    salary,
    AVG(salary) OVER (PARTITION BY department) as dept_avg,
    salary - AVG(salary) OVER (PARTITION BY department) as diff_from_avg
FROM employees;

-- JSON Operations
SELECT
    id,
    data->>'name' as name,
    (data->>'age')::int as age,
    jsonb_array_elements(data->'hobbies') as hobby
FROM users
WHERE (data->>'age')::int > 25;

-- Full Text Search
SELECT
    title,
    ts_rank_cd(to_tsvector('english', content), query) as rank
FROM articles, plainto_tsquery('english', 'your search terms') query
WHERE to_tsvector('english', content) @@ query
ORDER BY rank DESC;

-- Recursive Queries
WITH RECURSIVE subordinates AS (
    -- Base case
    SELECT employee_id, manager_id, name, 1 as level
    FROM employees
```

```

WHERE employee_id = 1

UNION ALL

-- Recursive case
SELECT e.employee_id, e.manager_id, e.name, s.level + 1
FROM employees e
INNER JOIN subordinates s ON s.employee_id = e.manager_id
)
SELECT * FROM subordinates;

```

Advanced Features

Partitioning and Inheritance

```

-- Create partitioned table
CREATE TABLE measurements (
    city_id      int not null,
    logdate      date not null,
    peaktemp     int,
    unitsales    int
) PARTITION BY RANGE (logdate);

-- Create partition
CREATE TABLE measurements_y2025m01 PARTITION OF measurements
FOR VALUES FROM ('2025-01-01') TO ('2025-02-01');

-- Show partitions
SELECT
    parent.relname as table_name,
    child.relname as partition_name,
    pg_get_expr(child.relpartbound, child.oid) as partition_expression
FROM pg_inherits
JOIN pg_class parent ON pg_inherits.inhparent = parent.oid
JOIN pg_class child ON pg_inherits.inhrelid = child.oid
WHERE parent.relname = 'measurements';

```

Materialized Views

```
-- Create materialized view
CREATE MATERIALIZED VIEW sales_summary AS
SELECT
    date_trunc('month', order_date) as month,
    product_category,
    SUM(amount) as total_sales,
    COUNT(*) as order_count
FROM orders
GROUP BY 1, 2
WITH DATA;

-- Refresh materialized view
REFRESH MATERIALIZED VIEW CONCURRENTLY sales_summary;

-- Show materialized view status
SELECT
    schemaname,
    matviewname,
    matviewowner,
    ispopulated
FROM pg_matviews;
```

Extensions and Custom Functions

```
-- List available extensions
SELECT * FROM pg_available_extensions;

-- Create custom function
CREATE OR REPLACE FUNCTION calculate_age(birthdate date)
RETURNS integer AS $$%
BEGIN
    RETURN extract(year from age(current_date, birthdate));
END;
$$ LANGUAGE plpgsql;

-- Show function definitions
SELECT
    p.proname as function_name,
    pg_get_functiondef(p.oid) as definition
FROM pg_proc p
```

```
JOIN pg_namespace n ON p.pronamespace = n.oid  
WHERE n.nspname = 'public';
```

Best Practices

Always use appropriate indexes

Create indexes for frequently queried columns

Monitor index usage

Remove unused indexes

Regular maintenance

Schedule VACUUM and ANALYZE

Monitor table bloat

Keep statistics up to date

Query optimization

Use EXPLAIN ANALYZE

Avoid SELECT *

Use appropriate JOIN types

Security

Regular security audits

Principle of least privilege

Strong password policies

Further Reading

[PostgreSQL Official Documentation](#)

[TheDBAdmin PostgreSQL Tutorials](#)

[PostgreSQL Performance Tuning Guide](#)

[Advanced PostgreSQL Security](#)

Remember to always test these queries in a development environment first, as some may be resource-intensive or require specific permissions. Adjust them according to your specific needs and database structure.

Share this article



Tags

oracle (1) postgresql (13) aws-dms (1) ora2pg (1) database-migration (1)

cloud-migration (1) database-administration (11) ai (6) machine-learning (4)
dba (2) career-development (2) data-science (2) mlops (1) qwen (1)
local-deployment (2) llm (2) ollama (2) deepseek (3) nosql (1)
career-guide (1) mongodb (2) cassandra (1) redis (2) docker (2)
docker-compose (2) containerization (2) monitoring (1) performance (5)
devops (1) backup (1) recovery (1) disaster-recovery (2) caching (1)
database (3) high-availability (1) replication (1) failover (1) scalability (1)
high-concurrency (1) database-tuning (1) optimization (1) tuning (1)
security (1) encryption (1) authentication (1) sql (1) tutorials (1)
automation (1) productivity (1) openai (1) claude (1) comparison (1)
technology (1) career (1) certification (1) welcome (1) introduction (1)

Recent Posts

[Oracle to PostgreSQL Migration: Comprehensive Guide Using AWS DMS & Ora2Pg](#)
Jan 31, 2025

[PostgreSQL DBA's Role in AI & ML: A Comprehensive Guide to the Future](#)
Jan 31, 2025

[Running Qwen 2.5 Locally: Complete Implementation Guide](#)
Jan 30, 2025

[Running DeepSeek Models Locally with Ollama: Complete Guide](#)
Jan 30, 2025

[How to Become a NoSQL Database Administrator: Complete Career Guide](#)
Jan 30, 2025



Expert database administration and consulting services for modern enterprises.

**Services**[Database Services](#)[Courses](#)[Blog](#)**Company**[About Us](#)[Contact](#)**Support**[Documentation](#)[Help Center](#)[Status](#)**Stay Updated**

Subscribe to our newsletter for the latest updates and insights.

[Subscribe](#)

© 2025 TheDBAdmin. All rights reserved.

[Privacy Policy](#) [Terms of Service](#) [Cookie Policy](#)