

[Back to Blog](#)

# 100+ PostgreSQL DBA Interview Questions and Answers [2025 Guide]

TheDBAdmin Team • February 12, 2025

[Contact](#)[postgresql performance tuning](#)[postgresql backup recovery](#)[postgresql high availability](#)[postgresql security](#)[postgresql monitoring](#)[postgresql aws rds](#)[postgresql azure](#)[postgresql pgvector](#)[postgresql cloud](#)

This comprehensive guide presents 100+ essential PostgreSQL Database Administrator (DBA) interview questions and answers, organized by topic. Whether you're preparing for a junior or senior PostgreSQL DBA interview, studying for PostgreSQL certification, or looking to assess your database administration expertise, these questions cover all key aspects of PostgreSQL administration and management.

## Want to master PostgreSQL?

Check out our comprehensive [PostgreSQL DBA Course](#)! This course covers everything from basic administration to advanced topics, helping you become a proficient PostgreSQL Database Administrator. Perfect for both beginners and experienced DBAs looking to enhance their skills.

# Basic Concepts (1-10)

---

## What is PostgreSQL?

Open-source RDBMS emphasizing extensibility and SQL compliance

Object-relational database supporting both relational and non-relational features

## What are the key features of PostgreSQL?

ACID compliance

Multi-Version Concurrency Control (MVCC)

Table inheritance

Custom functions

Table partitioning

Rich indexing options

## What is MVCC in PostgreSQL?

Mechanism allowing multiple transactions to access the same data simultaneously

Each transaction sees a snapshot of data from a particular point in time

## What are roles in PostgreSQL?

Database objects managing database access permissions

Can be users or groups

Support inheritance of permissions

## What is WAL in PostgreSQL?

Write-Ahead Logging

Ensures data integrity by logging changes before they're written to data files

## What is the postgresql.conf file?

Main configuration file containing database parameters

Controls memory, connections, replication, etc.

## What is the pg\_hba.conf file?

Host-Based Authentication configuration file

Controls client authentication methods and access

## What are tablespaces in PostgreSQL?

Physical locations where databases store their data

Allow distributing database objects across different storage devices

## What is the difference between TRUNCATE and DELETE?

TRUNCATE is faster, removes all rows at once

DELETE removes rows one by one, can use WHERE clause

**What is autovacuum in PostgreSQL?**

Automated process that removes dead tuples and updates statistics  
Essential for maintaining database performance

## Backup & Recovery (11-20)

---

**What is pg\_dump?**

Logical backup tool for single database  
Creates SQL script or archive file format

**What is pg\_basebackup?**

Physical backup tool for entire cluster  
Creates exact copy of database files

**What is Point-in-Time Recovery (PITR)?**

Ability to restore database to specific point in time  
Uses WAL archives and base backup

**What is continuous archiving?**

Method of archiving WAL files continuously  
Essential for PITR capability

**How do you take a hot backup?**

Use pg\_basebackup while database is running  
Ensure archive\_mode is enabled

**What is recovery.conf?**

In PostgreSQL <12: Recovery configuration file  
In PostgreSQL ≥12: Uses recovery.signal and standby.signal

**What is archive\_command?**

Configuration parameter specifying WAL archiving method  
Example: 'cp %p /archive/%f'

**What is restore\_command?**

Specifies how to retrieve archived WAL files during recovery  
Example: 'cp /archive/%f %p'

**What is recovery\_target\_time?**

Parameter specifying point-in-time for recovery  
Format: 'YYYY-MM-DD HH:MM:SS'

**What is pg\_dumpall?**

Backs up entire PostgreSQL cluster

Includes roles and tablespace definitions

## Performance Tuning (21-30)

---

**What is EXPLAIN ANALYZE?**

Shows query execution plan with actual timings

Helps identify performance bottlenecks

**What are statistics in PostgreSQL?**

Data about tables and columns

Used by query planner to create efficient plans

**What is shared\_buffers?**

Memory parameter for shared data cache

Typically 25% of total RAM

**What is work\_mem?**

Memory for sort operations and hash tables

Set per operation, not total

**What is effective\_cache\_size?**

Estimate of disk cache available

Influences query planner decisions

**What is vacuum?**

Reclaims storage from dead tuples

Updates statistics for query planner

**What is analyze?**

Updates statistics about table contents

Improves query planner estimates

**What are index types in PostgreSQL?**

B-tree (default)

Hash

GiST

SP-GiST

GIN

BRIN

**What is table partitioning?**

Splitting large tables into smaller pieces  
Improves query performance and manageability

**What is parallel query?**

Multiple workers executing parts of query  
Controlled by max\_parallel\_workers parameter

## High Availability (31-40)

---

**What is streaming replication?**

Real-time WAL shipping to standby servers  
Provides high availability

**What is synchronous replication?**

Ensures transaction commit waits for standby  
Provides stronger consistency guarantees

**What is Patroni?**

HA solution for PostgreSQL  
Manages failover and replication

**What is pg\_rewind?**

Tool to re-synchronize primary and standby  
Used after failback scenarios

**What is replication slot?**

Prevents removal of needed WAL files  
Ensures standby doesn't fall behind

**What is hot\_standby?**

Allows queries on standby server  
May impact replication performance

**What is promote\_trigger\_file?**

File triggering standby promotion  
Used in manual failover

**What is primary\_conninfo?**

Connection string to primary server  
Used by standby for replication

**What is application\_name in replication?**

Identifies standby servers

Used in monitoring and configuration

**What is pg\_stat\_replication?**

View showing replication status

Shows lag and connection details

## Security (41-50)

---

**What is SSL in PostgreSQL?**

Encrypts client/server communications

Configured in postgresql.conf

**What is Row-Level Security (RLS)?**

Fine-grained access control at row level

Implemented using policies

**What are GRANT/REVOKE?**

Commands managing permissions

Control access to database objects

**What is password encryption?**

Passwords stored using SCRAM-SHA-256

Previous methods: MD5, password

**What is client authentication?**

Methods: trust, password, md5, cert

Configured in pg\_hba.conf

**What is pgcrypto?**

Extension for cryptographic functions

Supports various encryption methods

**What are connection limits?**

max\_connections parameter

Per-role connection limits

**What is pg\_audit?**

Extension for detailed audit logging

Tracks database activities

**What is SSL certificate authentication?**

Client authentication using certificates

Requires SSL configuration

**What are default privileges?**

ALTER DEFAULT PRIVILEGES command

Sets permissions for new objects

## Monitoring (51-60)

---

**What is pg\_stat\_activity?**

View showing server process activities

Shows current queries and states

**What is pg\_stat\_statements?**

Extension tracking query execution statistics

Shows query performance metrics

**What are wait events?**

Show what processes are waiting for

Help identify bottlenecks

**What is log\_min\_duration\_statement?**

Logs queries exceeding specified duration

Helps identify slow queries

**What is pg\_stat\_bgwriter?**

Shows background writer statistics

Helps tune checkpoints

**What is pg\_stat\_database?**

Database-wide statistics

Shows transactions, tuples, cache hits

**What is track\_io\_timing?**

Enables I/O timing measurements

Shows detailed I/O statistics

**What is log\_checkpoints?**

Logs checkpoint activity

Helps monitor I/O activity

**What is pg\_stat\_user\_tables?**

Table access and update statistics

Shows vacuum and analyze activity

**What is pg\_locks?**

Shows current locks

Helps diagnose blocking issues

## Cloud & AWS RDS (61-70)

---

**What is RDS for PostgreSQL?**

AWS managed PostgreSQL service

Handles maintenance and backups

**What are parameter groups in RDS?**

Manage database configuration

Some parameters can't be modified

**What is Multi-AZ in RDS?**

Automatic failover capability

Synchronous replication to standby

**What is RDS read replica?**

Asynchronous replication copy

Offloads read operations

**What is Aurora PostgreSQL?**

AWS-optimized PostgreSQL variant

Claims better performance

**What is RDS automated backup?**

Daily backups and transaction logs

Enables point-in-time recovery

**What is RDS encryption?**

At-rest encryption using KMS

SSL for in-transit encryption

**What is RDS monitoring?**

Enhanced monitoring metrics

Performance Insights



**What is RDS proxy?**

Connection pooling service

Reduces connection management overhead

**What is RDS event notification?**

SNS notifications for RDS events

Monitors database changes

## Azure PostgreSQL (71-80)

---

**What is Azure Database for PostgreSQL?**

Microsoft's managed PostgreSQL service

Offers single server and flexible server

**What is Flexible Server?**

Newer deployment option

More control over configuration

**What is Azure high availability?**

Zone-redundant HA

Automatic failover

**What is Azure backup?**

Automated backups

Geo-redundant storage option

**What is Azure networking?**

Private Link

VNET integration

**What is Azure monitoring?**

Azure Monitor integration

Query Performance Insight

**What is Azure scaling?**

Vertical and horizontal scaling

Storage autogrowth

**What is Azure maintenance?**

Planned maintenance windows

Automatic minor version updates

**What is Azure firewall?**

Server-level firewall rules

VNET service endpoints

**What is Azure PostgreSQL pricing?**

Compute + Storage pricing model

Reserved capacity options

## Advanced Features (81-90)

---

**What is pgvector?**

Vector similarity search extension

Used for AI/ML applications

**What is pg\_cron?**

Job scheduling extension

Runs periodic SQL commands

**What is pg\_stat\_monitor?**

Advanced monitoring extension

Provides detailed query statistics

**What is logical replication?**

Replication at object level

Supports selective replication

**What is pg\_partman?**

Partition management extension

Automates partition creation

**What is pg\_bouncer?**

Connection pooler

Reduces connection overhead

**What is Foreign Data Wrapper?**

Connects to external data sources

Supports various data formats

**What is pg\_upgrade?**

Upgrades PostgreSQL version

Supports fast and link modes

**What is event trigger?**

Triggers on database events  
DDL operation notifications

**What is table sampling?**

TABLESAMPLE clause  
Quick data sampling methods

## Troubleshooting (91-100)

---

**How to handle connection issues?**

Check pg\_hba.conf  
Verify network connectivity  
Check max\_connections

**How to identify slow queries?**

Use pg\_stat\_statements  
Check log\_min\_duration\_statement  
EXPLAIN ANALYZE

**How to handle disk full?**

VACUUM FULL  
Remove unused indexes  
Archive old data

**How to fix high CPU usage?**

Check for missing indexes  
Optimize queries  
Adjust work\_mem

**How to handle replication lag?**

Monitor pg\_stat\_replication  
Check network bandwidth  
Adjust wal\_keep\_segments

**How to fix database bloat?**

Regular VACUUM  
Adjust autovacuum settings  
Monitor table/index bloat

**How to handle connection pooling?**

Use pg\_bouncer

Adjust pool\_size

Monitor pool statistics

### How to fix lock contention?

Monitor pg\_locks

Optimize transaction duration

Adjust lock timeout

### How to handle index corruption?

REINDEX command

Check system logs

Verify disk health

### How to improve checkpoint performance?

Adjust checkpoint\_timeout

Set checkpoint\_completion\_target

Monitor checkpoint timing

---

## Share this article



## Tags

postgresql interview questions (1)    postgresql dba interview (1)

database administrator interview (1)    postgresql administration (1)

postgresql troubleshooting (1)    postgresql performance tuning (1)

postgresql backup recovery (1)    postgresql high availability (1)    postgresql security (1)

postgresql monitoring (1)    postgresql aws rds (1)    postgresql azure (1)

postgresql pgvector (1)    postgresql cloud (1)    oracle (1)    postgresql (13)

aws-dms (1)    ora2pg (1)    database-migration (1)    cloud-migration (1)

database-administration (11)    ai (6)    machine-learning (4)    dba (2)

career-development (2)    data-science (2)    mlops (1)    qwen (1)

local-deployment (2)    llm (2)    ollama (2)    deepseek (3)    nosql (1)

career-guide (1)    mongodb (2)    cassandra (1)    redis (2)    docker (2)

docker-compose (2)    containerization (2)    monitoring (1)    performance (5)

devops (1)

backup (1)

recovery (1)

disaster-recovery (2)

caching (1)

database (3)

high-availability (1)

replication (1)

failover (1)

scalability (1)

high-concurrency (1)

database-tuning (1)

optimization (1)

tuning (1)

security (1)

encryption (1)

authentication (1)

sql (1)

tutorials (1)

automation (1)

productivity (1)

openai (1)

claude (1)

comparison (1)

technology (1)

career (1)

certification (1)

welcome (1)

introduction (1)

## Recent Posts

100+ PostgreSQL DBA Interview Questions and Answers [2025 Guide]

Feb 12, 2025

Oracle to PostgreSQL Migration: Comprehensive Guide Using AWS DMS & Ora2Pg

Jan 31, 2025

PostgreSQL DBA's Role in AI & ML: A Comprehensive Guide to the Future

Jan 31, 2025

Running Qwen 2.5 Locally: Complete Implementation Guide

Jan 30, 2025

Running DeepSeek Models Locally with Ollama: Complete Guide

Jan 30, 2025



Expert database administration and consulting services for modern enterprises.



Services

Database Services

Courses

Blog

Company

About Us

Contact

## Support

[Documentation](#)

[Help Center](#)

[Status](#)

## Stay Updated

Subscribe to our newsletter for the latest updates and insights.

Subscribe

---

© 2025 TheDBAdmin. All rights reserved.

[Privacy Policy](#) [Terms of Service](#) [Cookie Policy](#)