

PostgreSQL with Patroni: Installation and Configuration

In this document, we will set up a PostgreSQL Patroni Cluster step by step. Our goal is to create a highly available PostgreSQL cluster. The installation process will follow these steps:


1. Etcd Cluster Installation
2. PostgreSQL + Patroni Installation
3. HAProxy + Keepalived Installation
4. pgBackRest Backup Solution Setup



We will use the following server infrastructure:

- .190 → VIP
- .191 → HAProxy Primary Node
- .192 → HAProxy Secondary Node
- .193 → PostgreSQL Node 1
- .194 → PostgreSQL Node 2
- .195 → PostgreSQL Node 3
- .196 → Etcd Node 1
- .197 → Etcd Node 2
- .198 → Etcd Node 3
- .199 → Backup Node

High Available Postgres Architecture



Now, let's proceed with the first step: Etcd Cluster Installation.

Etcd Installation

Etcd is a key-value store service used by Patroni for distributed database management. We will set up a three-node etcd cluster.

Step 1: Install Etcd Binaries

Perform the following steps on all etcd nodes:

```
dnf -y install curl wget vim

ETCD_RELEASE=$(curl -s https://api.github.com/repos/etcd-io/etcd/releases/latest | grep
tag_name | cut -d '"' -f 4)
echo $ETCD_RELEASE
wget https://github.com/etcd-io/etcd/releases/download/${ETCD_RELEASE}/etcd-
${ETCD_RELEASE}-linux-amd64.tar.gz
tar xvf etcd-${ETCD_RELEASE}-linux-amd64.tar.gz

cd etcd-${ETCD_RELEASE}-linux-amd64

mv etcd* /usr/local/binls /usr/local/bin

etcd --version
etcdctl version
etcdutl version
```

Step 2: Create Directories and Users

```
mkdir -p /var/lib/etcd/
```

```
mkdir /etc/etcd
groupadd --system etcd
useradd -s /sbin/nologin --system -g etcd etcd

chown -R etcd:etcd /var/lib/etcd/ && chmod 0775 /var/lib/etcd/
```

Step 3: Systemd Service File

```
# vi /etc/systemd/system/etcd.service
[Unit]
Description=Etcd - Highly Available Key Value Store

Documentation=man:etcd
After=network.target
Wants=networkonline.target

[Service]
Environment=DAEMON_ARGS=

Environment=ETCD_NAME=%H
Environment=ETCD_DATA_DIR=/var/lib/etcd/default
EnvironmentFile=/etc/etcd/etcd.conf
Type=notify
User=etcd
PermissionsStartOnly=true
ExecStart=/usr/local/bin/etcd $DAEMON_ARGS
Restart=on-abnormal
LimitNOFILE=65536

[Install]
WantedBy=multi-user.target
```

Step 4: Etcd Configuration

Etcd Node 1 (.196)

```
# cat /etc/etcd/etcd.conf
```

[Member]

```
ETCD_LISTEN_PEER_URLS="http://*.196:2380,http://localhost:2380"
ETCD_LISTEN_CLIENT_URLS="http://localhost:2379,http://*.196:2379"
ETCD_DATA_DIR="/var/lib/etcd/default"
```

[Clustering]

```
ETCD_INITIAL_ADVERTISE_PEER_URLS="http://*.196:2380"
ETCD_ADVERTISE_CLIENT_URLS="http://*.196:2379"
ETCD_INITIAL_CLUSTER="etcd1=http://*.196:2380,etcd2=http://*.197:2380,etcd3=http://*.198:2380"
```

```
ETCD_INITIAL_CLUSTER_TOKEN="etcd-cluster"
ETCD_INITIAL_CLUSTER_STATE="new"
ETCD_NAME="etcd1"
ETCD_ELECTION_TIMEOUT=5000
ETCD_HEARTBEAT_INTERVAL=1000
ETCD_ENABLE_V2=true
```

Etcd Node 2 (.197)

```
# cat /etc/etcd/etcd.conf
```

[Member]

```
ETCD_LISTEN_PEER_URLS="http://**.**.197:2380,http://localhost:2380"
ETCD_LISTEN_CLIENT_URLS="http://localhost:2379,http://**.**.197:2379"
ETCD_DATA_DIR="/var/lib/etcd/default"
```

[Clustering]

```
ETCD_INITIAL_ADVERTISE_PEER_URLS="http://**.**.196:2380"
ETCD_ADVERTISE_CLIENT_URLS="http://**.**.197:2379"
ETCD_INITIAL_CLUSTER="etcd1=http://**.**.196:2380,etcd2=http://**.**.197:2380,etcd3=http://**.**.198:2380"
ETCD_INITIAL_CLUSTER_TOKEN="etcd-cluster" ETCD_INITIAL_CLUSTER_STATE="new"
ETCD_NAME="etcd2"
ETCD_ELECTION_TIMEOUT=5000
ETCD_HEARTBEAT_INTERVAL=1000
ETCD_ENABLE_V2=true
```

Etcd Node 3 (.198)

```
# cat /etc/etcd/etcd.conf
```

[Member]

```
ETCD_LISTEN_PEER_URLS="http://**.**.198:2380,http://localhost:2380"
ETCD_LISTEN_CLIENT_URLS="http://localhost:2379,http://**.**.198:2379"
ETCD_DATA_DIR="/var/lib/etcd/default"
```

[Clustering]

```
ETCD_INITIAL_ADVERTISE_PEER_URLS="http://**.**.198:2380"
ETCD_ADVERTISE_CLIENT_URLS="http://**.**.198:2379"
ETCD_INITIAL_CLUSTER="etcd1=http://**.**.196:2380,etcd2=http://**.**.197:2380,etcd3=http://**.**.198:2380"
ETCD_INITIAL_CLUSTER_TOKEN="etcd-cluster"
ETCD_INITIAL_CLUSTER_STATE="new"
ETCD_NAME="etcd3"
ETCD_ELECTION_TIMEOUT=5000
ETCD_HEARTBEAT_INTERVAL=1000
ETCD_ENABLE_V2=true
```

Step 4: Start Etcd Service

```
Copy systemctl enable etcd
systemctl start etcd
systemctl status etcd
```

Step 5: Verify Etcd Cluster Status

```
etcdctl member list
```

```
etcdctl --endpoints=*.*.*.196:2379,*.*.*.197:2379,*.*.*.198:2379 endpoint status
--write-out=table
```

Expected output:

ENDPOINT	ID	VERSION	DB SIZE	IS LEADER	IS LEARNER	RAFT TERM	RAFT INDEX	RAFT APPLIED INDEX	ERRORS
..*.196:2379	84ebce0bfc393c1	3.5.11	29 kB	false	false	40	100512574	100512574	
..*.197:2379	76d84aedaf38bba4	3.5.11	29 kB	true	false	40	100512574	100512574	
..*.198:2379	b3d60bce62f139c	3.5.11	29 kB	false	false	40	100512574	100512574	

Patroni Installation Guide

Step 1: Installing Required Packages

```
# apply each postgres node
dnf -y install https://dl.fedoraproject.org/pub/epel/epel-release-latest-9.noarch.rpm
#(https://yum.postgresql.org/)

dnf install -y https://download.postgresql.org/pub/repos/yum/repopms/EL-9-x86_64/pgdg-
redhat-repo-latest.noarch.rpm
dnf -qy module disable postgresq

# Install PostgreSQL:

dnf install -y postgresql13-server
dnf install -y postgresql13-contrib
dnf install -y postgresql13-devel# Install Patroni
yum install patroni
yum install python3-urllib3
yum install patroni-etcd.x86_64
```

```
cp -p /usr/share/doc/patroni/postgres0.yml /etc/patroni/patroni.yml
```

Step 2: Node 1 (.193) Configuration

```
#vi /etc/patroni/patroni.yml-----

scope: clusternamespace: /var/lib/pgsql/config/name: node1restapi:
  listen: **.**.**.193:8008
  connect_address: **.**.**.193:8008etcd:
  host: **.**.**.196:2379
  host: **.**.**.197:2379
  host: **.**.**.198:2379bootstrap:
dcs:
  ttl: 30
  loop_wait: 10
  retry_timeout: 10
  maximum_lag_on_failover: 1048576
  postgresql:
    use_pg_rewind: true
    use_slots: true
    parameters:
  initdb:
    - encoding: UTF8
    - data-checksums
    - waldir: /pg_wal
    - wal-segsize=512
  pg_hba:
    - host replication replicator 127.0.0.1/32 scram-sha-256
    - host replication replicator **.**.**.193/32 scram-sha-256
    - host replication replicator **.**.**.194/32 scram-sha-256
    - host replication replicator **.**.**.195/32 scram-sha-256
    - host all all **.**.**.*/24 scram-sha-256

# Please separate your database usage per node per IP block for system security
users:
  admin:
    password: admin
    options:
      - createrole
      - createdbpostgresql:
listen: **.**.**.193:5432
connect_address: **.**.**.193:5432
data_dir: /pg_data/data
bin_dir: /usr/pgsql-13/bin
pgpass: /tmp/pgpass
authentication:
  replication:
    username: replicator
    password: test123
  superuser:
    username: postgres
```

```
password: test123
create_replica_methods:
  - pgbackrest
  - basebackup
pgbackrest:
  command: pgbackrest --stanza=pg_backup restore --type=none
  keep_data: True
  no_params: True
basebackup:
  checkpoint: 'fast'tags:
  nofailover: false
  noloadbalance: false
  clonefrom: false
  nosync: false
```

Step 3: Node 2 (.194) Configuration

```
# vi /etc/patroni/patroni.yml-----

scope: cluster
namespace: /var/lib/pgsql/config/
name: node2

restapi:
  listen: *.*.*.194:8008
  connect_address: *.*.*.194:8008

etcd:
  host: *.*.*.196:2379
  host: *.*.*.197:2379
  host: *.*.*.198:2379

bootstrap:

dcs:
  ttl: 30
  loop_wait: 10
  retry_timeout: 10
  maximum_lag_on_failover: 1048576
  postgresql:
    use_pg_rewind: true
    use_slots: true
    parameters:
  initdb:
    - encoding: UTF8
    - data-checksums
    - waldir: /pg_wal
    - wal-segsize=512
  pg_hba:
    - host replication replicator 127.0.0.1/32 scram-sha-256
```

- host replication replicator **.**.193/32 scram-sha-256
- host replication replicator **.**.194/32 scram-sha-256
- host replication replicator **.**.195/32 scram-sha-256
- host all all **.**.*/24 scram-sha-256

Please separate your database usage per node per IP block for system security

```

users:
  admin:
    password: admin
    options:
      - createrole
      - createdbpostgresql:
listen: **.194:5432
connect_address: **.194:5432
data_dir: /pg_data/data
bin_dir: /usr/pgsql-13/bin
pgpass: /tmp/pgpass
authentication:
  replication:
    username: replicator
    password: test123
  superuser:
    username: postgres
    password: test123
create_replica_methods:
  - pgbackrest
  - basebackup
pgbackrest:
  command: pgbackrest --stanza=pg_backup restore --type=none
  keep_data: True
  no_params: True
basebackup:
  checkpoint: 'fast'tags:
  nofailover: false
  noloadbalance: false
  clonefrom: false
  nosync: false

```

Step 4: Node 3 (.195) Configuration

Copy# vi /etc/patroni/patroni.yml-----

```

scope: cluster
namespace: /var/lib/pgsql/config/
name: node3
restapi:
  listen: **.195:8008
  connect_address: **.195:8008

etcd:
  host: **.196:2379

```



```
host: *.*.*.197:2379
host: *.*.*.198:2379

bootstrap:
  dcs:
    ttl: 30
    loop_wait: 10
    retry_timeout: 10
    maximum_lag_on_failover: 1048576
  postgresql:
    use_pg_rewind: true
    use_slots: true
    parameters:
  initdb:
    - encoding: UTF8
    - data-checksums
    - waldir: /pg_wal
    - wal-segsize=512
  pg_hba:
    - host replication replicator 127.0.0.1/32 scram-sha-256
    - host replication replicator *.*.*.193/32 scram-sha-256
    - host replication replicator *.*.*.194/32 scram-sha-256
    - host replication replicator *.*.*.195/32 scram-sha-256
    - host all all *.*.*.**/24 scram-sha-256

# Please separate your database usage per node per IP block for system security
users:
  admin:
    password: admin
    options:
      - createrole
      - createdb
  postgresql:
listen: *.*.*.195:5432
connect_address: *.*.*.195:5432
data_dir: /pg_data/data
bin_dir: /usr/pgsql-13/bin
pgpass: /tmp/pgpass
authentication:
  replication:
    username: replicator
    password: test123
  superuser:
    username: postgres
    password: test123
create_replica_methods:
  - pgbackrest
  - basebackup
pgbackrest:
  command: pgbackrest --stanza=pg_backup restore --type=none
  keep_data: True
```

```
no_params: True
basebackup:
checkpoint: 'fast'tags:
nofailover: false
noloadbalance: false
clonefrom: false
nosync: false
```

Step 5: Start Patroni Service

```
systemctl enable patroni
systemctl start patroni
systemctl status patroni
```

Step 6: Check Cluster Status

```
patronictl -c /etc/patroni/patroni.yml list
```

Expected Output:

```
+ Cluster: cluster --+-----+-----+-----+
| Member | Host   | Role | State | TL | Lag in MB |
+-----+-----+-----+-----+-----+
| node1  | *.**.*.193 | Replica | streaming | 2 | 0 |
| node2  | *.**.*.194 | Leader  | running  | 2 |   |
| node3  | *.**.*.195 | Replica | streaming | 2 | 0 |
+-----+-----+-----+-----+-----+
```

HAProxy Installation

Step 1: Installing Required Package

```
dnf install haproxy.x86_64
```

Step 2: Configure HAProxy for Primary Node

```
# vim /etc/haproxy/haproxy.cfg

#for master haproxy node

global
    maxconn 1000
    log 127.0.0.1 local0
defaults
    log global
    mode tcp
    retries 2
    timeout client 20m
    timeout connect 4s
```

```

timeout server 20m
timeout check 5slisten stats
mode http
bind *:7000
stats enable
stats uri /
frontend a_listen_fe#bind *:5000#bind *:5001
acl is-read-service-dead nbsrv(standby) lt 1
use_backend postgres if is-read-service-dead
default_backend standby
listen postgres
    bind **.*.*.*.190:5000
    option httpchk OPTIONS/master
    http-check expect status 200
    default-server inter 3s fall 4 rise 3 on-marked-down shutdown-sessions
    server node1 **.*.*.*.193:5432 maxconn 1000 check port 8008
    server node2 **.*.*.*.194:5432 maxconn 1000 check port 8008
    server node3 **.*.*.*.195:5432 maxconn 1000 check port 8008
    server node4 **.*.*.*.199:5432 maxconn 1000 check backup port 8008listen standby
    bind **.*.*.*.190:5001
    option httpchk OPTIONS/replica
    http-check expect status 200
    default-server inter 3s fall 4 rise 3 on-marked-down shutdown-sessions
    server node1 **.*.*.*.193:5432 maxconn 1000 check port 8008
    server node2 **.*.*.*.194:5432 maxconn 1000 check port 8008
    server node3 **.*.*.*.195:5432 maxconn 1000 check port 8008

```

Step 3: Configure HAProxy for Secondary Node

```

#for Secondary haproxy node
# -----

global
    maxconn 100
    log 127.0.0.1 local0
defaults
    log global
    mode tcp
    retries 2
    timeout client 120m
    timeout connect 4s
    timeout server 120m
    timeout check 5s
listen stats
    mode http
    bind *:7000
    stats enable
    stats uri /
frontend a_listen_fe#bind *:5000#bind *:5000
acl is-read-service-dead nbsrv(standby) lt 1
use_backend postgres if is-read-service-dead

```

```

default_backend standby
listen postgres
    bind *.**.*.190:5000
    option httpchk OPTIONS/master
    http-check expect status 200
    default-server inter 3s fall 4 rise 3 on-marked-down shutdown-sessions
    server node1 *.**.*.193:5432 maxconn 1000 check port 8008
    server node2 *.**.*.194:5432 maxconn 1000 check port 8008
    server node3 *.**.*.195:5432 maxconn 1000 check port 8008
    server node4 *.**.*.199:5432 maxconn 1000 check backup port 8008
listen standby
    bind *.**.*.190:5001
    option httpchk OPTIONS/replica
    http-check expect status 200
    default-server inter 3s fall 4 rise 3 on-marked-down shutdown-sessions
    server node1 *.**.*.193:5432 maxconn 1000 check port 8008
    server node2 *.**.*.194:5432 maxconn 1000 check port 8008
    server node3 *.**.*.195:5432 maxconn 1000 check port 8008

```

Step 4: Ceck HAProxy Status

```

systemctl start haproxy.service
systemctl status haproxy.service
systemctl enable haproxy.service

```

Keepalived Installation

Step 1: Installing Required Package

```
dnf install keepalived.x86_64
```

Step 2: Configure Keepalived for Primary Node

```

#for Primary haproxy node
#-----
# vi /etc/keepalived/keepalived.conf

global_defs {}
vrrp_script chk_haproxy {
    script "killall -0 haproxy" # widely used idiom
    interval 2 # check every 2 seconds
    weight 2 # add 2 points of prio if OK}
vrrp_instance VI_1 {
    interface ens192
    state MASTER
    priority 101
    virtual_router_id 51
    authentication {
        auth_type PASS
        auth_pass test123
    }
    virtual_ipaddress {
        *.**.*.190/24
    }
}

```

```

}
unicast_src_ip *.*.*.*.191 # This node
unicast_peer {
*.*.*.*.192 # Other nodes
}
track_script {
chk_haproxy
}
}}
```

Step 3: Configure Keepalived for Secondary Node

```

#for Secondary haproxy node
#-----
# vi /etc/keepalived/keepalived.conf

global_defs {}
vrrp_script chk_haproxy {
script "killall -0 haproxy" # widely used idiom
interval 2 # check every 2 seconds
weight 2 # add 2 points of prio if OK}
vrrp_instance VI_1 {
interface ens192
state BACKUP
priority 99
virtual_router_id 51
authentication {
auth_type PASS
auth_pass test123
}
virtual_ipaddress {
*.*.*.*.190/24
}
unicast_src_ip *.*.*.*.192 # This node
unicast_peer {
*.*.*.*.191 # Other nodes
}
track_script {
chk_haproxy
}
}}
```

Step 4: Ceck Keepalived Status

```

systemctl start keepalived
systemctl status keepalived
systemctl enable keepalived
```

pgBackRest Installation

Step 1: Installing the Required Package

```

yum install pgbackrest
```

Step 2: Configure pgBackRest on Backup Node

```
vi /etc/pgbackrest.conf

[global]
repo1-path=/pg_backup/pg_backup_patroni
repo1-retention-full=14
repo1-retention-full-type=time
repo1-host-user=pgbackrest
archive-check=n
process-max=1
log-level-console=info
log-path=/pg_backup/pgbackrest/log
log-level-file=debug
start-fast=y
delta=y
compress-level=3
[pg_backup]
pg1-host=*.*.*.193
pg1-host-user=postgres
pg1-database=postgres
pg1-path=/pg_data/data
pg1-port=5432
pg2-host=*.*.*.194
pg2-host-user=postgres
pg2-database=postgres
pg2-path=/pg_data/data
pg2-port=5432
pg3-host=*.*.*.195
pg3-host-user=postgres
pg3-database=postgres
pg3-path=/pg_data/data
pg3-port=5432
```

Step 3: Configure PostgreSQL Nodes

For node1, edit the Patroni configuration file:

```
patronictl -c /etc/patroni/patroni.yml edit-config
# Add the following archive_command:
archive_command = 'pgbackrest --stanza=pg_backup archive-push %p && cp -i %p
/var/lib/pgsql/archive/%f'
```

Also, configure the pgBackRest settings for node1 in /etc/pgbackrest.conf:

```
# Repeat similar configurations for node2 and node3.

[global]
repo1-host=*.*.*.199
repo1-host-user=pgbackrest
log-level-console=info
```

```
log-level-file=debugdelta=y  
[pg_backup]  
pg1-path=/pg_data/data  
recovery-option=primary_conninfo=host=*.*.*.190 user=replication
```

Repeat similar configurations for `node2` and `node3`.

Step 4: Initialize the pgBackRest Stanza on the Backup Node

```
pgbackrest --stanza=pg_backup stanza-create  
pgbackrest --stanza=pg_backup check --log-level-console=info
```

Step 5: Run Full and Incremental Backups

```
pgbackrest --stanza=pg_backup --type=full backup  
# For incremental backups  
pgbackrest --stanza=pg_backup --type=incr backup
```

Step 6: Check the Status of pgBackRest Backups

```
pgbackrest --stanza=pg_backup info
```

Conclusion

In this guide, we've successfully walked through the setup of a highly available PostgreSQL cluster using Patroni. By configuring the Etcd cluster for distributed configuration management, setting up PostgreSQL with Patroni for automatic failover and replication, and implementing HAProxy and Keepalived for load balancing and high availability, we ensure a resilient and fault-tolerant system. Additionally, by integrating pgBackRest for backup solutions, we've set up a reliable mechanism for disaster recovery, which is crucial for maintaining data integrity and availability. This high availability setup ensures that even if one or more nodes fail, the cluster can continue to function seamlessly, minimizing downtime and ensuring business continuity.