Member-only story

10 Common Reasons Why Your PostgreSQL Queries Are Slow (And How to Fix Them)

Ajaymaurya · Follow

2 min read · Mar 26, 2025

Listen Share More

PostgreSQL is a powerful and feature-rich database, but slow queries can cripple performance if not optimized properly. Have you ever wondered why your queries

take forever to execute? In this guide, we'll explore 10 common reasons why PostgreSQL queries are slow and provide actionable solutions to fix them.

1. Missing or Inefficient Indexes

Problem:
Indexes speed up query performance, but if they're missing or not used properly, queries can take significantly longer.

Fix:

- Use **EXPLAIN ANALYZE** to check if an index is being used.
- Create appropriate indexes using `CREATE INDEX` based on your query patterns
- Consider **GIN, GiST, BRIN, or B-Tree indexes** for specific use cases.

Example:

```
sql

CREATE INDEX idx_users_email ON users(email);
```

2. Poorly Written Queries

Problem:
Bad query structure can lead to excessive table scans, unnecessary joins, and redundant operations.

Fix:

- Use **SELECT only the necessary columns** instead of `SELECT *`.
- Optimize joins by ensuring proper **indexing on join columns.**
- Avoid **subqueries** when a **JOIN** or **CTE (Common Table Expression)** can do the job.

3. Lack of VACUUM and ANALYZE

Problem:
PostgreSQL doesn't automatically reclaim storage from deleted/updated rows, causing performance degradation.

Fix:

- Run **VACUUM** to clean up dead tuples.
- Use **ANALYZE** to update statistics for the query planner.
- Enable **autovacuum** for automatic maintenance.

Example:

```
sql

VACUUM ANALYZE users;
```

4. Bloated Tables and Indexes

Problem:
Over time, tables and indexes accumulate unnecessary data, leading to increased query time.

Fix:

- Regularly run `VACUUM FULL` and `REINDEX`.
- Monitor table bloat using **pg_stat_user_tables.**

Example:

```
sql

REINDEX TABLE users;
```

5. Using the Wrong Data Types

Fix:

- Use **INTEGER instead of TEXT** for numerical values.
- Use **UUID** instead of large VARCHARs for IDs.
- Prefer **TIMESTAMP** over STRING for date-related fields.

6. Too Many Sequential Scans

Problem:
Sequential scans scan entire tables, slowing down large dataset queries.

Fix:

- Add proper indexes to enforce **Index Scans** instead.
- Use **SET enable_seqscan = OFF** for testing index usage.

Example:

```
sql

SET enable_seqscan = OFF;
```

7. Locking Issues and Deadlocks

Problem:
Long-running transactions cause table locks, making queries wait indefinitely.

Fix:

- Keep transactions short and commit early.
- Use **LOCK TIMEOUT** to prevent waiting indefinitely

Example:

```
sql

SET lock_timeout = '5s';
```

8. Lack of Connection Pooling

Problem:

Excessive database connections slow down performance.

Fix:

- Use connection pooling tools like **PgBouncer**.
- Limit max connections and use **persistent connections.**

9. Overuse of ORDER BY and DISTINCT

Problem:

Sorting large datasets requires significant memory and slows down queries.

Fix:

- Create indexes on sorted columns.
- Avoid **DISTINCT** if `GROUP BY` can achieve the same

10. Incorrect PostgreSQL Configuration

Problem:

Default PostgreSQL settings may not be optimized for large workloads.

Fix:

- Adjust `shared_buffers`, `work_mem`, and `effective_cache_size`.
- Use **pg_stat_statements** to monitor slow queries.

Postgresql Data Sql Database Index

Follow

Written by Ajaymaurya

208 Followers · 781 Following

Engineer in Electronics and IT

No responses yet

Gvadakte

What are your thoughts?

More from Ajaymaurya

Ajaymaurya

What Are the Best Indexing Strategies for High-Performance PostgreSQL Queries?

PostgreSQL is a powerful open-source relational database, but ensuring optimal performance requires well-planned indexing strategies. A...

Mar 15 1

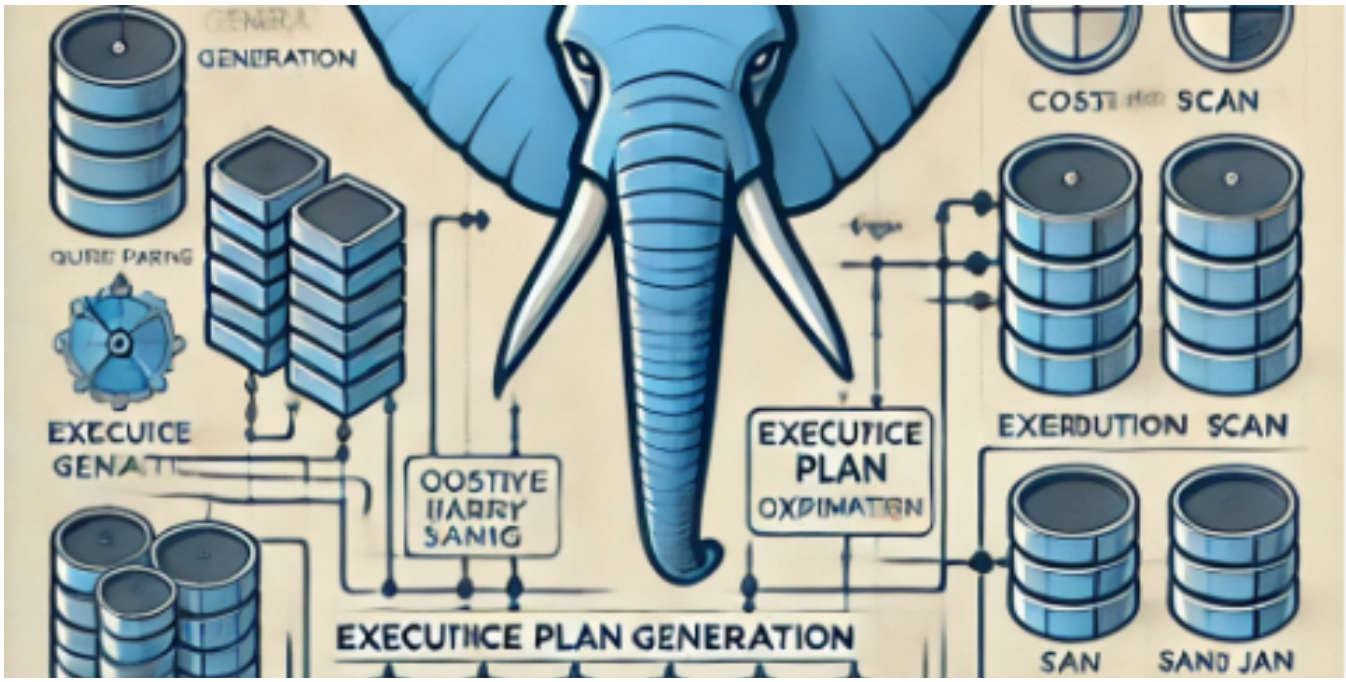

Ajaymaurya

PostgreSQL Internals: How the Query Planner and Optimizer Work

PostgreSQL is a powerful open-source database known for its efficiency and reliability. But have you ever wondered how it decides the best...

Mar 29 16

Ajaymaurya

How Often Should You Reindex Your PostgreSQL Database? A Data-Driven Approach

PostgreSQL is one of the most powerful and flexible relational databases available today. However, like any database system, it requires...

Mar 24

Ajaymaurya

How Do B-Tree, Hash, GIN, and BRIN Indexes Work in PostgreSQL?

When working with large databases, performance is a key concern. PostgreSQL, one of the most powerful relational databases, provides...

Mar 8

See all from Ajaymaurya

Recommended from Medium

In Databases by Sergey Egorenkov

Making SQL query 40x faster for 10 million rows table

Make your SQL query really fast using this approach

Mar 17 10

In Hack the Stack by Coders Stop

9 Database Optimization Tricks SQL Experts Are Hiding From You

Most developers learn enough SQL to get by—SELECT, INSERT, UPDATE, DELETE, and maybe a few JOINs. They might even know how to create...

Mar 27 186 5

Ajaymaurya

What Are the Best Indexing Strategies for High-Performance PostgreSQL Queries?

PostgreSQL is a powerful open-source relational database, but ensuring optimal performance requires well-planned indexing strategies. A...

Mar 15 1

at it Means	Best Used For
e directly in the row	Simple data like INT
e in the row (unless large)	Larger types, but tri
press + store out-of-row	Long texts, large ob
e out-of-row, no compression	When compression

Udbhav Singh

Storages in PostgreSQL

What Are Storage Types in PostgreSQL — And Why Should You Care?

6d ago 18

Tihomir Manushev

Vector Search with pgvector in PostgreSQL

Simple AI-powered similarity search

Mar 9

Mateus Trentz

7 Amazing DBeaver Tips and Tricks to Improve Your SQL Workflow

Straight-to-the-point tips for the best SQL IDE

Mar 18 45 2

See more recommendations