

# PostgreSQL Performance Tuning: A Comprehensive Guide

## Table of Contents

1. Query Optimization
2. Indexing Strategies
3. Configuration Tuning
4. Memory Management
5. Monitoring and Analysis
6. Maintenance Operations

## Query Optimization

### ● EXPLAIN ANALYZE

```
-- Basic EXPLAIN ANALYZE
EXPLAIN ANALYZE
SELECT * FROM orders
WHERE order_date >= '2025-01-01'
AND customer_id IN (
    SELECT id FROM customers WHERE country = 'USA'
);

-- Show buffers and timing information
EXPLAIN (ANALYZE, BUFFERS, TIMING)
SELECT * FROM orders
WHERE order_date >= '2025-01-01';
```

### ● Query Optimization Techniques

```
-- Use EXISTS instead of IN for better performance
-- Before
SELECT * FROM orders
WHERE customer_id IN (SELECT id FROM customers WHERE country = 'USA');

-- After
SELECT * FROM orders o
WHERE EXISTS (
    SELECT 1 FROM customers c
    WHERE c.id = o.customer_id
    AND c.country = 'USA'
);
```

```

-- Use JOIN instead of correlated subqueries
-- Before
SELECT *,
    (SELECT COUNT(*) FROM order_items oi WHERE oi.order_id = o.id)
FROM orders o;

-- After
SELECT o.*, COUNT(oi.id)
FROM orders o
LEFT JOIN order_items oi ON oi.order_id = o.id
GROUP BY o.id;

```

## Indexing Strategies

### ● Index Types

#### ● -- B-tree index (default)

```
CREATE INDEX idx_orders_date ON orders(order_date);
```

#### ● -- Partial index

```
CREATE INDEX idx_orders_status ON orders(status)
WHERE status IN ('pending', 'processing');
```

#### ● -- Multi-column index

```
CREATE INDEX idx_orders_customer_date ON orders(customer_id, order_date);
```

#### ● -- Expression index

```
CREATE INDEX idx_lower_email ON customers(LOWER(email));
```

#### ● -- BRIN index for sequential data

```
CREATE INDEX idx_orders_date_brin ON orders USING BRIN(order_date);
```

### ● Index Maintenance

```

-- Find unused indexes
SELECT
    schemaname || '.' || tablename as table_name,
    indexname,
    idx_scan,
    idx_tup_read,
    idx_tup_fetch
FROM pg_stat_user_indexes
WHERE idx_scan = 0
AND schemaname NOT IN ('pg_catalog', 'pg_toast')
ORDER BY pg_relation_size(indexrelid) DESC;

-- Reindex table
REINDEX TABLE orders;

-- Concurrent reindex (no lock)
CREATE INDEX CONCURRENTLY idx_new_index ON orders(column_name);
DROP INDEX CONCURRENTLY idx_old_index;

```

## Configuration Tuning

### ● Memory Settings

```
# postgresql.conf

# Memory Configuration
shared_buffers = 2GB      # 25% of RAM for dedicated servers
work_mem = 16MB           # Depends on max_connections
maintenance_work_mem = 256MB # For maintenance operations
effective_cache_size = 6GB # 75% of RAM for dedicated servers

# Query Planning
random_page_cost = 1.1    # For SSD storage
effective_io_concurrency = 200 # For SSD storage
```

### ● Connection Settings

```
# postgresql.conf

# Connection Settings
max_connections = 100
superuser_reserved_connections = 3

# Statement Timeout
statement_timeout = '1min'
lock_timeout = '10s'
idle_in_transaction_session_timeout = '1min'
```

## Memory Management

### ● Vacuum Settings

```
# postgresql.conf

# Autovacuum Configuration
autovacuum = on
autovacuum_vacuum_scale_factor = 0.1
autovacuum_analyze_scale_factor = 0.05
autovacuum_vacuum_cost_delay = 2ms
autovacuum_vacuum_cost_limit = 200
ini
```

## Buffer Cache Management

```
-- Check buffer cache hit ratio
SELECT
    sum(heap_blks_read) as heap_read,
```

```

sum(heap_blks_hit) as heap_hit,
sum(heap_blks_hit) / (sum(heap_blks_hit) + sum(heap_blks_read))::float as ratio
FROM pg_statio_user_tables;

-- Find tables with low cache hit ratio
SELECT
    schemaname,
    relname,
    heap_blks_read,
    heap_blks_hit,
    heap_blks_hit::float / (heap_blks_read + heap_blks_hit) as hit_ratio
FROM pg_statio_user_tables
WHERE heap_blks_read + heap_blks_hit > 0
ORDER BY hit_ratio ASC;

```

## Monitoring and Analysis

### ● Performance Monitoring

```

-- Monitor active queries
SELECT
    pid,
    age(clock_timestamp(), query_start) as duration,
    username,
    query
FROM pg_stat_activity
WHERE state != 'idle'
AND query NOT ILIKE '%pg_stat_activity%'
ORDER BY duration DESC;

-- Find slow queries
SELECT
    substring(query, 1, 50) as short_query,
    round(total_time::numeric, 2) as total_time,
    calls,
    round(mean_time::numeric, 2) as mean_time,
    round((100 * total_time / sum(total_time::numeric) over ()):numeric, 2) as percentage
FROM pg_stat_statements
ORDER BY total_time DESC
LIMIT 10;

```

### ● Table Statistics

```

-- Table size and bloat
SELECT
    schemaname,
    tablename,
    pg_size_pretty(pg_total_relation_size(schemaname || '.' || tablename)) as total_size,
    pg_size_pretty(pg_table_size(schemaname || '.' || tablename)) as table_size,
    pg_size_pretty(pg_indexes_size(schemaname || '.' || tablename)) as index_size,
    pg_size_pretty(
        pg_total_relation_size(schemaname || '.' || tablename) -

```

```
pg_table_size(schemaname || '.' || tablename)
) as bloat_size
FROM pg_tables
WHERE schemaname NOT IN ('pg_catalog', 'information_schema')
ORDER BY pg_total_relation_size(schemaname || '.' || tablename) DESC;
```

sql

## Maintenance Operations

### ● Regular Maintenance Tasks

```
-- Analyze tables
ANALYZE VERBOSE;
```

-- Update table statistics

```
ANALYZE VERBOSE mytable;
```

-- VACUUM tables

```
VACUUM (VERBOSE, ANALYZE) mytable;
```

-- Reindex database

```
REINDEX DATABASE mydb;
```

sql

### ● Maintenance Schedule

```
-- Create maintenance function
CREATE OR REPLACE FUNCTION perform_maintenance()
RETURNS void AS $$
BEGIN
    -- Vacuum analyze all tables
    VACUUM (ANALYZE, VERBOSE);

    -- Update statistics
    ANALYZE VERBOSE;

    -- Reindex specific tables if needed
    REINDEX TABLE frequently_updated_table;
END;
$$ LANGUAGE plpgsql;
```

-- Schedule maintenance (using cron or similar)

```
SELECT perform_maintenance();
```

sql

## Performance Tuning Checklist

### ● Query Optimization

1. Use EXPLAIN ANALYZE
2. Optimize JOIN operations
3. Use appropriate subquery types

## ● Indexing

1. Create necessary indexes
2. Remove unused indexes
3. Use appropriate index types

## ● Configuration

1. Optimize memory settings
2. Configure autovacuum
3. Set appropriate timeouts

## ● Monitoring

1. Track slow queries
2. Monitor cache hit ratios
3. Check for bloat

## ● Maintenance

1. Regular VACUUM
2. Update statistics
3. Reindex when needed