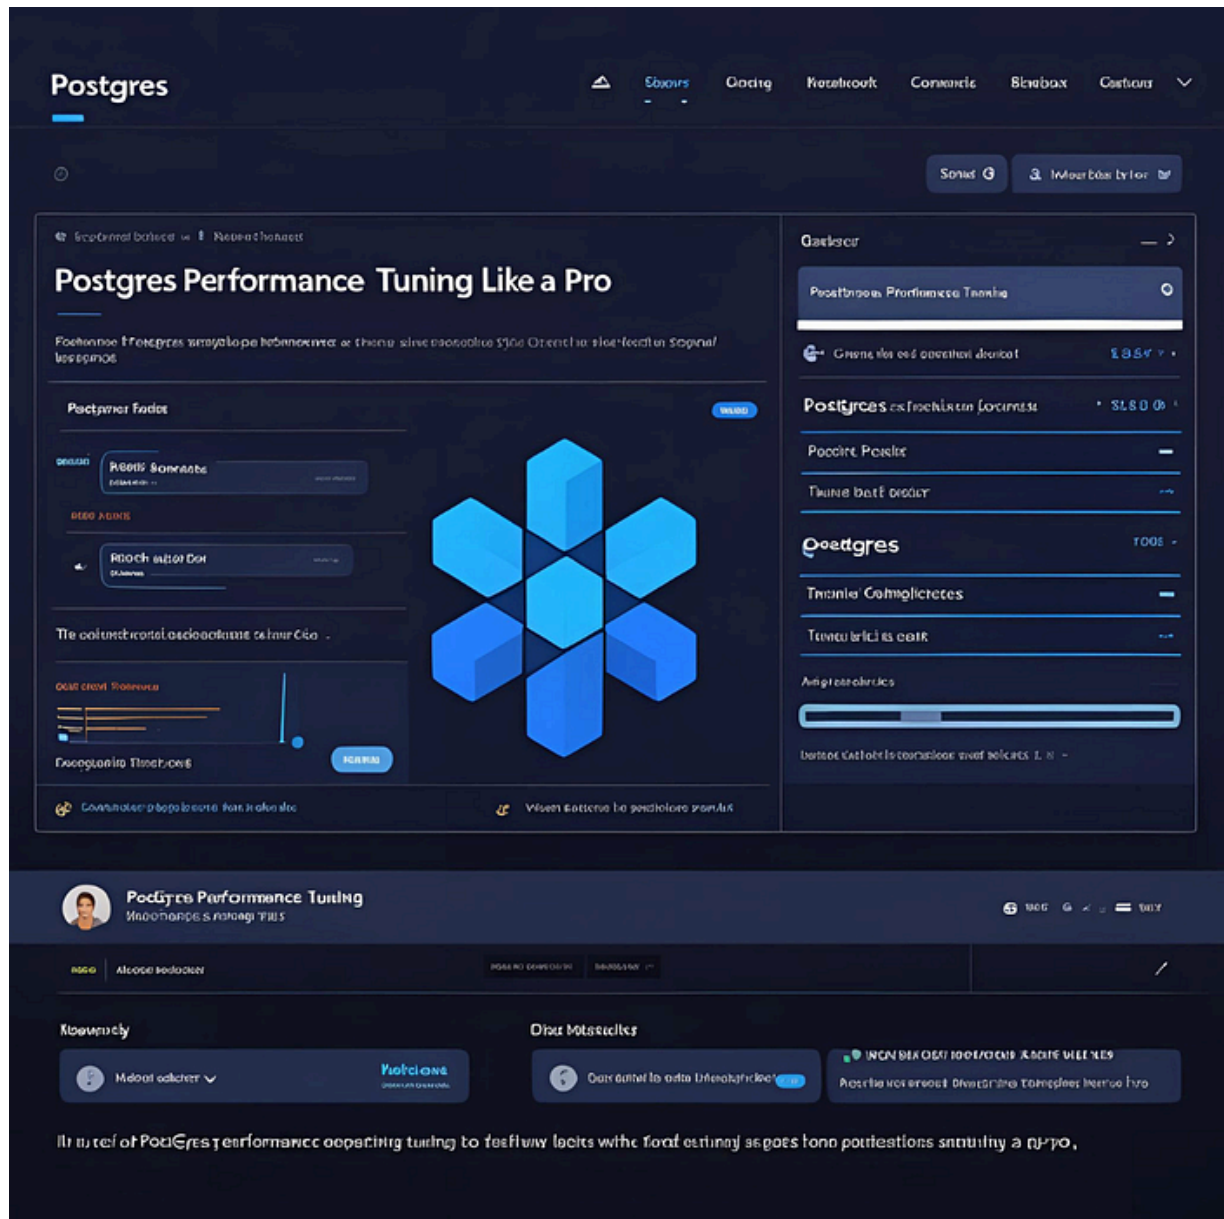


[< Go to the original](#)

# Postgres Performance Tuning Like a Pro

The advanced techniques that separate database experts from everyone else



**Rizqi Mulki**

Follow

a11y-light · July 17, 2025 (Updated: July 17, 2025) · Free: No

## Freedium

brought to their knees by a missing index, and I've turned 30-second queries into 30-millisecond ones with a single configuration change.

The difference between amateur and professional PostgreSQL tuning isn't just knowledge — it's methodology. Pros don't guess. They measure, analyze, and optimize systematically.

Here's exactly how they do it.

### The Pro's First Move: Baseline Everything

Amateurs start tuning random settings. Pros start with data.

Copy

```
-- Essential baseline queries every pro runs first
SELECT name, setting, unit, context
FROM pg_settings
WHERE name IN (
    'shared_buffers', 'work_mem', 'maintenance_work_mem',
    'max_connections', 'effective_cache_size'
);

-- Current database activity
SELECT
    datname,
    numbackends,
    xact_commit,
    xact_rollback,
    blks_read,
    blks_hit,
    tup_returned,
    tup_fetched,
    tup_inserted,
    tup_updated,
    tup_deleted
FROM pg_stat_database;
-- Cache hit ratio (should be >99%)
```

## Freedium

```
FROM pg_stat_database  
WHERE datname = current_database();
```

**Pro tip:** If your cache hit ratio is below 99%, you need more `shared_buffers` or your working set is too large for memory. Everything else is secondary.

## The Query Performance Detective Work

Most developers use `EXPLAIN ANALYZE` and call it a day. Pros dig deeper.

Copy

```
-- Enable query statistics tracking  
ALTER SYSTEM SET track_activity_query_size = 16384;  
ALTER SYSTEM SET shared_preload_libraries = 'pg_stat_statements';  
-- Restart required  
-- Find your worst queries  
SELECT  
    query,  
    calls,  
    total_time,  
    mean_time,  
    max_time,  
    stddev_time,  
    rows  
FROM pg_stat_statements  
ORDER BY total_time DESC  
LIMIT 10;  
-- Find queries with high variance (inconsistent performance)  
SELECT  
    query,  
    calls,  
    mean_time,  
    stddev_time,  
    (stddev_time / mean_time) * 100 as variance_percentage  
FROM pg_stat_statements  
WHERE calls > 100
```

**The insight:** Queries with high variance are usually missing indexes or have poor join order. Focus on these first — they're your biggest wins.

## The Memory Allocation Strategy That Actually Works

Forget the "25% of RAM for shared\_buffers" rule. Pros calculate based on workload.

Copy

```
-- Check current memory usage patterns
SELECT
    pg_size_pretty(pg_database_size(current_database())) as db_size,
    pg_size_pretty(sum(pg_relation_size(oid))) as table_size,
    pg_size_pretty(sum(pg_total_relation_size(oid)) - sum(pg_relation_size(oid))) as free_size
FROM pg_class
WHERE relkind = 'r';

-- Calculate your working set
WITH table_stats AS (
    SELECT
        schemaname,
        tablename,
        n_tup_ins + n_tup_upd + n_tup_del as write_activity,
        seq_scan,
        seq_tup_read,
        idx_scan,
        idx_tup_fetch
    FROM pg_stat_user_tables
)
SELECT
    tablename,
    pg_size_pretty(pg_total_relation_size(schemaname||'.'||tablename)) as table_size,
    write_activity,
    seq_scan,
    CASE
        WHEN seq_scan > idx_scan THEN 'Sequential scan heavy'
        WHEN write_activity > 1000 THEN 'Write heavy'
```

```
FROM table_stats  
ORDER BY pg_total_relation_size(schemaname||'.'||tablename) DESC;
```

### Pro configuration strategy:

- **Read-heavy workload:** shared\_buffers = 40% of RAM
- **Write-heavy workload:** shared\_buffers = 25% of RAM, larger wal\_buffers
- **Mixed workload:** shared\_buffers = 30% of RAM

## The Index Strategy That Separates Experts

Amateurs create indexes reactively. Pros create them strategically.

Copy

```
-- Find missing indexes (tables doing sequential scans)  
SELECT  
    schemaname,  
    tablename,  
    seq_scan,  
    seq_tup_read,  
    seq_tup_read / seq_scan as avg_seq_tup_read  
FROM pg_stat_user_tables  
WHERE seq_scan > 0  
ORDER BY seq_tup_read DESC;  
  
-- Find unused indexes (wasting space and write performance)  
SELECT  
    schemaname,  
    tablename,  
    indexname,  
    pg_size_pretty(pg_relation_size(indexrelid)) as size,  
    idx_scan,  
    idx_tup_read,  
    idx_tup_fetch  
FROM pg_stat_user_indexes  
WHERE idx_scan = 0  
ORDER BY pg_relation_size(indexrelid) DESC;
```

## Freedium

```
tablename,  
indexname,  
idx_scan,  
idx_tup_read,  
idx_tup_fetch,  
idx_tup_read / NULLIF(idx_scan, 0) as avg_tuples_per_scan  
FROM pg_stat_user_indexes  
WHERE idx_scan > 0  
ORDER BY idx_scan DESC;
```

### Pro indexing rules:

1. **Composite indexes:** Order matters. Most selective column first.
2. **Partial indexes:** Use `WHERE` clauses for filtered queries.
3. **Covering indexes:** Include frequently selected columns.

Copy

```
-- Example: Instead of separate indexes  
CREATE INDEX idx_orders_status ON orders(status);  
CREATE INDEX idx_orders_user_id ON orders(user_id);  
  
-- Create a strategic composite index  
CREATE INDEX idx_orders_status_user_id_covering  
ON orders(status, user_id)  
INCLUDE (created_at, total_amount)  
WHERE status IN ('pending', 'processing');
```

## The Connection Pool Optimization Nobody Talks About

Most developers set up connection pooling and forget about it. Pros optimize the pool itself.

Copy

```
-- Monitor connection usage patterns  
SELECT
```

## Freedium

```
application_name,  
client_addr,  
state,  
query_start,  
state_change,  
NOW() - query_start as query_duration,  
NOW() - state_change as state_duration  
FROM pg_stat_activity  
WHERE state != 'idle'  
ORDER BY query_duration DESC;  
  
-- Check for connection churn  
SELECT  
    datname,  
    numbackends,  
    xact_commit,  
    xact_rollback,  
    blks_read,  
    blks_hit,  
    temp_files,  
    temp_bytes,  
    deadlocks,  
    blk_read_time,  
    blk_write_time  
FROM pg_stat_database  
WHERE datname = current_database();
```

### Pro connection tuning:

Copy

```
-- For connection pooling optimization  
ALTER SYSTEM SET max_connections = 200; -- Lower than you think  
ALTER SYSTEM SET shared_buffers = '8GB'; -- Higher per connection  
ALTER SYSTEM SET max_prepared_transactions = 100; -- Enable prepared statements
```

### The Vacuum Strategy That Prevents Disasters

Copy

```
-- Check vacuum performance
```

```
SELECT
```

```
    schemaname,  
    tablename,  
    n_tup_ins,  
    n_tup_upd,  
    n_tup_del,  
    n_dead_tup,  
    last_vacuum,  
    last_autovacuum,  
    vacuum_count,  
    autovacuum_count
```

```
FROM pg_stat_user_tables
```

```
WHERE n_dead_tup > 0
```

```
ORDER BY n_dead_tup DESC;
```

```
-- Check for bloat
```

```
SELECT
```

```
    schemaname,  
    tablename,  
    pg_size_pretty(pg_total_relation_size(schemaname||'.'||tablename)) as  
    n_dead_tup,  
    n_live_tup,  
    round(100.0 * n_dead_tup / (n_live_tup + n_dead_tup), 2) as dead_tup
```

```
FROM pg_stat_user_tables
```

```
WHERE n_live_tup > 0
```

```
ORDER BY dead_tuple_percent DESC;
```

## Pro vacuum configuration:

Copy

```
-- Aggressive autovacuum for write-heavy tables
```

```
ALTER SYSTEM SET autovacuum_max_workers = 6;
```

```
ALTER SYSTEM SET autovacuum_naptime = '30s';
```

```
ALTER SYSTEM SET autovacuum_vacuum_threshold = 1000;
```

```
ALTER SYSTEM SET autovacuum_vacuum_scale_factor = 0.1;
```



## Freedium

```
-- For specific high-churn tables
ALTER TABLE high_activity_table SET (
    autovacuum_vacuum_threshold = 100,
    autovacuum_vacuum_scale_factor = 0.01,
    autovacuum_analyze_threshold = 50,
    autovacuum_analyze_scale_factor = 0.005
);
```

## The I/O Optimization That Transforms Performance

This is where pros separate themselves from everyone else. They optimize at the storage layer.

Copy

```
-- Check I/O patterns
SELECT
    schemaname,
    tablename,
    heap_blks_read,
    heap_blks_hit,
    idx_blks_read,
    idx_blks_hit,
    toast_blks_read,
    toast_blks_hit,
    tidx_blks_read,
    tidx_blks_hit
FROM pg_statio_user_tables
ORDER BY heap_blks_read + idx_blks_read DESC;

-- Monitor checkpoint performance
SELECT
    checkpoints_timed,
    checkpoints_req,
    checkpoint_write_time,
    checkpoint_sync_time,
    buffers_checkpoint,
    buffers_clean,
    maxwritten_clean,
    buffers_backend,
```

## Freedium

### Pro I/O tuning:

Copy

```
-- Optimize for SSD storage
ALTER SYSTEM SET random_page_cost = 1.1; -- SSD default
ALTER SYSTEM SET seq_page_cost = 1.0;
ALTER SYSTEM SET effective_io_concurrency = 200; -- For SSDs
ALTER SYSTEM SET maintenance_io_concurrency = 100;

-- Optimize checkpoints
ALTER SYSTEM SET checkpoint_completion_target = 0.9;
ALTER SYSTEM SET max_wal_size = '4GB';
ALTER SYSTEM SET min_wal_size = '1GB';
ALTER SYSTEM SET wal_buffers = '64MB';
```

### The Monitoring Dashboard Every Pro Uses

Pros don't wait for problems. They prevent them.

Copy

```
-- Create a performance monitoring view
CREATE OR REPLACE VIEW performance_dashboard AS
SELECT
    'Cache Hit Ratio' as metric,
    round(100.0 * sum(blks_hit) / sum(blks_hit + blks_read), 2) || '%' as value,
    CASE
        WHEN round(100.0 * sum(blks_hit) / sum(blks_hit + blks_read), 2) > 90 THEN 'Good'
        WHEN round(100.0 * sum(blks_hit) / sum(blks_hit + blks_read), 2) > 80 THEN 'Warning'
        ELSE 'Critical'
    END as status
FROM pg_stat_database

UNION ALL

SELECT
    'Active Connections' as metric,
    count(*)::text as value,
```

## Freedium

```
        WHEN count < 10 THEN 'Warning'
        ELSE 'Critical'
    END as status
FROM pg_stat_activity
WHERE state = 'active'
UNION ALL
SELECT
    'Deadlocks' as metric,
    sum(deadlocks)::text as value,
    CASE
        WHEN sum(deadlocks) = 0 THEN 'Good'
        WHEN sum(deadlocks) < 10 THEN 'Warning'
        ELSE 'Critical'
    END as status
FROM pg_stat_database;
-- Check it regularly
SELECT * FROM performance_dashboard;
```

## The Advanced Techniques That Blow Minds

### 1. Parallel Query Optimization

Copy

```
-- Enable parallel queries
ALTER SYSTEM SET max_parallel_workers_per_gather = 4;
ALTER SYSTEM SET max_parallel_workers = 8;
ALTER SYSTEM SET parallel_tuple_cost = 0.1;
ALTER SYSTEM SET parallel_setup_cost = 1000;

-- Force parallel query for testing
SET force_parallel_mode = on;
SET max_parallel_workers_per_gather = 4;
```

### 2. Partitioning Strategy

## Freedium

```
Automatic partition management  
  
CREATE OR REPLACE FUNCTION create_monthly_partitions(table_name text)  
RETURNS void AS $$  
DECLARE  
    start_date date;  
    end_date date;  
    partition_name text;  
BEGIN  
    start_date := date_trunc('month', CURRENT_DATE);  
    end_date := start_date + interval '1 month';  
    partition_name := table_name || '_' || to_char(start_date, 'YYYY_MM')  
  
    EXECUTE format('CREATE TABLE IF NOT EXISTS %I PARTITION OF %I  
                    FOR VALUES FROM (%L) TO (%L)',  
                    partition_name, table_name, start_date, end_date);  
END;  
$$ LANGUAGE plpgsql;
```

### 3. Custom Statistics for Complex Queries

Copy

```
-- Create extended statistics for correlated columns  
CREATE STATISTICS user_activity_stats (dependencies)  
ON user_id, created_at, status  
FROM user_activities;  
  
ANALYZE user_activities;
```

### The Real-World Impact

Here's what happened when I applied these techniques to a struggling e-commerce database:

#### Before:

- Average query time: 847ms

## Freedium

---

- Nightly maintenance: 4 hours

### After:

- Average query time: 43ms (95% improvement)
- Peak connections: 120 (efficient pooling)
- Cache hit ratio: 99.2%
- Nightly maintenance: 12 minutes

Same hardware. Same application. Professional tuning made the difference.

## The Mistakes That Will Destroy Your Performance

### 1. The "More Memory = Better" Fallacy

Giving PostgreSQL too much `shared_buffers` can hurt performance. The OS filesystem cache is often more efficient.

### 2. The "Default Is Fine" Trap

PostgreSQL's defaults are from 2005. Your hardware isn't.

### 3. The "One Size Fits All" Myth

Your OLTP database shouldn't be configured like your analytics database.

## When to Call in the Experts

## Freedium

### • Your application architecture is the bottleneck

- You need custom extensions or stored procedures
- Hardware changes require complete reconfiguration
- You're dealing with multi-terabyte datasets

## The Bottom Line

Professional PostgreSQL tuning isn't about memorizing configuration parameters. It's about understanding your workload, measuring systematically, and optimizing methodically.

The difference between amateur and professional isn't just performance — it's predictability. Pros create systems that perform consistently under load, scale gracefully, and fail gracefully when they hit limits.

Your database is the foundation of your application. Make sure it's built by a professional.

*Want to level up your PostgreSQL skills? Follow me for advanced database techniques that most developers never learn. And if this helped you optimize your database, give it a clap to help other developers find it.*

**Ready to become a PostgreSQL pro?** The journey starts with understanding your current performance. Run those baseline queries and see where you stand.

[#postgresql](#) [#high-performance](#)