

[Back to Blog](#)**TheDBAdmin**[Contact](#)

PostgreSQL Performance Tuning: A Comprehensive Guide

TheDBAdmin Team • January 30, 2025

[postgresql](#)[performance](#)[database-administration](#)[optimization](#)[tuning](#)

Performance tuning is crucial for maintaining a healthy and responsive PostgreSQL database. This comprehensive guide covers essential techniques and best practices for optimizing your PostgreSQL database performance.

Table of Contents

[Query Optimization](#)[Indexing Strategies](#)[Configuration Tuning](#)[Memory Management](#)[Monitoring and Analysis](#)[Maintenance Operations](#)

Query Optimization

EXPLAIN ANALYZE

```
-- Basic EXPLAIN ANALYZE
EXPLAIN ANALYZE
SELECT * FROM orders
WHERE order_date >= '2025-01-01'
```

```
AND customer_id IN (  
    SELECT id FROM customers WHERE country = 'USA'  
);  
  
-- Show buffers and timing information  
EXPLAIN (ANALYZE, BUFFERS, TIMING)  
SELECT * FROM orders  
WHERE order_date >= '2025-01-01';
```

Query Optimization Techniques

```
-- Use EXISTS instead of IN for better performance  
-- Before  
SELECT * FROM orders  
WHERE customer_id IN (SELECT id FROM customers WHERE country = 'USA');  
  
-- After  
SELECT * FROM orders o  
WHERE EXISTS (  
    SELECT 1 FROM customers c  
    WHERE c.id = o.customer_id  
    AND c.country = 'USA'  
);  
  
-- Use JOIN instead of correlated subqueries  
-- Before  
SELECT *,  
    (SELECT COUNT(*) FROM order_items oi WHERE oi.order_id = o.id)  
FROM orders o;  
  
-- After  
SELECT o.*, COUNT(oi.id)  
FROM orders o  
LEFT JOIN order_items oi ON oi.order_id = o.id  
GROUP BY o.id;
```

Indexing Strategies

Index Types


```
-- B-tree index (default)
CREATE INDEX idx_orders_date ON orders(order_date);

-- Partial index
CREATE INDEX idx_orders_status ON orders(status)
WHERE status IN ('pending', 'processing');

-- Multi-column index
CREATE INDEX idx_orders_customer_date ON orders(customer_id, order_date)

-- Expression index
CREATE INDEX idx_lower_email ON customers(LOWER(email));

-- BRIN index for sequential data
CREATE INDEX idx_orders_date_brin ON orders USING BRIN(order_date);
```



Index Maintenance

```
-- Find unused indexes
SELECT
    schemaname || '.' || tablename as table_name,
    indexname,
    idx_scan,
    idx_tup_read,
    idx_tup_fetch
FROM pg_stat_user_indexes
WHERE idx_scan = 0
AND schemaname NOT IN ('pg_catalog', 'pg_toast')
ORDER BY pg_relation_size(indexrelid) DESC;

-- Reindex table
REINDEX TABLE orders;

-- Concurrent reindex (no lock)
CREATE INDEX CONCURRENTLY idx_new_index ON orders(column_name);
DROP INDEX CONCURRENTLY idx_old_index;
```

Configuration Tuning

Memory Settings

```
# postgresql.conf

# Memory Configuration
shared_buffers = 2GB          # 25% of RAM for dedicated servers
work_mem = 16MB               # Depends on max_connections
maintenance_work_mem = 256MB # For maintenance operations
effective_cache_size = 6GB     # 75% of RAM for dedicated servers

# Query Planning
random_page_cost = 1.1        # For SSD storage
effective_io_concurrency = 200 # For SSD storage
```

Connection Settings

```
# postgresql.conf

# Connection Settings
max_connections = 100
superuser_reserved_connections = 3

# Statement Timeout
statement_timeout = '1min'
lock_timeout = '10s'
idle_in_transaction_session_timeout = '1min'
```

Memory Management

Vacuum Settings

```
# postgresql.conf

# Autovacuum Configuration
autovacuum = on
autovacuum_vacuum_scale_factor = 0.1
autovacuum_analyze_scale_factor = 0.05
```

```
autovacuum_vacuum_cost_delay = 2ms
```

```
autovacuum_vacuum_cost_limit = 200
```

Buffer Cache Management

```
-- Check buffer cache hit ratio
```

```
SELECT
```

```
    sum(heap_blks_read) as heap_read,
```

```
    sum(heap_blks_hit)  as heap_hit,
```

```
    sum(heap_blks_hit) / (sum(heap_blks_hit) + sum(heap_blks_read))::float
```

```
FROM pg_statio_user_tables;
```

```
-- Find tables with low cache hit ratio
```

```
SELECT
```

```
    schemaname,
```

```
    relname,
```

```
    heap_blks_read,
```

```
    heap_blks_hit,
```

```
    heap_blks_hit::float / (heap_blks_read + heap_blks_hit) as hit_ratio
```

```
FROM pg_statio_user_tables
```

```
WHERE heap_blks_read + heap_blks_hit > 0
```

```
ORDER BY hit_ratio ASC;
```

Monitoring and Analysis

Performance Monitoring

```
-- Monitor active queries
```

```
SELECT
```

```
    pid,
```

```
    age(clock_timestamp(), query_start) as duration,
```

```
    username,
```

```
    query
```

```
FROM pg_stat_activity
```

```
WHERE state != 'idle'
```

```
AND query NOT ILIKE '%pg_stat_activity%'
```

```
ORDER BY duration DESC;
```

```
-- Find slow queries
```

```

SELECT
    substring(query, 1, 50) as short_query,
    round(total_time::numeric, 2) as total_time,
    calls,
    round(mean_time::numeric, 2) as mean_time,
    round((100 * total_time / sum(total_time::numeric) over ())::numeric
FROM pg_stat_statements
ORDER BY total_time DESC
LIMIT 10;

```

Table Statistics

```

-- Table size and bloat
SELECT
    schemaname,
    tablename,
    pg_size_pretty(pg_total_relation_size(schemaname || '.' || tablename)
    pg_size_pretty(pg_table_size(schemaname || '.' || tablename)) as tab_size,
    pg_size_pretty(pg_indexes_size(schemaname || '.' || tablename)) as index_size,
    pg_size_pretty(
        pg_total_relation_size(schemaname || '.' || tablename) -
        pg_table_size(schemaname || '.' || tablename)
    ) as bloat_size
FROM pg_tables
WHERE schemaname NOT IN ('pg_catalog', 'information_schema')
ORDER BY pg_total_relation_size(schemaname || '.' || tablename) DESC;

```

Maintenance Operations

Regular Maintenance Tasks

```

-- Analyze tables
ANALYZE VERBOSE;

-- Update table statistics
ANALYZE VERBOSE mytable;

-- VACUUM tables

```

```
VACUUM (VERBOSE, ANALYZE) mytable;
```

```
-- Reindex database
```

```
REINDEX DATABASE mydb;
```

Maintenance Schedule

```
-- Create maintenance function
```

```
CREATE OR REPLACE FUNCTION perform_maintenance()
```

```
RETURNS void AS $$
```

```
BEGIN
```

```
    -- Vacuum analyze all tables
```

```
    VACUUM (ANALYZE, VERBOSE);
```

```
    -- Update statistics
```

```
    ANALYZE VERBOSE;
```

```
    -- Reindex specific tables if needed
```

```
    REINDEX TABLE frequently_updated_table;
```

```
END;
```

```
$$ LANGUAGE plpgsql;
```

```
-- Schedule maintenance (using cron or similar)
```

```
SELECT perform_maintenance();
```

Performance Tuning Checklist

Query Optimization

- ☐ Use EXPLAIN ANALYZE
- ☐ Optimize JOIN operations
- ☐ Use appropriate subquery types

Indexing

- ☐ Create necessary indexes
- ☐ Remove unused indexes
- ☐ Use appropriate index types

Configuration

- ☐ Optimize memory settings
- ☐ Configure autovacuum
- ☐ Set appropriate timeouts

Monitoring

- ☐ Track slow queries

- ☐ Monitor cache hit ratios
- ☐ Check for bloat

Maintenance

- ☐ Regular VACUUM
- ☐ Update statistics
- ☐ Reindex when needed

Further Reading

[PostgreSQL Official Documentation](#)

[TheDBAdmin PostgreSQL Security Guide](#)

[50 Essential PostgreSQL Queries](#)

[PostgreSQL DBA Course](#)

Remember to test performance changes in a development environment first and monitor the impact of each optimization.

Share this article



Tags

oracle (1) postgresql (13) aws-dms (1) ora2pg (1) database-migration (1)
cloud-migration (1) database-administration (11) ai (6) machine-learning (4)
dba (2) career-development (2) data-science (2) mlops (1) qwen (1)
local-deployment (2) llm (2) ollama (2) deepseek (3) nosql (1)
career-guide (1) mongodb (2) cassandra (1) redis (2) docker (2)
docker-compose (2) containerization (2) monitoring (1) performance (5)
devops (1) backup (1) recovery (1) disaster-recovery (2) caching (1)
database (3) high-availability (1) replication (1) failover (1) scalability (1)
high-concurrency (1) database-tuning (1) optimization (1) tuning (1)
security (1) encryption (1) authentication (1) sql (1) tutorials (1)
automation (1) productivity (1) openai (1) claude (1) comparison (1)
technology (1) career (1) certification (1) welcome (1) introduction (1)

Recent Posts

Oracle to PostgreSQL Migration: Comprehensive Guide Using AWS DMS & Ora2Pg

Jan 31, 2025

PostgreSQL DBA's Role in AI & ML: A Comprehensive Guide to the Future

Jan 31, 2025

Running Qwen 2.5 Locally: Complete Implementation Guide

Jan 30, 2025

Running DeepSeek Models Locally with Ollama: Complete Guide

Jan 30, 2025

How to Become a NoSQL Database Administrator: Complete Career Guide

Jan 30, 2025



TheDBAdmin

Database Administration Services

Expert database administration and consulting services for modern enterprises.



Services

[Database Services](#)

[Courses](#)

[Blog](#)

Support

[Documentation](#)

[Help Center](#)

[Status](#)

Company

[About Us](#)

[Contact](#)

Stay Updated

Subscribe to our newsletter for the latest updates and insights.

Subscribe

© 2025 TheDBAdmin. All rights reserved.

[Privacy Policy](#) [Terms of Service](#) [Cookie Policy](#)