

Setting up replication in PostgreSQL involves configuring a primary (master) server and one or more standby (replica) servers. PostgreSQL supports two main types of replication: **physical replication** (streaming replication) and **logical replication**. In this guide, we'll focus on **physical replication**, which is commonly used for high availability (HA) setups.

## Steps for Setting Up Streaming (Physical) Replication in PostgreSQL

### 1. Prepare the Primary (Master) Server

#### a. Edit `postgresql.conf` on the Primary

On the primary server, you need to adjust some configuration settings to enable replication.

##### 1. Open the `postgresql.conf` file:

```
[postgres@localhost pg_data]$ ls -lrt postgresql.conf
-rw-----. 1 postgres postgres 29686 Oct 20 13:09 postgresql.conf
[postgres@localhost pg_data]$ pwd
/postgres/pg_data
[postgres@localhost pg_data]$ vi /postgres/pg_data/postgresql.conf
```

##### 2. Enable WAL archiving and configure replication settings: Add or modify the following parameters:

```
wal_level = replica           # Required for physical replication
max_wal_senders = 5           # Number of replication connections allowed
wal_keep_size = 64            # Amount of WAL files (in MB) to retain for
replicas
archive_mode = on             # Enable archiving of WAL segments (optional)
archive_command = 'cp %p /path_to_wal_archive/%f' # Command to archive WAL
(optional)
listen_addresses = '*'        # Allow connections from any IP (or specify an
IP)
```

- `wal_level = replica`: This is required for streaming replication.
- `max_wal_senders`: This sets the number of standby servers that can connect for replication.
- `wal_keep_size`: This defines the amount of WAL logs to keep to prevent standby servers from falling behind.
- `listen_addresses`: This allows remote connections to the primary server.

```
postgres=# show wal_level;
wal_level
-----
replica
(1 row)
```

```
postgres=# show archive_mode;
archive_mode
-----
on
(1 row)
```

```
postgres=# show max_wal_senders ;
max_wal_senders
-----
```

```
10
(1 row)
```

```
postgres=# show max_replication_slots;
max_replication_slots
-----
10
(1 row)
```

### 3. Reload or restart PostgreSQL to apply these changes:

```
systemctl restart postgresql
```

#### b. Edit `pg_hba.conf` on the Primary

The `pg_hba.conf` file controls client authentication, and you need to allow replication connections from the standby server.

##### 1. Open the `pg_hba.conf` file:

```
Vi /postgres/pg_data/pg_hba.conf

[postgres@localhost archivelogs]$ hostname -I
192.168.30.9 192.168.122.1
[postgres@localhost archivelogs]$
```

##### 2. Add a replication entry: Add a line that allows the standby server to connect for replication. Replace `standby_ip_address` with the IP address of the standby server:

```
host      replication      replication_user 192.168.30.9/32      md5
host      replication      all             192.168.30.9/32      md5
```

This allows the standby to connect using an IP address and `md5` authentication.

### 3. Reload PostgreSQL:

```
systemctl reload postgresql
```

#### c. Create a Replication Role on the Primary

You need a user that has replication privileges.

##### 1. Create the replication role:

```
postgres=# create user replication_user password 'postgres' replication;
CREATE ROLE
postgres=#
```

##### 2. Grant the role appropriate privileges, if needed.

## 2. Prepare the Standby (Replica) Server

#### a. Base Backup from Primary

To set up replication, the standby needs a copy of the data directory from the primary server.

## 1. Stop the standby server if it's running:

```
systemctl stop postgresql
```

## 2. Use the **pg\_basebackup** command to copy the data from the primary to the standby server:

```
pg_basebackup -h primary_ip_address -D /home/PG/backup/standby -U  
replication_user -P -R
```

```
[postgres@localhost rep]$ pg_basebackup -h localhost -D  
/home/PG/backup/rep/ -U replication_user -P -R
```

```
Password:
```

```
38300/38300 kB (100%), 1/1 tablespace
```

```
[postgres@localhost rep]$ ls -lrt
```

- **-h primary\_ip\_address**: IP address of the primary server.
- **-D /postgres/pg\_data/standby/data\_directory**: Data directory of the standby server.
- **-U replication\_user**: User created for replication.
- **-P**: Show progress.
- **-R**: Automatically creates a recovery.conf file on the standby, which is needed for replication.

```
[postgres@localhost rep]$ ls -lrt
```

```
total 284
```

```
-rw-----. 1 postgres postgres 225 Oct 20 17:39 backup_label  
drwx-----. 3 postgres postgres 60 Oct 20 17:39 pg_wal  
drwx-----. 7 postgres postgres 59 Oct 20 17:39 base  
drwx-----. 2 postgres postgres 6 Oct 20 17:39 pg_snapshots  
drwx-----. 2 postgres postgres 6 Oct 20 17:39 pg_serial  
drwx-----. 2 postgres postgres 6 Oct 20 17:39 pg_replslot  
drwx-----. 2 postgres postgres 6 Oct 20 17:39 pg_notify  
drwx-----. 4 postgres postgres 36 Oct 20 17:39 pg_multixact  
drwx-----. 4 postgres postgres 68 Oct 20 17:39 pg_logical  
drwx-----. 2 postgres postgres 6 Oct 20 17:39 pg_dynshmem  
drwx-----. 2 postgres postgres 6 Oct 20 17:39 pg_commit_ts  
drwx-----. 2 postgres postgres 32 Oct 20 17:39 log  
-rw-----. 1 postgres postgres 29686 Oct 20 17:39 postgresql.conf  
-rw-----. 1 postgres postgres 465 Oct 20 17:39 postgresql.auto.conf  
drwx-----. 2 postgres postgres 18 Oct 20 17:39 pg_xact  
-rw-----. 1 postgres postgres 3 Oct 20 17:39 PG_VERSION  
drwx-----. 2 postgres postgres 6 Oct 20 17:39 pg_twophase  
drwx-----. 2 postgres postgres 6 Oct 20 17:39 pg_tblspc  
drwx-----. 2 postgres postgres 6 Oct 20 17:39 pg_subtrans  
drwx-----. 2 postgres postgres 6 Oct 20 17:39 pg_stat_tmp  
drwx-----. 2 postgres postgres 6 Oct 20 17:39 pg_stat  
-rw-----. 1 postgres postgres 2640 Oct 20 17:39 pg_ident.conf  
-rw-----. 1 postgres postgres 5636 Oct 20 17:39 pg_hba.conf  
-rw-----. 1 postgres postgres 30 Oct 20 17:39 current_logfiles  
drwx-----. 2 postgres postgres 4096 Oct 20 17:39 global  
-rw-----. 1 postgres postgres 0 Oct 20 17:39 standby.signal  
-rw-----. 1 postgres postgres 225003 Oct 20 17:39 backup_manifest
```

```
[postgres@localhost rep]$ cat postgresql.auto.conf
```

```
# Do not edit this file manually!
```

```
# It will be overwritten by the ALTER SYSTEM command.
```

```
archive_mode = 'on'
```

```
archive_command = 'cp %p /home/PG/archive/logs/%f'
```

```
primary_conninfo = 'user=replication_user password=postgres
channel_binding=prefer host=localhost port=5432 sslmode=prefer
sslcompression=0 sslcertmode=allow sslssl_min_protocol_version=TLSv1.2
gssencmode=prefer krbsrvname=postgres gssdelegation=0
target_session_attrs=any load_balance_hosts=disable'
```

## b. Modify the Standby's Configuration

1. **Edit postgresql.conf:** You need to set up replication parameters on the standby server.

```
vi /home/PG/backup/rep/postgresql.conf
```

Add or modify the following lines:

```
hot_standby = on                # Enable read queries on the standby
```

2. **Edit recovery.conf:** The pg\_basebackup command will automatically create a recovery.conf file in the data directory of the standby server if you used the -R option. However, you should verify its content.

o

```
[postgres@localhost rep]$ cat postgresql.auto.conf
# Do not edit this file manually!
# It will be overwritten by the ALTER SYSTEM command.
archive_mode = 'on'
archive_command = 'cp %p /home/PG/archivelogs/%f'
primary_conninfo = 'user=replication_user password=postgres
channel_binding=prefer host=localhost port=5432 sslmode=prefer
sslcompression=0 sslcertmode=allow
sslssl_min_protocol_version=TLSv1.2 gssencmode=prefer
krbsrvname=postgres gssdelegation=0 target_session_attrs=any
load_balance_hosts=disable'
[postgres@localhost rep]$ vi postgresql.conf
port=5434 → port no changed to 5434
```

- o standby\_mode = 'on': Ensures the server stays in standby mode.
- o primary\_conninfo: Contains the connection information for the primary server.

3. **Start the standby server:**

```
systemctl start postgresql
```

or

```
[postgres@localhost rep]$ pg_ctl -D /home/PG/backup/rep/ start
waiting for server to start....2024-10-20 17:43:36.223 IST [27196] LOG:
redirecting log output to logging collector process
2024-10-20 17:43:36.223 IST [27196] HINT: Future log output will appear in
directory "log".
done
server started
[postgres@localhost rep]$ psql
psql (16.4)
Type "help" for help.

postgres=#
```

```
[postgres@localhost archivelogs]$ psql postgres --port=5434
psql (16.4)
Type "help" for help.

postgres=# \l+
```

### 3. Verify Replication Setup

#### a. Check Replication Status on the Primary

On the primary server, you can check if the standby is connected using the following SQL command:

```
SELECT * FROM pg_stat_replication;
```

```
[postgres@localhost rep]$ psql postgres --port=5434
psql (16.4)
Type "help" for help.
```

```
postgres=# SELECT * FROM pg_stat_replication;
```

```
 pid | usesysid | username | application_name | client_addr | client_hostname |
client_port | backend_start | backend_xmin | state | sent_lsn | write_lsn | flush_l
sn | replay_lsn | write_lag | flush_lag | replay_lag | sync_priority | sync_state |
reply_time
-----+-----+-----+-----+-----+-----+-----+-----+-----+-----+-----+-----+-----
(0 rows)

postgres=#
```

This should show an entry for each standby server connected.

#### b. Check Replication Lag (Optional)

You can monitor the replication lag using:

```
SELECT pg_last_wal_receive_lsn(), pg_last_wal_replay_lsn();
```

```
postgres=# SELECT pg_last_wal_receive_lsn(), pg_last_wal_replay_lsn();
```

```
 pg_last_wal_receive_lsn | pg_last_wal_replay_lsn
-----+-----
 0/53EC89D0              | 0/53EC89D0
(1 row)

postgres=#
```

The closer these two values are, the smaller the lag between the primary and standby.

### 4. Optional: Enable Read-Only Queries on the Standby

By default, the standby server in streaming replication mode is read-only. You can connect to the standby and execute read-only queries, which is useful for load balancing read operations.

### Important Considerations:

- **WAL Retention:** Ensure the primary retains enough WAL files to allow the standby to catch up in case it falls behind. This is controlled by `wal_keep_size`.
- **Failover:** In case of primary failure, you can promote the standby to a primary by running:

```
pg_ctl promote -D /postgres/pg_data/standby/data_directory
```

- **Monitoring:** It's important to monitor the replication status, especially in production environments, to ensure that the standby is not falling behind.

This is the basic setup for **streaming replication**. You can extend this with **asynchronous/synchronous replication** or implement **failover management tools** like `repmgr` or `Patroni` for automatic failover handling.