

★ Member-only story

# Setting Up Streaming Replication in PostgreSQL on RHEL 9



Oz · Following

2 min read · Jul 10, 2023



Listen




Share



More

Streaming replication in PostgreSQL allows you to create one or more standby servers that continuously update their data from a primary (master) server. Here's a step-by-step guide to setting up streaming replication between a primary and standby server on RHEL (Red Hat Enterprise Linux).



## Prerequisites

- Two RHEL servers: Primary (psql-primary) and Standby (psql-standby)
- PostgreSQL 12 installed on both servers

## Primary Server Configuration

### Host and IP Configuration

Edit the `/etc/hosts` and `/etc/hostname` files to set the hostname and IP addresses for both servers.

```
cd /etc
vi hosts
vi hostname
```

Add the following entries:

```
Primary server: psql-primary (Host and Hostname) IP Address: 10.6.128.139
Standby server: psql-standby (Host and Hostname) IP Address: 10.6.128.94
```

### pg\_hba.conf Configuration

Edit the `pg_hba.conf` file to allow replication connections.

```
nano /var/lib/pgsql/12/data/pg_hba.conf
```

Add the following lines:

```
# replication privilege.
local replication all peer
host replication all 10.6.128.94/32 trust
host replication all ::1/128 ident
```

## postgresql.conf Configuration

Edit the `postgresql.conf` file to configure replication settings.

```
vi /var/lib/pgsql/12/data/postgresql.conf
```

Add or modify the following settings:

```
# CONNECTIONS AND AUTHENTICATION
listen_addresses = '*'
```

```
# WRITE-AHEAD LOG
wal_level = replica
wal_log_hints = on

# REPLICATION
max_wal_senders = 10
wal_keep_size = 10
wal_receiver_status_interval = 5s
hot_standby_feedback = on
```

## Create Replication User

Create a replication user on the primary server.

```
sudo -u postgres createuser -U postgres repuser -P --replication
```

## Restart PostgreSQL

```
sudo systemctl restart postgresql-12
```

## Standby Server Configuration

## Stop PostgreSQL and Remove Data Directory

```
systemctl stop postgresql-12  
yes | rm -r /var/lib/pgsql/12/data/
```

## Create Data Directory and Set Permissions

```
mkdir /var/lib/pgsql/12/data/  
chown postgres:postgres /var/lib/pgsql/12/data/
```

## Initialize Standby Server from Primary

```
sudo -H -u postgres bash -c 'pg_basebackup --pgdata=/var/lib/pgsql/12/data/ --f
```

## Update postgresql.conf

```
vi /var/lib/pgsql/12/data/postgresql.conf
```

Add or modify the following settings:

[Open in app](#) ↗

Medium

🔍 Search



```
primary_conninfo = 'host=10.6.128.159 port=5432 user=repuser password=xxxxxx'  
primary_slot_name = 'standby1_slot'
```

## Restart PostgreSQL

```
sudo systemctl restart postgresql-12
```

## Monitoring Replication

### On Standby Server

```
su - postgres
psql
\x
select * from pg_stat_wal_receiver;
select * from pg_stat_replication;
```

### On Primary Server

bashCopy [code](#)

```
su - postgres
psql
\x
select * from pg_stat_replication;
```

## Conclusion

You should now have a working streaming replication setup between your primary and standby servers on RHEL. This setup provides high availability and fault tolerance for your PostgreSQL database.



Following

## Written by Oz

149 Followers · 13 Following

Database Administrator 

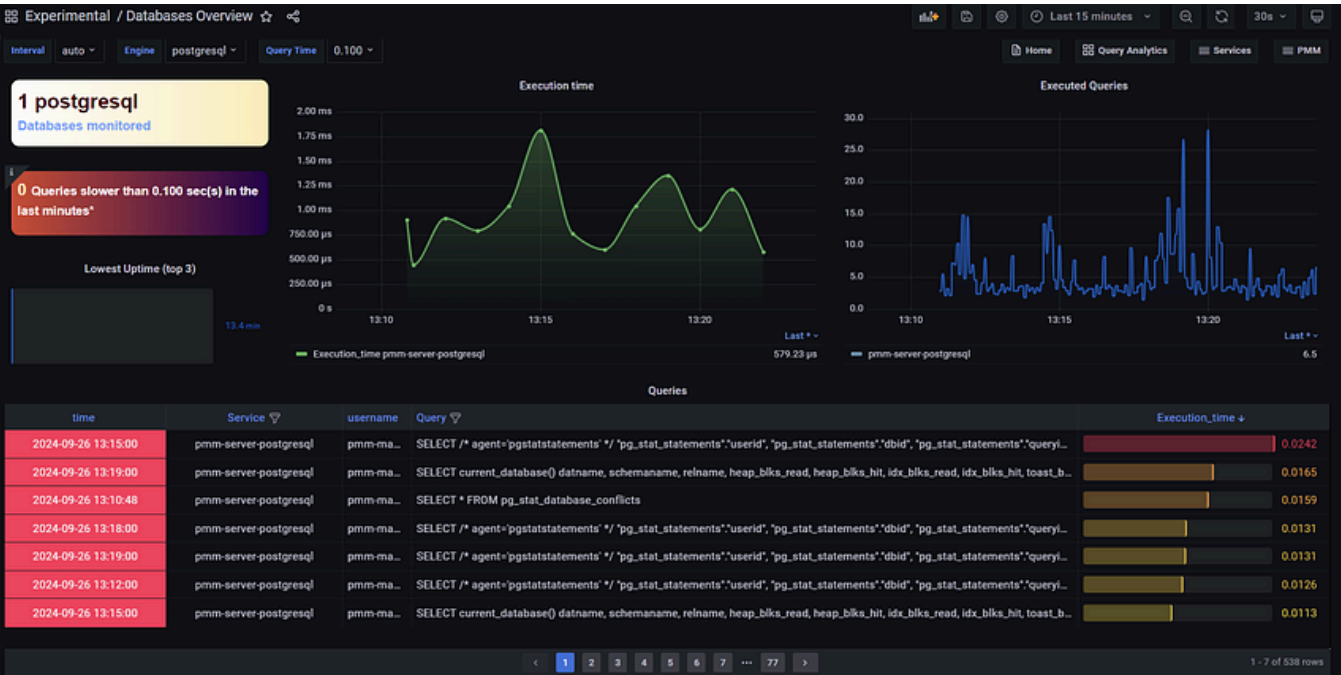
No responses yet



Gvadakte

What are your thoughts?

More from Oz



Installing Percona Monitoring & Management (PMM) with Postgres

Introduction:



Sep 26, 2024




54



1






Oz

## Pluggable Database Command

----- - create pluggable database pdb1 admin user root identified by test123; alter pluggable database...

✦ May 12, 2023



Oz

## RMAN Backup Basic Commands

rman target / rman target sys/password@YDKTST; backup database; backup database format '/backup/path/%d\_%t\_%s.rman'; backup tablespace...

★ May 11, 2023 🖱 1



 Oz

## delete jobs

★ May 8, 2023



See all from Oz

## Recommended from Medium





# podman

## with PostgreSQL

@mehmetozanguven



mehmetozanguven

### Running PostgreSQL with Podman

Instead of running PostgreSQL locally, we can easily run with Podman. Here are the basic steps you should follow.

Mar 28 🖱️ 2



## POSTGRES SQL PERFORMANCE OPTIMIZATION



# PostgreSQL



In Towards Dev by Nakul Mitra

### PostgreSQL Performance Optimization—Cleaning Dead Tuples & Reindexing

Performance optimization is crucial in PostgreSQL to ensure efficient query execution and minimal resource consumption.

Mar 28 🖱️ 1



```
b/postgresql/16/main/*

t patroni

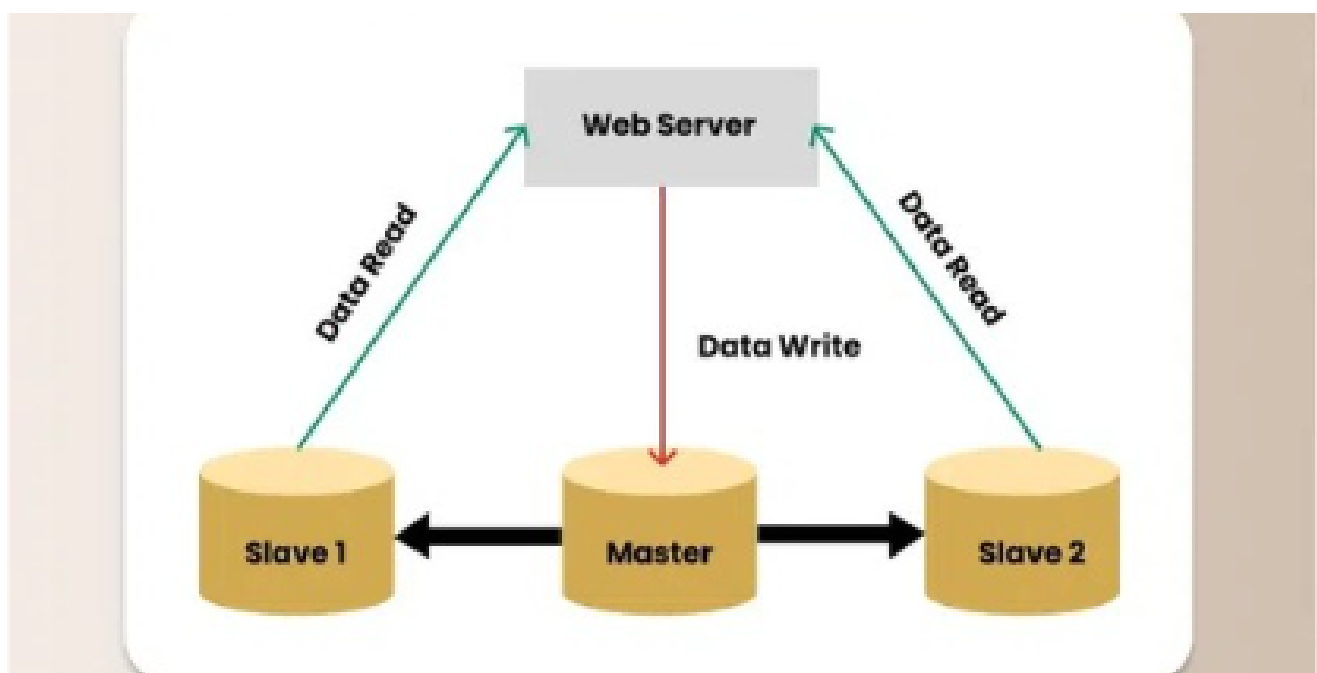
/etc/patroni.yml list
21665717) -----+-----+-----+
Role      | State      | TL | Lag in MB |
-----+-----+-----+-----+
Leader    | running    | 1  |           |
Replica   | streaming  | 1  | 0         |
Replica   | streaming  | 1  | 0         |
-----+-----+-----+-----+
```


Dickson Gathima

## Building a Highly Available PostgreSQL Cluster with Patroni, etcd, and HAProxy

Achieving high availability in PostgreSQL requires the right combination of tools and architecture.

Mar 14 🖱️ 4



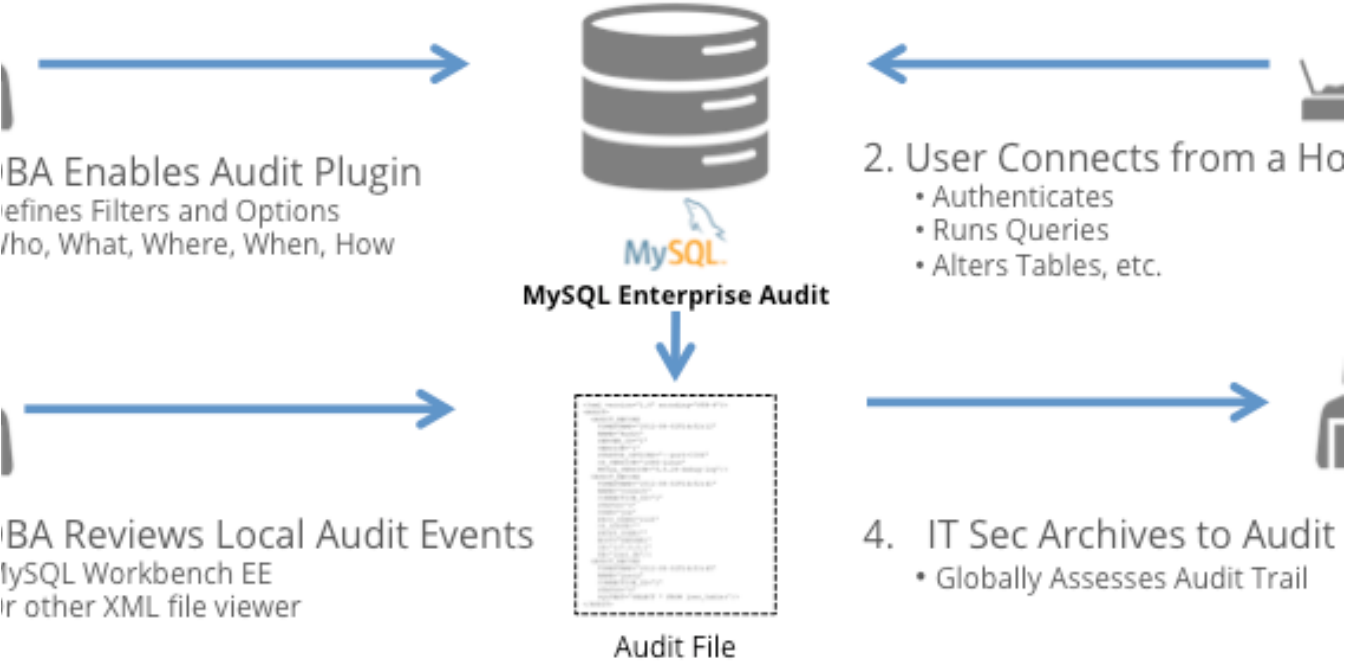
 Ramesh Rs


# How Bitbucket Solved Replication Lag in PostgreSQL

Bit bucket's Database Architecture

Feb 3

 24



 Hari


# How to Enable MYSQL DB audit logs

Step 1: Install the MySQL Enterprise Audit Plugin

★ Nov 12, 2024

 53





Josef Machytka

## Quick and Easy Data Exports to Parquet Format Using DuckDB

The Parquet format has become almost an industry standard for Data Lakes and Data Lakehouses, thanks to its efficiency and compact storage...

Dec 5, 2024   🖱️ 17   💬 1



See more recommendations