

Setup replication using repmgr

Streaming replication is a common PostgreSQL high-availability setup, consisting of a primary server (read-write) and one or more standby replicas (read-only). However, a key limitation of streaming replication is that it doesn't automatically failover to a replica if the primary server goes down, requiring manual intervention from a DBA to promote the replica to read-write mode. This also means the setup needs constant monitoring.

This is where `repmgr` comes into play, addressing this critical issue by automating failover to a replica in the event of a primary server failure, eliminating the need for DBA intervention

Installing repmgr and PostgreSQL

The setup below will be performed on two VMs running Ubuntu 22, with PostgreSQL 16 being used for the configuration.

hostname	ip address
postgresql-db01	10.217.10.6
postgresql-db02	10.217.10.7

Installing PostgreSQL 16

Perform the below steps in both DB servers

1. Import the repository signing key:

```
sudo apt install curl ca-certificates
sudo install -d /usr/share/postgresql-common/pgdg
sudo curl -o /usr/share/postgresql-common/pgdg/apt.postgresql.org.asc --fail
https://www.postgresql.org/media/keys/ACCC4CF8.asc
```

2. Create the repository configuration file:

```
sudo sh -c 'echo "deb [signed-by=/usr/share/postgresql-common/pgdg/apt.postgresql.org.asc]
https://apt.postgresql.org/pub/repos/apt $(lsb_release -cs)-pgdg main" >
/etc/apt/sources.list.d/pgdg.list'
```

3. Update the package lists:

```
sudo apt update
```

4. Install PostgreSQL 12

```
sudo apt -y install postgresql-16
```

baXterm by subscribing to the professional edition here: <https://mobaxterm.mobatek.net/>

```
sudo systemctl status postgresql@16-main.service
```

```
Sep 23 13:50:48 postgresql-db01 systemd[1]: Starting PostgreSQL Cluster 16-main...
Sep 23 13:50:50 postgresql-db01 systemd[1]: Started PostgreSQL Cluster 16-main.
lines 1-17/17 (END)
```

```
sudo nano /etc/postgresql/16/main/postgresql.conf
```

```
listen_addresses = '*'
max_wal_senders = 10

max_replication_slots = 10

wal_level = 'hot_standby' or 'replica' or 'logical'

hot_standby = on

archive_mode = on

archive_command = '/bin/true'

shared_preload_libraries = 'repmgr'
```

On the primary server, create another superuser called `repmgr` and create a database assigned to the `repmgr` user. This database will be used by `repmgr` for storing statistics.

```
create user repmgr;
ALTER USER repmgr WITH SUPERUSER;

create database repmgr with owner repmgr;
```

```
postgres=# create user repmgr;
CREATE ROLE
postgres=# create database repmgr with owner repmgr;
CREATE DATABASE
postgres=# \d
postgres=# \l
```

Name	Owner	Encoding	Locale Provider	Collate	Ctype	ICU Locale	ICU Rules	Access privileges
postgres	postgres	UTF8	libc	en_US.UTF-8	en_US.UTF-8			
repmgr	repmgr	UTF8	libc	en_US.UTF-8	en_US.UTF-8			
template0	postgres	UTF8	libc	en_US.UTF-8	en_US.UTF-8			=c/postgres +
template1	postgres	UTF8	libc	en_US.UTF-8	en_US.UTF-8			postgres=Ct/postgres +
								=c/postgres +
								postgres=Ct/postgres

(4 rows)

```
postgres=#
```

next update the `pg_hba.conf` file with connection allowed in both ipv4 and replication section as follow

```
sudo vi /etc/postgresql/16/main/pg_hba.conf
```

local	replication	repmgr		trust
host	replication	repmgr	127.0.0.1/32	trust
host	replication	repmgr	10.217.10.0/24	trust
local	repmgr	repmgr		trust
host	repmgr	repmgr	127.0.0.1/32	trust

```
host      repmgr      repmgr      10.217.10.0/24      trust
```

```
# Database administrative login by Unix domain socket
local all postgres peer
# TYPE DATABASE USER ADDRESS METHOD
# "local" is for Unix domain socket connections only
local all all peer
local repmgr repmgr trust
# IPv4 local connections:
host all all 127.0.0.1/32 scram-sha-256
host repmgr repmgr 10.217.10.0/24 trust
host repmgr repmgr 127.0.0.1/32 trust
# IPv6 local connections:
host all all ::1/128 scram-sha-256
# Allow replication connections from localhost, by a user with the
# replication privilege.
local replication all peer
host replication all 127.0.0.1/32 scram-sha-256
host replication all ::1/128 scram-sha-256
host replication repmgr 10.217.10.0/24 trust
-- INSERT --
```

Restart PostgreSQL services to get the all configuration loaded

```
sudo systemctl restart postgresql@16-main.service
```

Create a `repmgr.conf` on the master server with the following entire you can place the file in `/etc` directory

```
sudo nano /etc/repmgr.conf
```

```
cluster='failovertest'

node_id=1

node_name=node1

conninfo='host=10.217.10.6 user=repmgr dbname=repmgr connect_timeout=2'

data_directory='/var/lib/postgresql/16/main/'

failover=automatic

promote_command='/usr/bin/repmgr standby promote -f /etc/repmgr.conf --log-to-file'

follow_command='/usr/bin/repmgr standby follow -f /etc/repmgr.conf --log-to-file --upstream-node-id=%n'
```

"Now we will start by registering the primary server with `repmgr` using the following command:

```
repmgr -f /etc/repmgr.conf primary register
```

```

dba@postgresql-db01:~$ repmgr -f /etc/repmgr.conf primary register
WARNING: the following problems were found in the configuration file:
  parameter "cluster" is deprecated and will be ignored
INFO: connecting to primary database...
NOTICE: attempting to install extension "repmgr"
NOTICE: "repmgr" extension successfully installed
NOTICE: primary node record (ID: 1) registered
dba@postgresql-db01:~$

```

Check the status of the cluster by running the following command

```
repmgr -f /etc/repmgr.conf cluster show
```

```

dba@postgresql-db01:~$ repmgr -f /etc/repmgr.conf cluster show
WARNING: the following problems were found in the configuration file:
  parameter "cluster" is deprecated and will be ignored
ID | Name | Role | Status | Upstream | Location | Priority | Timeline | Connection string
-----+-----+-----+-----+-----+-----+-----+-----+-----
1 | node1 | primary | * running | | default | 100 | 1 | host=10.217.10.6 user=repmgr dbname=repmgr connect_timeout=2
dba@postgresql-db01:~$

```

Configuring replica server

1. Create the `repmgr.conf` file and populate it with the following parameters. I have placed the file in the `/etc/` directory.

```
sudo nano /etc/repmgr.conf
```

```

node_id=2

node_name=node2

conninfo='host=10.217.10.7 user=repmgr dbname=repmgr connect_timeout=2'
data_directory='/var/lib/postgresql/16/main/'
failover=automatic

promote_command='/usr/bin/repmgr standby promote -f /etc/repmgr.conf --log-to-file'

follow_command='/usr/bin/repmgr standby follow -f /etc/repmgr.conf --log-to-file --upstream-node-id=%n

```

- 2 . Stop PostgreSQL services in replica server and go to the data directory and remove all the files

```

sudo systemctl stop postgresql@16-main.service
sudo -i

cd /var/lib/postgresql/16/main
rm -rf *

```

```
dba@postgresql-db02:~$ sudo systemctl stop postgresql@16-main.service
dba@postgresql-db02:~$ sudo systemctl status postgresql@16-main.service
o postgresql@16-main.service - PostgreSQL Cluster 16-main
   Loaded: loaded (/lib/systemd/system/postgresql.service; enabled; vendor preset: enabled)
   Active: inactive (dead) since Mon 2024-09-23 15:08:55 UTC; 6s ago
   Process: 5572 ExecStop=/usr/bin/pg_ctlcluster --skip-systemctl-redirect -m fast 16-main stop (code=exited, status=0/SUCCESS)
   Main PID: 5067 (code=exited, status=0/SUCCESS)
   CPU: 705ms

Sep 23 14:47:41 postgresql-db02 systemd[1]: Starting PostgreSQL Cluster 16-main...
Sep 23 14:47:43 postgresql-db02 systemd[1]: Started PostgreSQL Cluster 16-main.
Sep 23 15:08:55 postgresql-db02 systemd[1]: Stopping PostgreSQL Cluster 16-main...
Sep 23 15:08:55 postgresql-db02 systemd[1]: postgresql@16-main.service: Deactivated successfully.
Sep 23 15:08:55 postgresql-db02 systemd[1]: Stopped PostgreSQL Cluster 16-main.
dba@postgresql-db02:~$ sudo -i
root@postgresql-db02:~# cd /var/lib/postgresql/16/main
root@postgresql-db02:/var/lib/postgresql/16/main# rm -rf *
root@postgresql-db02:/var/lib/postgresql/16/main# ls
root@postgresql-db02:/var/lib/postgresql/16/main#
```

3. now we can perform dry run which test our standby server configuration before we can add it to the cluster

```
sudo su - postgres
```

```
repmgr -h 10.217.10.6 -U repmgr -d repmgr -f /etc/repmgr.conf standby clone --dry-run
```

```
root@postgresql-db02:/var/lib/postgresql/16/main# sudo su - postgres
postgres@postgresql-db02:~$ repmgr -h 10.217.10.6 -U repmgr -d repmgr -f /etc/repmgr.conf standby clone --dry-run
NOTICE: destination directory "/var/lib/postgresql/16/main" provided
INFO: connecting to source node
DETAIL: connection string is: host=10.217.10.6 user=repmgr dbname=repmgr
DETAIL: current installation size is 29 MB
INFO: "repmgr" extension is installed in database "repmgr"
INFO: replication slot usage not requested; no replication slot will be set up for this standby
INFO: parameter "max_wal_senders" set to 10
NOTICE: checking for available walsenders on the source node (2 required)
INFO: sufficient walsenders available on the source node
DETAIL: 2 required, 10 available
NOTICE: checking replication connections can be made to the source server (2 required)
INFO: required number of replication connections could be made to the source server
DETAIL: 2 replication connections required
WARNING: data checksums are not enabled and "wal_log_hints" is "off"
DETAIL: pg_rewind requires "wal_log_hints" to be enabled
NOTICE: standby will attach to upstream node 1
HINT: consider using the -c/--fast-checkpoint option
INFO: would execute:
pg_basebackup -l "repmgr base backup" -D /var/lib/postgresql/16/main -h 10.217.10.6 -p 5432 -U repmgr -X stream
INFO: all prerequisites for "standby clone" are met
postgres@postgresql-db02:~$
```

4. Start cloning of the data directory from primary server by running the below command

```
repmgr -h 10.217.10.6 -U repmgr -d repmgr -f /etc/repmgr.conf standby clone
```

```
postgres@postgresql-db02:~$ repmgr -h 10.217.10.6 -U repmgr -d repmgr -f /etc/repmgr.conf standby clone
NOTICE: destination directory "/var/lib/postgresql/16/main" provided
INFO: connecting to source node
DETAIL: connection string is: host=10.217.10.6 user=repmgr dbname=repmgr
DETAIL: current installation size is 29 MB
INFO: replication slot usage not requested; no replication slot will be set up for this standby
NOTICE: checking for available walsenders on the source node (2 required)
NOTICE: checking replication connections can be made to the source server (2 required)
WARNING: data checksums are not enabled and "wal_log_hints" is "off"
DETAIL: pg_rewind requires "wal_log_hints" to be enabled
INFO: checking and correcting permissions on existing directory "/var/lib/postgresql/16/main"
NOTICE: starting backup (using pg_basebackup)...
HINT: this may take some time; consider using the -c/--fast-checkpoint option
INFO: executing:
pg_basebackup -l "repmgr base backup" -D /var/lib/postgresql/16/main -h 10.217.10.6 -p 5432 -U repmgr -X stream
NOTICE: standby clone (using pg_basebackup) complete
NOTICE: you can now start your PostgreSQL server
HINT: for example: pg_ctl -D /var/lib/postgresql/16/main start
HINT: after starting the server, you need to register this standby with "repmgr standby register"
postgres@postgresql-db02:~$
```

5. start PostgreSQL 16 services

```
systemctl start postgresql@16-main.service
```

- Update the `postgresql.conf` and `pg_hba.conf` files to match the configuration parameters used on the primary server.
- Register the standby server to the cluster

```
repmgr -f /etc/repmgr.conf standby register
```

```
postgres@postgresql-db02:~$ repmgr -f /etc/repmgr.conf standby register
INFO: connecting to local node "node2" (ID: 2)
INFO: connecting to primary database
WARNING: --upstream-node-id not supplied, assuming upstream node is primary (node ID: 1)
INFO: standby registration complete
NOTICE: standby node "node2" (ID: 2) successfully registered
postgres@postgresql-db02:~$
```

```
postgres@postgresql-db02:~$ repmgr -f /etc/repmgr.conf cluster show
```

ID	Name	Role	Status	Upstream	Location	Priority	Timeline	Connection string
1	node1	primary	* running		default	100	1	host=10.217.10.6 user=repmgr dbname=repmgr connect_timeout=2
2	node2	standby	running	node1	default	100	1	host=10.217.10.7 user=repmgr dbname=repmgr connect_timeout=2

```
postgres@postgresql-db02:~$
```

- check the replication by creating testdb in primary and check weather it get replicated to the standby server

```
sudo -u postgres psql
```

```
create database testdb;
```

```
dba@postgresql-db01:~$ sudo -u postgres psql
[sudo] password for dba:
psql (16.4 (Ubuntu 16.4-1.pgdg22.04+1))
Type "help" for help.

postgres=# create database testdb;
CREATE DATABASE
postgres=# \l
```

Name	Owner	Encoding	Locale Provider	Collate	Ctype	ICU Locale	ICU Rules	Access privileges
postgres	postgres	UTF8	libc	en_US.UTF-8	en_US.UTF-8			
repmgr	repmgr	UTF8	libc	en_US.UTF-8	en_US.UTF-8			
template0	postgres	UTF8	libc	en_US.UTF-8	en_US.UTF-8			=c/postgres +
template1	postgres	UTF8	libc	en_US.UTF-8	en_US.UTF-8			postgres=CTc/postgres +
testdb	postgres	UTF8	libc	en_US.UTF-8	en_US.UTF-8			=c/postgres +
								postgres=CTc/postgres

```
(5 rows)
postgres=#
```

```
postgres@postgresql-db02:~$ psql
psql (16.4 (Ubuntu 16.4-1.pgdg22.04+1))
Type "help" for help.

postgres=# \l
```

Name	Owner	Encoding	Locale Provider	Collate	Ctype	ICU Locale	ICU Rules	Access privileges
postgres	postgres	UTF8	libc	en_US.UTF-8	en_US.UTF-8			
repmgr	repmgr	UTF8	libc	en_US.UTF-8	en_US.UTF-8			
template0	postgres	UTF8	libc	en_US.UTF-8	en_US.UTF-8			=c/postgres +
template1	postgres	UTF8	libc	en_US.UTF-8	en_US.UTF-8			postgres=CTc/postgres +
testdb	postgres	UTF8	libc	en_US.UTF-8	en_US.UTF-8			=c/postgres +
								postgres=CTc/postgres

```
(5 rows)
postgres=#
```

I would like to extend my thanks to EDB for their excellent guide. You can view it yourself by visiting the following URL

<https://www.enterprisedb.com/postgres-tutorials/how-implement-repmgr-postgresql-automatic-failover>